
Artificial Intelligence
A Modern Approach
Third Edition
Stuart Russell & Peter Norvig

Chapter 3
Solving Problems
by Searching

3.1 Problem-Solving Agents

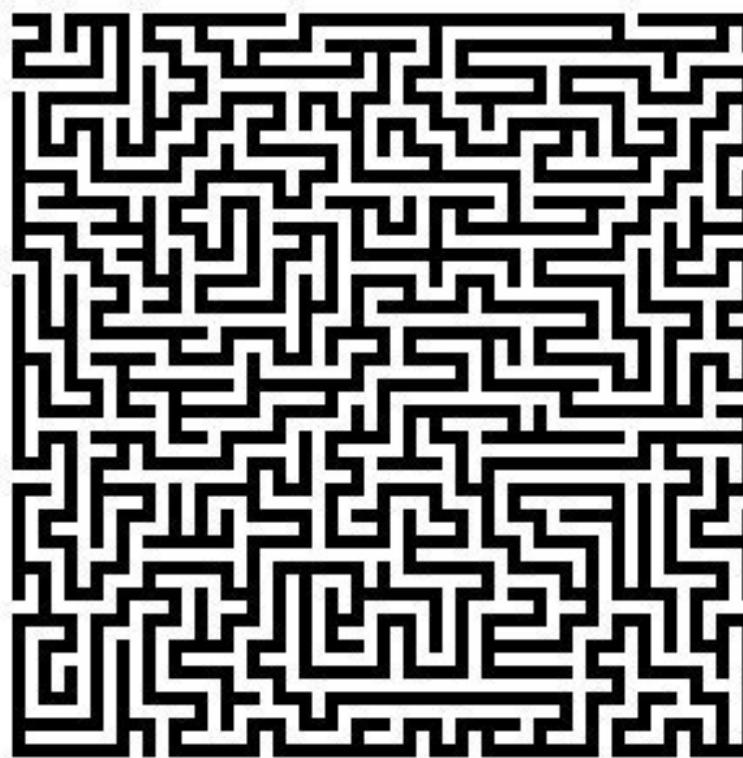
- ❖ Formulate a real-world problem to a **search** problem
- ❖ A **simple problem-solving** agent
 - ❖ **Formulate** a goal and a problem
 - ❖ **Search** for a **sequence of actions** that solves the problem
 - ❖ **Execute** the **actions** one by one

function SIMPLE-PROBLEM-SOLVING-AGENT(*percept*) **returns** an action
persistent: *seq*, an action sequence, initially empty
 state, some description of the current world state
 goal, a goal, initially null
 problem, a problem formulation

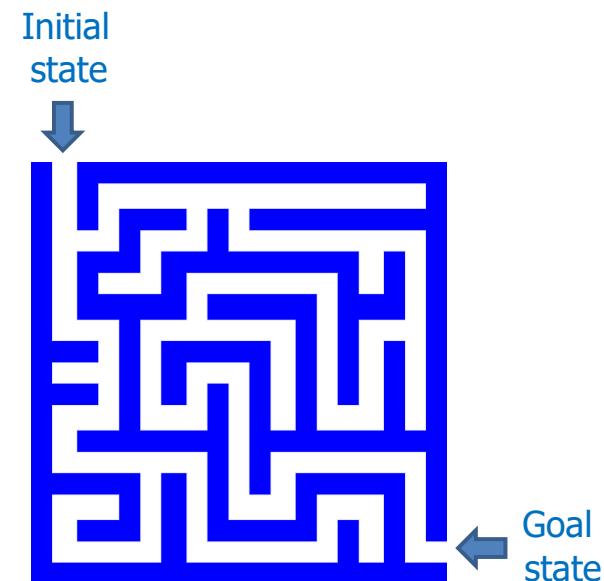
state \leftarrow UPDATE-STATE(*state*, *percept*)
if *seq* is empty **then**
 goal \leftarrow FORMULATE-GOAL(*state*) from “current state” to the “goal state”
 problem \leftarrow FORMULATE-PROBLEM(*state*, *goal*)
 seq \leftarrow SEARCH(*problem*)
 if *seq* = failure **then return** a null action
action \leftarrow FIRST(*seq*) action sequence = {*a*₁, *a*₂, *a*₃, ..., *a*_{*n*}}
seq \leftarrow REST(*seq*) action sequence = {*a*₂, *a*₃, ..., *a*_{*n*}}
return *action*

Search

- ❖ Solving problems by **searching**



- ❖ **Search**
 - ❖ Consider the problem of designing **goal-based agents** in **observable, deterministic, discrete environments**
 - ❖ **Solution** : a fixed sequence of actions
 - ❖ **Search** : the process of looking for the sequence of actions that reaches the goal



- ❖ Search problem components

- ❖ States
- ❖ Initial state
- ❖ Actions
- ❖ Transition model $f: S \times A \rightarrow S$

- ❖ What is the **result** of performing a given action in a given state?

- ❖ Goal state

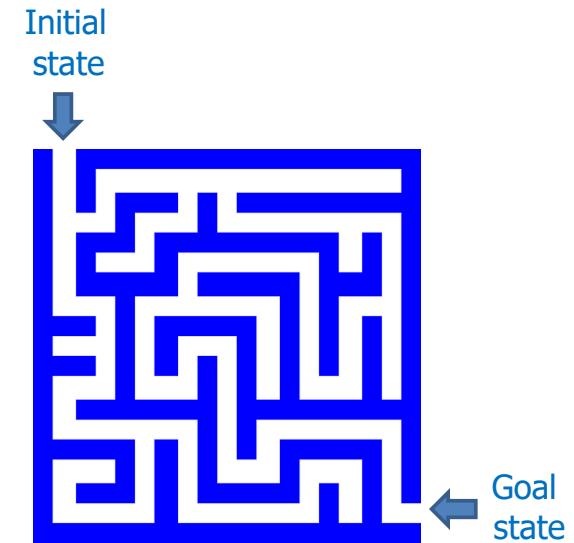
May be one goal or multiple goals

- Find an optimal path to reach the one goal
- Find an optimal goal of the multiple goals

- ❖ Path cost

- ❖ Assume that it is a **sum of nonnegative step costs**

- ❖ The **optimal solution** is the sequence of actions that gives the **lowest path cost** for reaching the goal

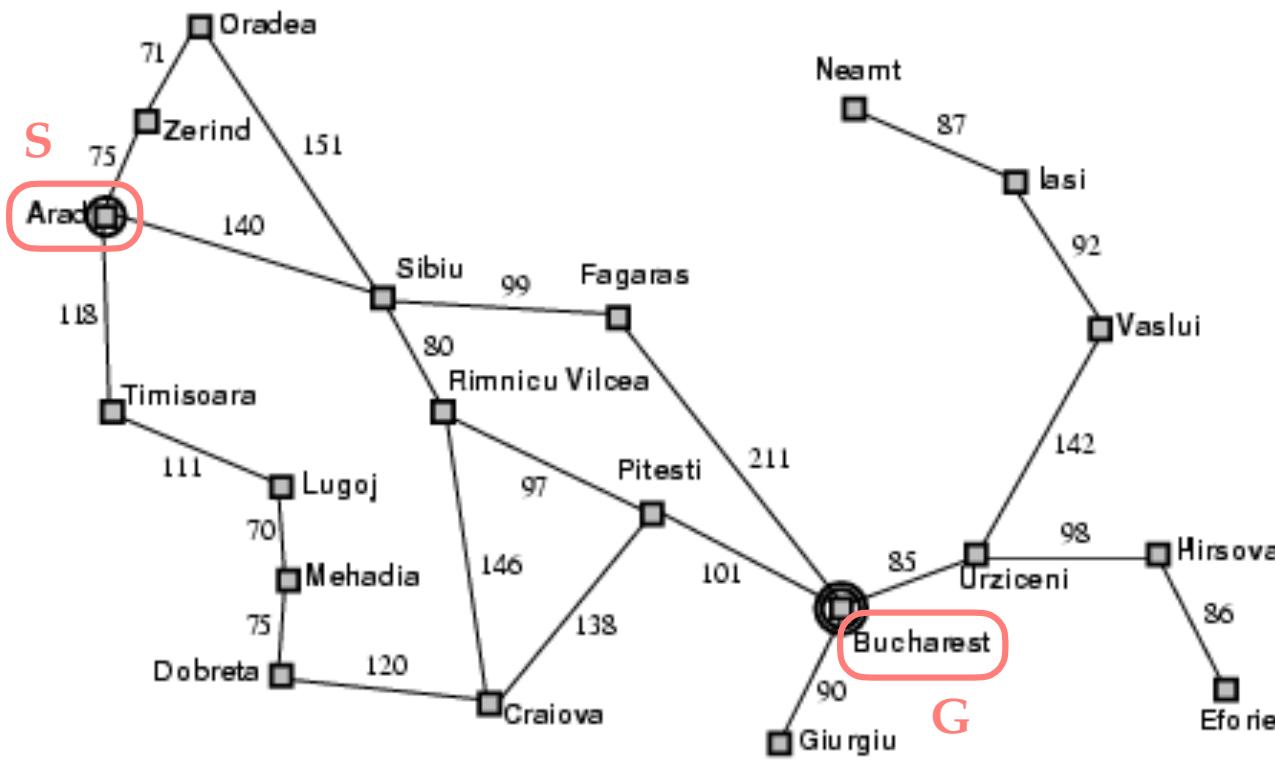


3.2 Example Problems

- ❖ Example: Romania
- ❖ Example: Vacuum world
- ❖ Example: 8-puzzle
- ❖ Example: 8-Queen Puzzle
- ❖ Real-world problems

Example: Romania

- ❖ On vacation in Romania
 - ❖ Currently in Arad
 - ❖ Flight leaves tomorrow from Bucharest



- ❖ **Formulate goal**

- ❖ Be in Bucharest

- ❖ **Formulate problem**

- ❖ States: various cities with initial state: Arad

- ❖ ***Actions (s)*: set of actions available in state s** [go from one city to another]

- ❖ Transition model

$$f: S \times A \rightarrow S$$

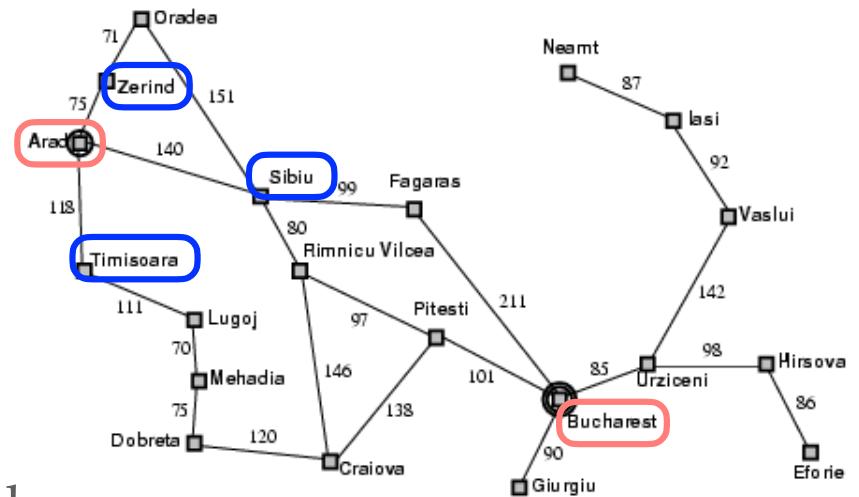
- ❖ ***Result(s,a)* = state that results from action a in state s**
- ❖ Alternatively: **successor function: $S(s) =$ set of action-state pairs**

e.g., $S(Arad) = \{<Arad \rightarrow Zerind, Sibiu, Timisoara>, \dots\}$

- ❖ [If you go from city A to city B, you end up in city B]

- ❖ **Goal state (or Goal test: set of states)**

- ❖ e.g., $s = \text{"at Bucharest"}$





- ❖ **Path cost** (additive)

step cost: c_1, c_2, c_3
path cost: $c_1 + c_2 + c_3$

- ❖ e.g., sum of **distances** (edge costs), number of actions executed, etc.
- ❖ $c(x,a,y)$ is the **step cost**, assumed to be ≥ 0 , step cost is summed to yield path cost

- ❖ **Find solution**

- ❖ **solution** = sequence of actions leading from initial state to a goal state
- ❖ e.g. sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

- ❖ **State space**

- ❖ The **initial state, actions, and transition model** define the **state space** of the problem
 - ❖ The **set of all states reachable** from initial state by any sequence of actions
 - ❖ Can be represented as a **directed graph** where the **nodes** are **states** and **links** between **nodes** are **actions**

- ❖ Select a state space

- ❖ Real world is absurdly complex

- ❖ State space must be **abstracted** for problem solving

- ❖ (Abstract) state \leftarrow set of **real** states

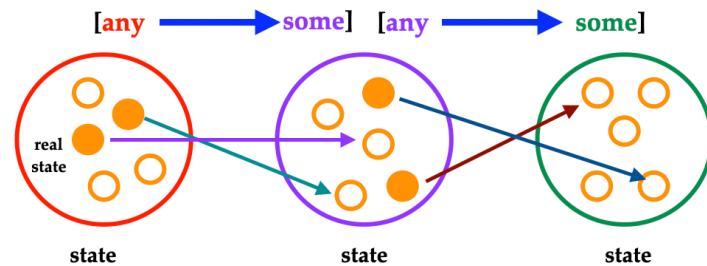
- ❖ (Abstract) action \leftarrow complex combination of **real** actions

- ❖ e.g., "Arad \rightarrow Zerind" represents a complex set of possible routes, detours, rest stops, etc.

- ❖ For guaranteed **realizability**, **any** real state "in Arad" must get to **some** real state "in Zerind"

- ❖ (Abstract) solution \leftarrow set of **real** paths that are solutions in the real world

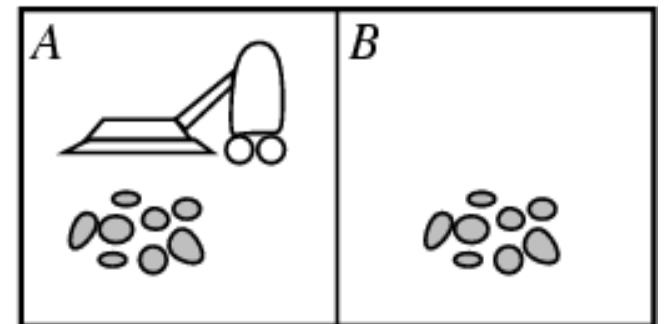
- ❖ Each abstract action should be "easier" than the original problem



- ❖ Environments types
 - ❖ Static / dynamic
 - ❖ Static: no attention to changes in environment
 - ❖ Observable / partially observable / unobservable
 - ❖ Observable: it knows initial state
 - ❖ Deterministic / stochastic
 - ❖ Deterministic: no new percepts were necessary, we can predict the future perfectly given our actions
 - ❖ Discrete / continuous
 - ❖ Discrete: we can enumerate all possibilities

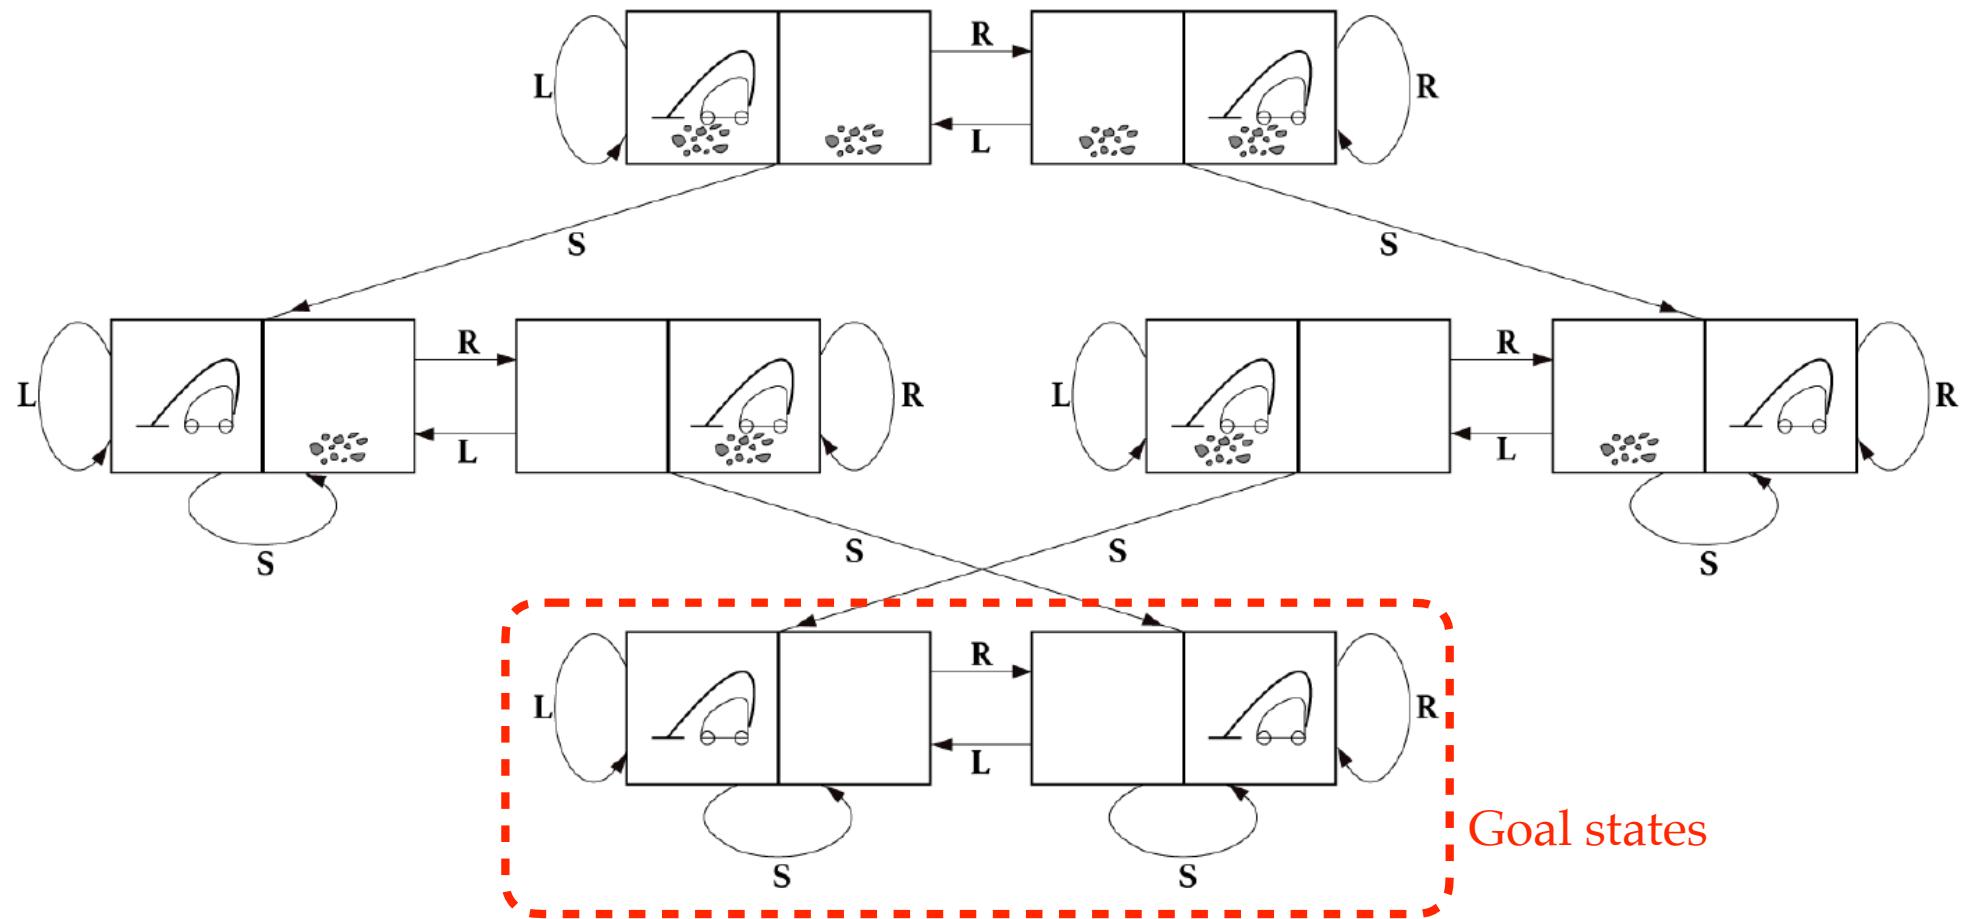
Example: Vacuum World

- ❖ **States:** dirt and robot location
- ❖ **Initial state:** any one of the above states. (ignore dirt amounts etc.)
- ❖ **Actions:** Left, Right, Suck, NoOp
- ❖ **Transition model:** the next figure
- ❖ **Goal test:** no dirt at all locations
- ❖ **Path cost:** 1 per action (0 for NoOp)



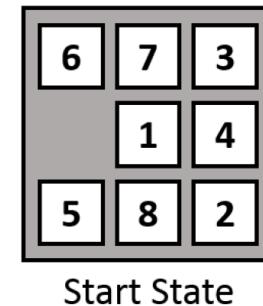
❖ Vacuum World State Space Graph

L: Left
R: Right
S: Suck

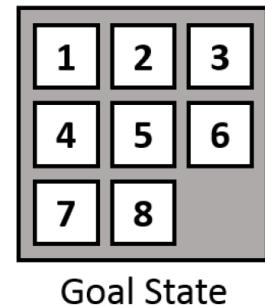


Example: 8-puzzle

- ❖ **Initial state:** the left figure (given)
- ❖ **States** (locations of tiles)
 - ❖ 8-puzzle (3×3 board): 362,880 states ($= 9!$)
 - ❖ 15-puzzle (4×4 board): 2.6 trillion states ($= 16!$)
 - ❖ solved optimally in a few milliseconds
 - ❖ 24-puzzle (5×5 board): 10^{25} trillion states ($= 25!$)
 - ❖ solved optimally in several hours or less time



Start State

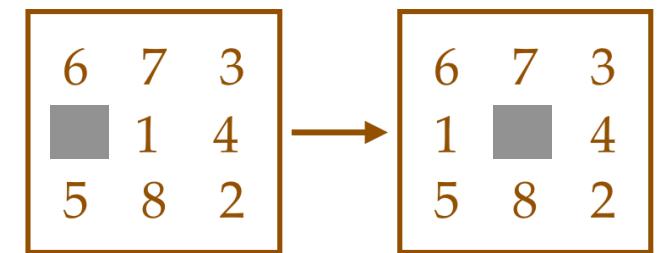


Goal State

6	7	3
1	4	
5	8	2

Start State

- ❖ **Actions:** move blank left, right, up, down (instead of 1 up, 6 down, ...)
- ❖ **Transition model:** common sense
- ❖ **Goal test:** the right figure (given)
- ❖ **Path cost:** 1 per move
- ❖ Note: optimal solution of sliding-block n -puzzle is **NP -hard**



1	2	3
4	5	6
7	8	

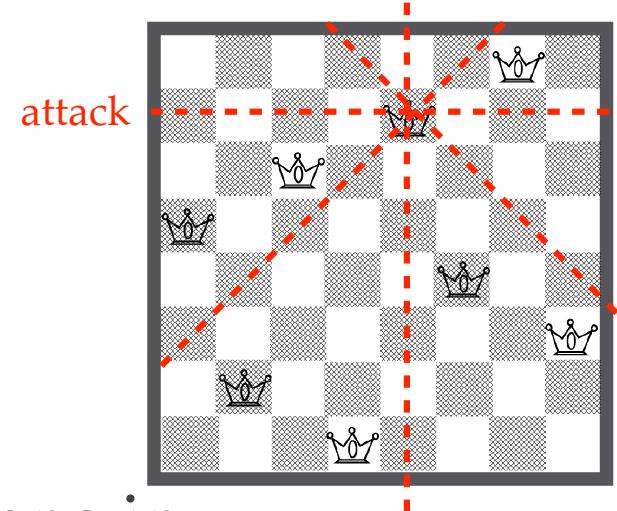
Goal State

NP-hardness (non-deterministic polynomial-time hardness): In computational complexity theory, NP-Hard is the defining property of a class of problems that are informally "at least as hard as the hardest problems in NP".

Example: 8-Queen Puzzle

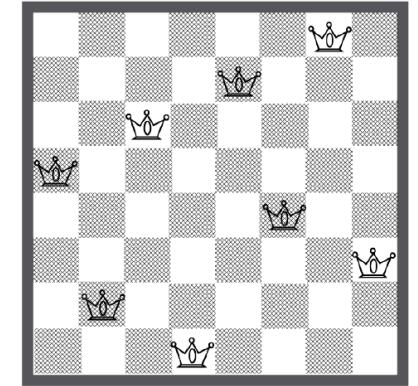
- ❖ **States:**

- ❖ [Type1] Any arrangement of $n \leq 8$ queens, or
 - ❖ [Type2] Arrangements of $n \leq 8$ queens in leftmost n columns, 1 per column, such that no queen attacks any other (BETTER!!)
- ❖ **Initial state:** no queen on the board

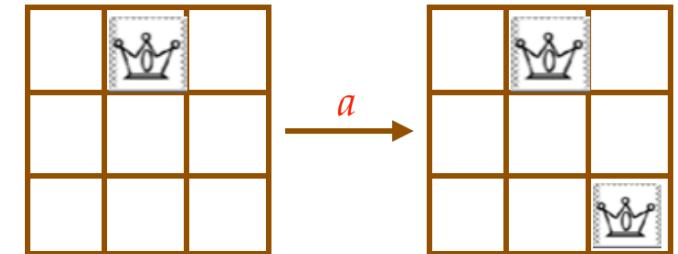


❖ Actions:

- ❖ [Type1] Add queen to **any empty square**
- ❖ [Type2] Add queen to **each column** from left to right (BETTER!)
- ❖ [Type3] Add queen to **leftmost empty square**, such that it is not attacked by other queens (BEST!!).



❖ Transition model:



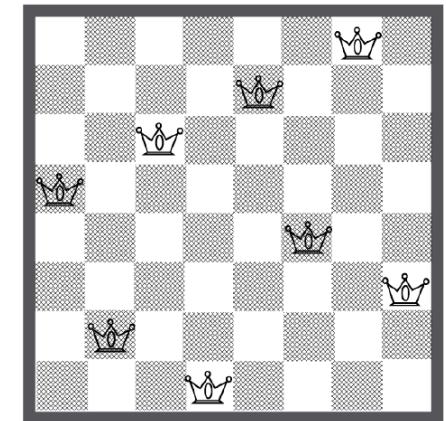
- ❖ Returns the board with a queen added to the specified square
- ❖ Goal test: 8 queens are on the board, none attacked
- ❖ Path cost: number of trials (1 per move)

❖ [Type1] States: any 0~8 queens on the board

0 queen 8 queens on the board

❖ State space: $C_0^{64} + C_1^{64} + C_2^{64} + \dots + C_8^{64} \simeq 5.1 \times 10^9$

❖ Solution space: $64 \cdot 63 \cdot \dots \cdot 57 \simeq 1.8 \times 10^{14}$

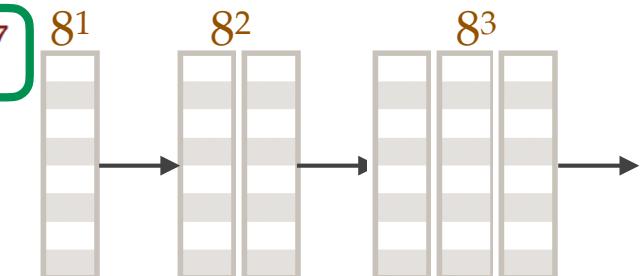


❖ [Type2] States: one queen per column (left → right)

0 queen 8 queens on the left n columns

❖ State space: $8^0 + 8^1 + 8^2 + \dots + 8^8 \simeq 1.9 \times 10^7$

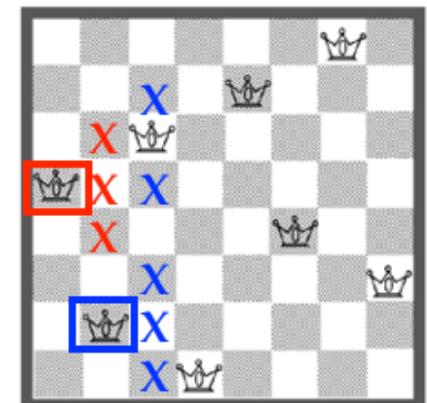
❖ Solution space: $8^8 \simeq 1.6 \times 10^7$



❖ [Type3] States: all possible arrangements of n ($0 \leq n \leq 8$) queens at leftmost n columns with no queens attacked

❖ Actions: Add a queen to the next column with no queens attacked, or backtrack if attacked

❖ State space: 2057



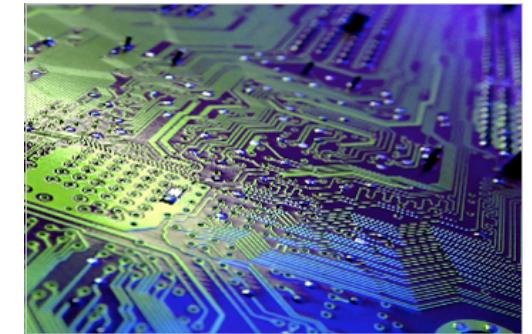
$8 \times (8-3) \times (8-5) \times \dots$

Real-World Problems

- ❖ **Rout finding**
 - ❖ Routing video streams in **computer networks**
 - ❖ **Military** operations planning
 - ❖ **Airline** travel-planning systems
- ❖ **Touring problems**
 - ❖ Closely related to route-finding problems, but with an important difference, e.g. **visit every city at least once**, starting and ending in the same location.

❖ VLSI layout

- ❖ Position millions of components (transistors) and connections on a chip to
 - ❖ **Minimize** area
 - ❖ **Minimize** circuit delays
 - ❖ **Minimize** stray capacitances
 - ❖ **Maximize** manufacturing yield



- ❖ **Robot navigation**
 - ❖ Rather than following a discrete set of routes, a robot can move in a **continuous space** with an infinite set of possible actions and states
- ❖ **Automatic assembly sequencing**
 - ❖ Find an **order** in which to assemble the parts of some object

3.3 Searching for Solutions

- ❖ Search can be classified into **uninformed** search and **informed** search
 - ❖ **Uninformed search**
 - ❖ Doesn't predict **what might happen** if I do action A and action B?
 - ❖ The only information is **how much does it cost** if I do action A, and how much does it cost if I do action B?
 - ❖ **Informed search or heuristic search**
 - ❖ Use **problem-specific knowledge** beyond the definition of the problem itself — can find solutions more efficiently
 - ❖ A **heuristic** is a technique designed for solving a problem more quickly when classic methods are too slow, or for finding an **approximate solution** when classic methods fail to find any exact solution

❖ Tree search algorithms

- Tree search doesn't check repeated states (may have infinite states)
- Graph search checks repeated states and is more efficient (should have sufficient space)

1. Begin at the **start node**

2. Explore state space by generating a list of **all possible successors** of already-explored states (a.k.a. **expanding states**)

3. Maintain a **frontier** or a list of **unexpanded states**

4. At each step, **pick a state** from the **frontier** to expand

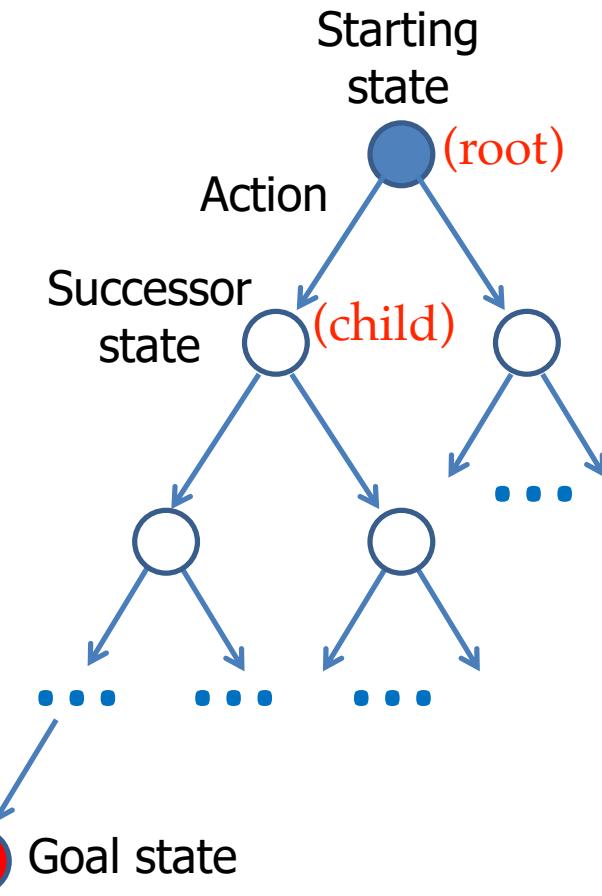
5. Keep going and **evaluate** every generated state

Is it a goal state?

6. Try to expand as **few states** as possible

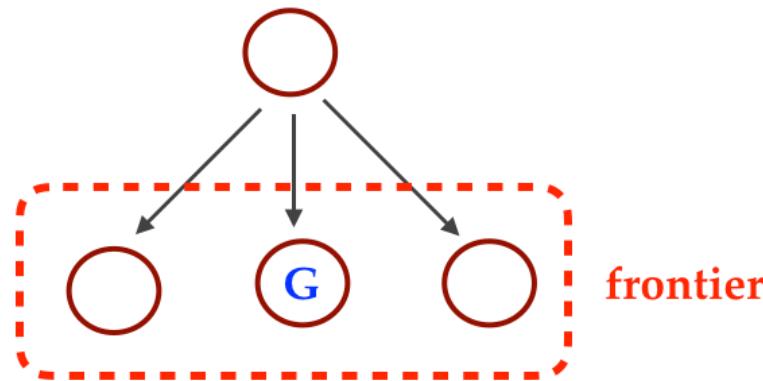
❖ Search tree

- ❖ “What if” tree of possible **actions** and **outcomes**
- ❖ The **root node** corresponds to the **starting state**
- ❖ The **children** of a node correspond to the **successor states** of that node’s state
- ❖ A **path** through the tree corresponds to a **sequence of actions**
- ❖ A **solution** is a **path ending** in the **goal state**



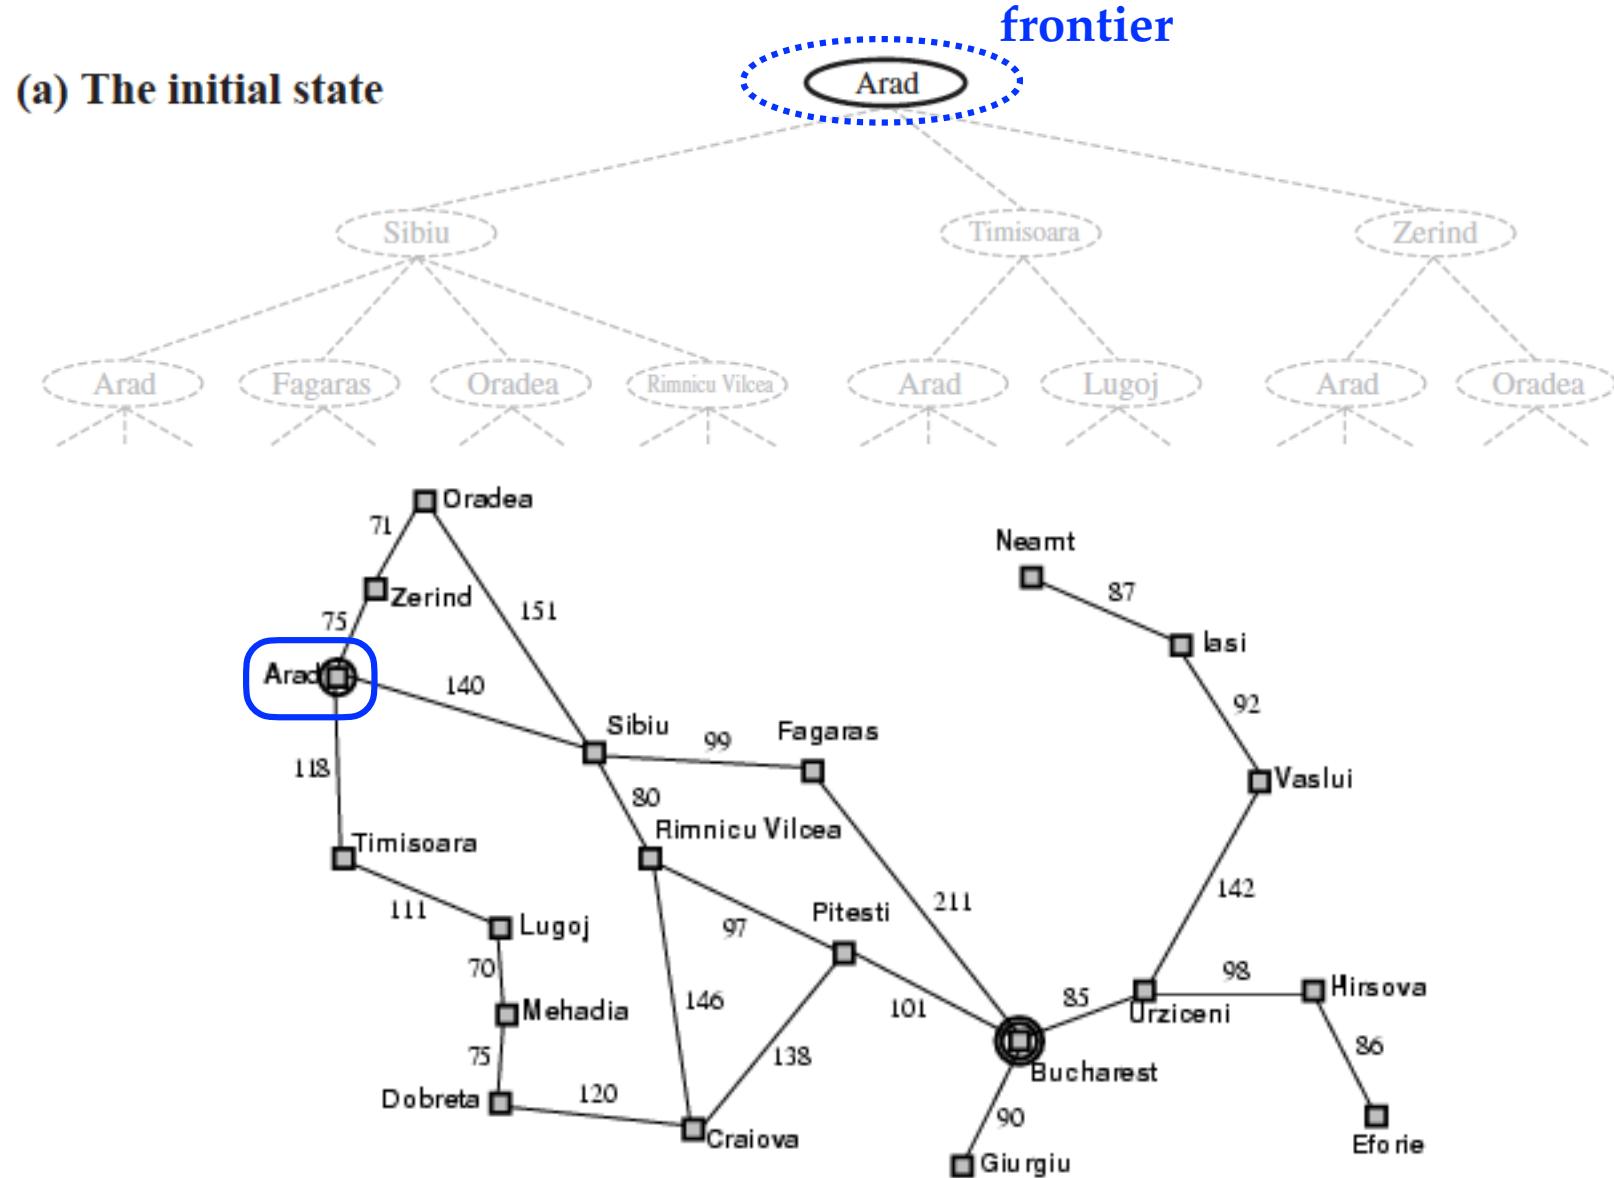
```

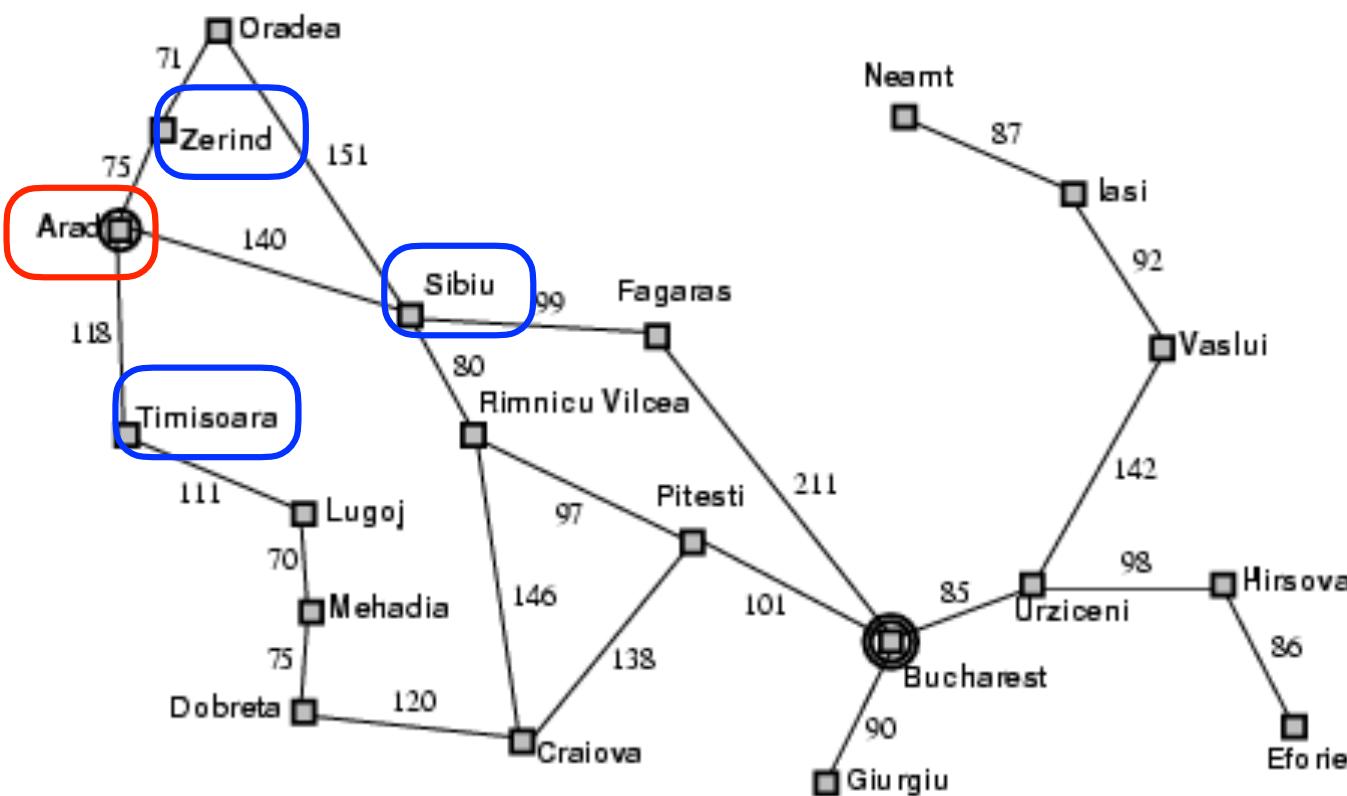
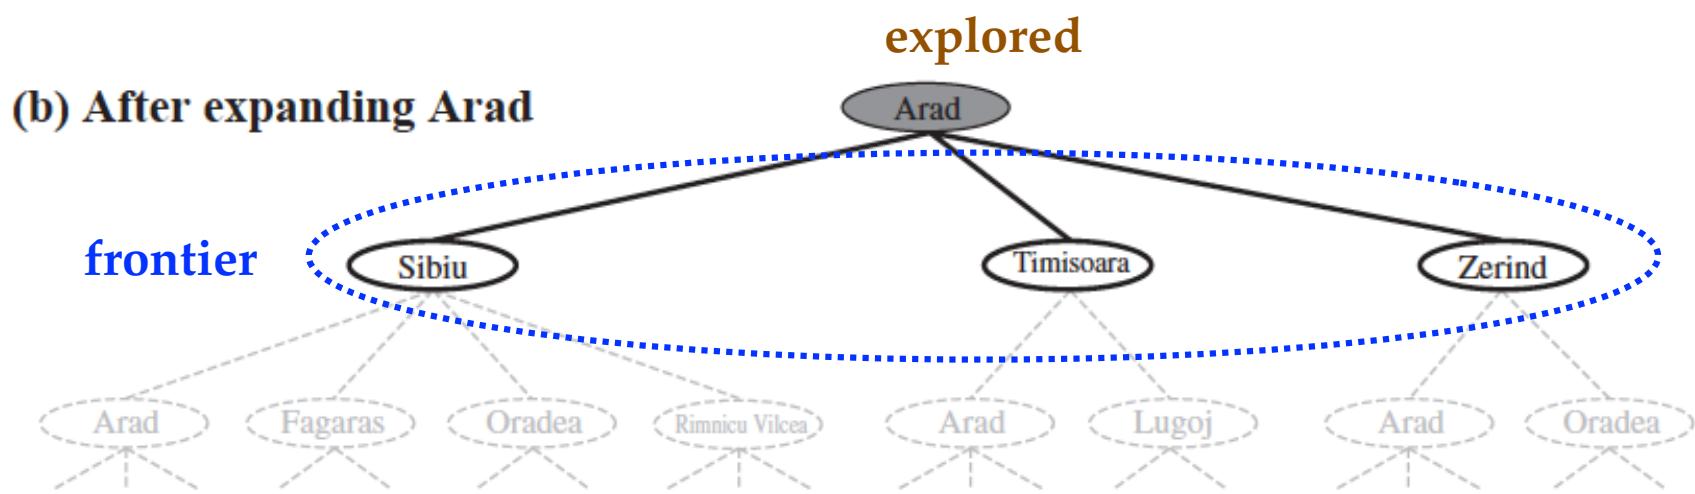
function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier
  
```



- Usually, we check whether a node is a “goal” (goal test) when it is **popped out** from the frontier
- This is because sometimes we cannot guarantee **optimal solution** if we do the “goal test” at the time we put a state into a frontier

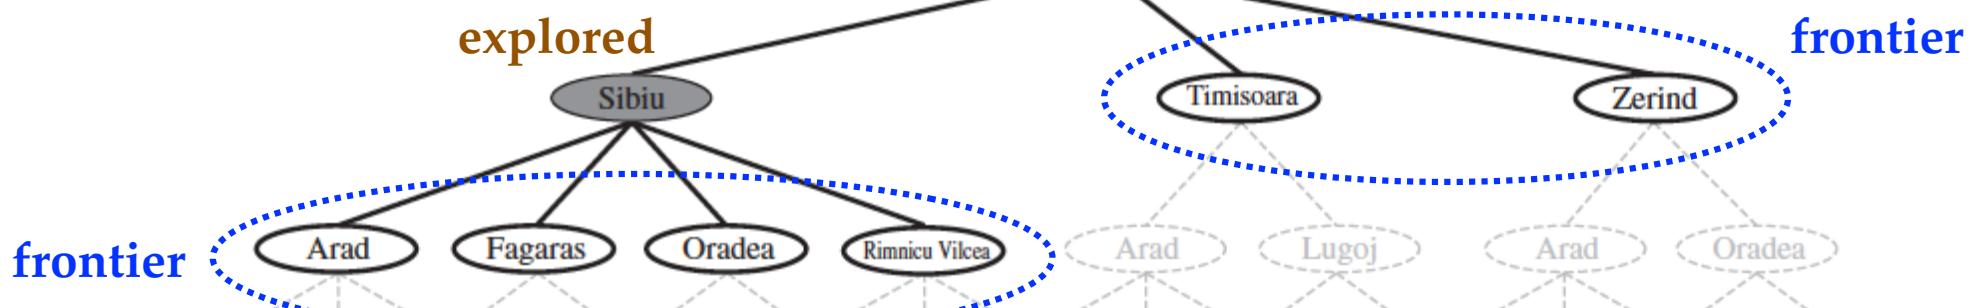
Tree Search Example



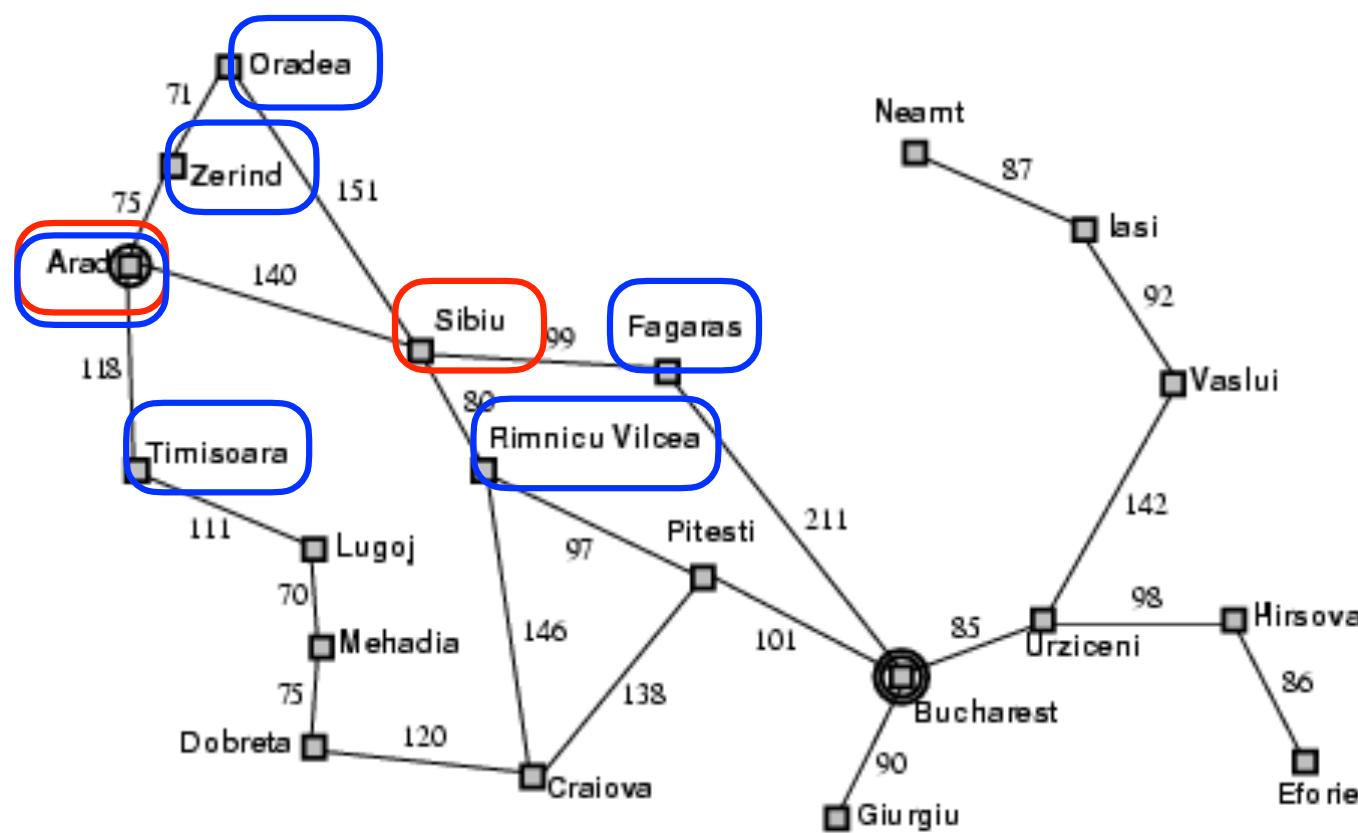


explored

(c) After expanding Sibiu

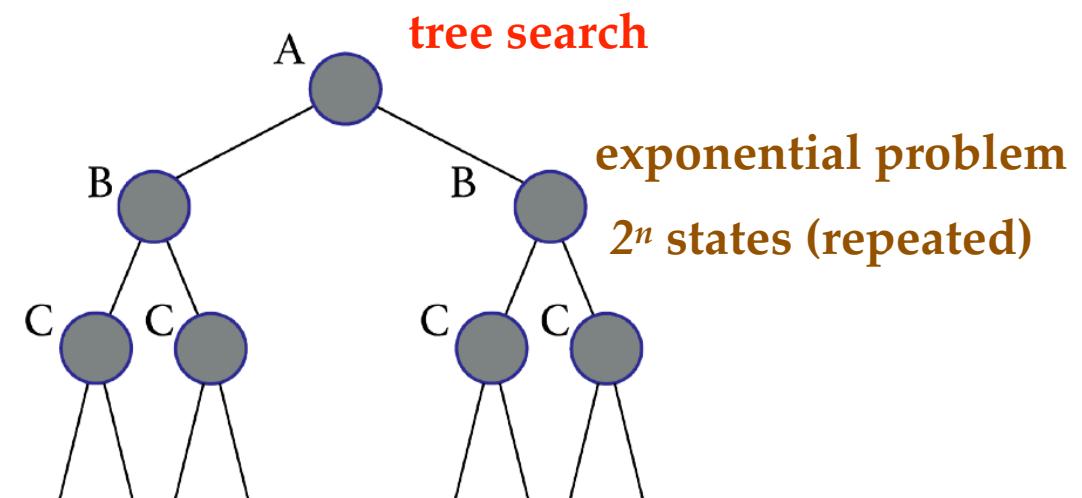
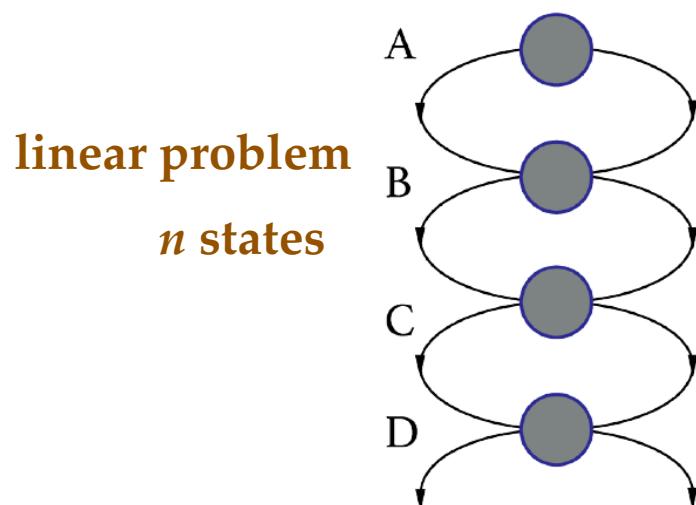


frontier



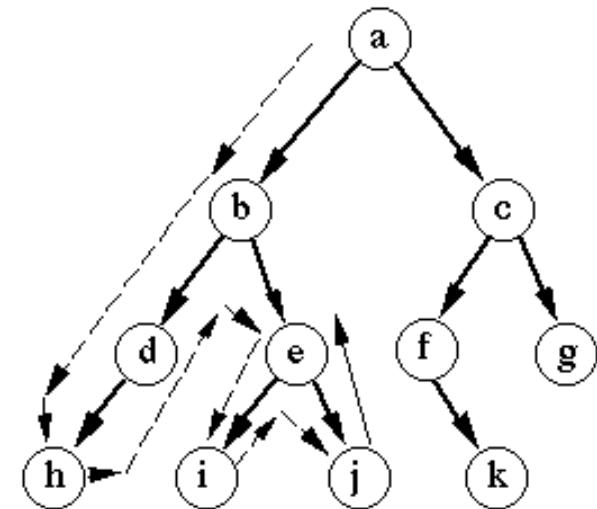
Repeated States in Tree Search

- ❖ Failure to detect repeated states can turn a **linear** problem into an **exponential** problem
- ❖ Use a **queue** to record **explored** states
- ❖ For fast detection of repeated states, **hashing** techniques are usually adopted (**hash table**)

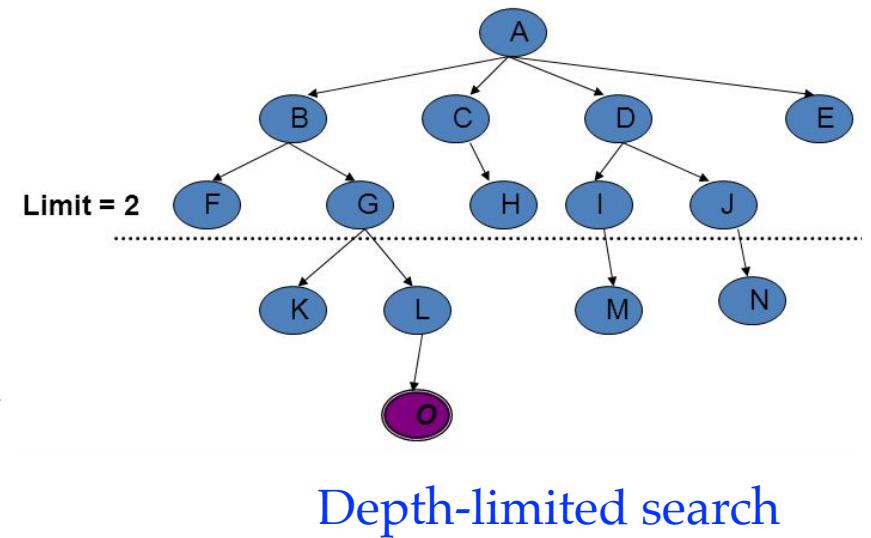


Solutions to Repeated States

- ❖ Graph search
 - ❖ Faster (efficient), but memory inefficient
 - ❖ Never generate a state generated before
 - ❖ Must keep track of **all possible states** (use a lot of memory)
 - ❖ e.g., 8-puzzle problem has $9! = 362,880$ states
 - ❖ Approximation for **DFS** (Depth First Search): only **avoid repeated states** in its **(limited) memory**
 - ❖ Approximation for **DLS** (Depth Limited Search): **avoid infinite loops** by checking path back to root

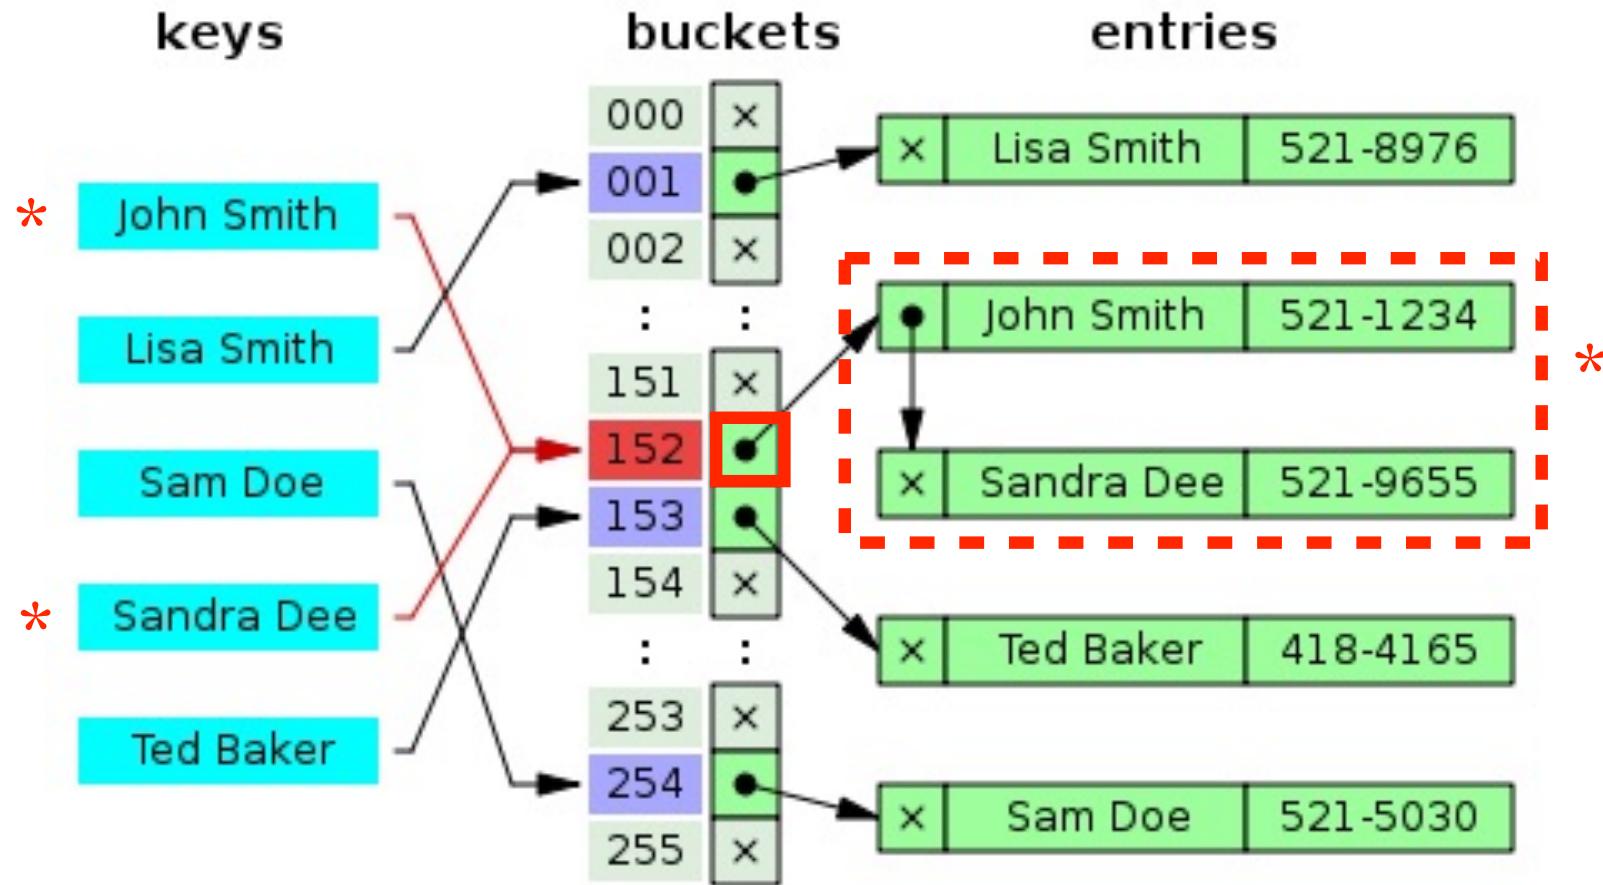


Depth-first search



Depth-limited search

- ❖ “Visited (repeated)?” test usually implemented as a hash table



“John Smith” and “Sandra Dee” are different but have the same hash value

function GRAPH-SEARCH(*problem*) **returns** a solution, or failure

initialize the frontier using the initial state of *problem*

initialize the explored set to be empty

loop do

if the frontier is empty **then return** failure

choose a leaf node and remove it from the frontier

if the node contains a goal state **then return** the corresponding solution

add the node to the explored set (explored set is used to check if there is

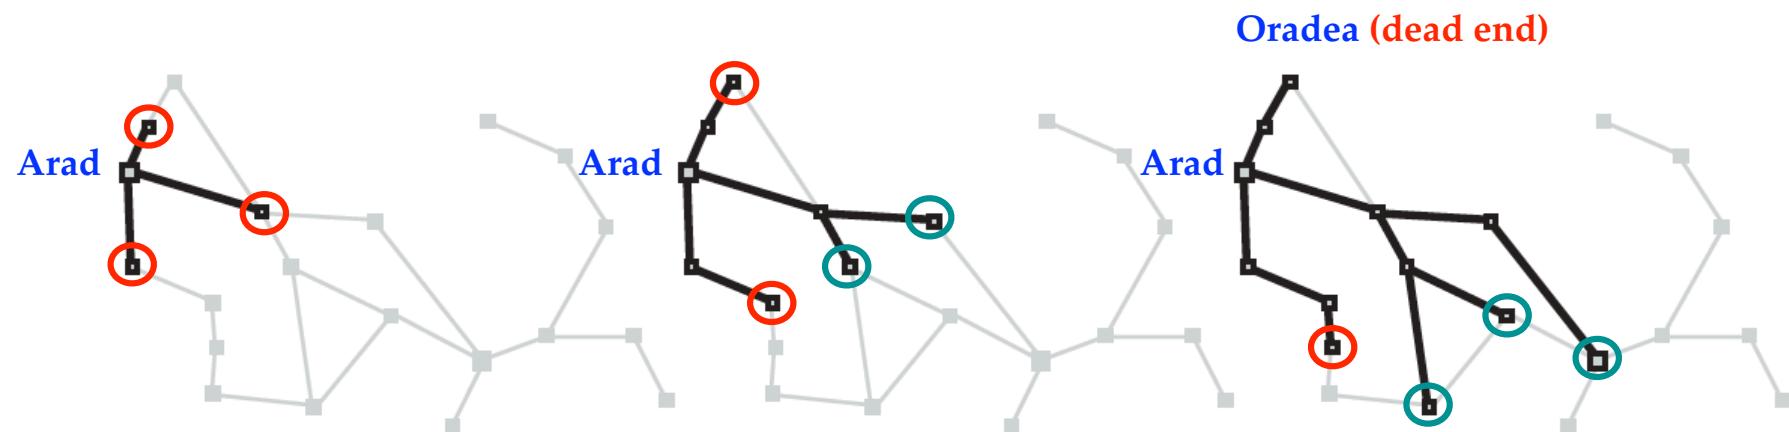
expand the chosen node. (repeated state)

if not in the frontier or explored set

add the resulting nodes to the frontier

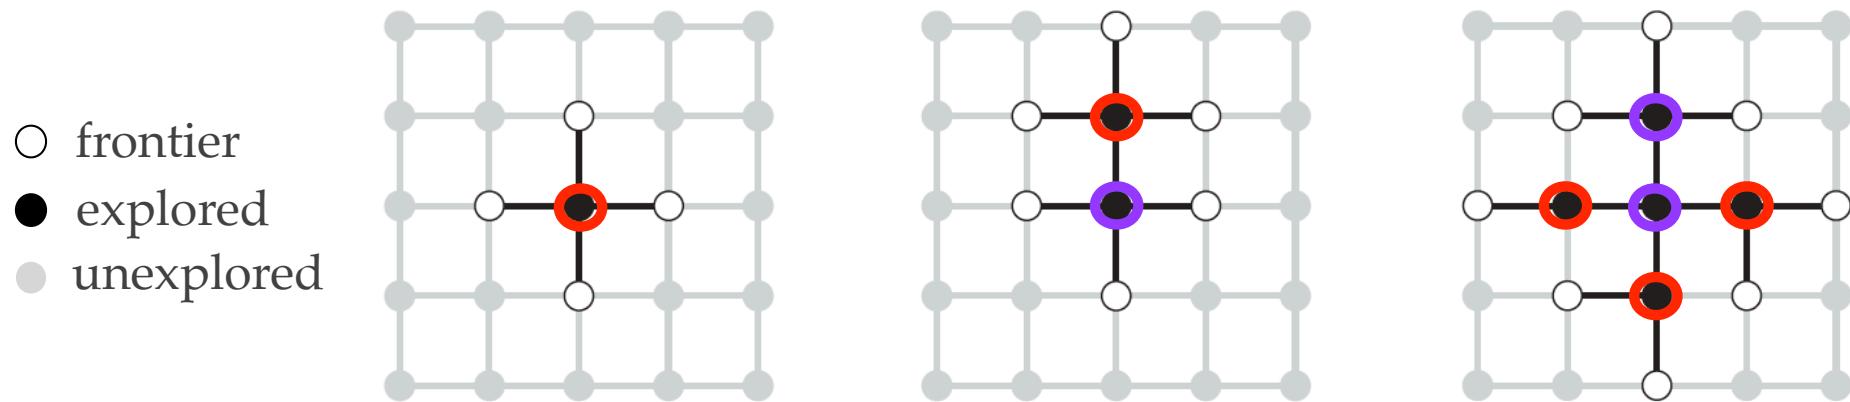
Search Tree

- ❖ A sequence of **search trees** generated by a graph search on the Romania problem
- ❖ At each stage, we **extend** each path by one step
- ❖ At the **third** stage, the northernmost city (**Oradea**) has become a **dead end**

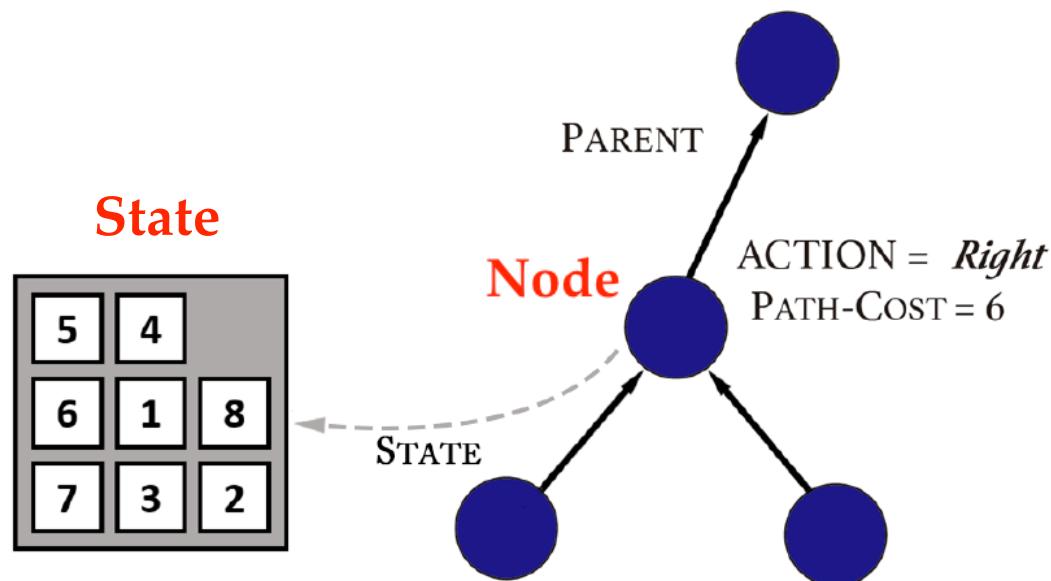


Graph Search & Frontier Separation

- ❖ The frontier **separates** the state space into **explored** and **unexplored** regions (**loop invariant proof**)



- ❖ State vs. node (map state to node)
 - ❖ State: a representation of a physical configuration
 - ❖ Node: a data structure constituting part of a search tree contains much more info. such as
 - ❖ state, parent node, action, depth, path cost $g(x)$



function CHILD-NODE(*problem, parent, action*) **returns** a node

return a node with

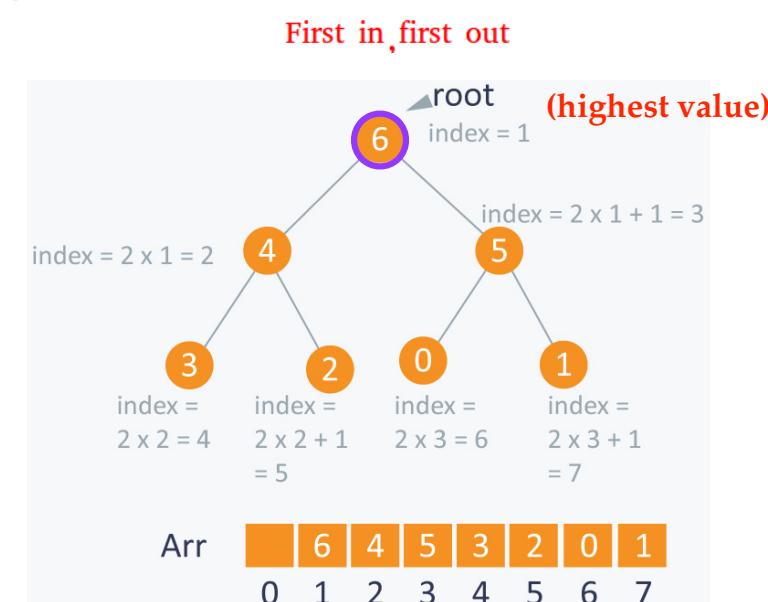
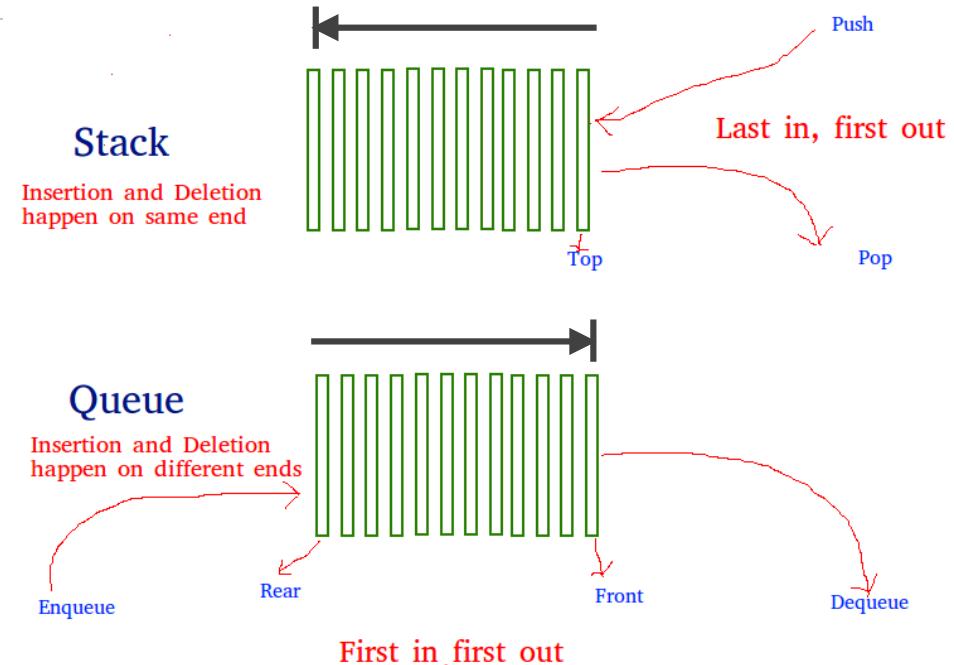
s' STATE = *problem.RESULT(parent.STATE, action)*,

PARENT = *parent*, ACTION = *action*,

PATH-COST = *parent.PATH-COST + problem.STEP-COST(parent.STATE, action)*

Infrastructure for Search Algorithms

- ❖ The appropriate data structure to maintain the frontier is a **queue** (memory buffer)
- ❖ Can be
 - ❖ **LIFO** (a.k.a. stack)
 - ❖ **FIFO** (a.k.a. queue)
 - ❖ **Priority queue (heap)**

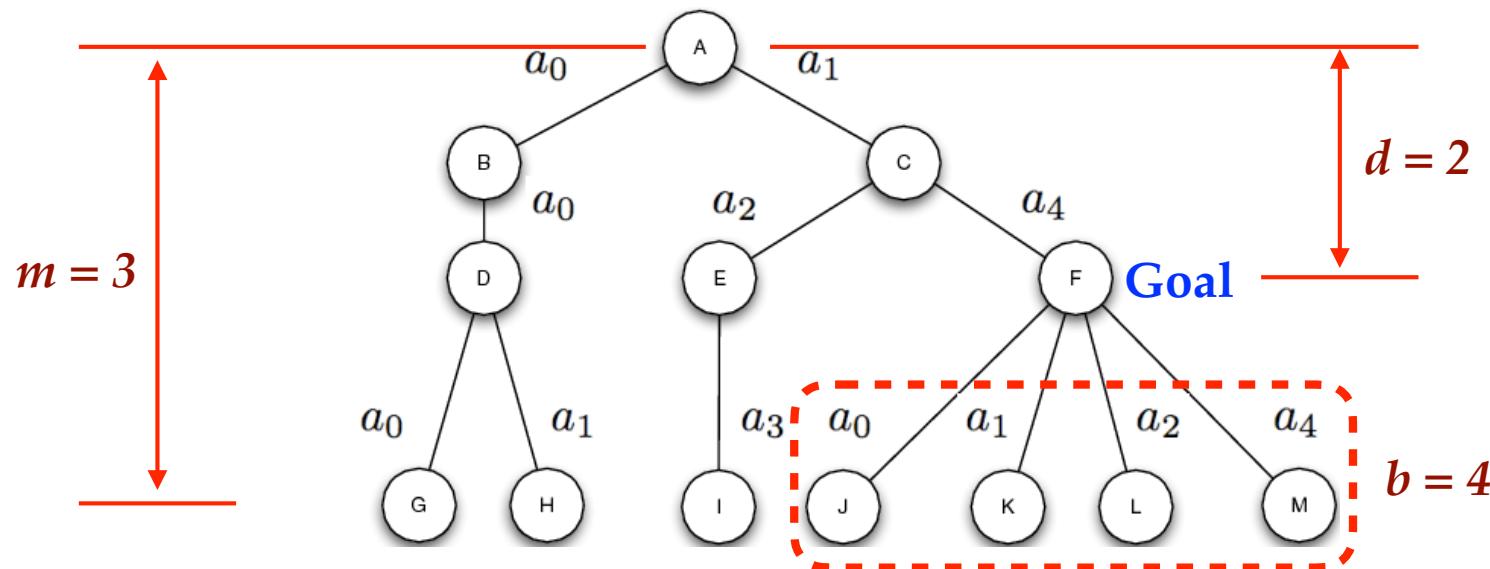


Tree Search Algorithms

- ❖ A **strategy** is defined by picking the **order** of node expansion
- ❖ Strategies are evaluated along the following dimensions
 - ❖ **Completeness** - does it **always find a solution** if one exists?
 - ❖ **Optimality** - does it always find a **least-cost solution** (given multiple goals)?
 - ❖ **Time complexity** - **number of nodes** generated/expanded
 - ❖ **Space complexity** - maximum number of nodes in **memory**

AI algorithms usually should concern more space complexity than time complexity.

- ❖ Time and space complexity are measured in terms of
 - ❖ **b** - max branching factor of the search tree
 - ❖ **d** - depth of the least-cost solution
 - ❖ **m** - max depth of the state space (may be ∞)

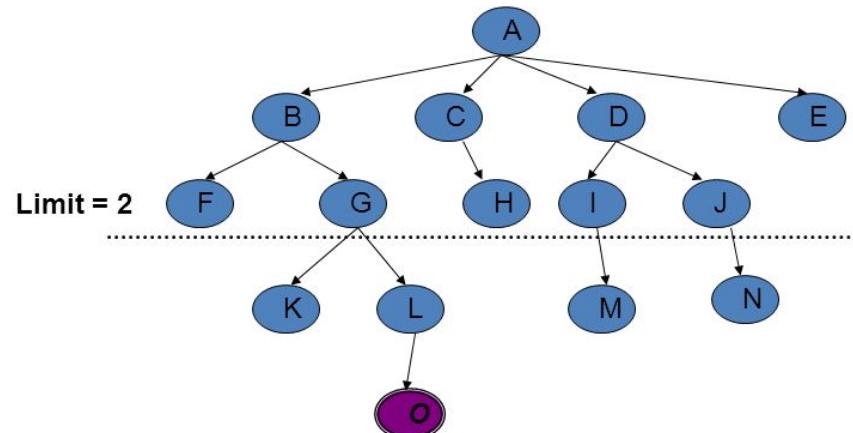
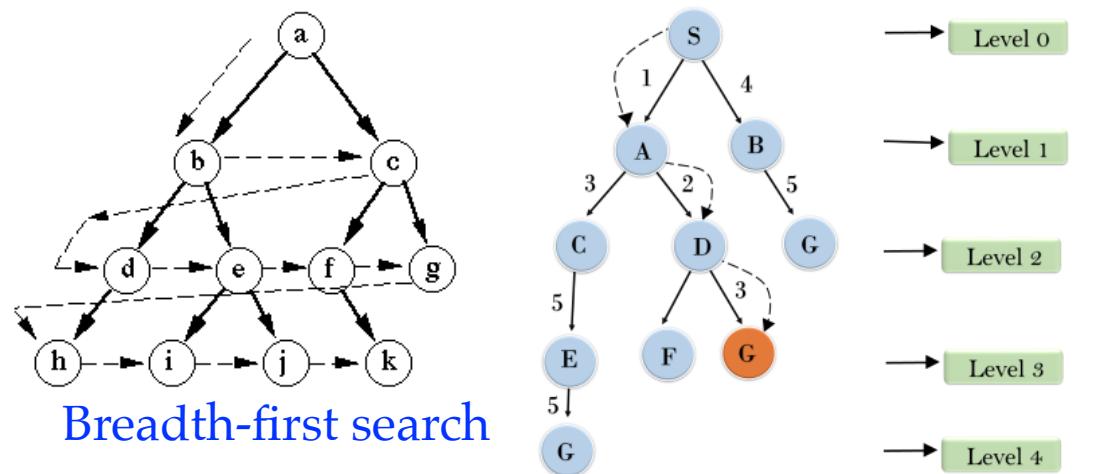


3.4 Uninformed Search Strategies (Blind Search)

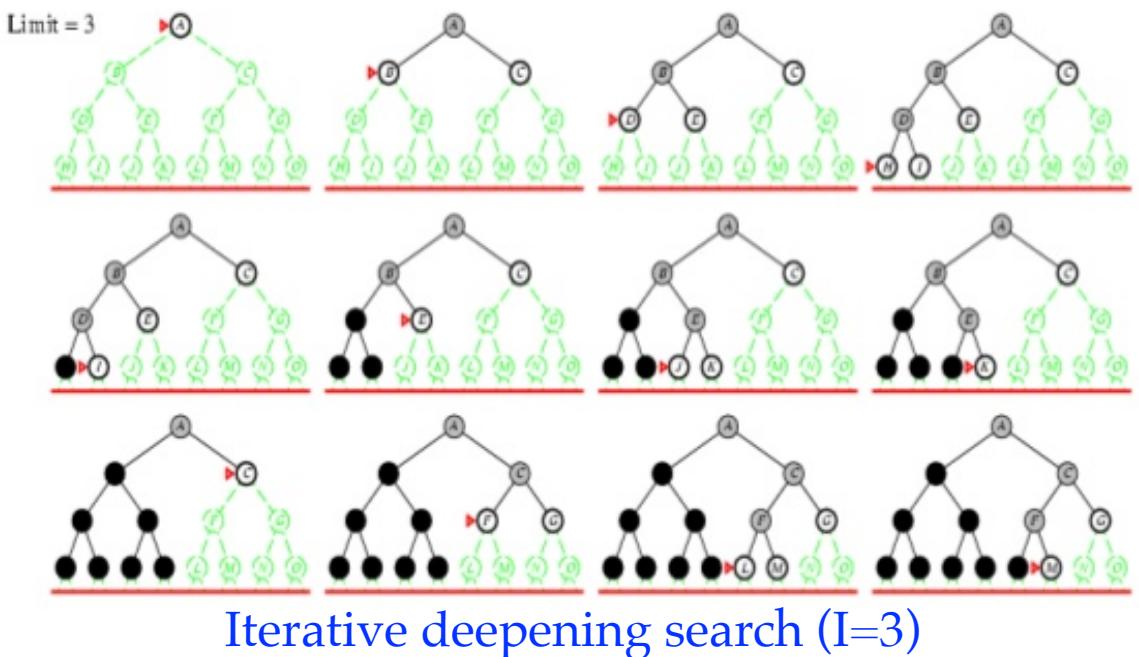
- ❖ **Uninformed (blind)**
 - ❖ No clue whether one non-goal state (current exploration) is better than any other. Your search is **blind**
 - ❖ Uninformed strategies use only the **information available** in the **problem definition**, e.g.
 - ❖ Initial state, goal state, transition model, goal test, path cost, step cost

❖ Blind searches

- ❖ Breadth-first search
- ❖ Uniform-cost search
- ❖ Depth-first search
- ❖ Depth-limited search
- ❖ Iterative deepening search



Depth-limited search

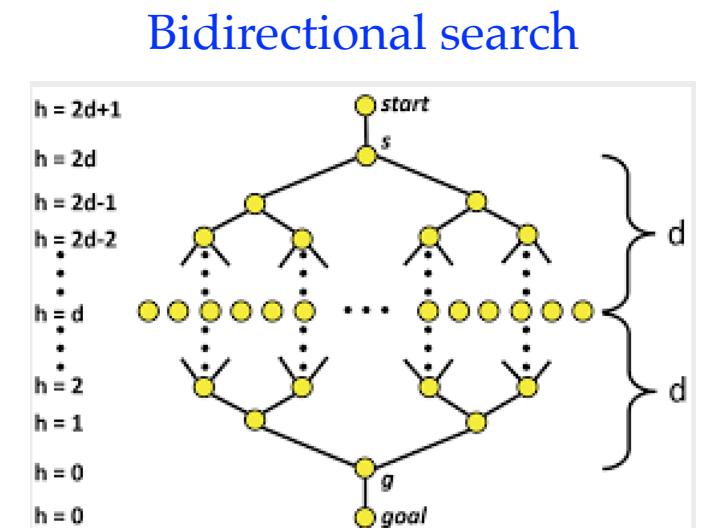


Uninformed Search Design Choices

- ❖ **Queue for frontier**
 - ❖ FIFO? LIFO? priority queue?
- ❖ **Goal-test**
 - ❖ Do goal-test when node **inserted** into frontier?
 - ❖ Do goal-test when node **removed** from frontier?
- ❖ **Tree search or graph search**
 - ❖ Forget expanded nodes? [tree search]
 - ❖ Remember expanded nodes? [graph search]

Queue for Frontier

- ❖ FIFO (First In, First Out)
 - ❖ Breadth-First Search (BFS)
- ❖ LIFO (Last In, First Out)
 - ❖ Depth-First Search (DFS)
- ❖ Priority queue sorted by path cost so far
 - ❖ Uniform Cost Search (UCS)
- ❖ Iterative Deepening Search (IDS) uses Depth-First Search
- ❖ Bidirectional Search can use either Breadth-First or Uniform Cost Search

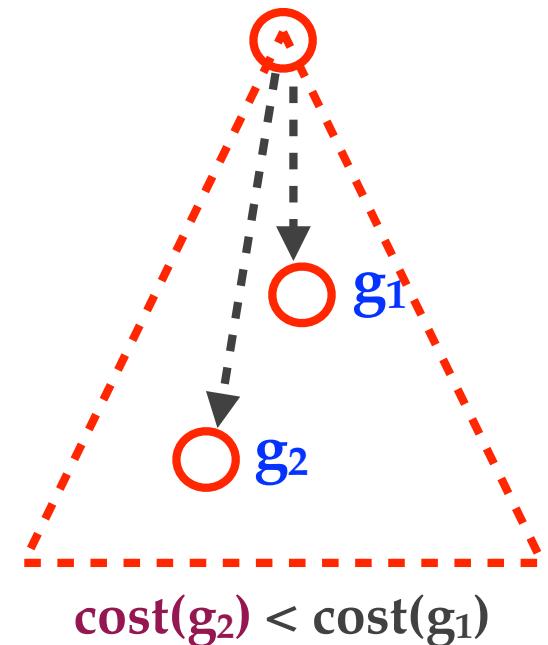


When to do Goal-Test? When Generated? When Popped?

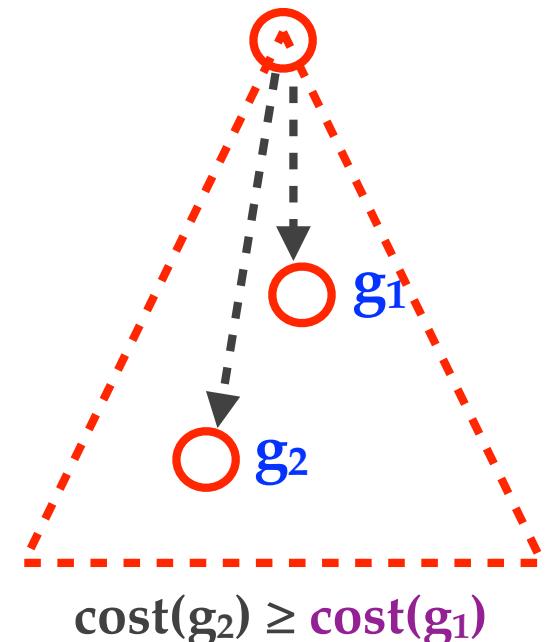
- ❖ Do Goal-Test when node is popped from queue

IF you care about finding the optimal path AND your search space may have both short expensive and long cheap paths to a goal

- ❖ Guard against a short expensive goal
- ❖ e.g.,
 - ❖ Uniform Cost Search with variable step costs



- ❖ Otherwise, do Goal-Test when **node inserted**
 - ❖ e.g.,
 - ❖ Breadth First Search (BFS)
 - ❖ Depth First Search (DFS)
 - ❖ Uniform Cost Search when **cost** is a **non-decreasing function of depth only**
 - ❖ Depth-Limited Search (DLS)
 - ❖ Iterative-Deepening Search (IDS)



Goal Test for General Tree Search

```
function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
```

Goal test after pop

```
function EXPAND( node, problem ) returns a set of nodes (successors)
  successors  $\leftarrow$  the empty set
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do
    s  $\leftarrow$  a new NODE
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1
    add s to successors
  return successors
```

Goal Test for General Graph Search

```
function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed  $\leftarrow$  an empty set (closed: set of repeated states)
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    if STATE[node] is not in closed then
      add STATE[node] to closed (set of repeated states)
      fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
```

Goal test after pop

Goal Test for Breadth-First Graph Search

```
function BREADTH-FIRST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  frontier  $\leftarrow$  a FIFO queue with node as the only element
  explored  $\leftarrow$  an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the shallowest node in frontier */
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)
        frontier  $\leftarrow$  INSERT(child, frontier)
```

Goal test before push

Goal Test for Uniform Cost Search (Sort by g)

- ❖ A^* is identical but uses $f = g + h$

Greedy best-first (GBFS) is identical but uses h

```
function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
    node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
    frontier  $\leftarrow$  a priority queue ordered by PATH-COST, with node as the only element
    explored  $\leftarrow$  an empty set
    loop do
        if EMPTY?(frontier) then return failure
        node  $\leftarrow$  POP(frontier) /* chooses the lowest-cost node in frontier */
        if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
        add node.STATE to explored
        for each action in problem.ACTIONS(node.STATE) do
            child  $\leftarrow$  CHILD-NODE(problem, node, action)
            if child.STATE is not in explored or frontier then
                frontier  $\leftarrow$  INSERT(child, frontier)
            else if child.STATE is in frontier with higher PATH-COST then
                replace that frontier node with child
        Goal test after popped
```

DLS: Depth-Limited Search

IDS: Iterative-Deepening Search

Goal Test for DLS & IDS

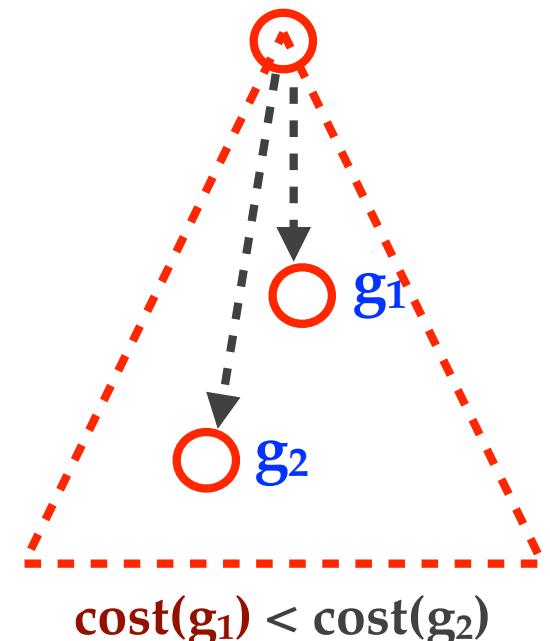
```
function DEPTH-LIMITED-SEARCH( problem, limit) returns soln/fail/cutoff
    RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)
function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
    cutoff-occurred?  $\leftarrow$  false
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    else if DEPTH[node] = limit then return cutoff
    else for each successor in EXPAND(node, problem) do
        result  $\leftarrow$  RECURSIVE-DLS(successor, problem, limit)
        if result = cutoff then cutoff-occurred?  $\leftarrow$  true
        else if result  $\neq$  failure then return result
    if cutoff-occurred? then return cutoff else return failure
```

Goal test before push

```
function ITERATIVE-DEEPENING-SEARCH( problem) returns a solution, or fail-
ure
    inputs: problem, a problem
    for depth  $\leftarrow$  0 to  $\infty$  do
        result  $\leftarrow$  DEPTH-LIMITED-SEARCH( problem, depth)
        if result  $\neq$  cutoff then return result
```

Summary: When to do Goal-Test?

- ❖ For **BFS**, **DFS**, **DLS**, and **IDS**, the goal test is done when the **child node is generated**
 - ❖ These are **not optimal searches** in the general case
 - ❖ BFS and IDS are **optimal if cost is a (increasing) function of depth only**; then, **optimal goals are also shallowest goals** and so will be found first



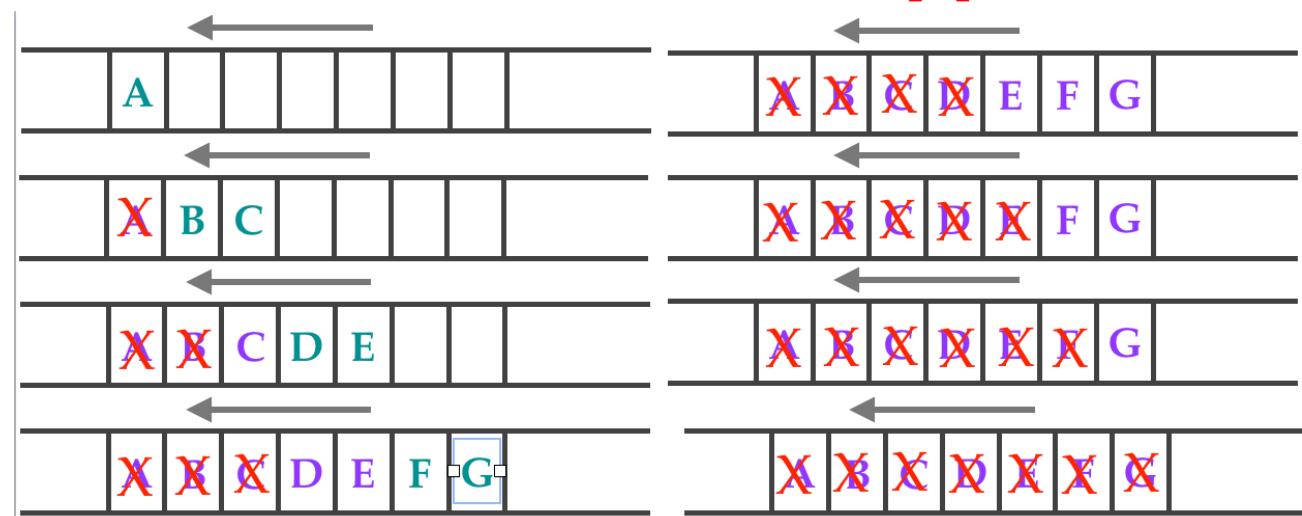
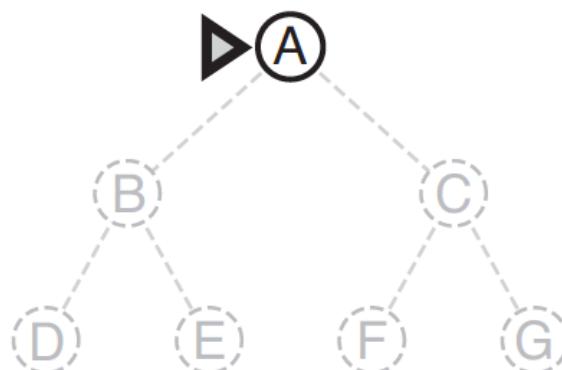
- ❖ For **UCS** (Uniform Cost Search) and **A*** the goal test is done when the node is **removed** from the queue
 - ❖ This precaution **avoids** finding a short expensive path before a long cheap path

Breadth-First Search (BFS)

- ❖ Expand **shallowest unexpanded node**
- ❖ **Frontier** (or fringe): nodes in queue to be explored
- ❖ Frontier is a **first-in-first-out (FIFO)** queue, i.e., new successors go at end of the queue
- ❖ **Goal-Test when inserted**

Initial state = A
Is A a goal state?

Put A at **end** of queue.
frontier = [A]

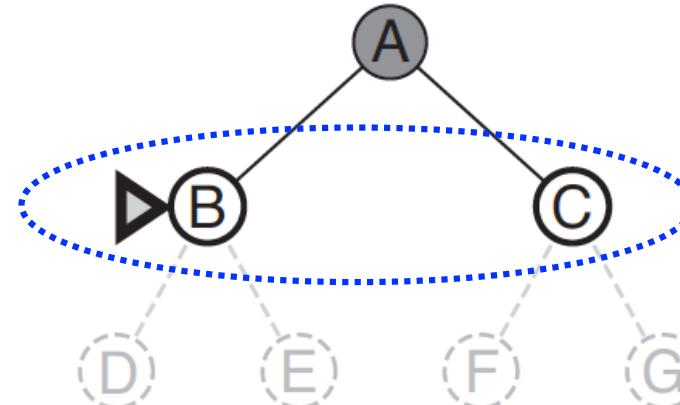


Expand A to B, C.

Is B or C a goal state?

Put B, C at end of queue.

frontier = [B,C]

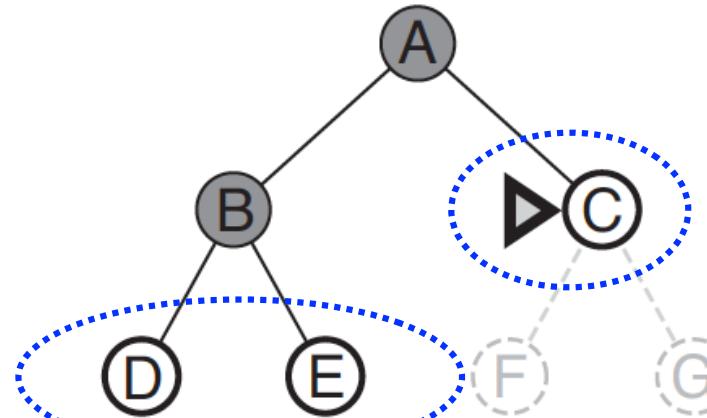


Expand B to D, E

Is D or E a goal state?

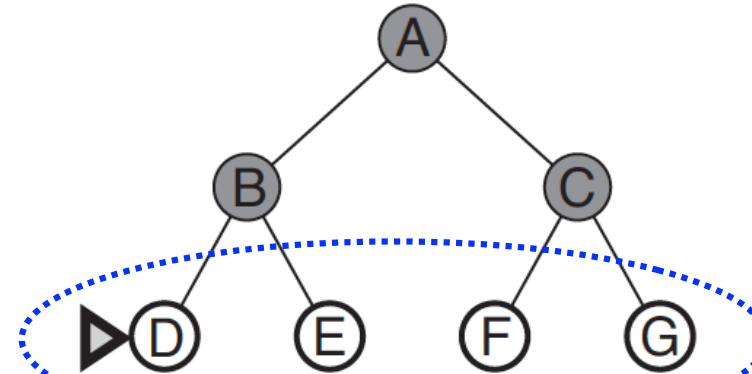
Put D, E at end of queue

frontier=[C,D,E]



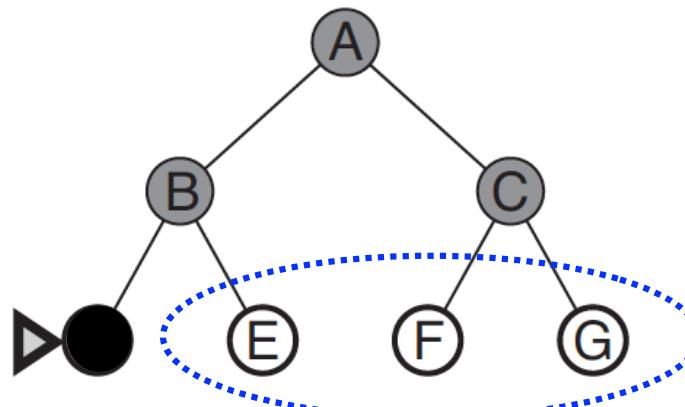
Expand C to F, G.
Is F or G a goal state?

Put F, G at end of queue.
frontier = [D,E,F,G]



Expand D to no children.
Forget D.

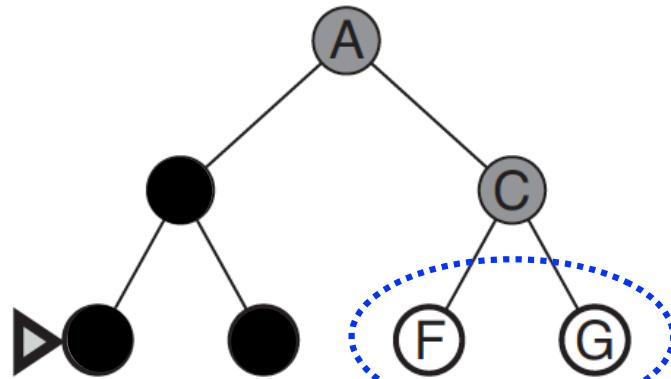
frontier = [E,F,G]



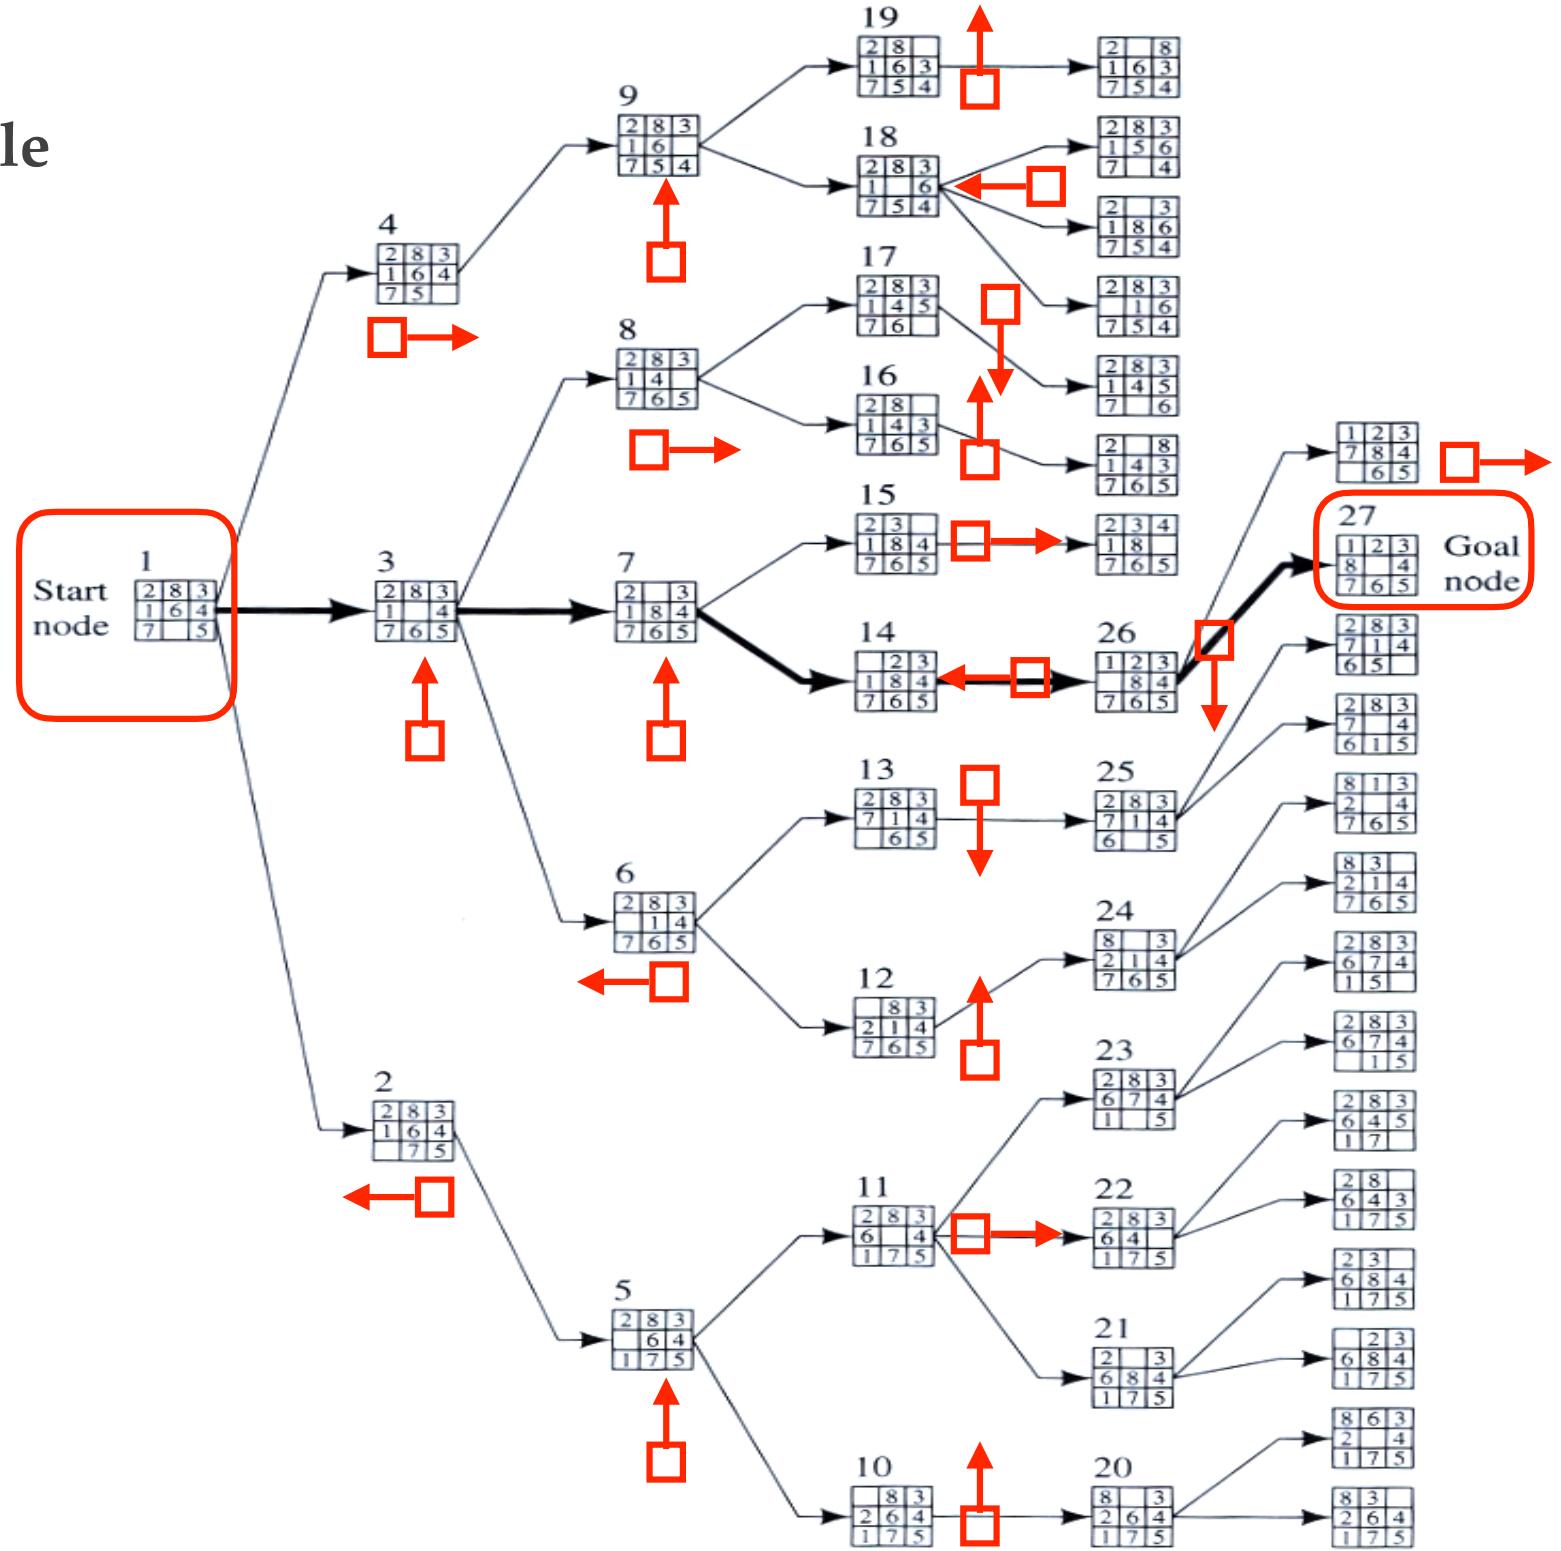
Expand E to no children.

Forget B,E.

frontier = [F,G]



BFS for 8-puzzle

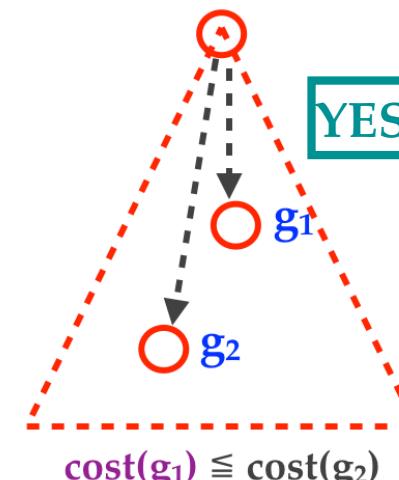
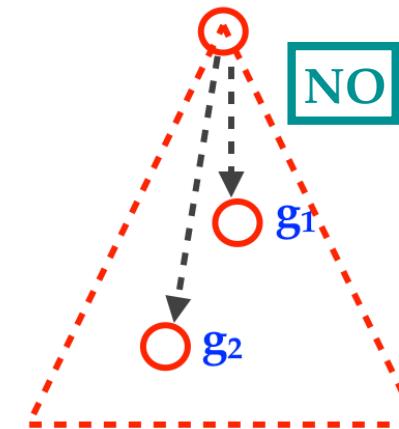
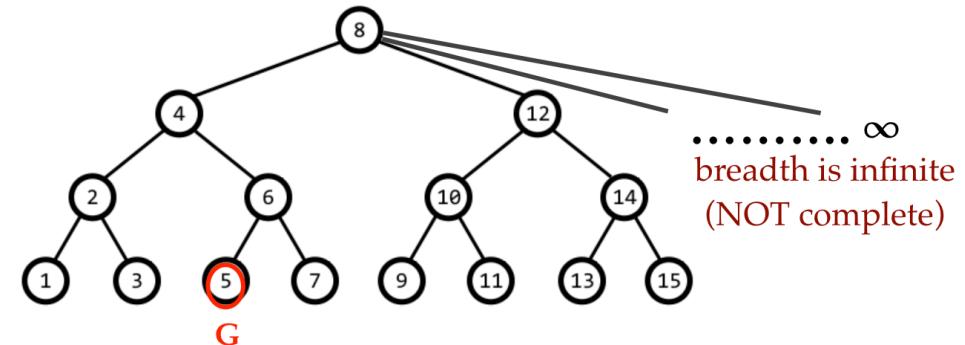


❖ Properties of BFS

- ❖ **Completeness:** Yes, it always reaches a goal (if b is finite)

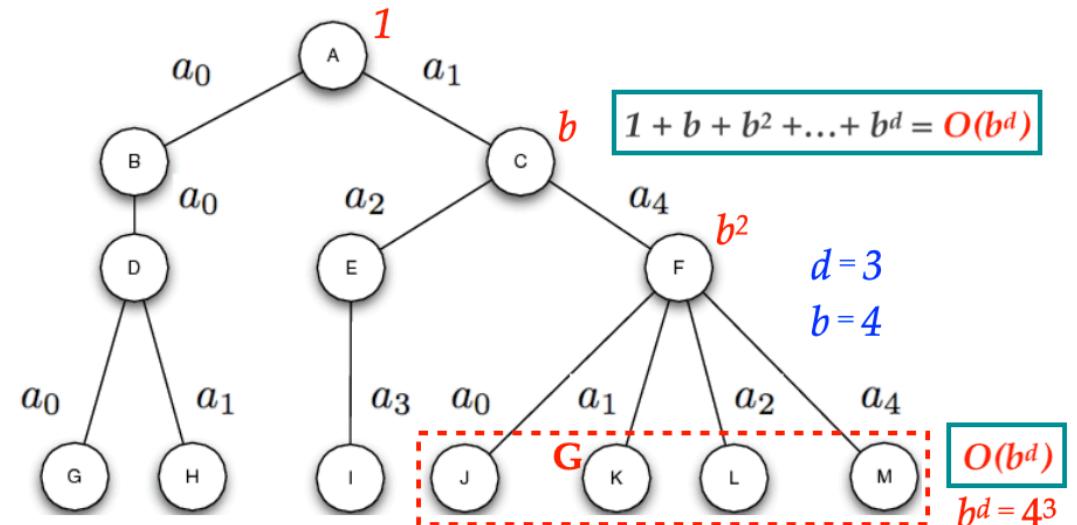
❖ Optimality

- ❖ No, for general cost functions
- ❖ Yes only if the path cost is a non-decreasing function of depth $f(d) \geq f(d-1)$, e.g., step-cost = constant, as in the 8-puzzle
 - ❖ All optimal goal nodes occur on the same level
 - ❖ Optimal goals are always shallower than non-optimal goals
 - ❖ An optimal goal will be found before any non-optimal goal



d - depth of the least-cost solution

- ❖ **Time complexity:** $1 + b + b^2 + \dots + b^d = O(b^d)$
 - ❖ Or $O(b^{d+1})$ if goal test is applied after expansion (this is the number of nodes we generate)



- ❖ **Space complexity:** $O(b^d)$
 - ❖ Keep each node in frontier in memory (i.e., it is dominated by the size of the frontier)
 - ❖ e.g. assume branching factor $b = 7$, if 10^6 nodes/sec, 10^3 bytes/node, $d = 7$ then 1GB/sec and 24hrs = 86.4TB

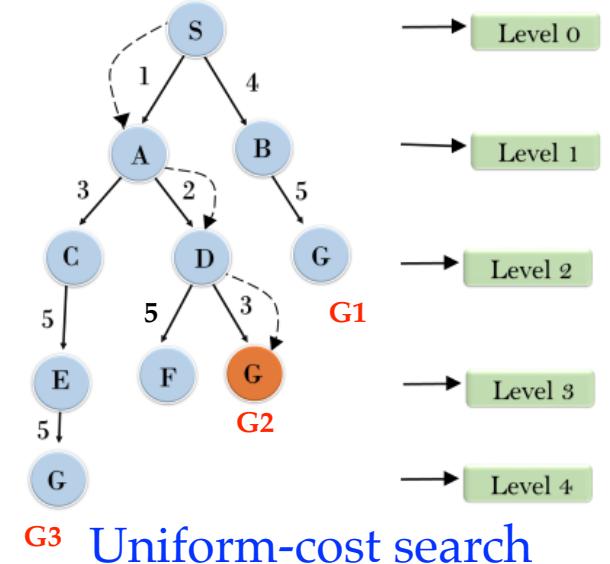
Space is the big problem!

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6 <i>d=7</i>	10^6	1.1 seconds	1 gigabyte 10GB
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes

- ❖ Time and memory requirements for breadth-first search
- ❖ The numbers shown assume branching factor $b = 7$;
 10^6 nodes/sec, 10^3 bytes/node

Uninformed-Cost Search

- ❖ Expand the unexpanded node with the **lowest path cost $g(n)$**
- ❖ Equivalent to BFS if step costs are all equal
- ❖ Frontier is a **priority queue**, i.e., new successors are merged into the queue sorted by $g(n)$
 - ❖ Can remove successors already on queue w/ higher $g(n)$
 - ❖ Saves memory, costs time
 - ❖ Goal-test when node is popped off queue



Explored	Frontier
S A:1, B:4	A, B
S, A B:4, C:4, D:3	B, C, D
S, A, D B:4, C:4, F:8, G2:6	B, C, F, G2
S, A, D, B C:4, F:8, G2:6, G1:9	C, F, G2, G1
S, A, D, B, C G2:6, G1:9	F, G2, G1, E
S, A, D, B, C, G2	G2, G1, E

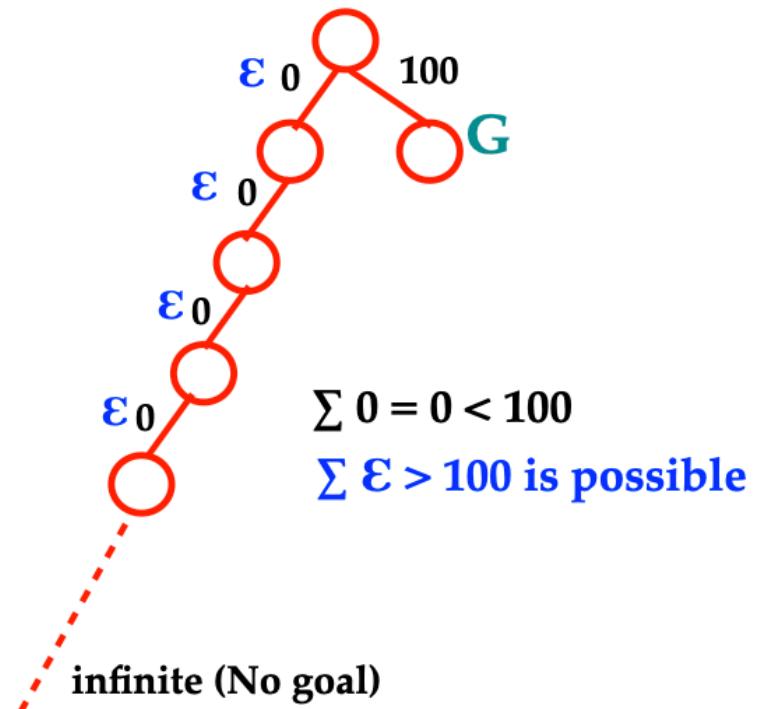
function UNIFORM-COST-SEARCH(*problem*) **returns** a solution, or failure

$node \leftarrow$ a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0
 $frontier \leftarrow$ a priority queue ordered by PATH-COST, with *node* as the only element
 $explored \leftarrow$ an empty set

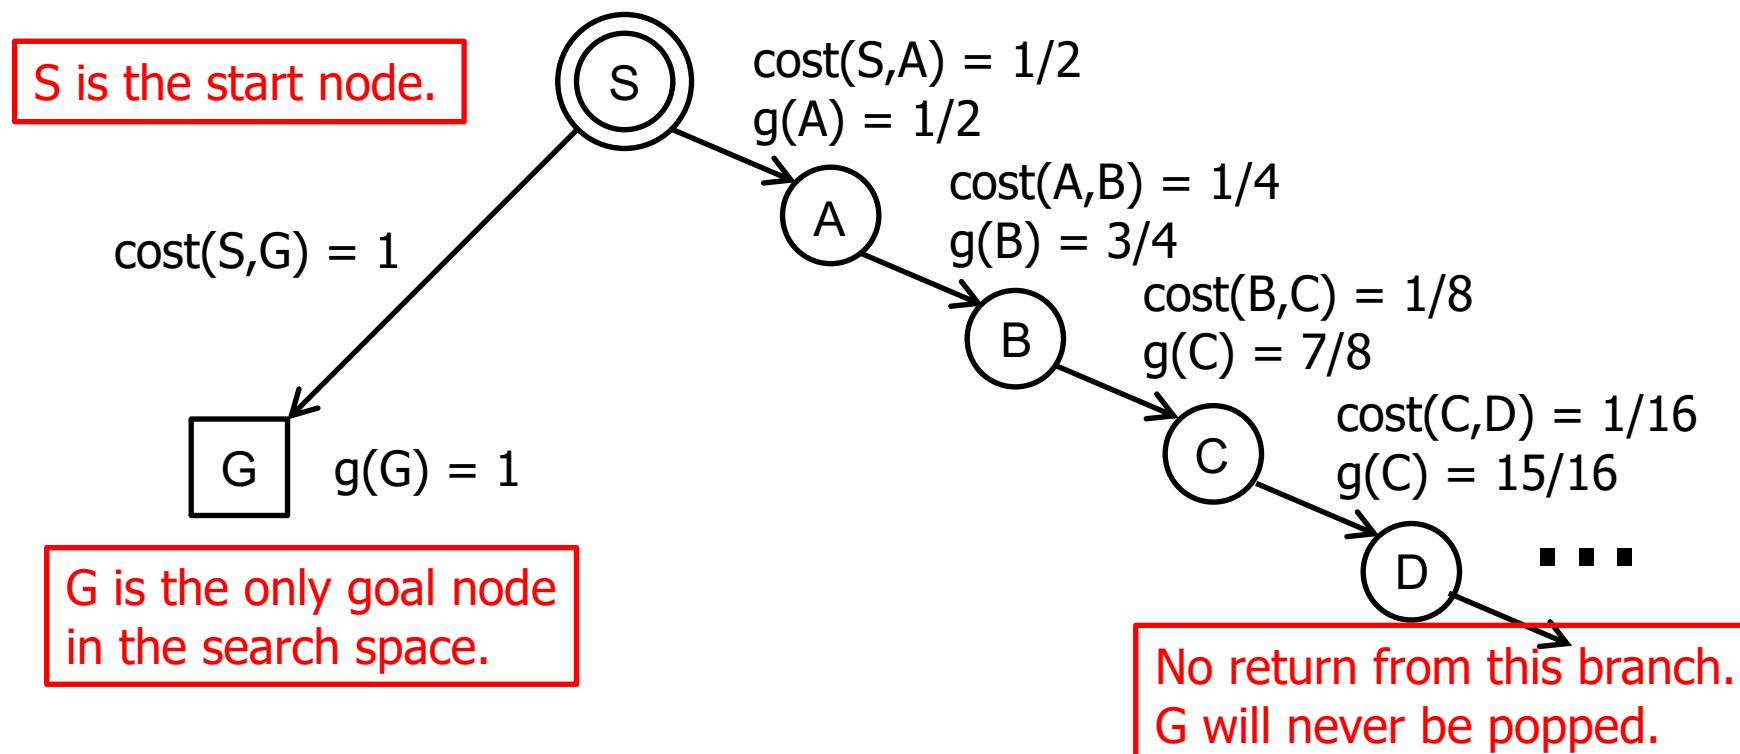
loop do

- if** EMPTY?(*frontier*) **then return** failure
- $node \leftarrow \text{POP}(frontier)$ /* chooses the lowest-cost node in *frontier* */
- if** *problem.GOAL-TEST(node.STATE)* **then return** SOLUTION(*node*)
- add *node.STATE* to *explored* **Goal test after pop**
- for each** *action* **in** *problem.ACTIONS(node.STATE)* **do**
- child* \leftarrow CHILD-NODE(*problem*, *node*, *action*)
- if** *child.STATE* is not in *explored* or *frontier* **then**
- frontier* \leftarrow INSERT(*child*, *frontier*)
- else if** *child.STATE* is in *frontier* with higher PATH-COST **then**
- replace that *frontier* node with *child*

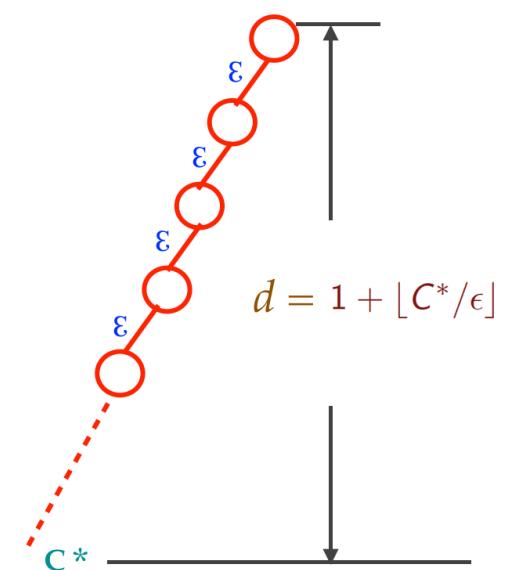
- ❖ **Implementation:** Frontier = queue ordered by path cost (priority queue)
- ❖ **Properties of Uniform-Cost Search**
 - ❖ **Completeness:** Yes, if b is finite and step cost $\geq \varepsilon > 0$
(otherwise it can get stuck in infinite loops, e.g, $\sum 0 = 0 < 100$)
 - ❖ **Optimality:** Yes - nodes expanded in increasing order of $g(n)$
(for any step cost $\geq \varepsilon > 0$)



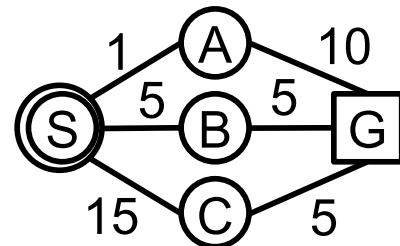
- ❖ Why require **step cost** $\geq \varepsilon > 0$?
 - ❖ Otherwise, an **infinite regress** is possible
 - ❖ [Step cost = 0] e.g, $\sum 0 = 0 < 100$
 - ❖ [Step cost has no lower bound $\geq \varepsilon > 0$]
 - e.g. $\sum_{n=1 \text{ to } \infty} 2^{-n} = 1.$ ($1, \frac{1}{2}, \frac{1}{4}, \dots$)



- ❖ **Time complexity:** # of nodes with $g \leq \text{cost of optimal solution}$
- ❖ Maximum depth is given by $1 + \lfloor C^*/\epsilon \rfloor$, where C^* is the cost of the optimal solution
- ❖ Time complexity is $1 + b + b^2 + \dots + b^d \approx O(b^{1+d})$
 $= O(b^{1+\lfloor C^*/\epsilon \rfloor})$ where $d = \lfloor C^*/\epsilon \rfloor$
- ❖ **Space complexity:** # of nodes with $g \leq \text{cost of optimal solution}$, $O(b^{1+\lfloor C^*/\epsilon \rfloor}) \approx O(b^{1+d})$

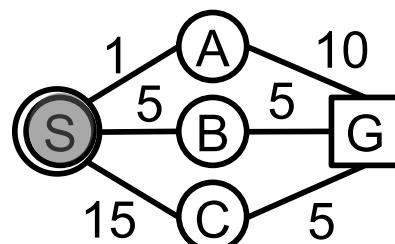
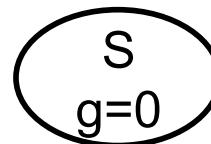


❖ Example hand-simulated search: **Search tree** method



Route finding problem.
Steps labeled w/cost.

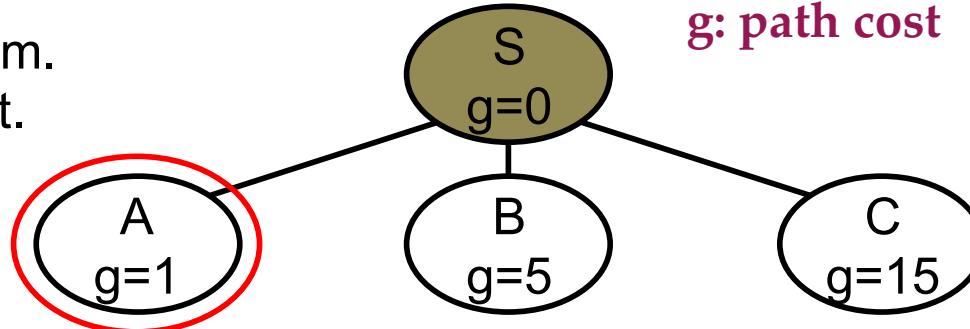
Order of node expansion: _____
Path found: _____ Cost of path found: _____

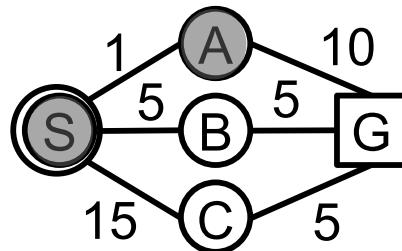


Route finding problem.
Steps labeled w/cost.

Order of node expansion: **S** _____
Path found: _____ Cost of path found: _____

g: path cost





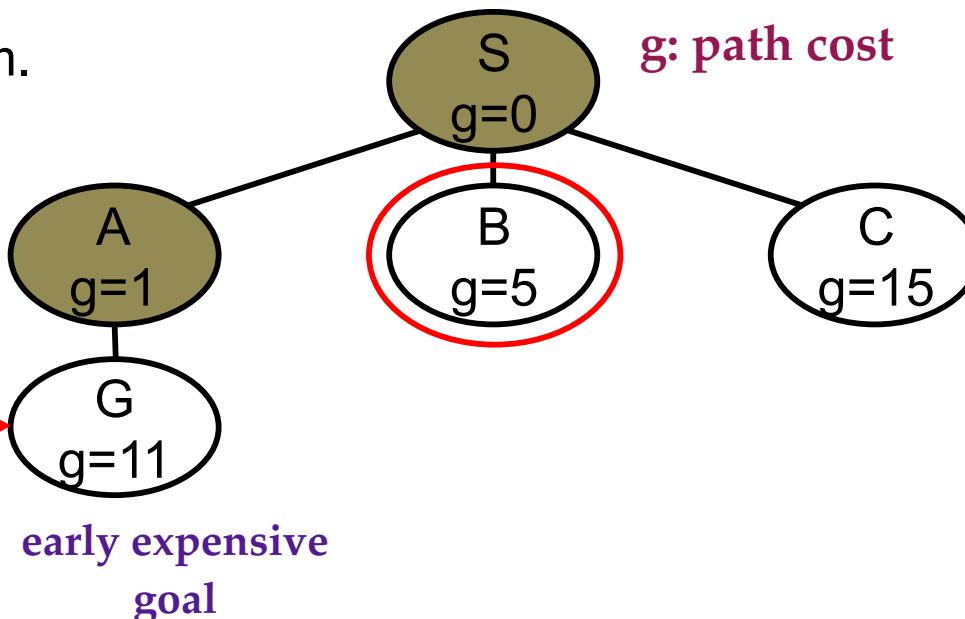
Order of node expansion: S A

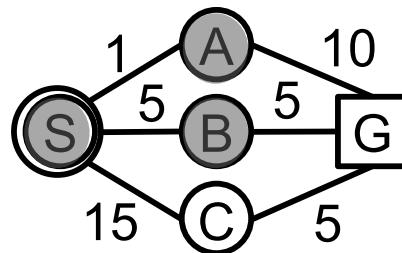
Path found: _____

Cost of path found: _____

Route finding problem.
Steps labeled w/cost.

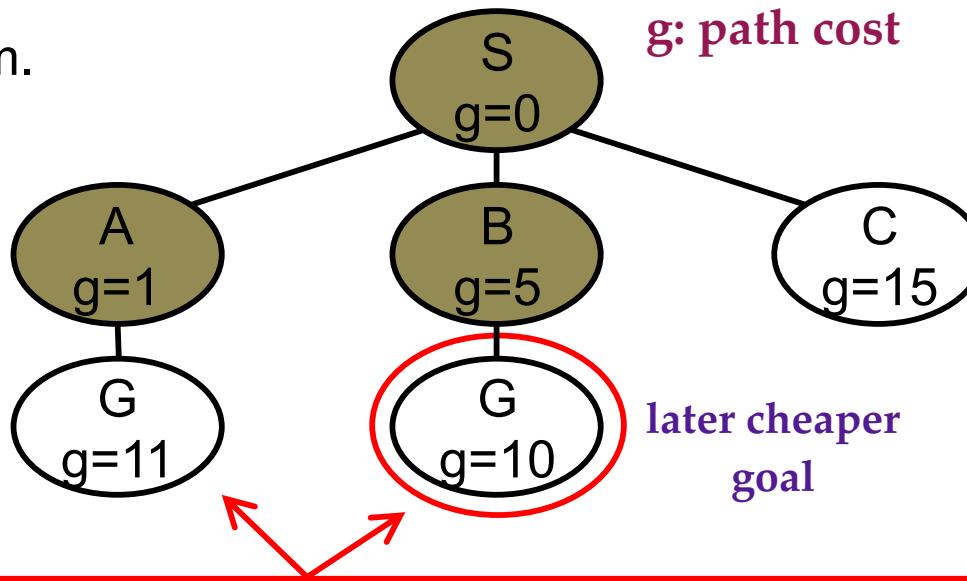
This early
expensive goal
node will go
back onto the
queue until
after the later
cheaper goal
is found.



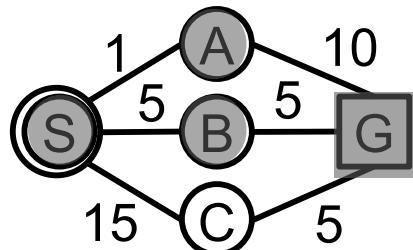


Order of node expansion: S A B
 Path found: _____ Cost of path found: _____

Route finding problem.
 Steps labeled w/cost.

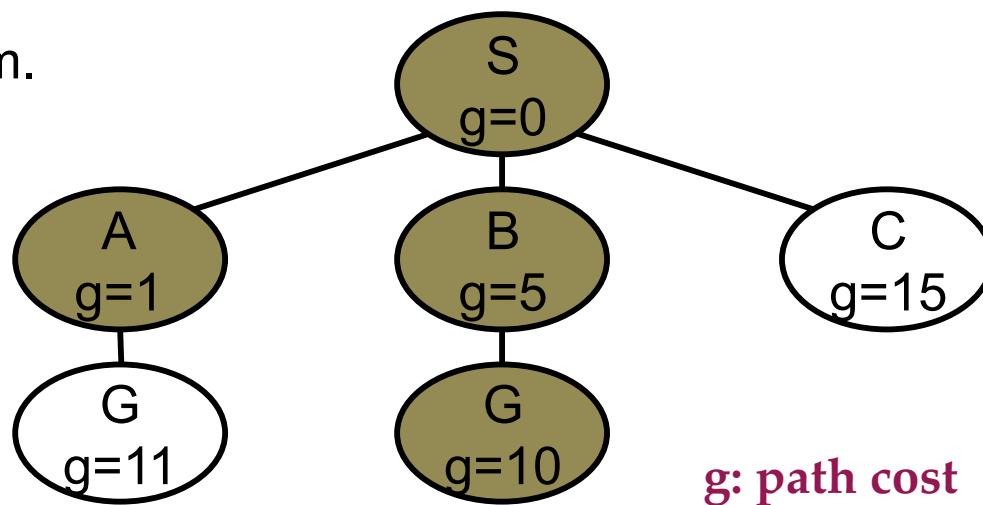


If we were doing graph search we would remove the higher-cost of identical nodes and save memory. However, UCS is optimal even with tree search, since lower-cost nodes sort to the front.



Order of node expansion: S A B G
 Path found: S B G Cost of path found: 10

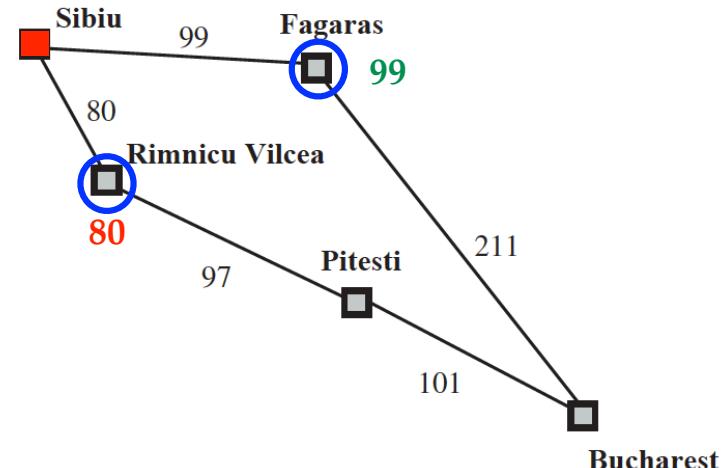
Route finding problem.
 Steps labeled w/cost.



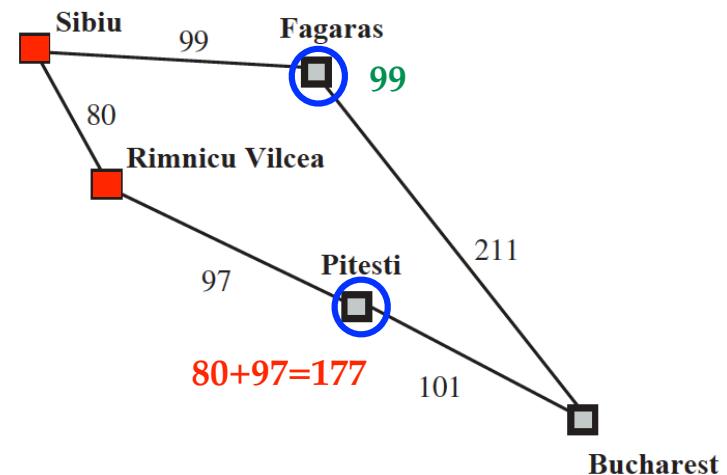
Technically, the goal node is not really expanded, because we do not generate the children of a goal node. It is listed in "Order of node expansion" only for your convenience, to see explicitly where it was found.

Romania example

- The successors of Sibiu are Rimnicu Vilcea (80) and Fagaras (99)



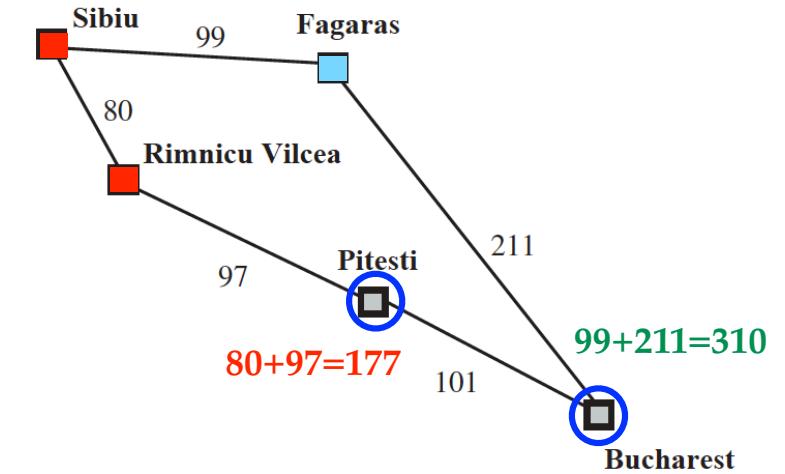
- The least-cost node, Rimnicu Vilcea (80), is expanded next, adding Pitesti (97) with cost $80 + 97 = 177$



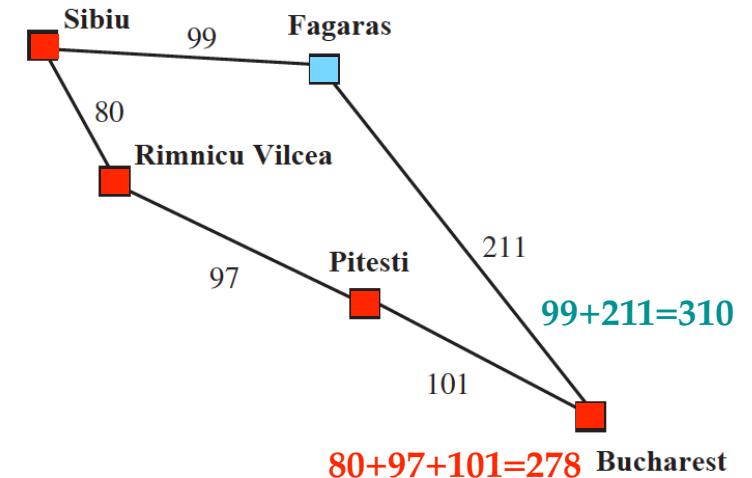
- The least-cost node is now Fagaras (99), so it is expanded, adding Bucharest (211) with cost $99+211=310$

Goal-test when node is popped off queue

- Now a goal node has been generated, but uniform-cost search keeps going, choosing Pitesti for expansion and adding a second path to Bucharest (101) with cost $80+97+101=278$



- Now the algorithm checks to see if this new path is better than the old one; it is, so the old one is discarded
- Bucharest, now with g-cost 278, is selected for expansion and the solution is returned



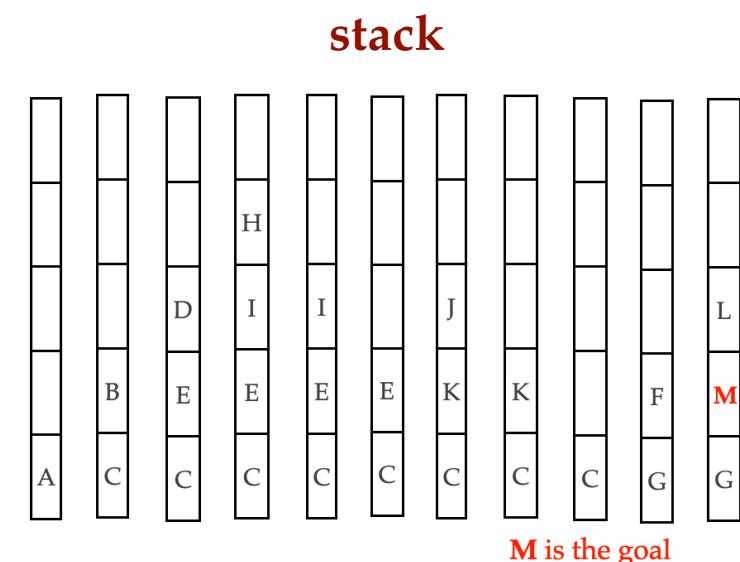
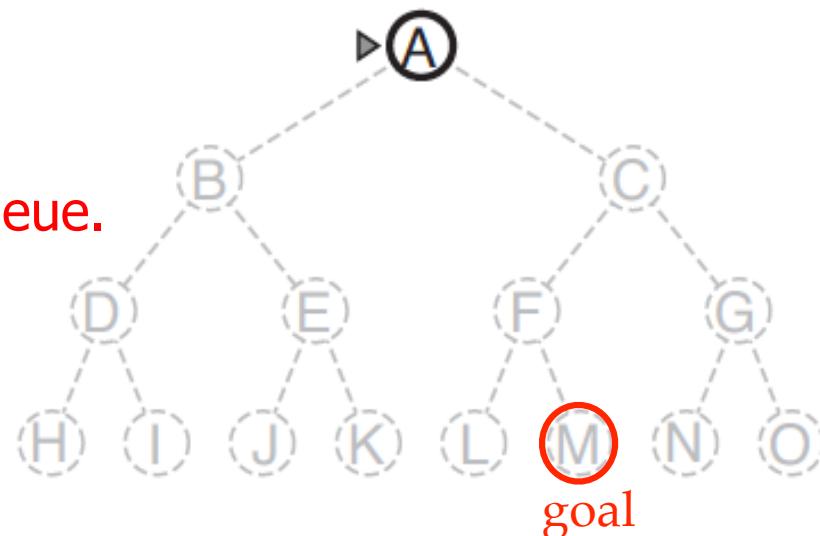
Depth-First Search

- ❖ Expand deepest unexpanded node
- ❖ Frontier = Last In First Out (LIFO) queue, i.e., new successors go at the front of the queue
- ❖ Goal-Test when inserted

Usually use “recursive” or “stack” to implement.

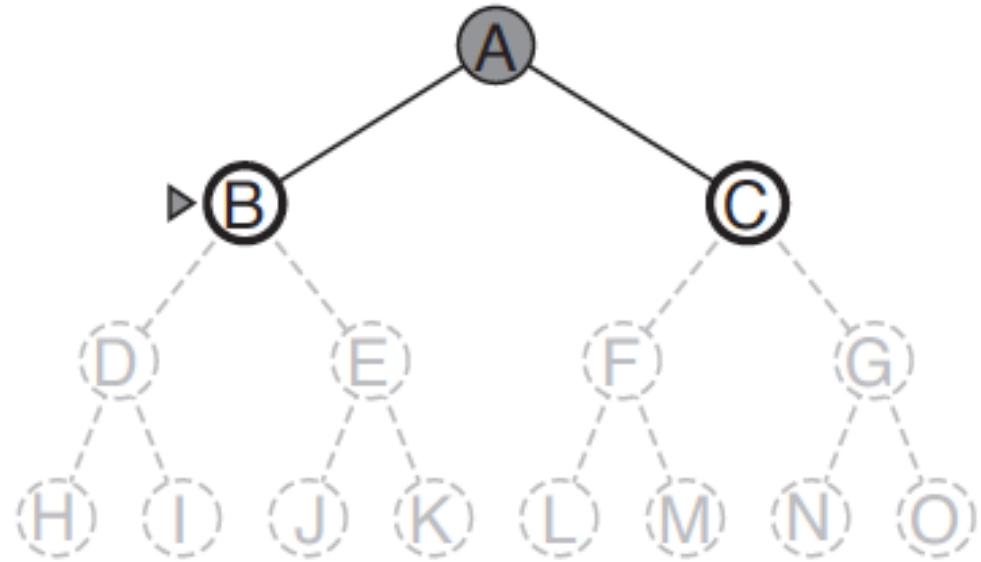
Initial state = A
Is A a goal state?

Put A at front of queue.
frontier = [A]



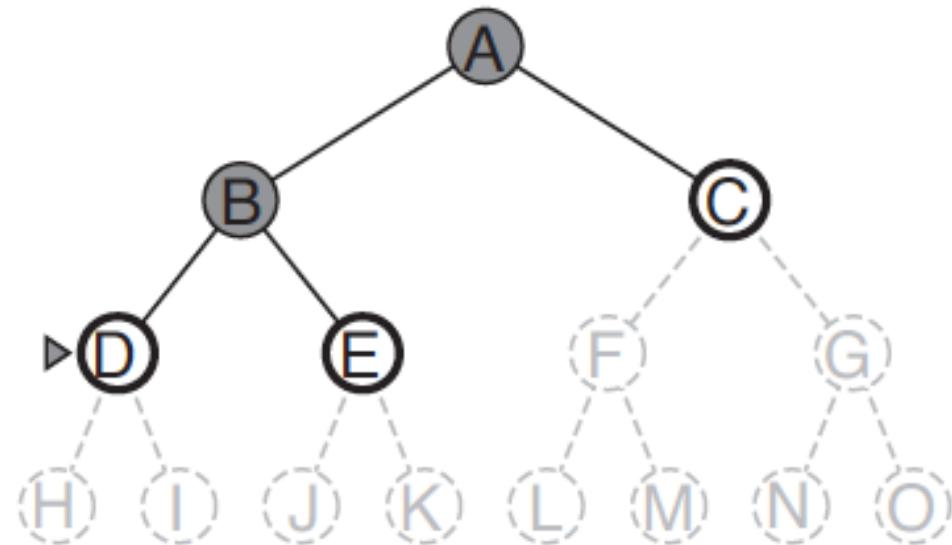
Expand A to B, C.
Is B or C a goal state?

Put B, C at front of queue.
frontier = [B,C]



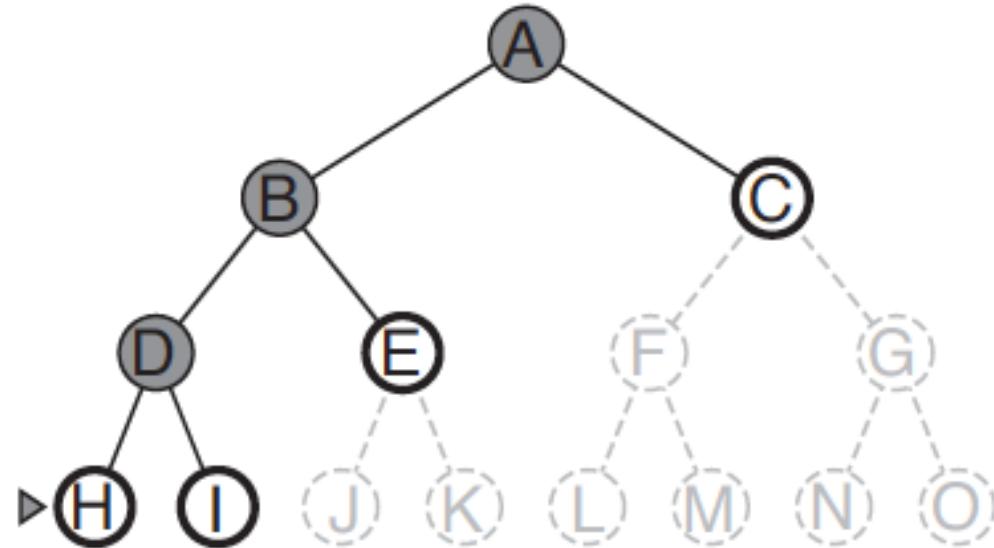
Expand B to D, E.
Is D or E a goal state?

Put D, E at front of queue.
frontier = [D,E,C]



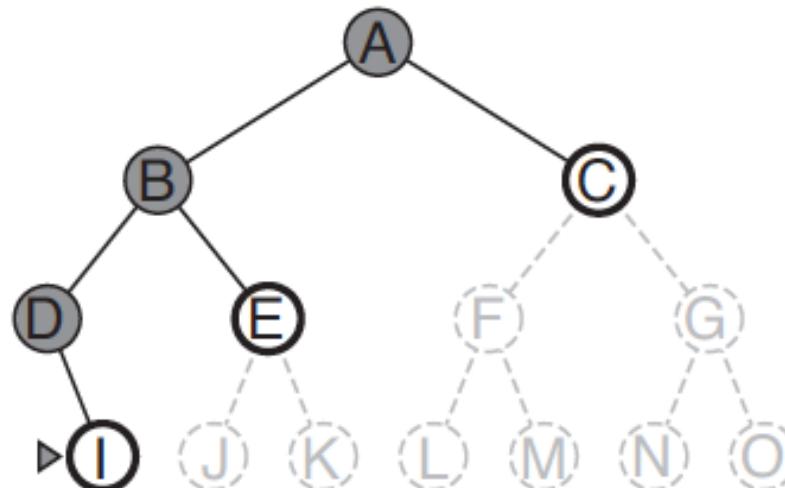
Expand D to H, I.
Is H or I a goal state?

Put H, I at front of queue.
frontier = [H,I,E,C]



Expand H to no children.
Forget H.

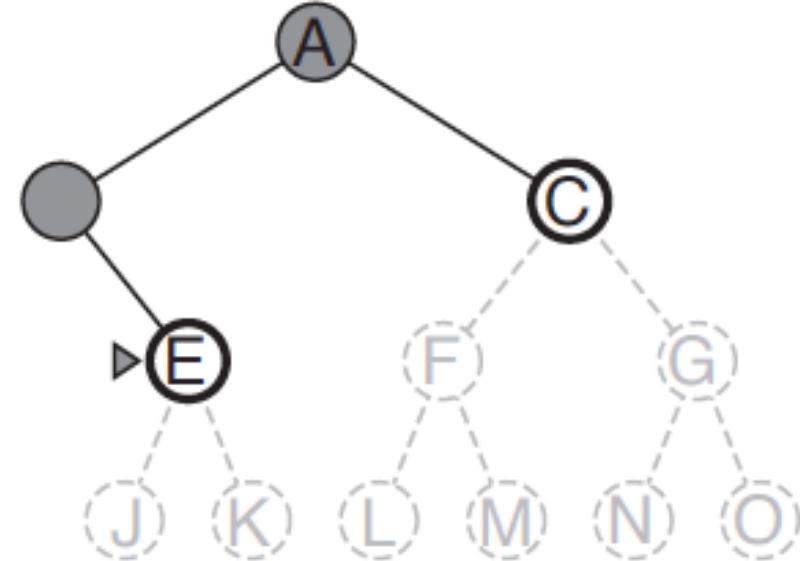
frontier = [I,E,C]



Expand I to no children.

Forget D, I.

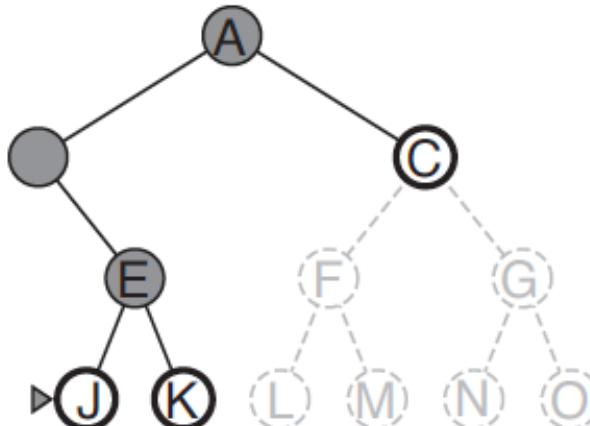
frontier = [E,C]



Expand E to J, K.

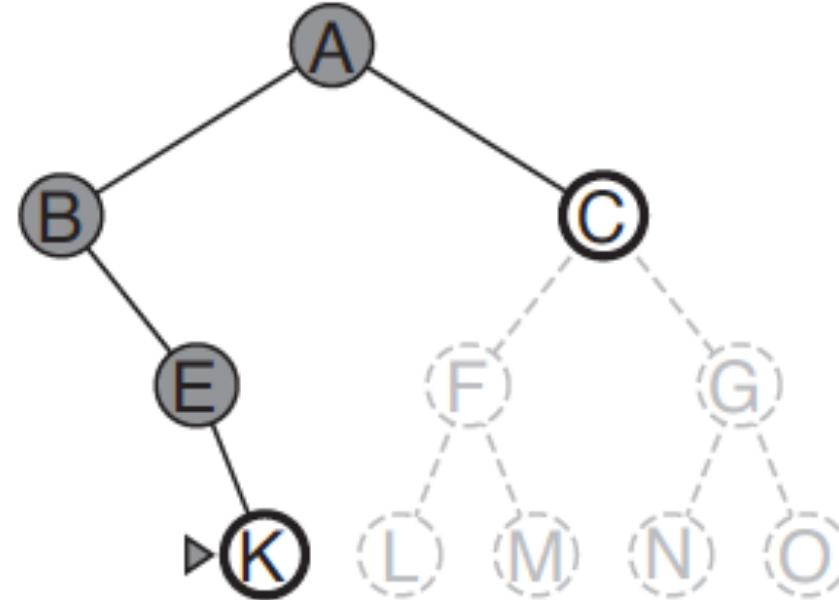
Is J or K a goal state?

Put J, K at front of queue.
frontier = [J,K,C]



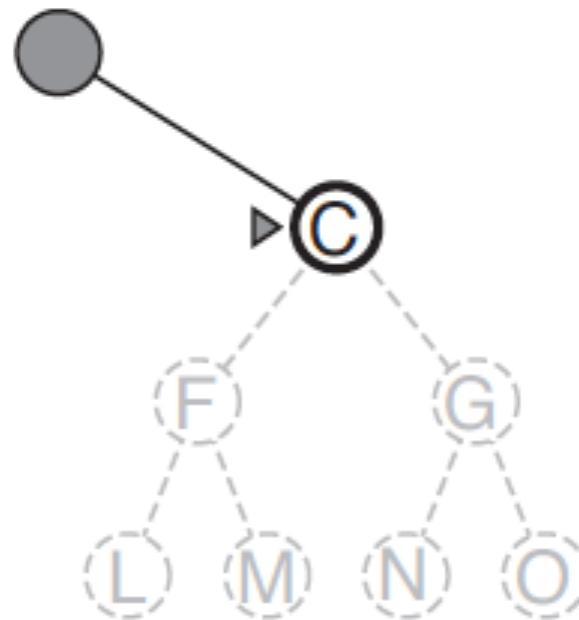
Expand J to no children.
Forget J.

frontier = [K,C]



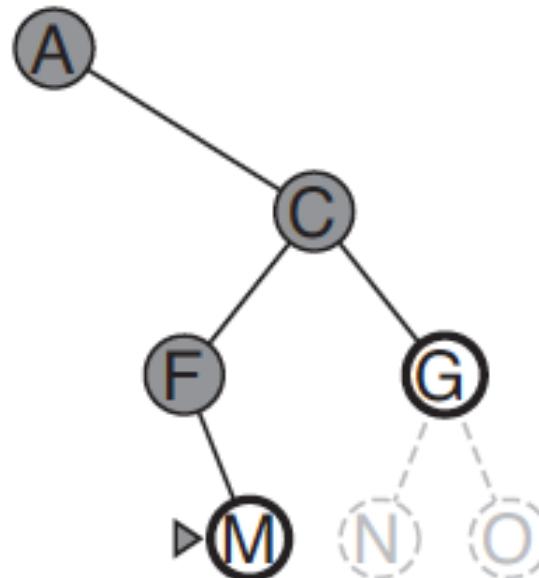
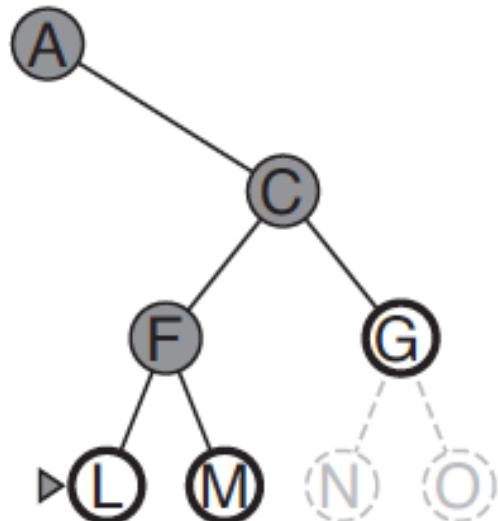
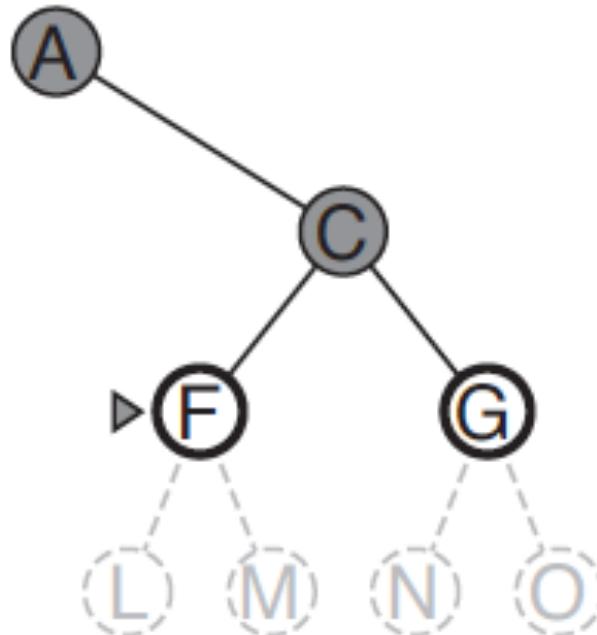
Expand K to no children.
Forget B, E, K.

frontier = [C]



Expand C to F, G.
Is F or G a goal state?

Put F, G at front of queue.
frontier = [F,G]

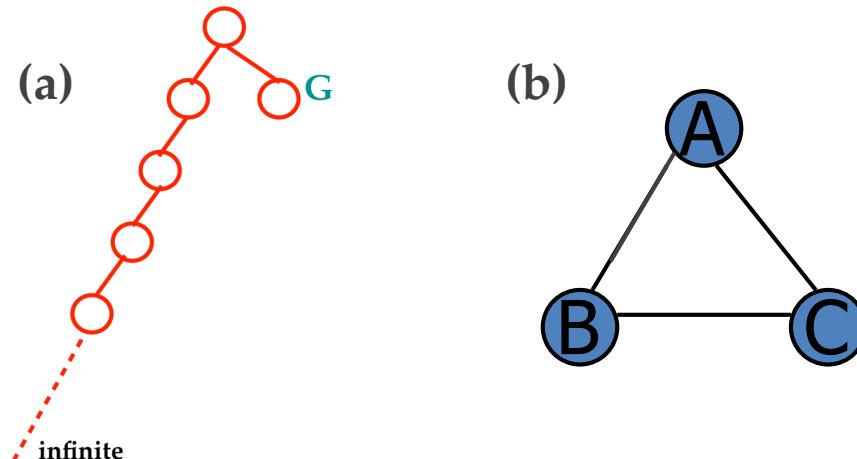


M is the only goal node.

- ❖ Properties of DFS

- ❖ Completeness

- ❖ No, fails in



- ❖ (a) infinite-depth spaces (may miss goal entirely)

- ❖ (b) spaces with loops

- ❖ (b) can be solved by modifying to avoid loops / repeated states along path

- ❖ Check if current nodes occurred before on path to root

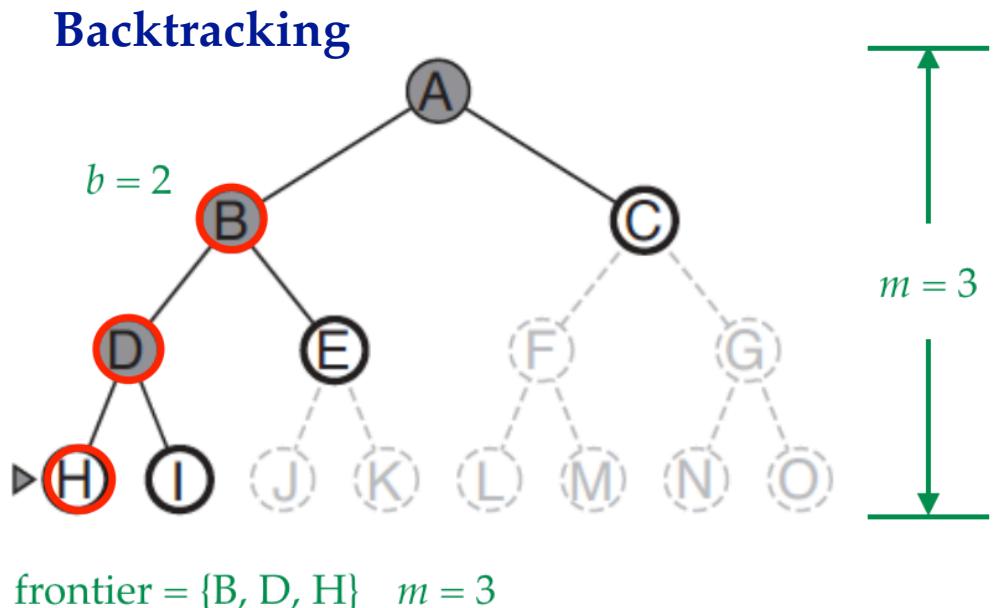
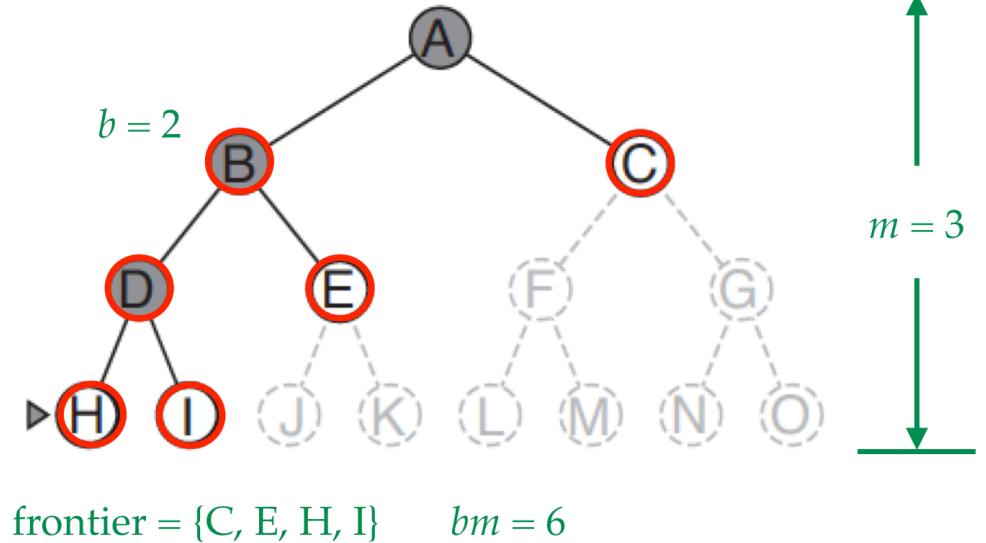
- ❖ Can use graph search (remember all nodes ever seen)

- ❖ Problem with graph search: space is exponential, not linear

- ❖ **Optimality: No.** It may find a non-optimal goal first
- ❖ **Time complexity**
 - ❖ $O(b^m)$ with $m = \max$ depth of space
 - ❖ **Terrible** if m is much larger than d (m may be infinite)
 - ❖ If solutions are **dense**, may be **much faster** than BFS

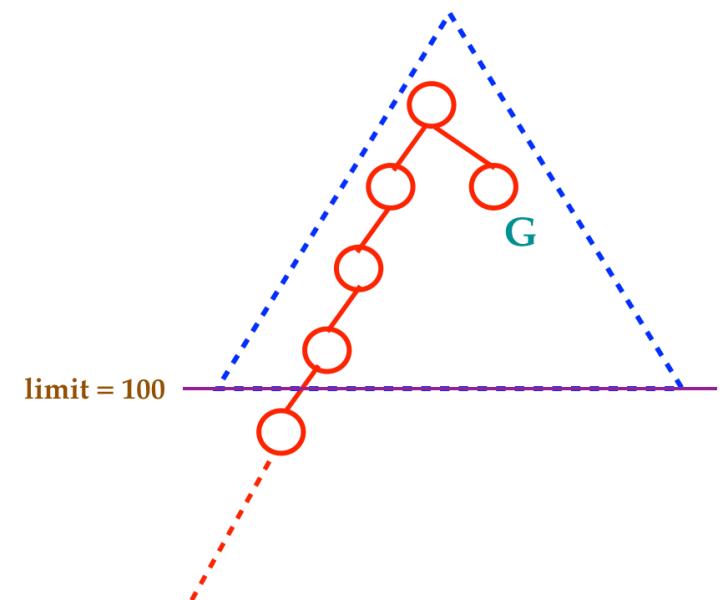
❖ Space complexity

- ❖ $O(bm)$, linear space! [only have to remember the nodes in frontier]
- ❖ Backtracking technique only generate **one successor** instead of all successors
→ space complexity is reduced to **$O(m)$**
(remember only a **single path + expanded unexplored nodes**)
→ Backtracking should trade off **higher time complexity**



Depth-Limited Search (DLS)

- ❖ DLS combines the advantages of BFS (better time complexity) and DFS (better space complexity)
- ❖ DFS never terminates if $m \rightarrow \infty$.
- ❖ DLS = DFS with **depth limit l**
- ❖ Nodes at depth l have no successors
- ❖ **Recursive implementation**



```
function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff  
  return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)
```

```
function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff  
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)  
  else if limit = 0 then return cutoff  
  else  
    cutoff_occurred?  $\leftarrow$  false  
    for each action in problem.ACTIONS(node.STATE) do  
      child  $\leftarrow$  CHILD-NODE(problem, node, action)  
      result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)  
      if result = cutoff then cutoff_occurred?  $\leftarrow$  true  
      else if result  $\neq$  failure then return result  
    if cutoff_occurred? then return cutoff else return failure
```

❖ Properties of DLS

❖ Completeness:

❖ Not complete if $l < d$ (G')

❖ Complete otherwise (G)

❖ Optimality: Not optimal in general (even if $l > d$)

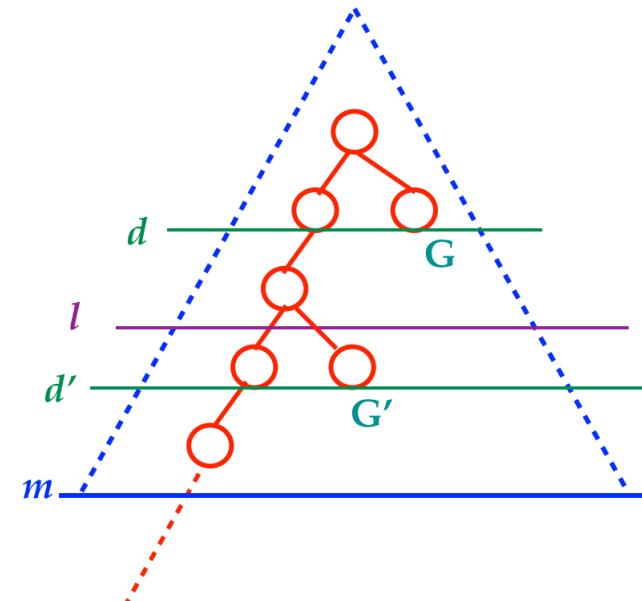
❖ Time complexity: $O(b^l)$

❖ Space complexity: $O(bl)$, linear space

❖ Terminate with two kinds of failure

❖ Failure: no solution

❖ Cutoff: no solution within the depth limit



❖ DFS

❖ Time complexity: $O(b^m)$

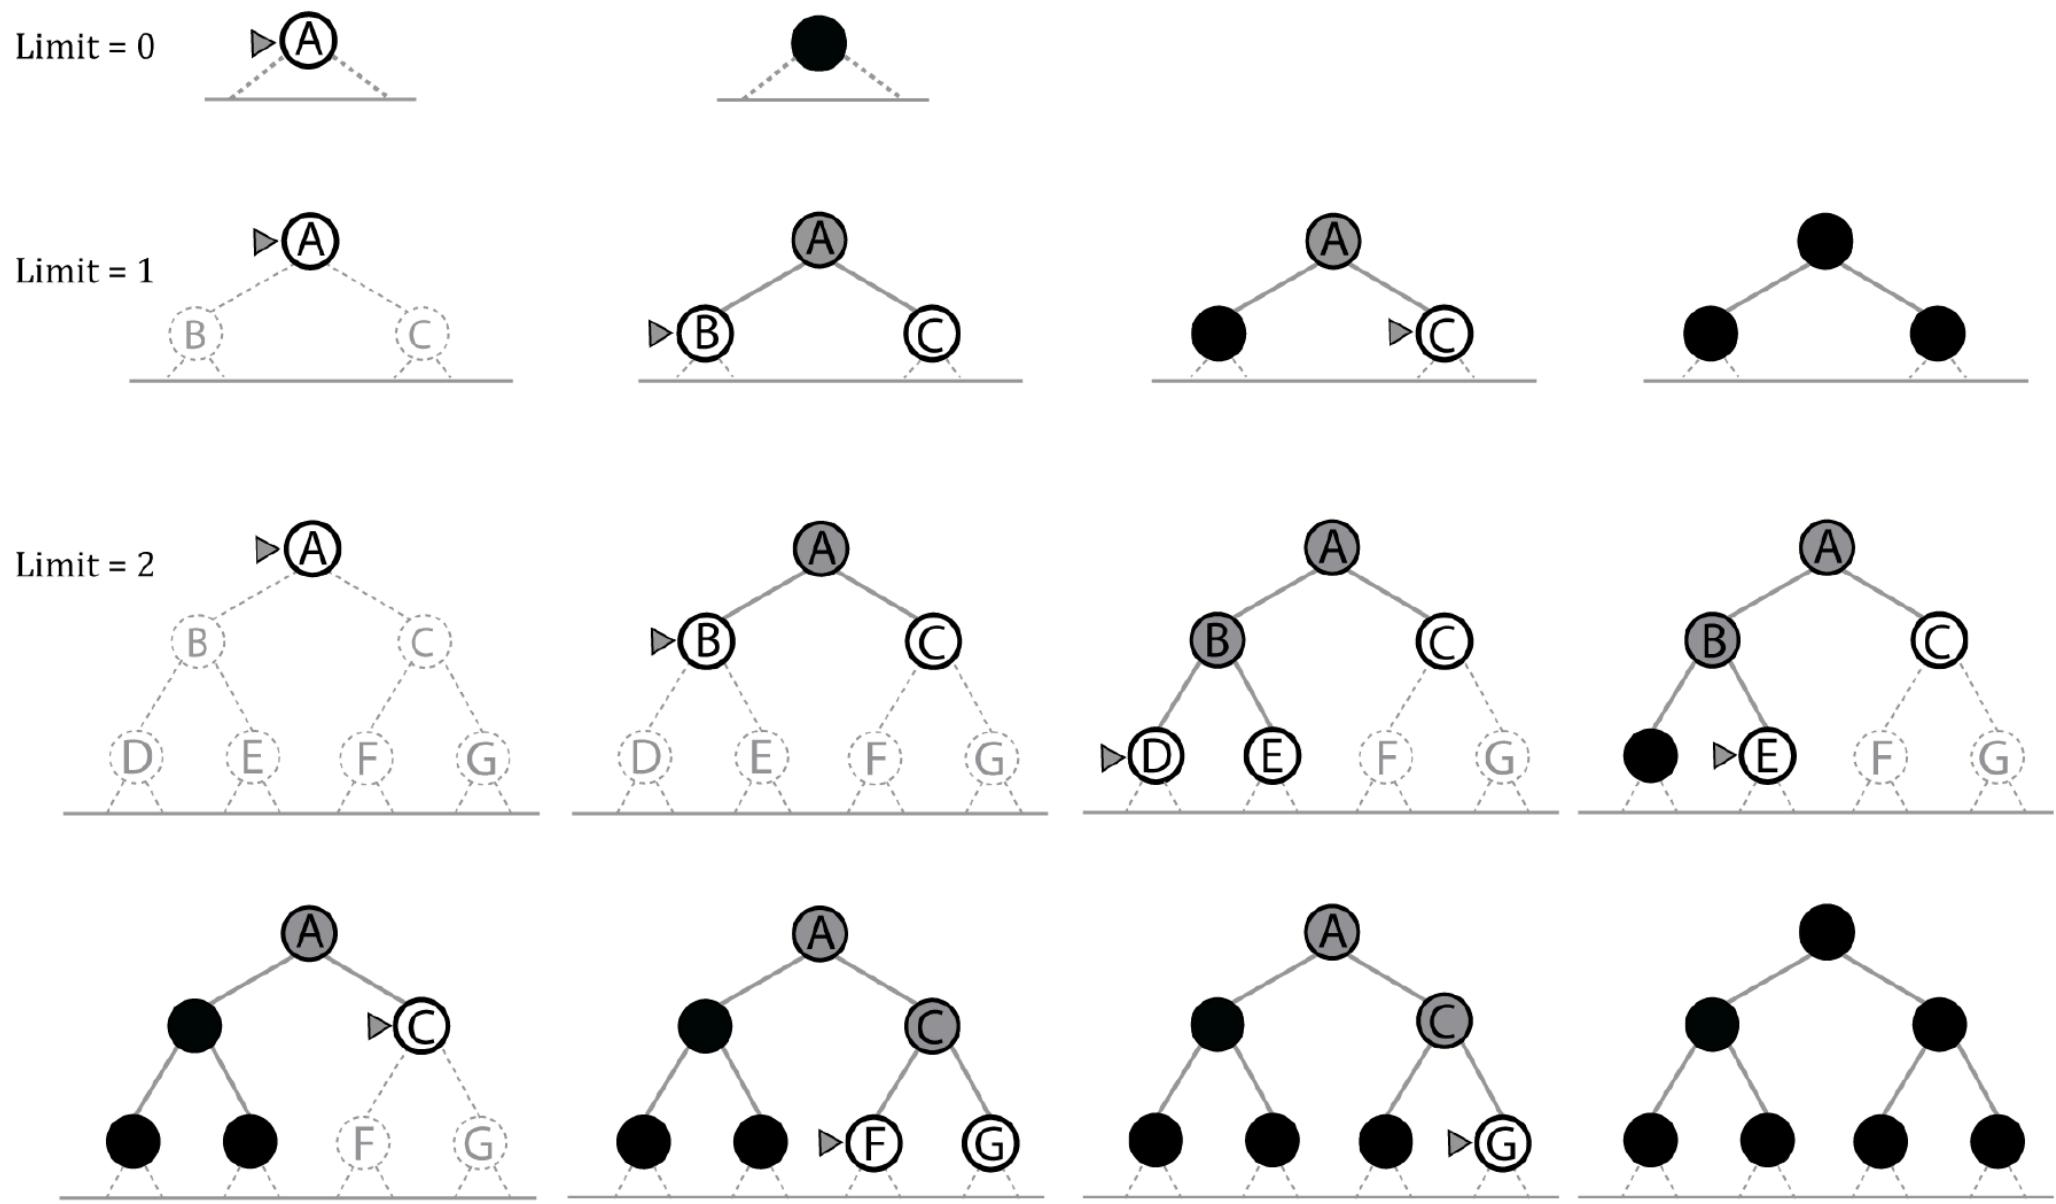
❖ Space complexity: $O(bm)$

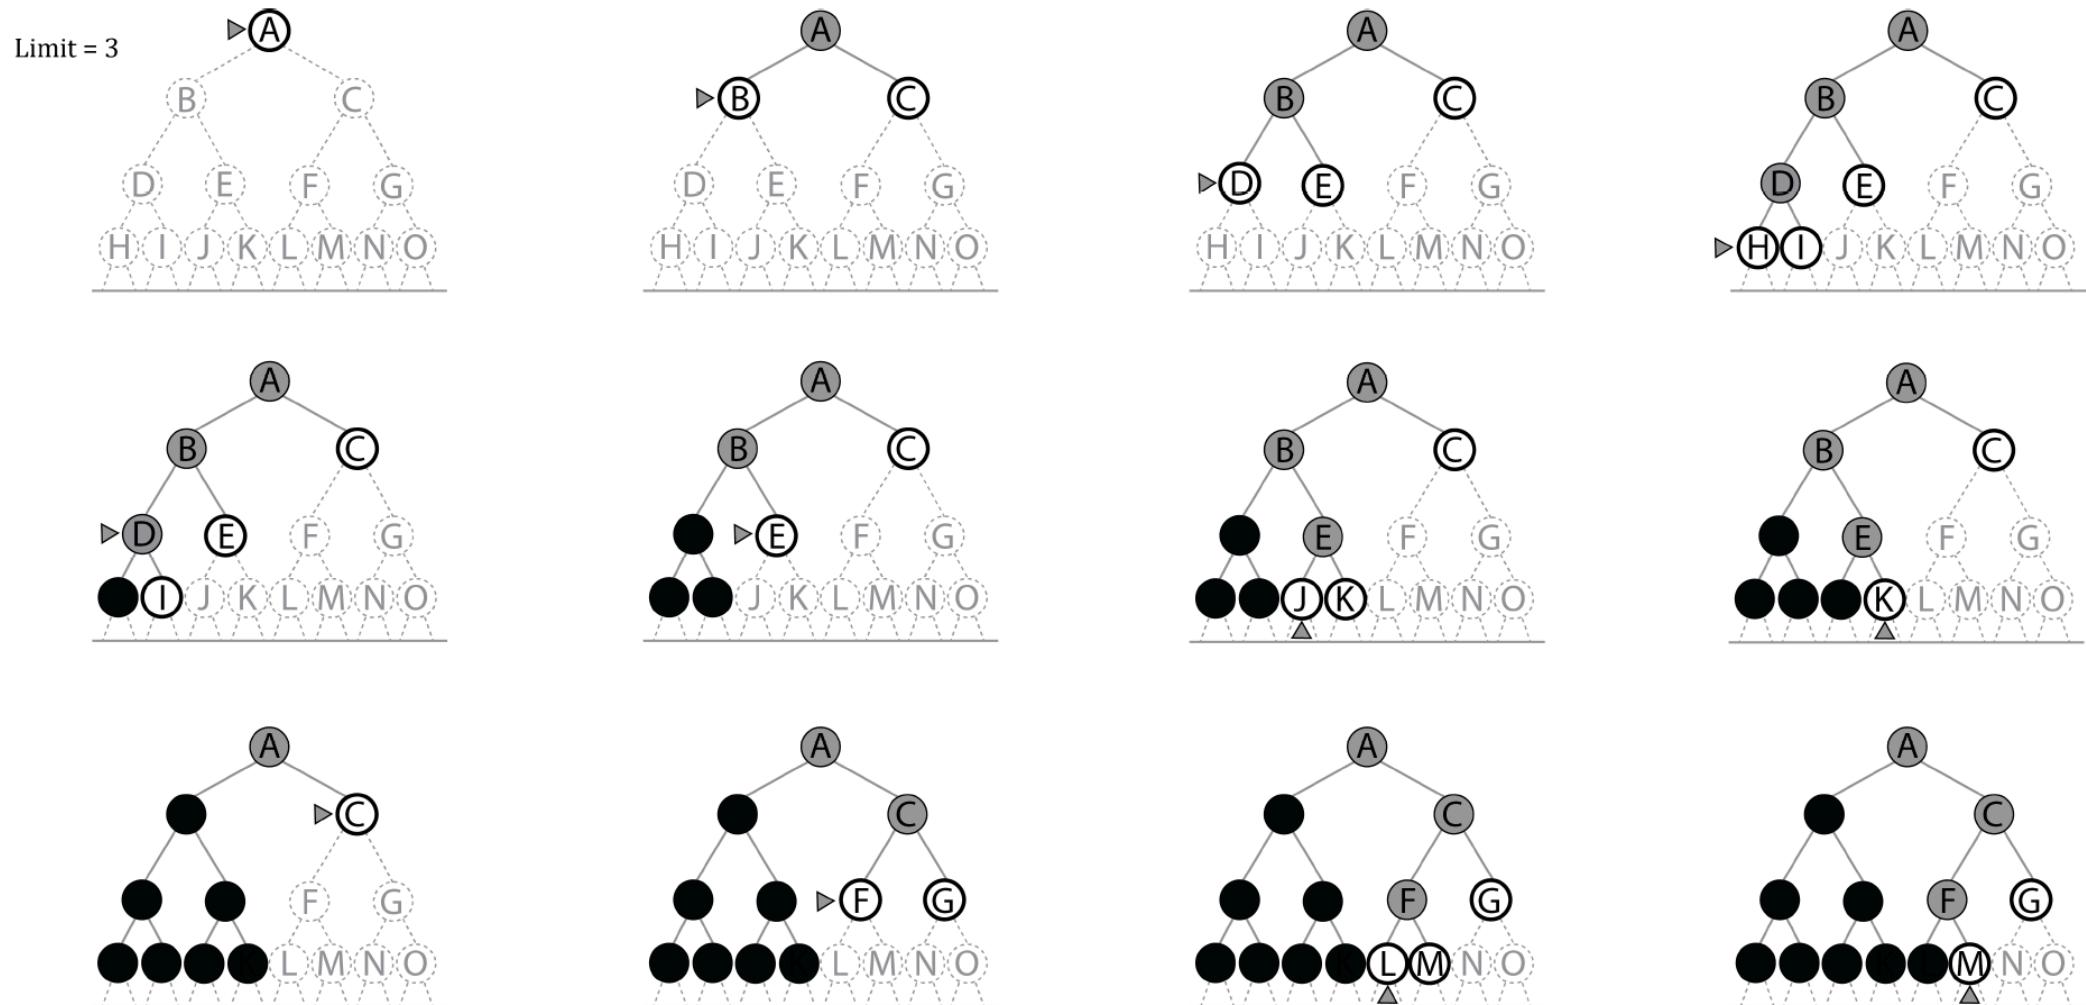
Iterative-Deepening Search (IDS)

- ❖ IDS is an improvement of DLS
- ❖ Combine the benefits of BFS and DFS
- ❖ Call DLS **iteratively** with **increasing depth limit**
- ❖ Use DFS as a subroutine
 - ❖ Check the root
 - ❖ Do a DFS searching for a path of length 1
 - ❖ If there is no path of length 1, do a DFS searching for a path of length 2
 - ❖ If there is no path of length 2, do a DFS searching for a path of length 3...

- ❖ IDS terminates when a **solution is found** or if the **depth limited search returns failure**, meaning that **no solution exists**

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure  
  for depth = 0 to  $\infty$  do  
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)  
    if result  $\neq$  cutoff then return result
```





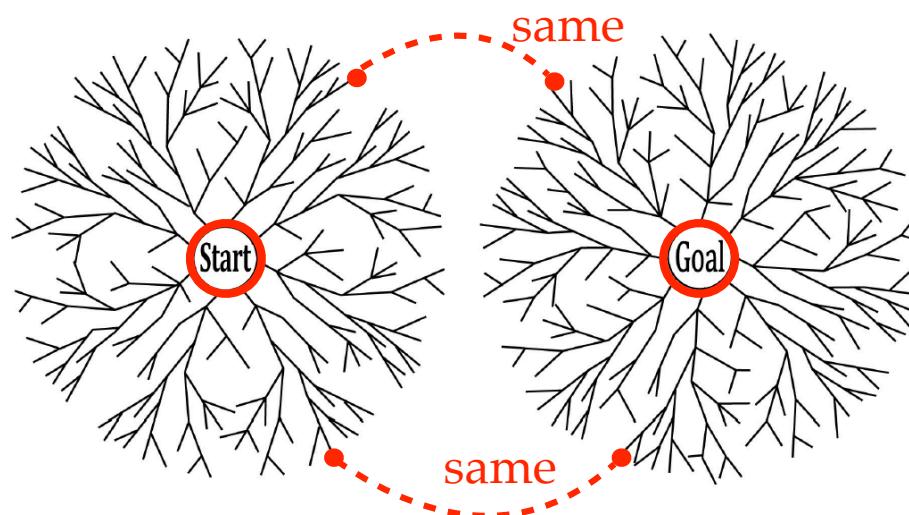
- ❖ Properties of Iterative Deepening Search
 - ❖ Completeness: Yes [$l = 0$ to ∞]
 - ❖ Optimality
 - ❖ No, for general cost functions
 - ❖ Yes, only if step cost = 1
 - ❖ Can be modified to explore uniform-cost tree (optimal)
 - ❖ e.g. Firstly search the nodes with path cost ≤ 10

- ❖ **Time complexity**
 - ❖ The nodes on the **bottom level** (depth d) are generated **once** — $(1)b^d$
 - ❖ Those on the **next-to-bottom level** are generated **twice**, and so on — $(2)b^{(d-1)}$
 - ❖ The children of the **root**, which are generated d times. — $(d)b^1$
 - ❖ Total: $O(b^d)$ is asymptotically the same as BFS
$$(d)b^1 + (d - 1)b^2 + \dots + (2)b^{(d-1)} + (1)b^d = O(b^d)$$
- ❖ **Space complexity:** $O(bd)$

- ❖ Numerical comparison (time complexity) for $b = 10$, $d = 5$, solution at **far right leaf**
 - ❖ $N(IDS) = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$ $(d)b^1 + (d - 1)b^2 + \dots + (2)b^{(d-1)} + (1)b^d = O(b^d)$
 - ❖ $N(BFS) = 10 + 100 + 1,000 + 10,000 + 100,000 = 111,100$ $b + b^2 + \dots + b^d = O(b^d)$
- ❖ Repeated search in IDS is **not severe**
- ❖ **IDS is preferred when search space is large and depth is unknown**
- ❖ More advantages of IDS in **adversarial search**

Bidirectional Search

- ❖ Idea
 - ❖ Simultaneously search **forward** from S and **backward** from G
 - ❖ Stop when both “**meet in the middle**”
 - ❖ Need to keep track of the **intersection** of 2 **open sets** of nodes



$$2 * b^{d/2}$$

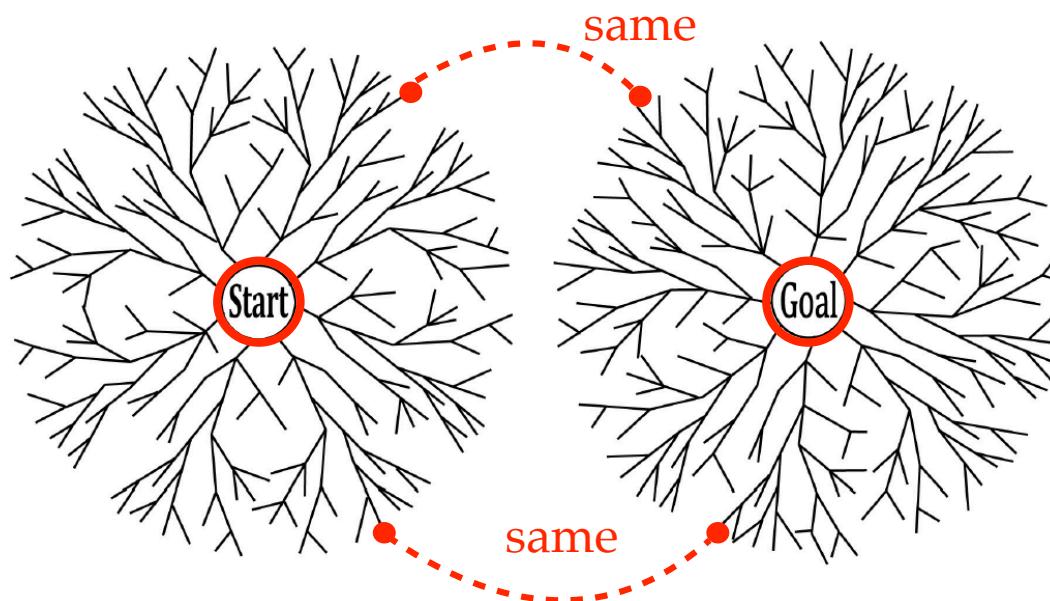
- ❖ **Time complexity**: reduce from $O(b^d)$ to $O(b^{d/2})$
 - ❖ Though the reduction is attractive, **how to search backward?**
 - ❖ Need a way to specify the **predecessors** of G
 - ❖ This can be difficult
 - ❖ e.g., predecessors of checkmate in chess?
 - ❖ Which to take if there are **multiple goal states**?
 - ❖ 8-puzzle : only one goal
 - ❖ 8-queen puzzle : many goals (Which are the predecessors?)

Problem:



e.g. Robot falls down and has no way to climb up

- ❖ **Space complexity**: increase from $O(bd)$ to $O(b^{d/2})$ [the space spent in saving the **matching** of the states from two ends - space complexity] as well, can be problematic



Succeed when a branch from the start node **meets** a branch from the goal node.

Comparison of Uninformed Search Strategies

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(bm)$	$O(bl)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

- ❖ $\textcolor{red}{b}$: the branching factor
- ❖ $\textcolor{red}{d}$: the depth of the shallowest solution
- ❖ $\textcolor{red}{m}$: the max depth of the search tree
- ❖ $\textcolor{red}{l}$: the depth limit

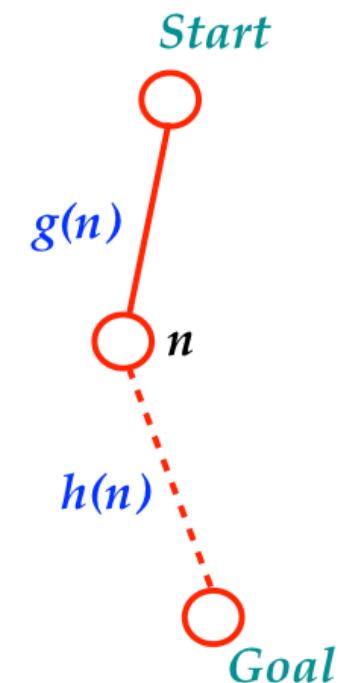
❖ *Superscript*

- ❖ $\textcolor{red}{a}$: complete if b is finite
- ❖ $\textcolor{red}{b}$: complete if step costs $\geq \epsilon$ for positive ϵ
- ❖ $\textcolor{red}{c}$: optimal if step costs are all identical
- ❖ $\textcolor{red}{d}$: if both directions use BFS

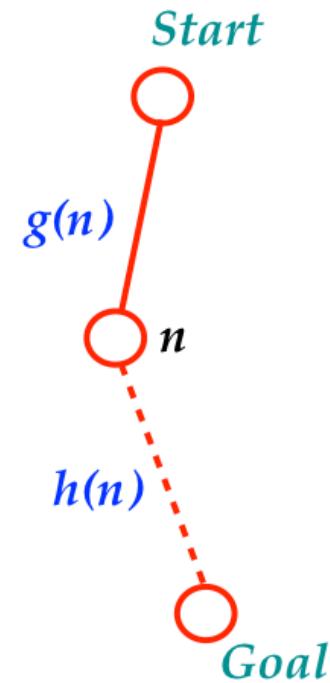
3.5 Informed (Heuristic) Search Strategies

- ❖ **Idea**
 - ❖ Give the algorithm “**hints**” about the **desirability** of different **states**
 - ❖ Use an **evaluation function** to **rank** nodes and select the **most promising one** for expansion
- ❖ **Heuristic**
 - ❖ **Definition:** a **rule of thumb** (or set of rules, e.g. straight-line distance) intended to increase the probability of solving some problem

- ❖ Heuristic function
 - ❖ $h(n)$: estimates the (optimal) **remaining cost** of reaching goal from node n
 - ❖ Defined using only the **state** of node n
 - ❖ $h(n) = 0$ if n is a **goal** node



- ❖ **Best-First Search**
 - ❖ Greedy search ($h(n)$)
 - ❖ A* search ($g(n) + h(n)$)
 - ❖ Optimality of A*
- ❖ **Memory Bounded Search**
 - ❖ Iterative deepening A*
 - ❖ Recursive best-first search
 - ❖ Simplified memory-bounded A*
- ❖ **Heuristic**
 - ❖ Performance
 - ❖ Generating heuristics



Best-First Search

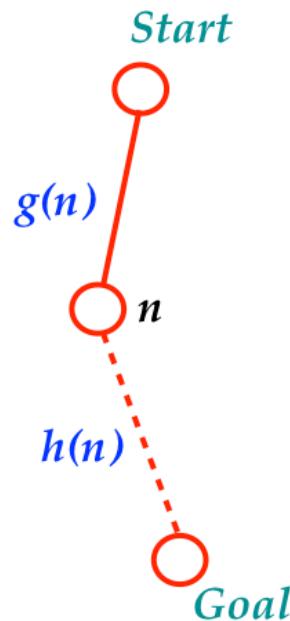
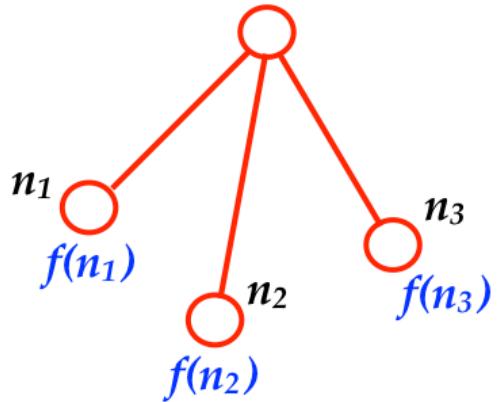
- ❖ **Informed** search, a.k.a. **heuristic** search
- ❖ **Idea:** Use an **evaluation function** for each node to estimate the **desirability**
 - ❖ Expand the most desirable unexpanded node
- ❖ The **evaluation function** is called **heuristic**, denoted as $h(n)$
 - ❖ It estimates the cost from node n to the **closest (best) goal**

- ❖ Special cases
 - ❖ Greedy search, $f(n) = h(n)$
 - ❖ Uniform Cost Search, $f(n) = g(n)$
path cost
 - ❖ A* search, $f(n) = g(n) + h(n)$

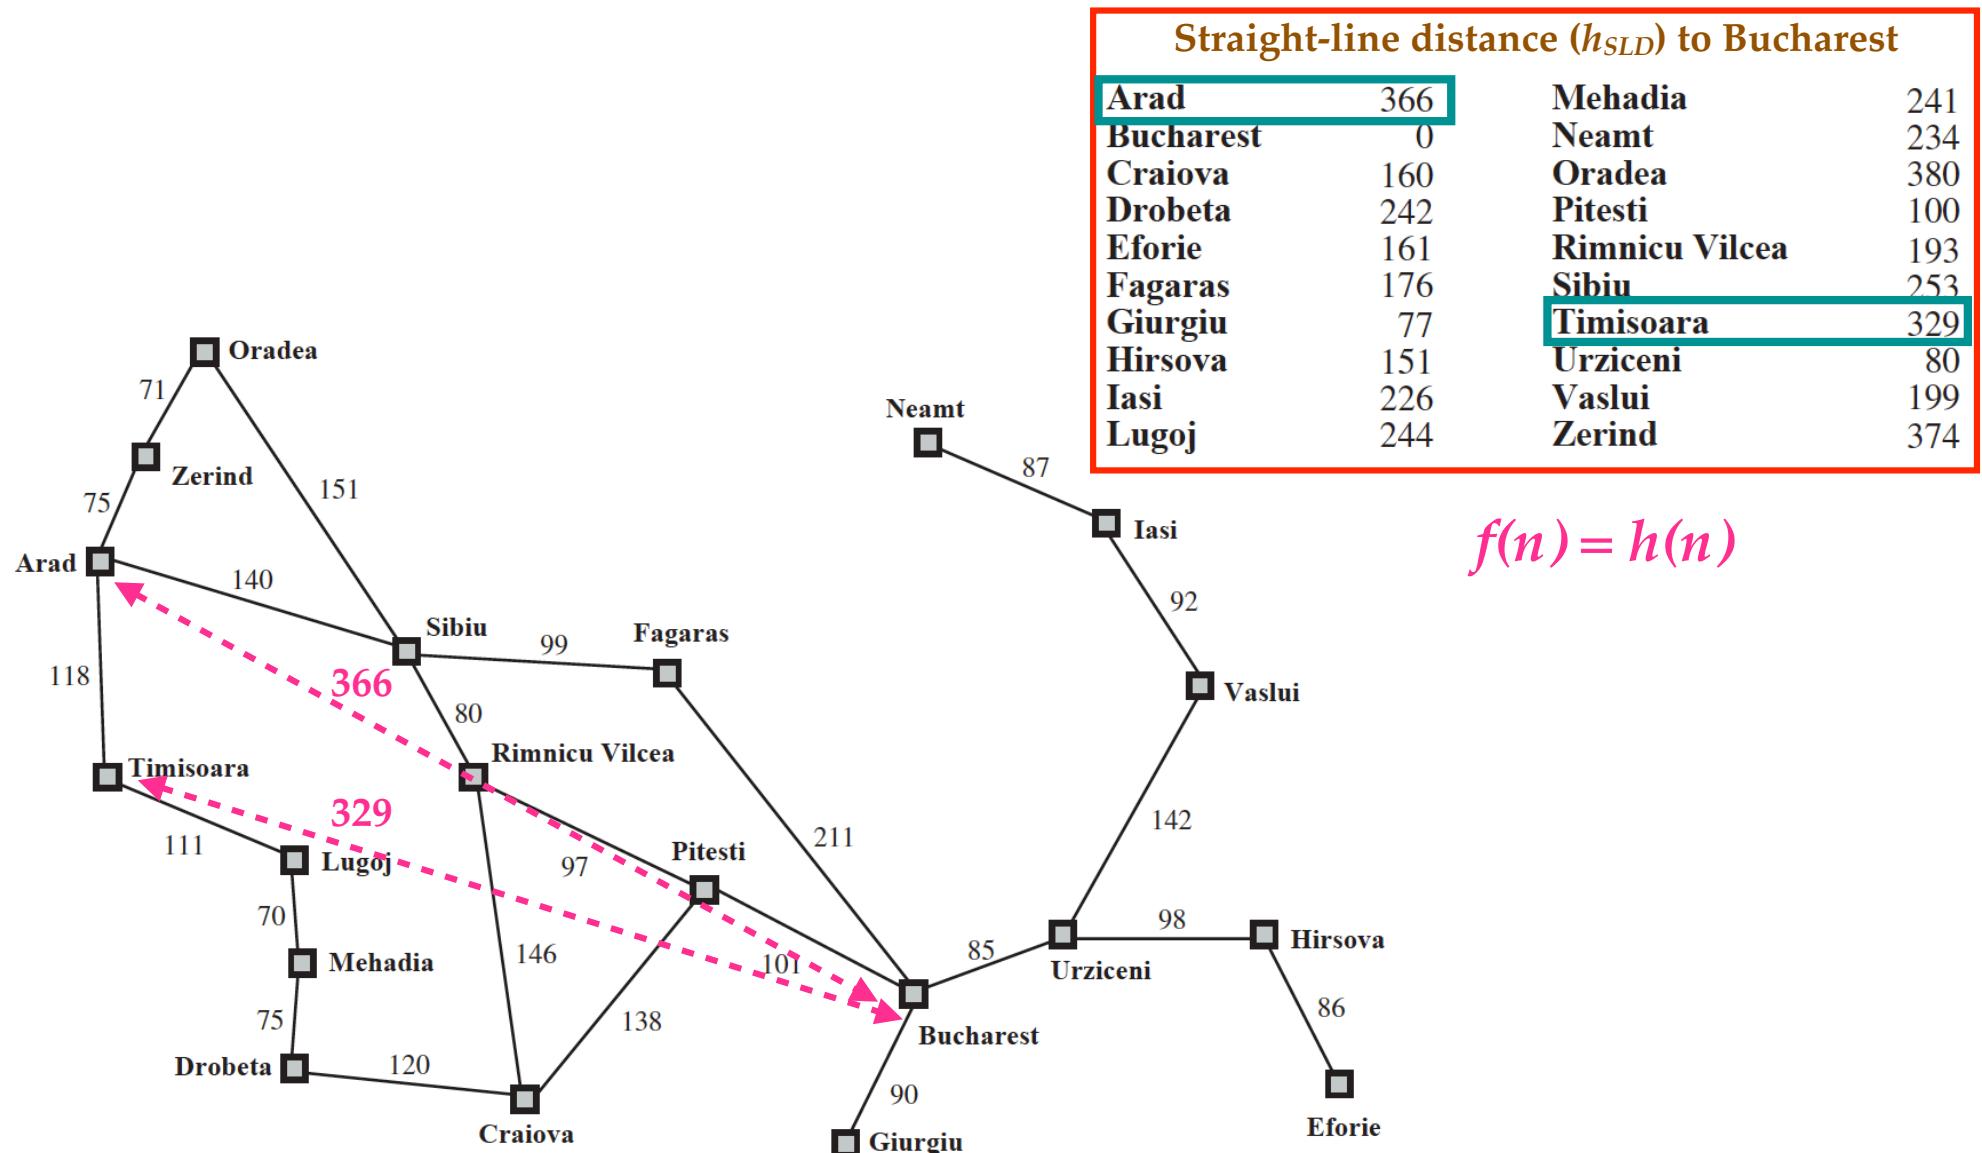
A* Search \approx Uniform Cost Search [$g(n)$] + Greedy Search [$h(n)$]

may be optimal

well-designed $h(n)$ may be sooner
to reach the goal

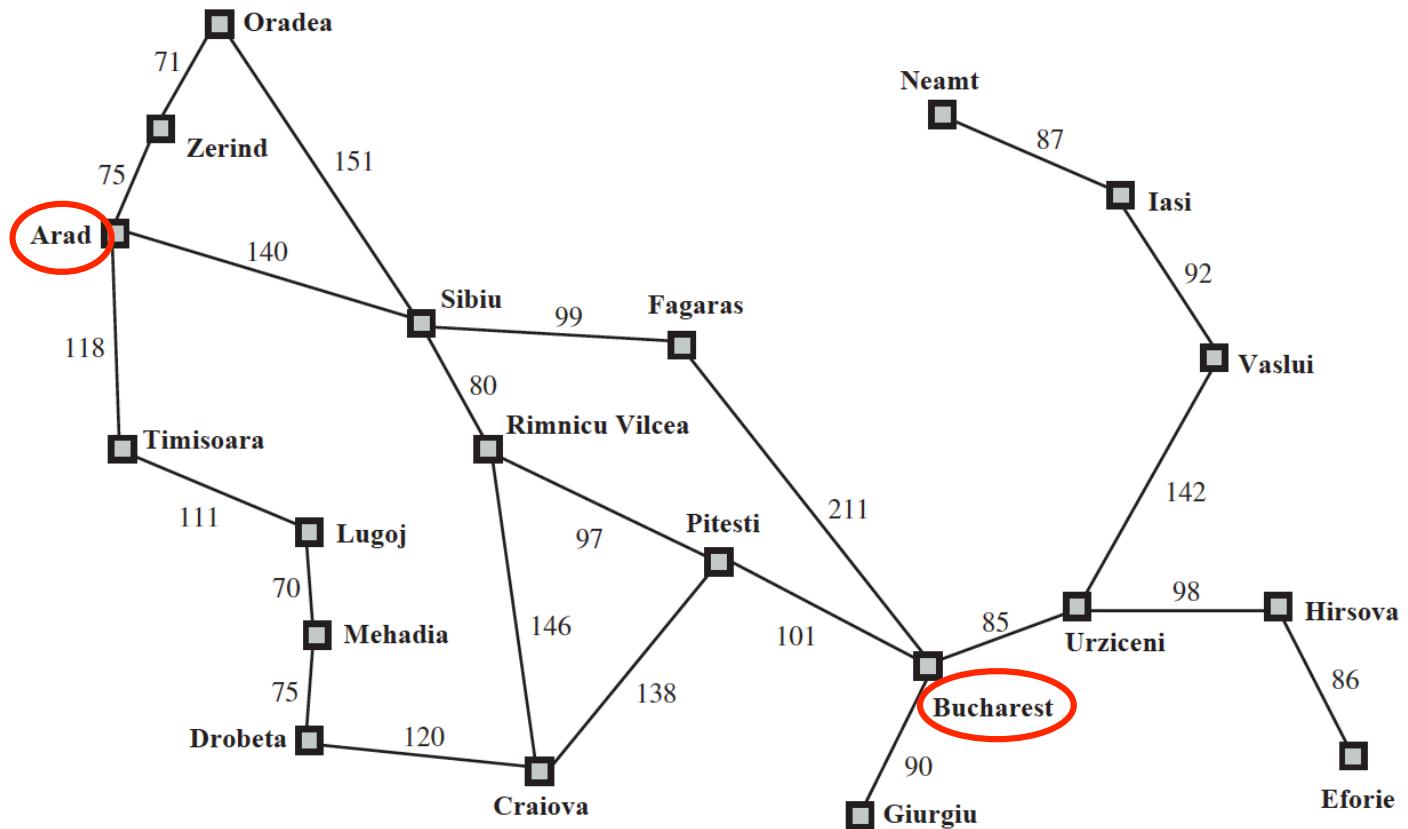


❖ Heuristic ($h(n)$) for the Romania problem



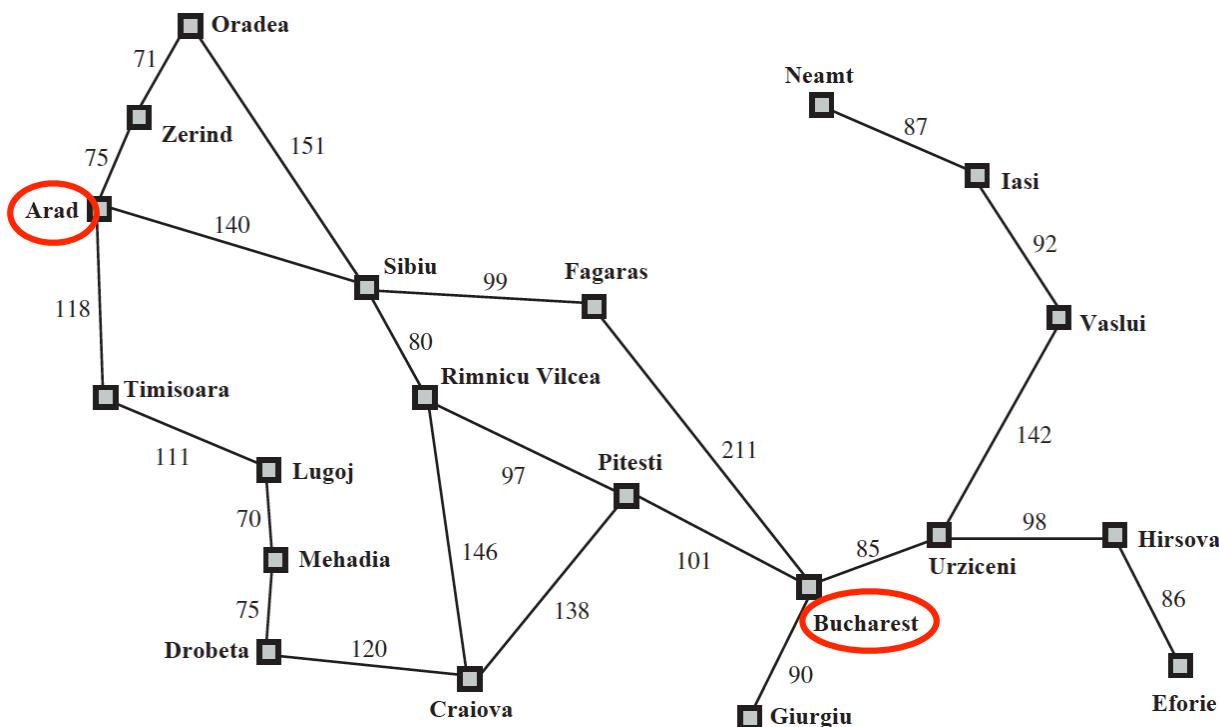
Greedy Best-First Search

- ❖ Expand the node that has the **lowest** value of the heuristic function $h(n)$
- ❖ Example



- ❖ Stages in a greedy best-first tree search for Bucharest with the **straight-line distance** heuristic h_{SLD}
- ❖ Nodes are labeled with their h -values

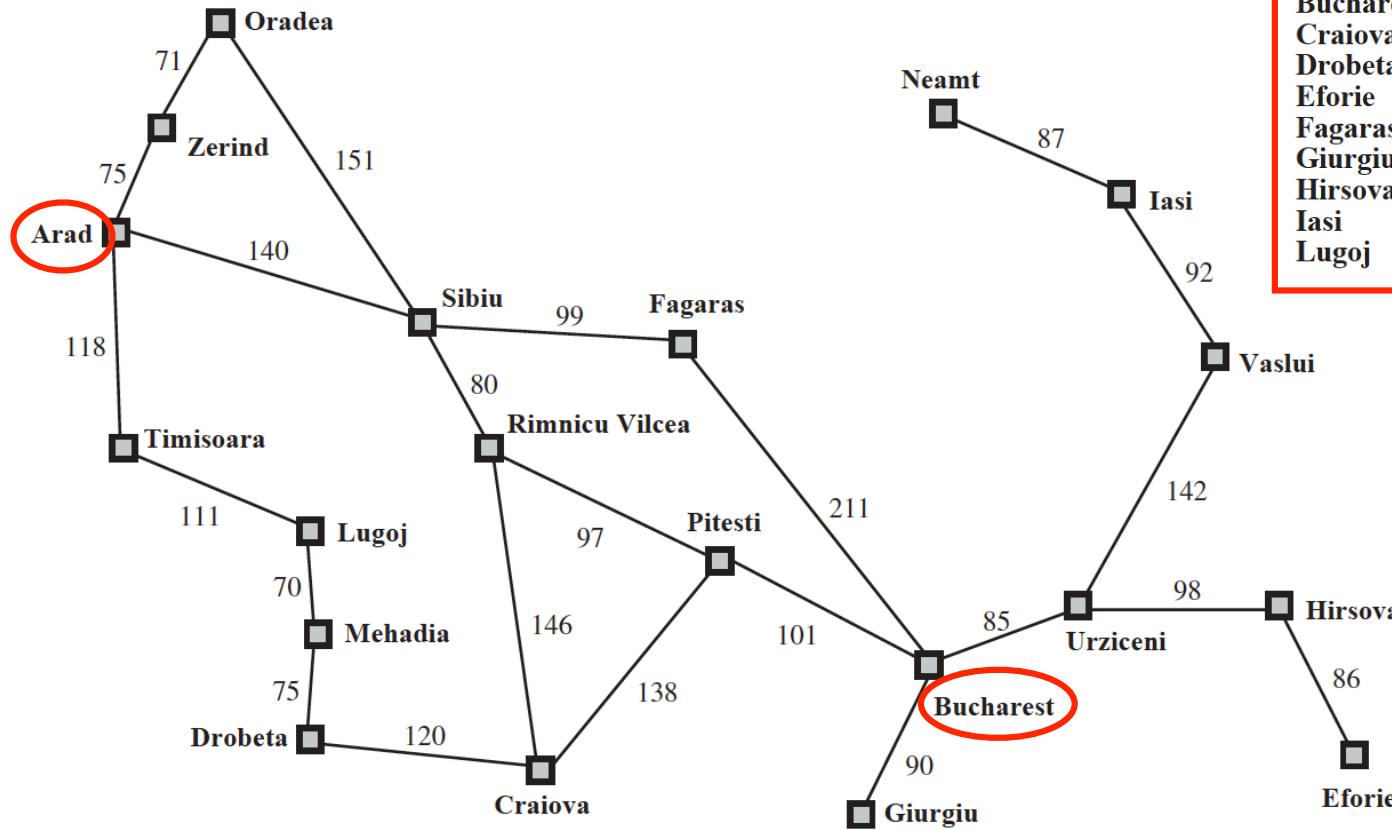
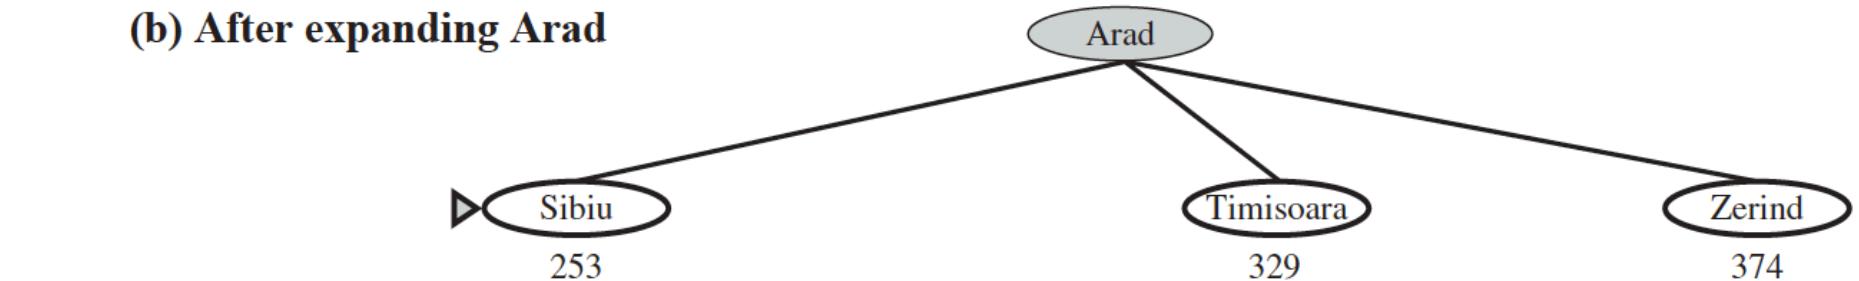
(a) The initial state



Straight-line distance (h_{SLD}) to Bucharest

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

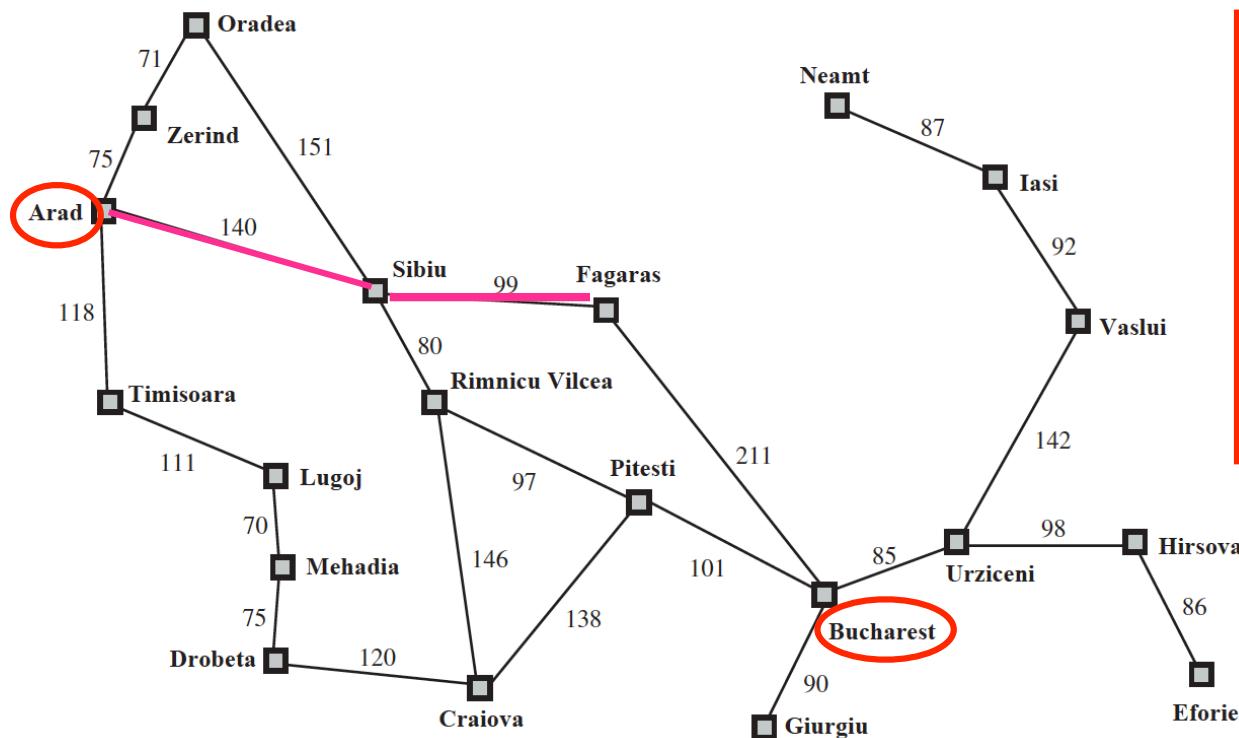
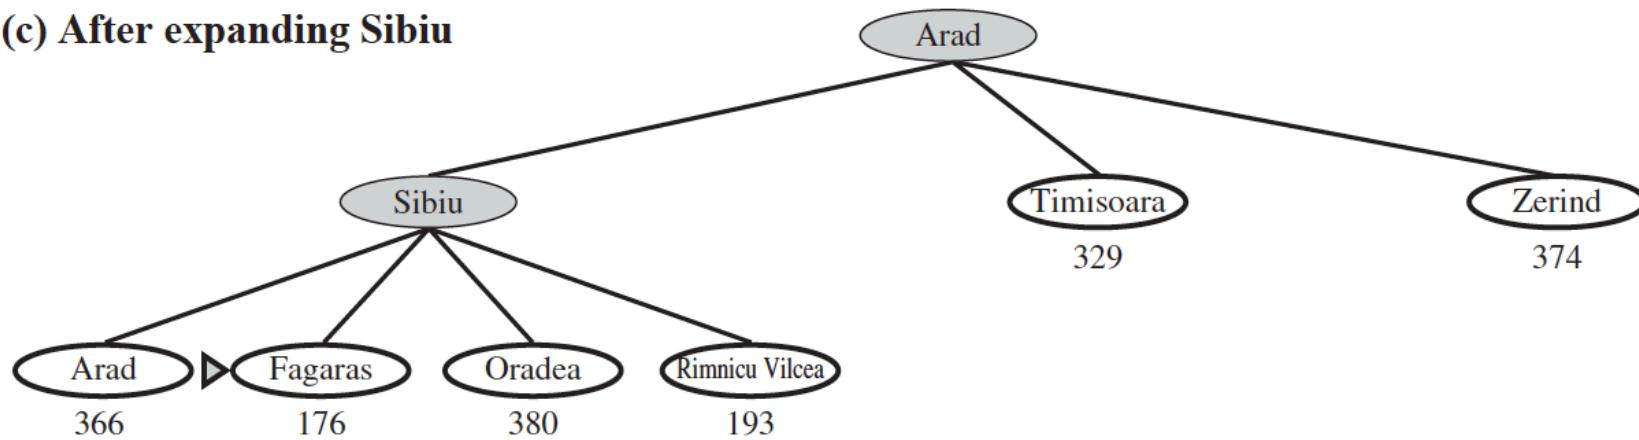
(b) After expanding Arad



Straight-line distance (h_{SLD}) to Bucharest

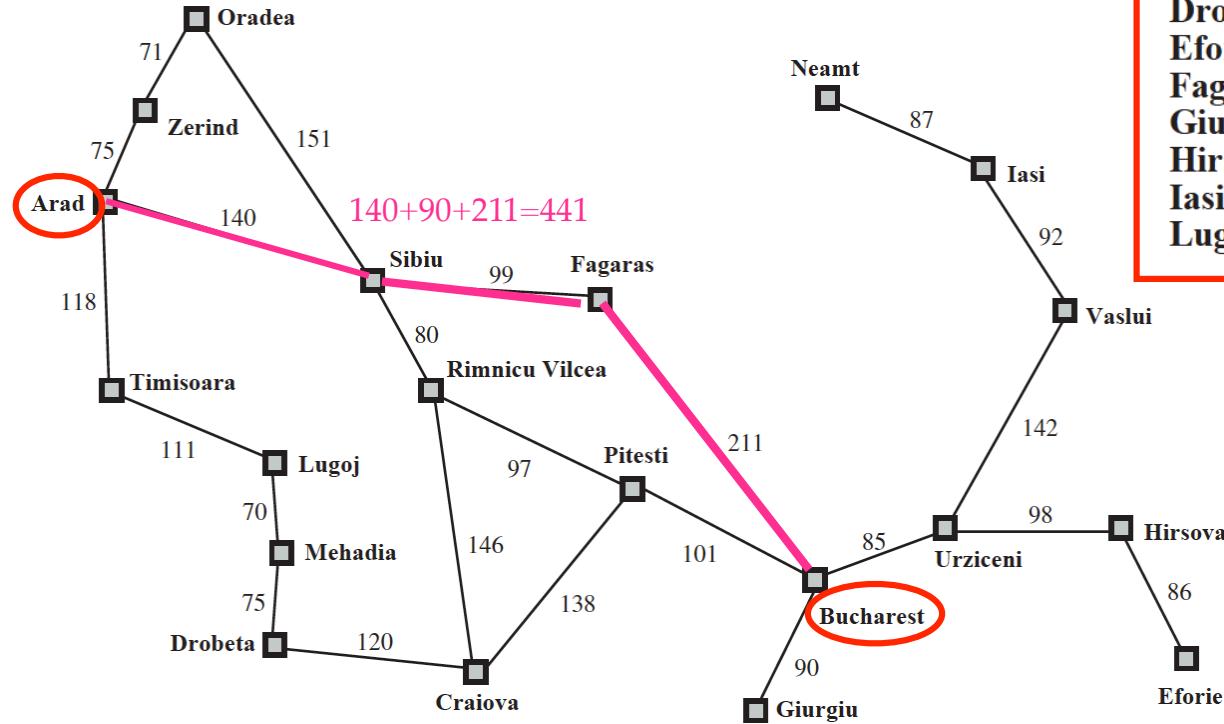
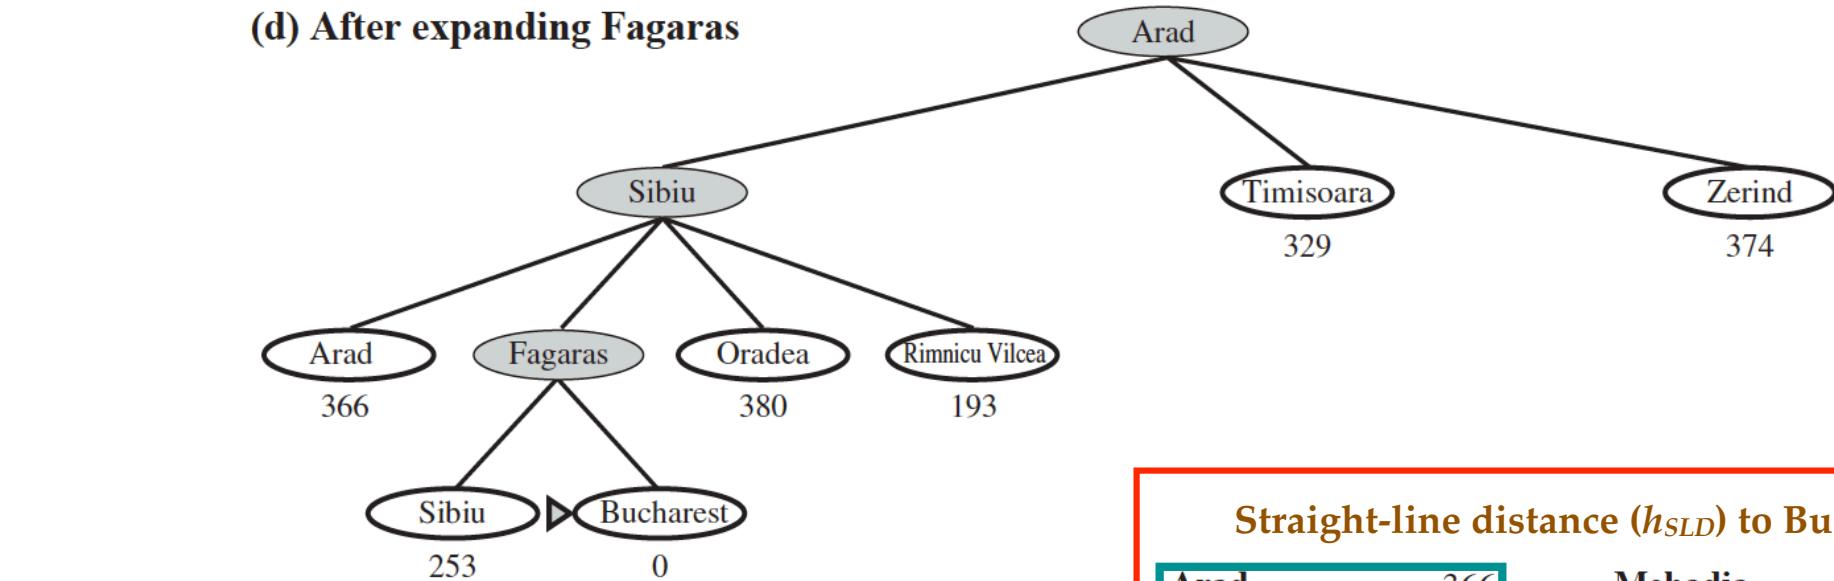
Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

(c) After expanding Sibiu



Straight-line distance (h_{SLD}) to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Drobeta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

(d) After expanding Fagaras



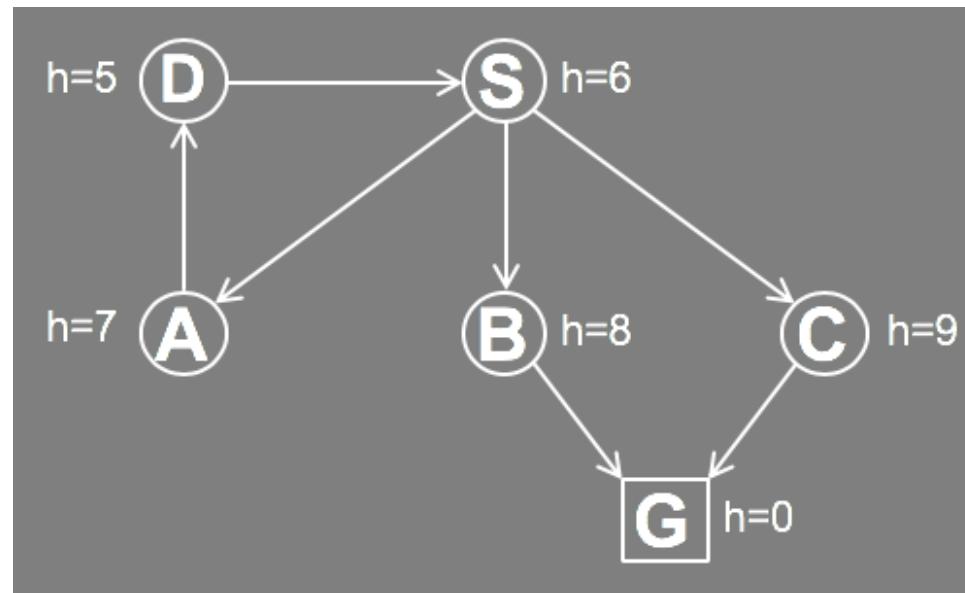
Straight-line distance (h_{SLD}) to Bucharest

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

- ❖ Properties of Greedy Search
 - ❖ Completeness: No
 - ❖ Tree-Search (no repeated states check) may get stuck in loops and never reach any goal even infinite state spaces

Order of node expansion: S A D S A D S A D...

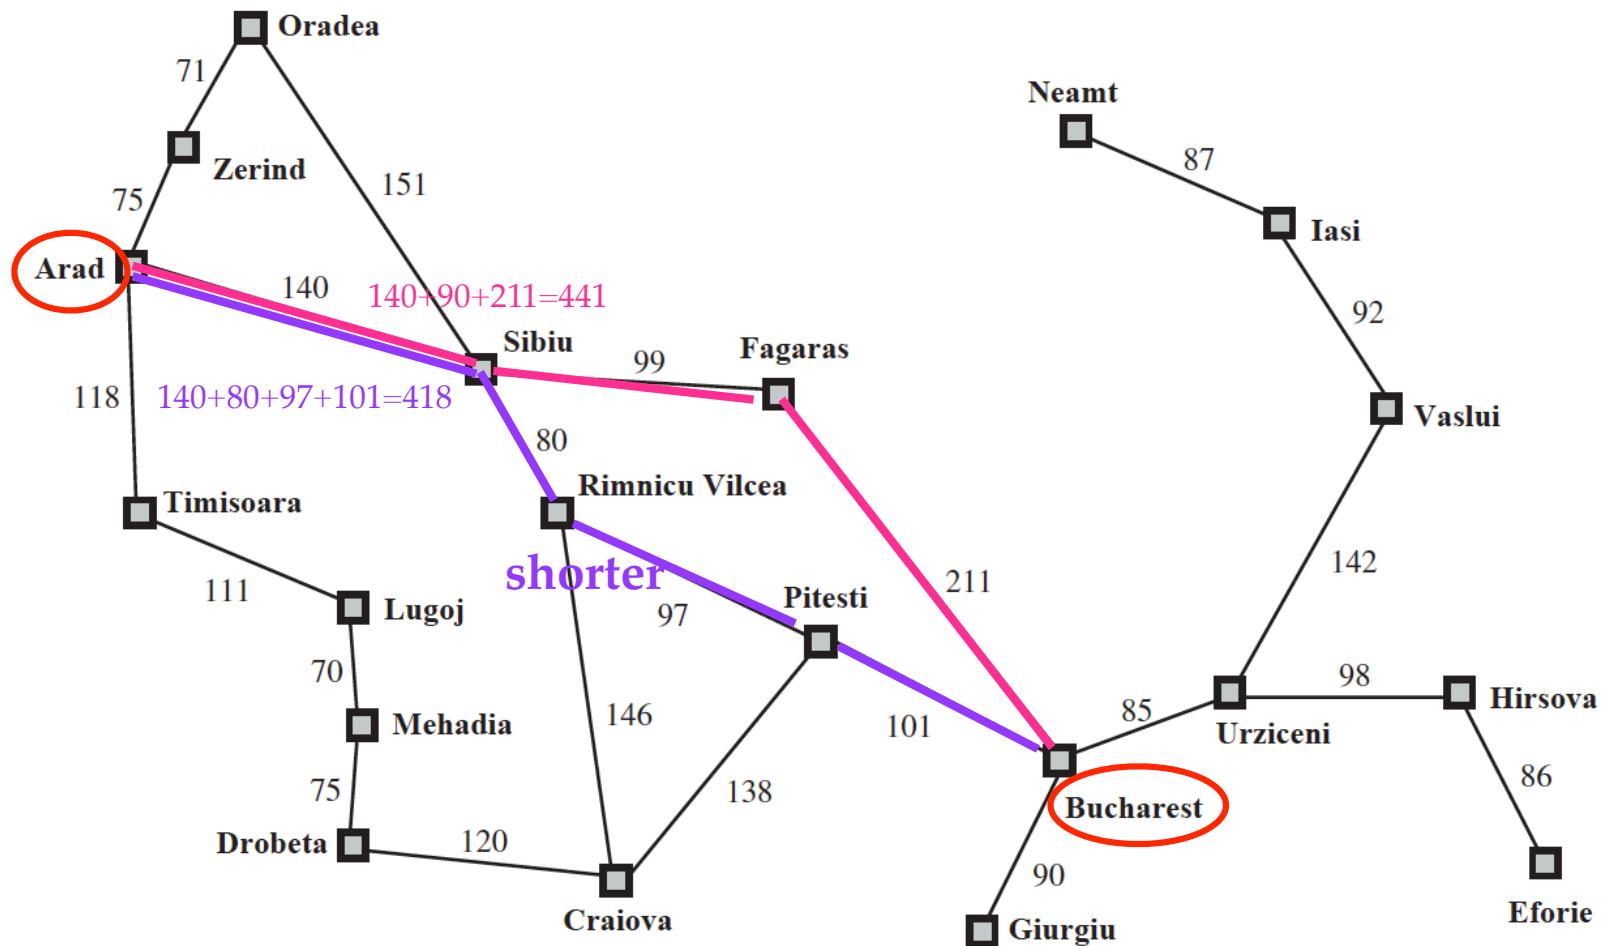
Path found: none Cost of path found: none



- ❖ Graph-Search (check repeated states) is complete in finite spaces, but not complete in infinite ones

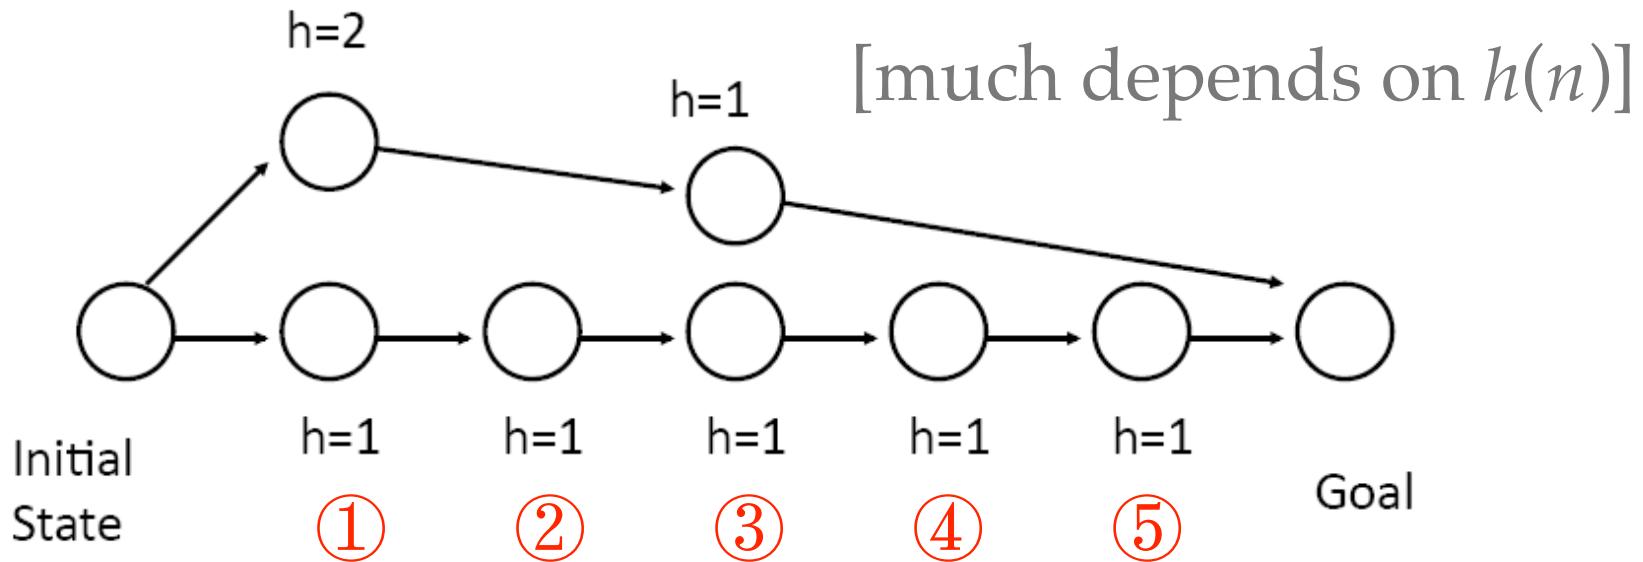
❖ Optimality: No

❖ E.g., Arad \rightarrow Sibiu \rightarrow Rimnicu Vilcea \rightarrow Pitesti \rightarrow Bucharest is shorter!



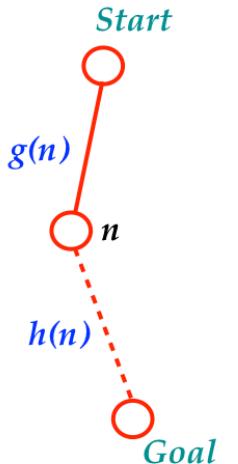
- ❖ **Time complexity** [much depends on $h(n)$]
 - ❖ Worst case: $O(b^m)$
 - ❖ Best case: $O(bd)$ – If $h(n)$ is 100% accurate (cheating?)
 - ❖ i.e., time complexity highly depends on the $h(n)$ function
- ❖ **Space complexity**: $O(b^m)$, since it keeps all nodes in memory

- ❖ How can we fix the greedy problem?



→ A* search helps to fix the greedy problem.

A* Search



- ❖ **Objective:** minimizing the **total estimated solution cost**
- ❖ **Idea:** Avoid expanding paths that are **already expensive**
- ❖ **Generally, A* is the preferred simple heuristic search**
- ❖ Evaluation function: $f(n) = g(n) + h(n)$
 - ❖ [Uniform Cost] $g(n)$: cost so far to reach n (path cost)
 - ❖ [Greedy Best First] $h(n)$: estimated (optimal) cost to goal from node n (heuristic)
 - ❖ [A*] $f(n)$: estimated total cost from the starting node to goal through n

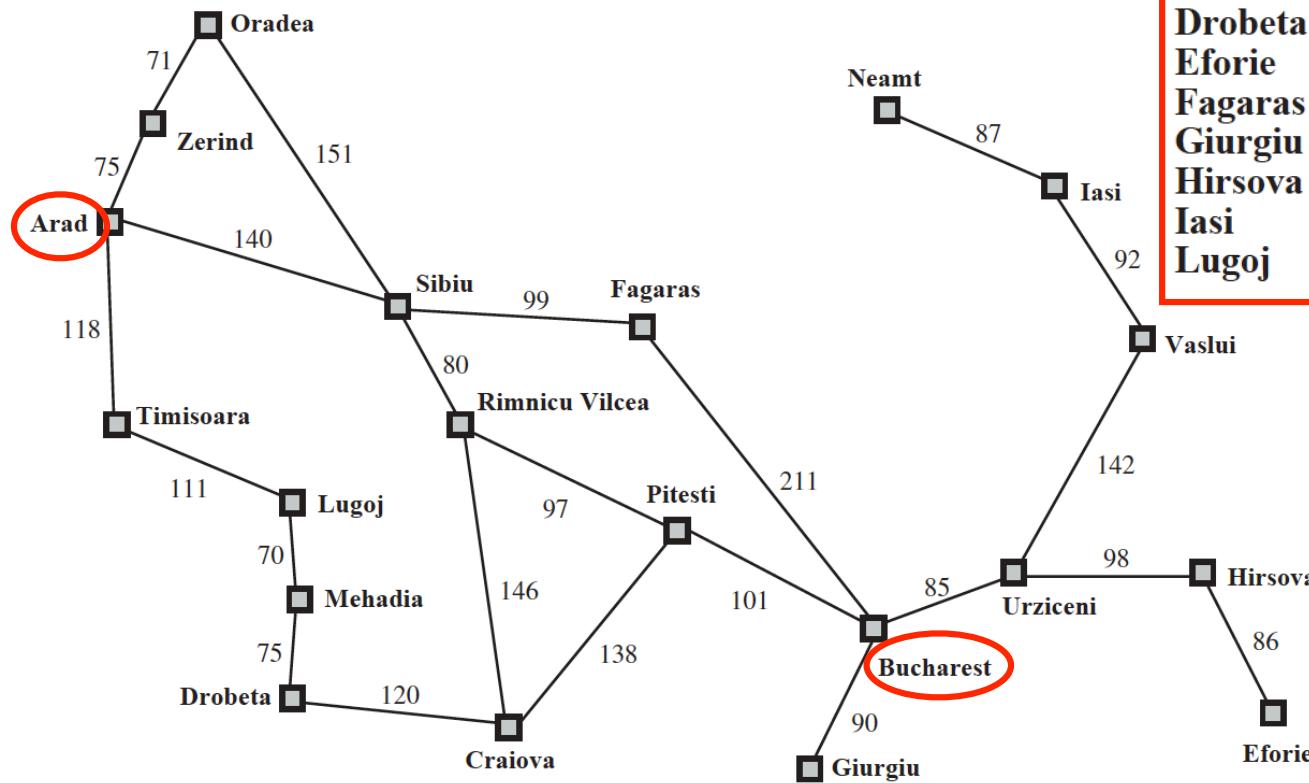
- ❖ A* search combines the advantages of the uniform-cost search and the greedy search

A* Search \approx Uniform Cost Search [$g(n)$] + Greedy Search [$h(n)$]

- ❖ Memory-efficient versions of A* are available
 - ❖ Recursive Best-First Search (RBFS)
 - ❖ Simplified Memory-Bounded A* (SMA*)

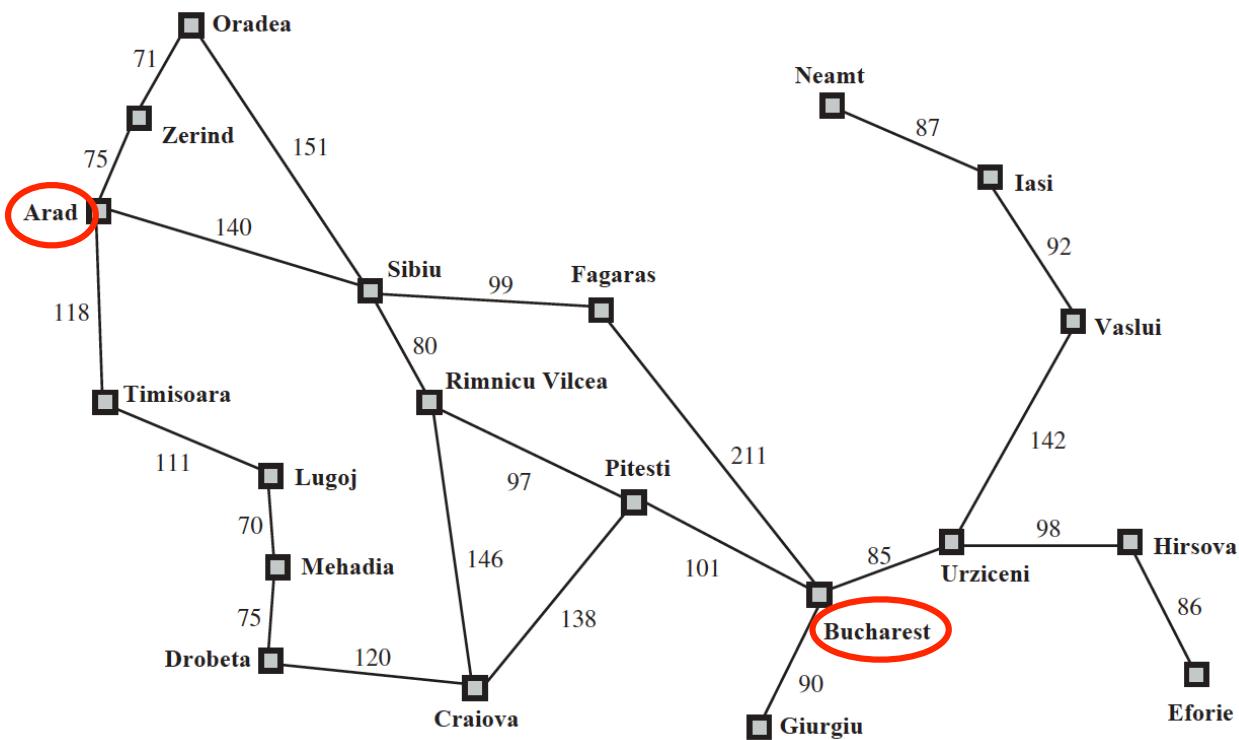
❖ Example

- ❖ Stages in an A* search for Bucharest
- ❖ Nodes are labeled with $f = g + h$
- ❖ The h values are the straight-line distances to Bucharest



Straight-line distance (h_{SLD})			
Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

(a) The initial state

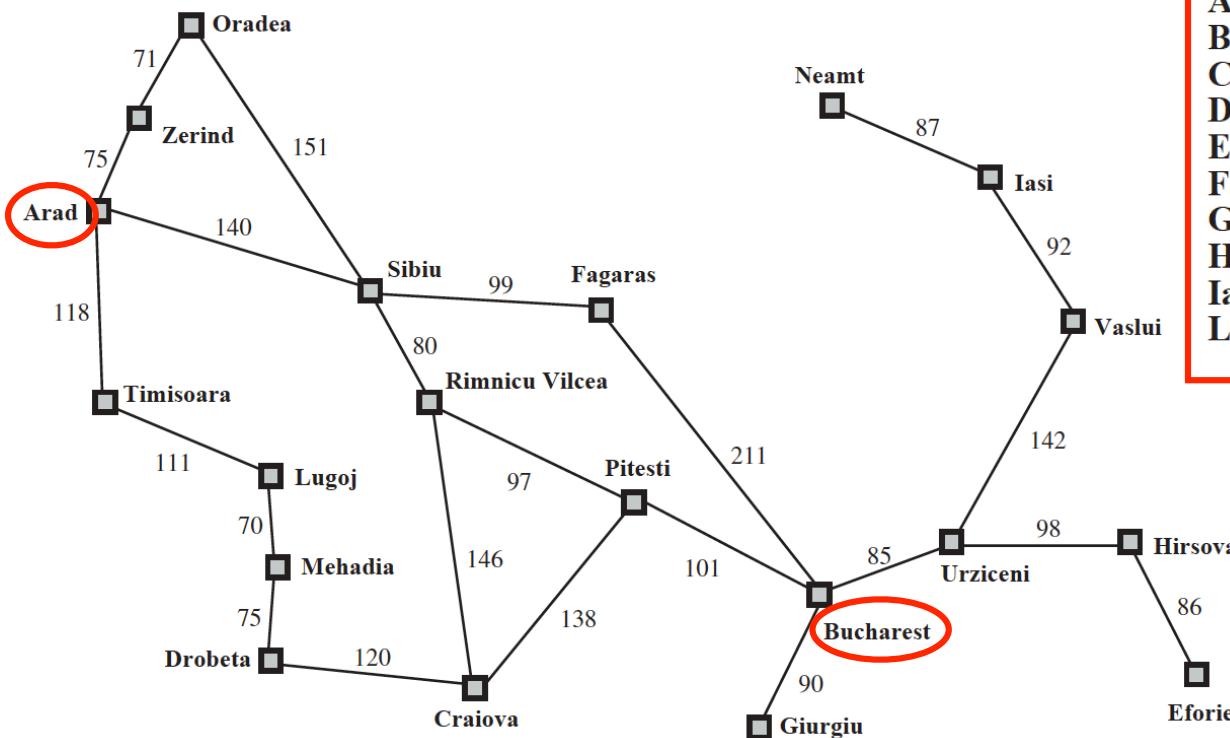
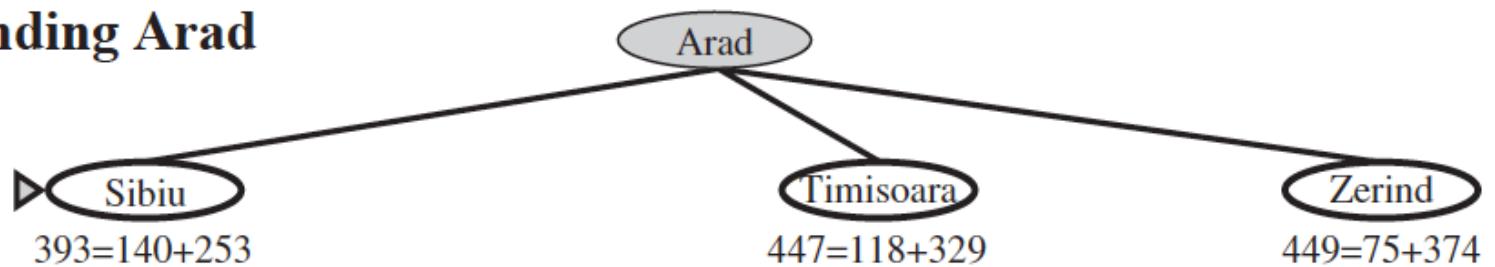


► Arad
366=0+366

Straight-line distance (h_{SLD})

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

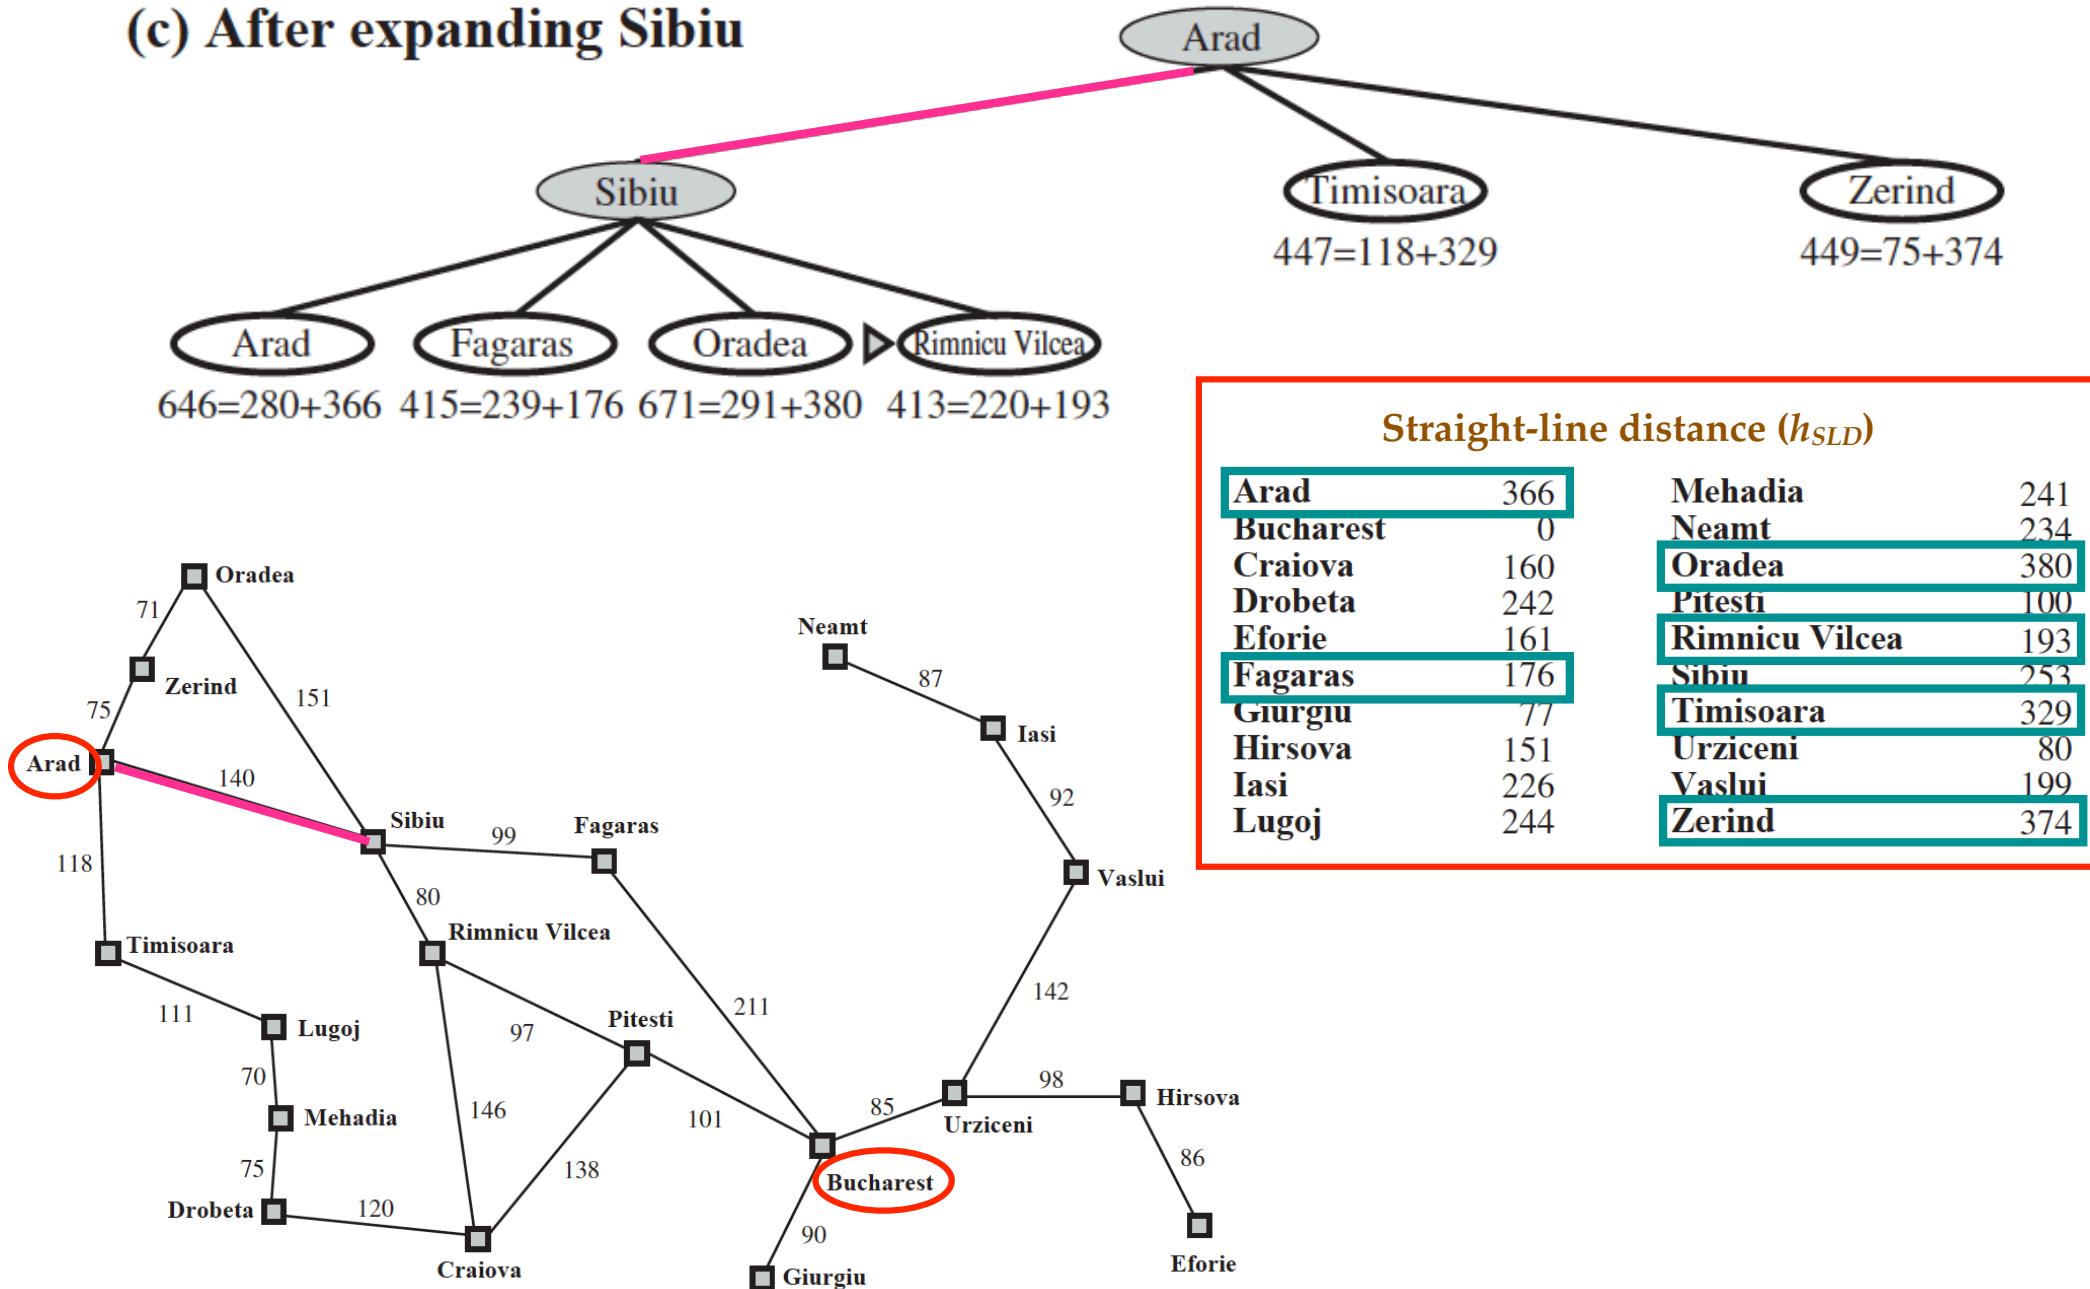
(b) After expanding Arad



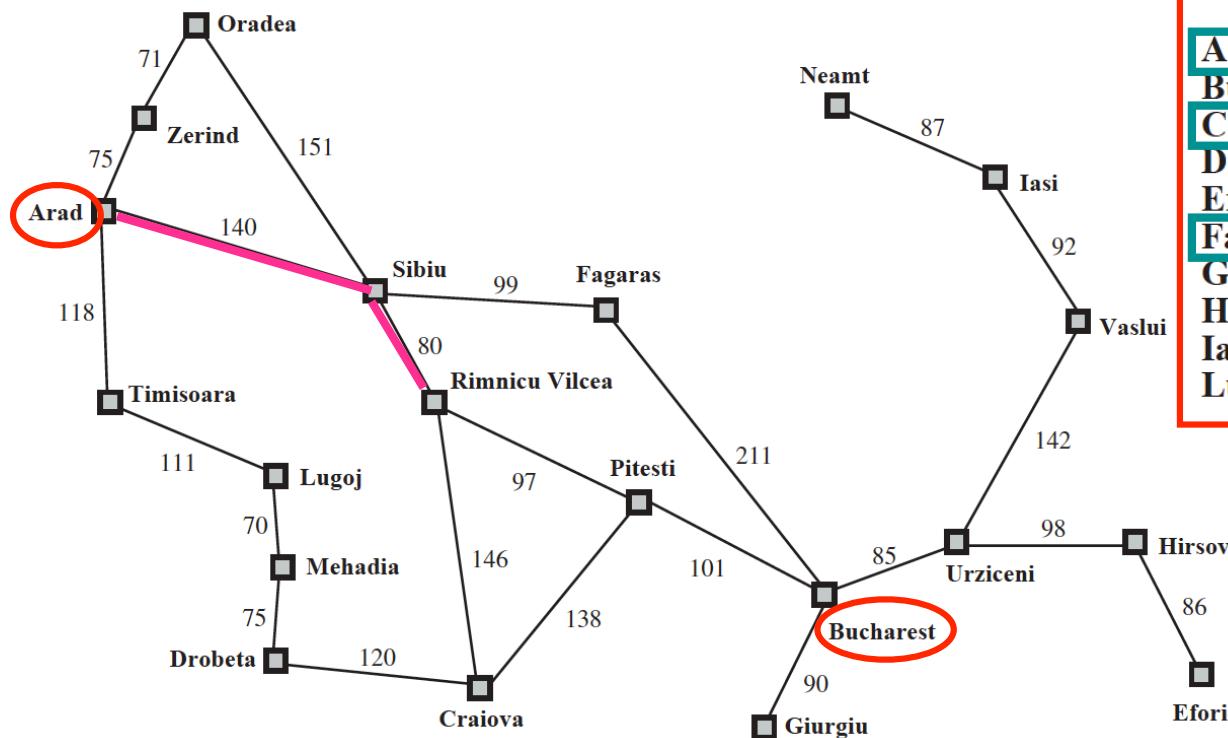
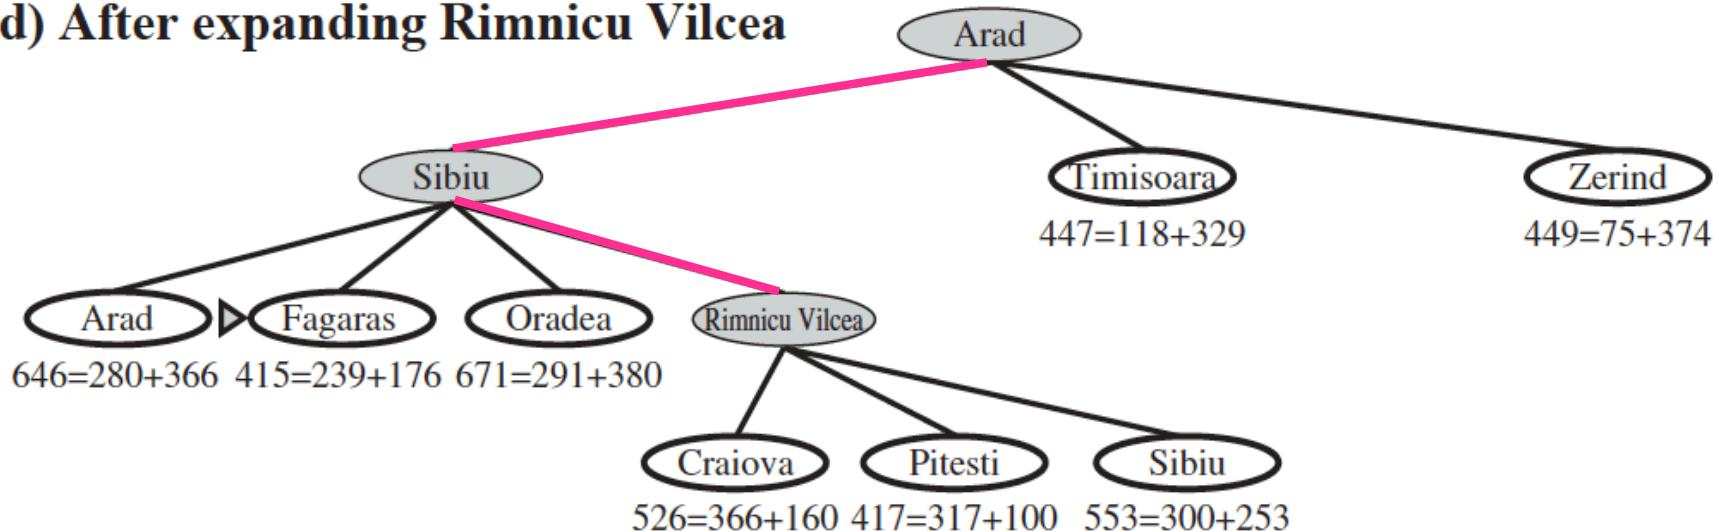
Straight-line distance (h_{SLD})

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

(c) After expanding Sibiu

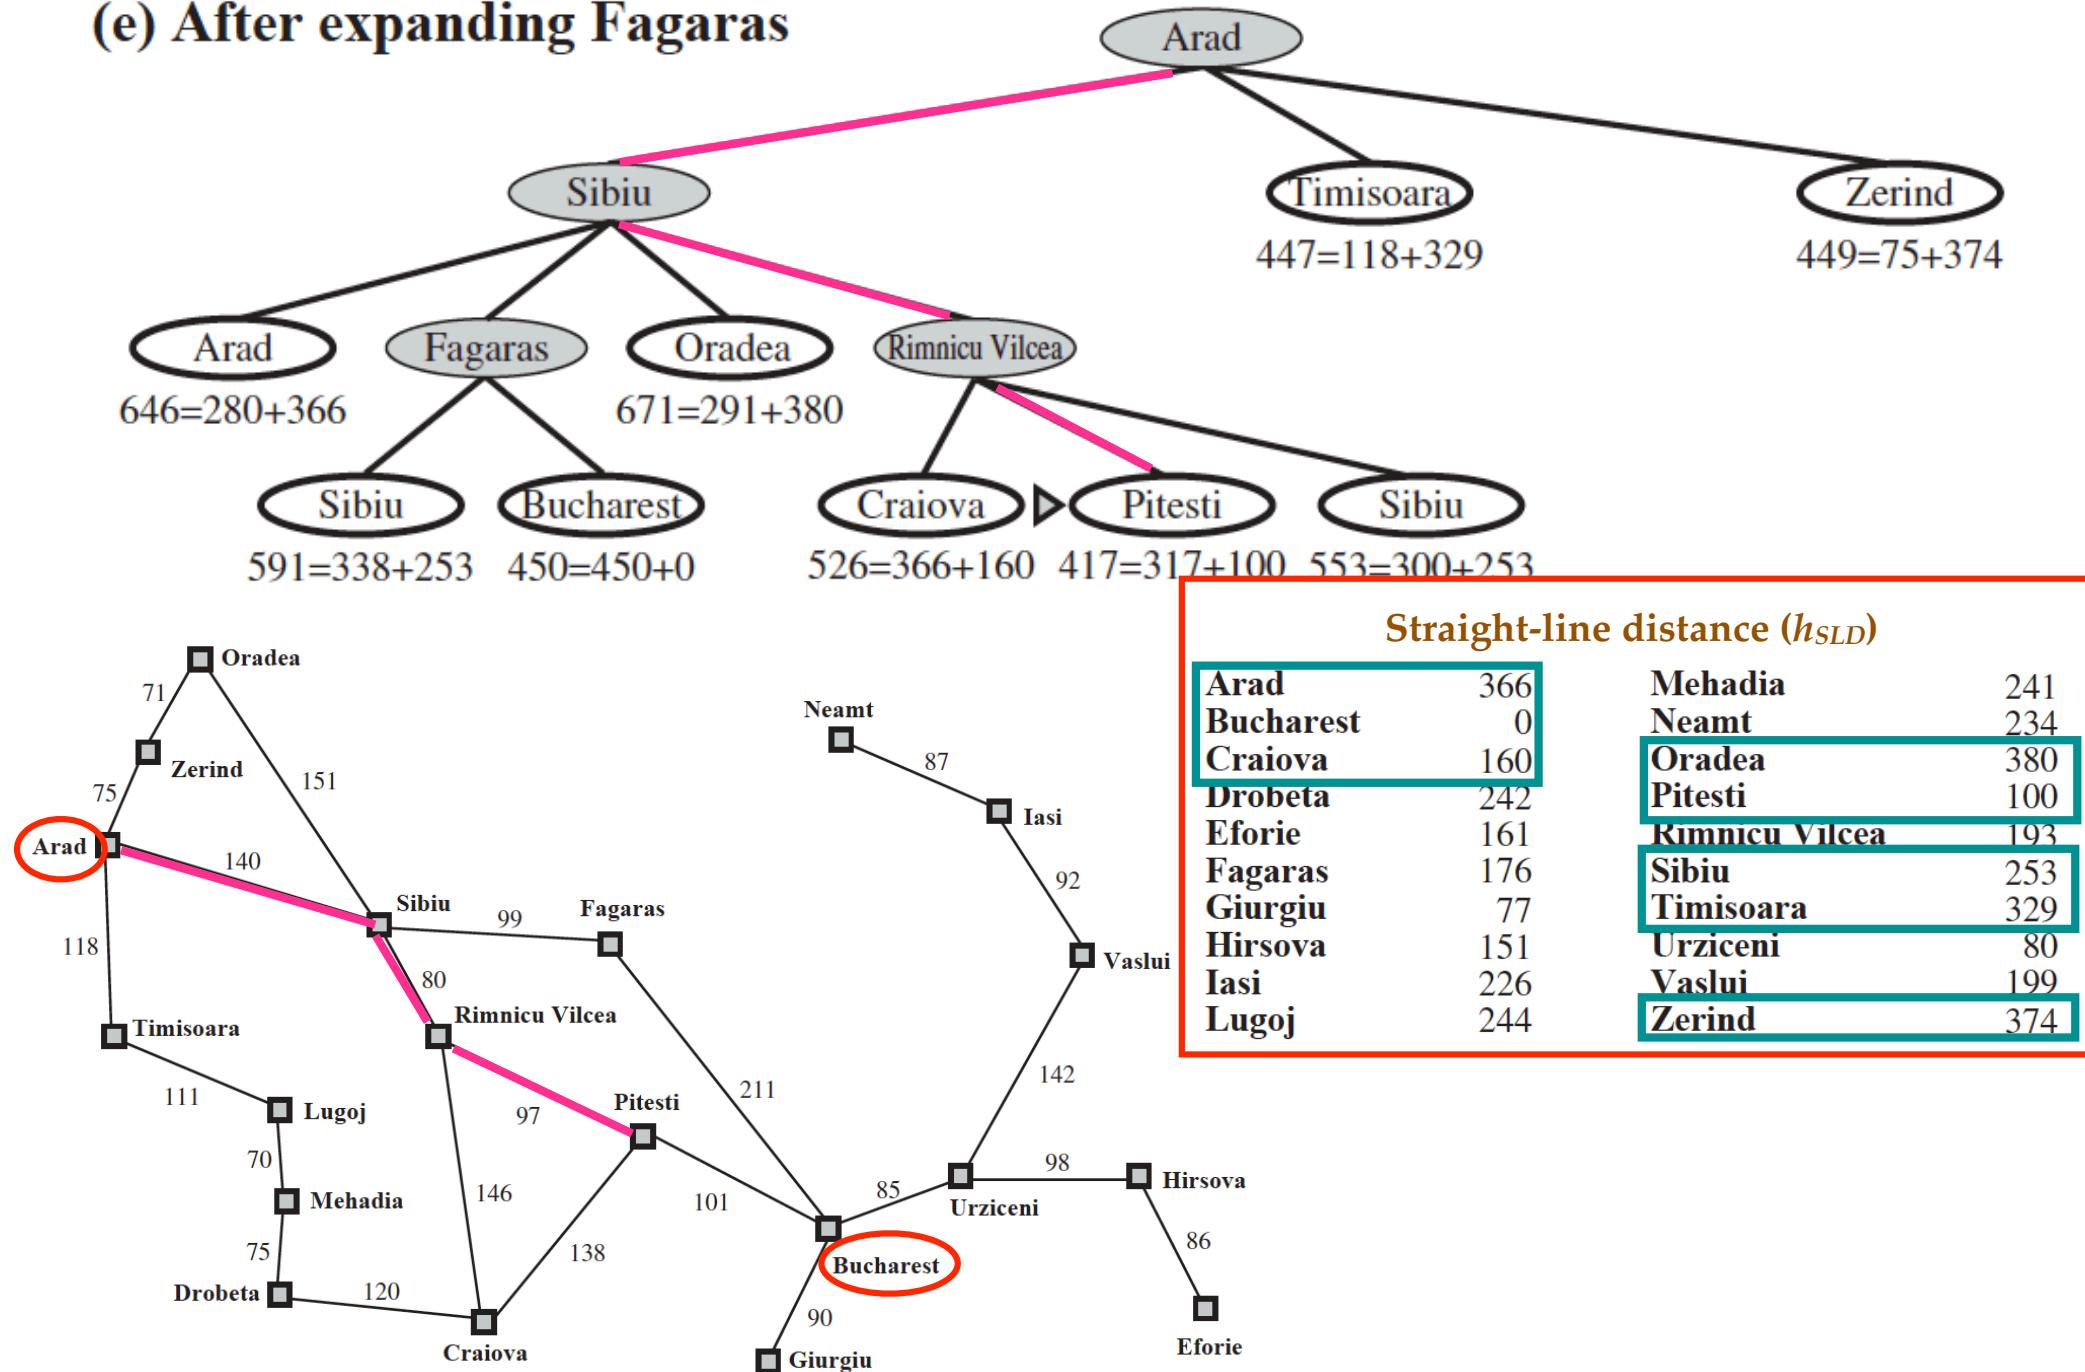


(d) After expanding Rimnicu Vilcea

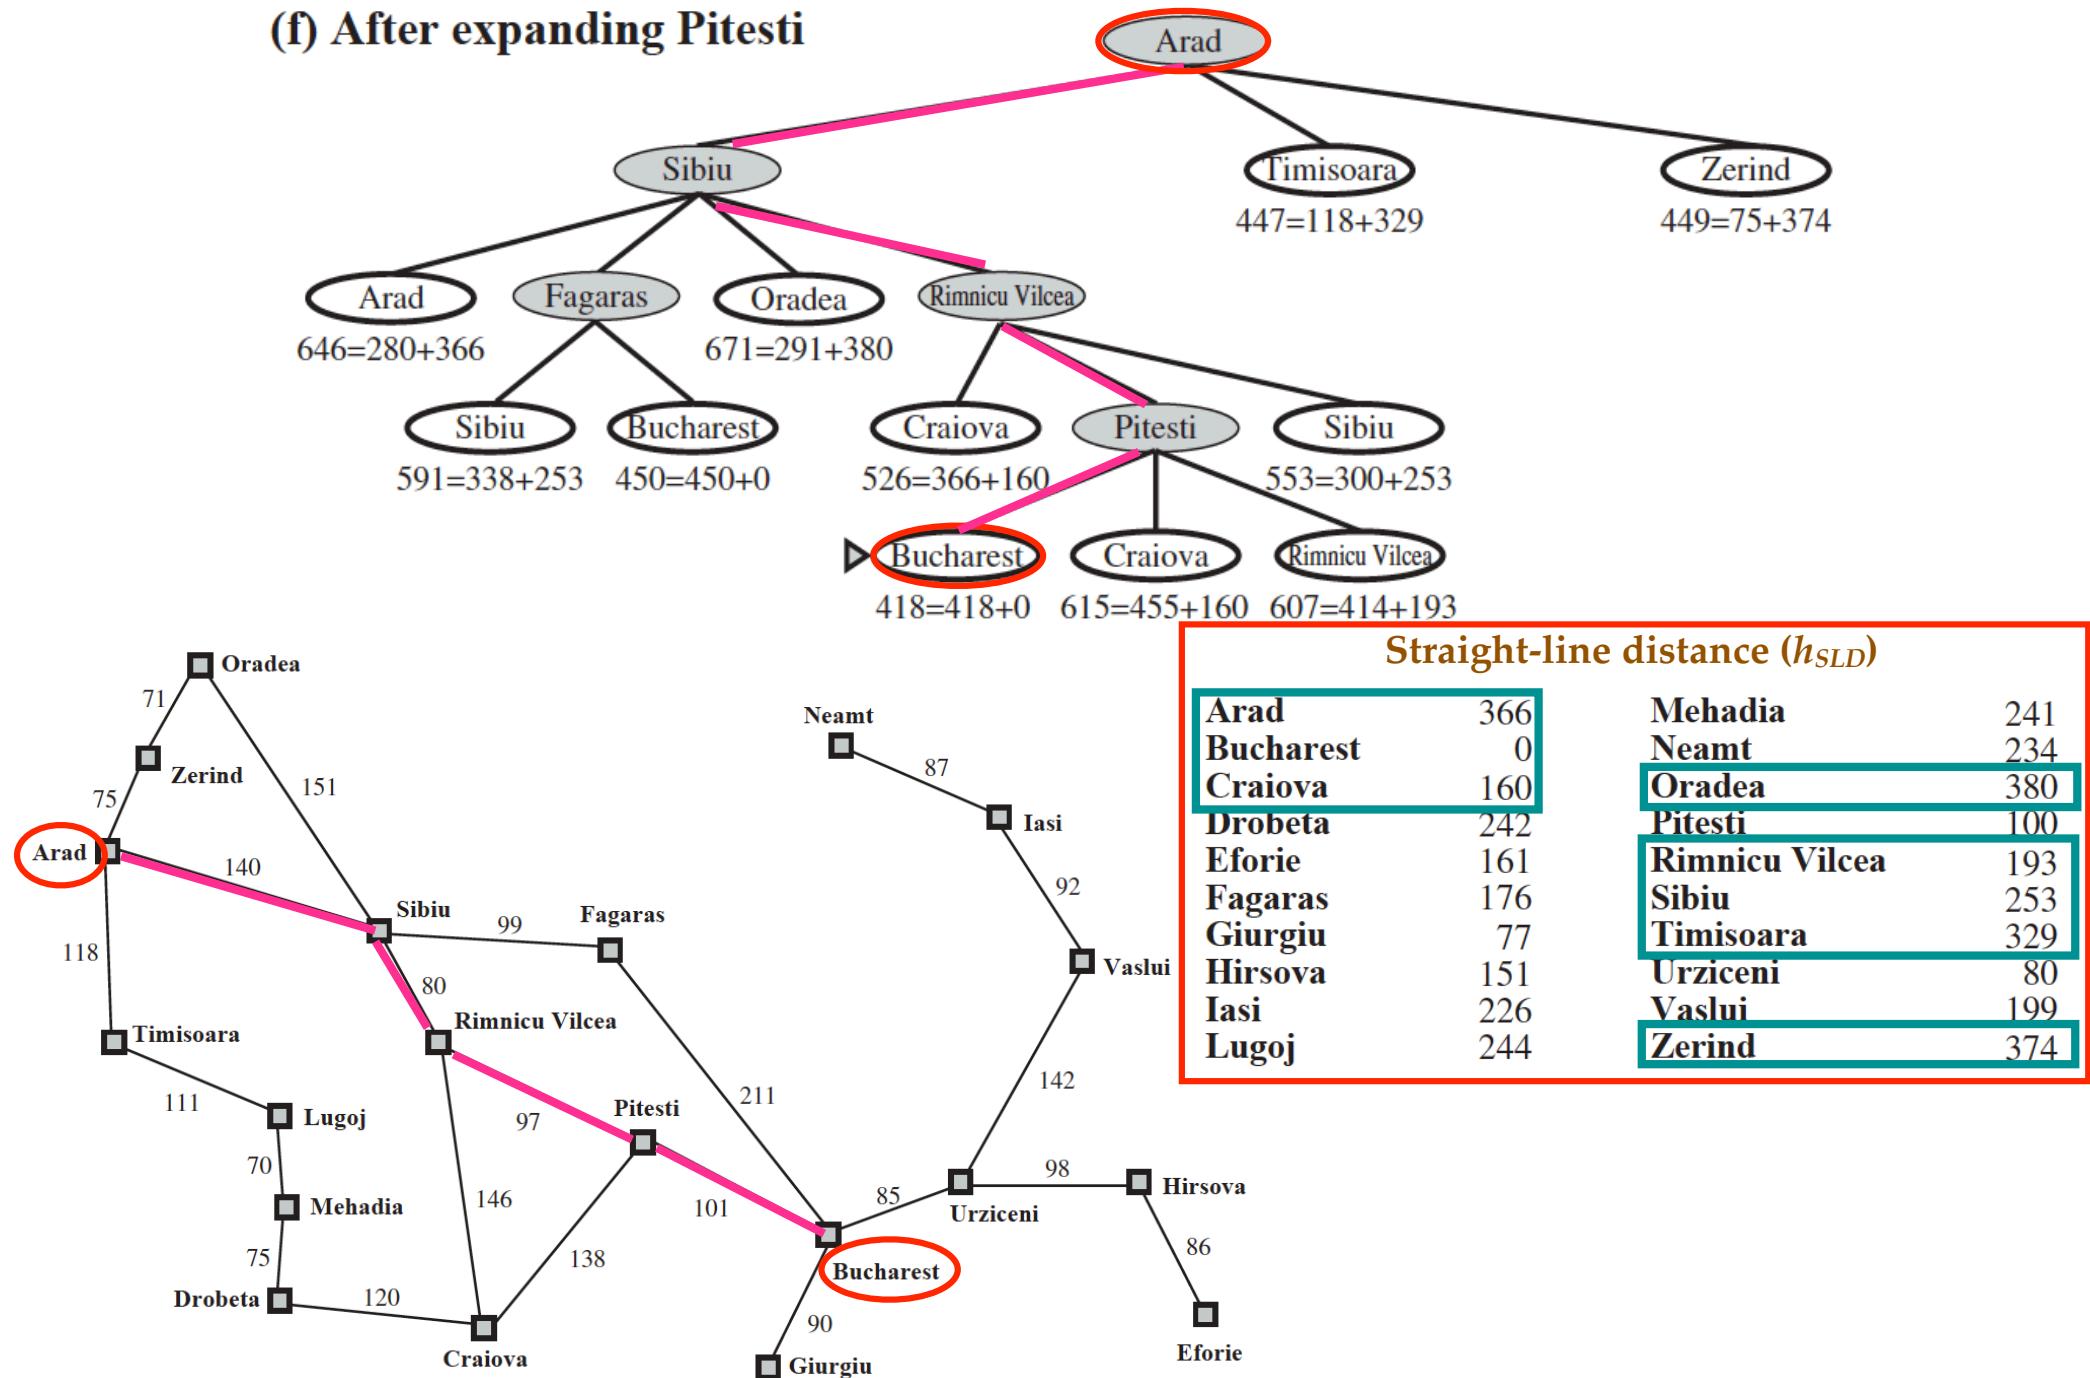


Straight-line distance (h_{SLD})		
Arad	366	
Bucharest	0	
Craiova	160	
Drobeta	242	
Eforie	161	
Fagaras	176	
Giurgiu	77	
Hirsova	151	
Iasi	226	
Lugoj	244	
Mehadia		241
Neamt		234
Oradea		380
Pitesti		100
Rimnicu Vilcea		193
Sibiu		253
Timisoara		329
Urziceni		80
Vaslui		199
Zerind		374

(e) After expanding Fagaras



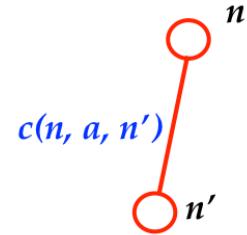
(f) After expanding Pitesti



❖ Properties of A* Search

❖ Completeness

$$f_k, f_{k+1}, \dots, f_\infty \leq f(goal)$$



❖ Yes. Unless infinite nodes with $f \leq f(goal)$

[most time $c(n, a, n') \geq \epsilon > 0$]

❖ Optimality

$$= g(goal) + h(goal)$$

\downarrow

❖ Depends on whether h is **admissible** (tree search)/**consistent** (graph search)

❖ Optimal on **trees** if h is **admissible** (tree has at most single path to goal)

❖ Optimal on **graphs** if h is **admissible** and **consistent** (graph may have more than one path to goal)

Admissible

- Never overestimates the **actual cost**.
- $\forall n, h(n) \leq h^*(n)$, where h^* is the actual cost.
- e.g., $h_{SLD}(n) \leq h^*(n)$.

Consistent

- A.k.a. **monotonicity**, **[nondecreasing]**
- \forall successor n' of any n generated by any action a , $h(n) \leq c(n, a, n') + h(n')$, where c is the **step cost**.

triangle inequality

- ❖ **Admissible heuristics**
 - ❖ A heuristic $h(n)$ is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true cost** to reach the goal state from n
 - ❖ An admissible heuristic **never overestimates** the cost to reach the goal
 - ❖ Example: $h_{SLD}(n) \leq h^*(n)$ (straight line distance never overestimates the actual road distance)
- ❖ **Theorem:** If $h(n)$ is **admissible**, A* using Tree Search is **optimal**

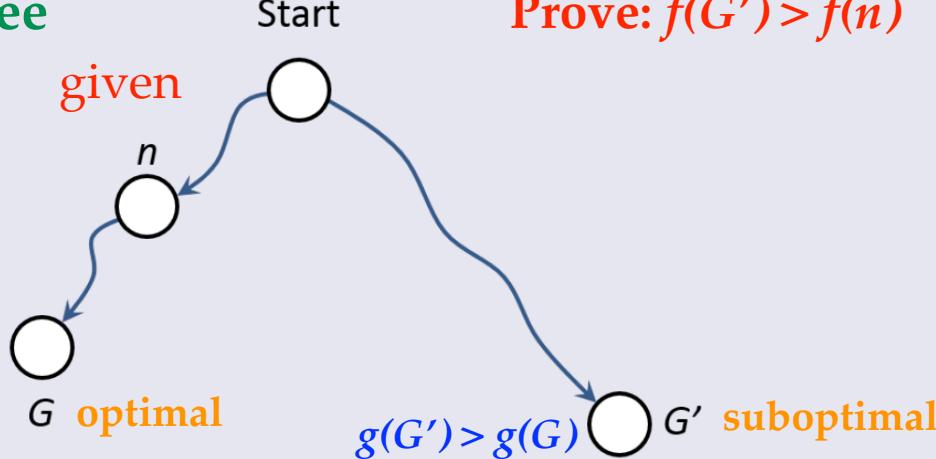
- ❖ **Time complexity**
 - ❖ $O(b^{\varepsilon d})$ for **constant** step costs, where $\varepsilon = (h^* - h) / h^*$ (relative error) and d is the solution depth
 - ❖ Time complexity much depends on the heuristic
 - ❖ e.g. $\varepsilon = 0$ when $h = h^*$, $\varepsilon = 1$ when $h = 0$
 - $O(b^0) = O(1)$ linear time
 - $O(b^{\varepsilon d}) = O(b^d)$ [same as BFS]
- ❖ Effective branch factor is b^ε
- ❖ **Space complexity**
 - ❖ $O(b^d)$, since it keeps all nodes in memory

- ❖ Optimality of A* on Trees [at most one path to the goal]
- ❖ Optimal on trees if h is admissible

Proof of A*'s optimality on trees.

Suppose some suboptimal goal G' has been generated and is in the queue. $g(G') > g(G)$
 Let n be an unexpanded node on a shortest path to an optimal goal G .

Tree

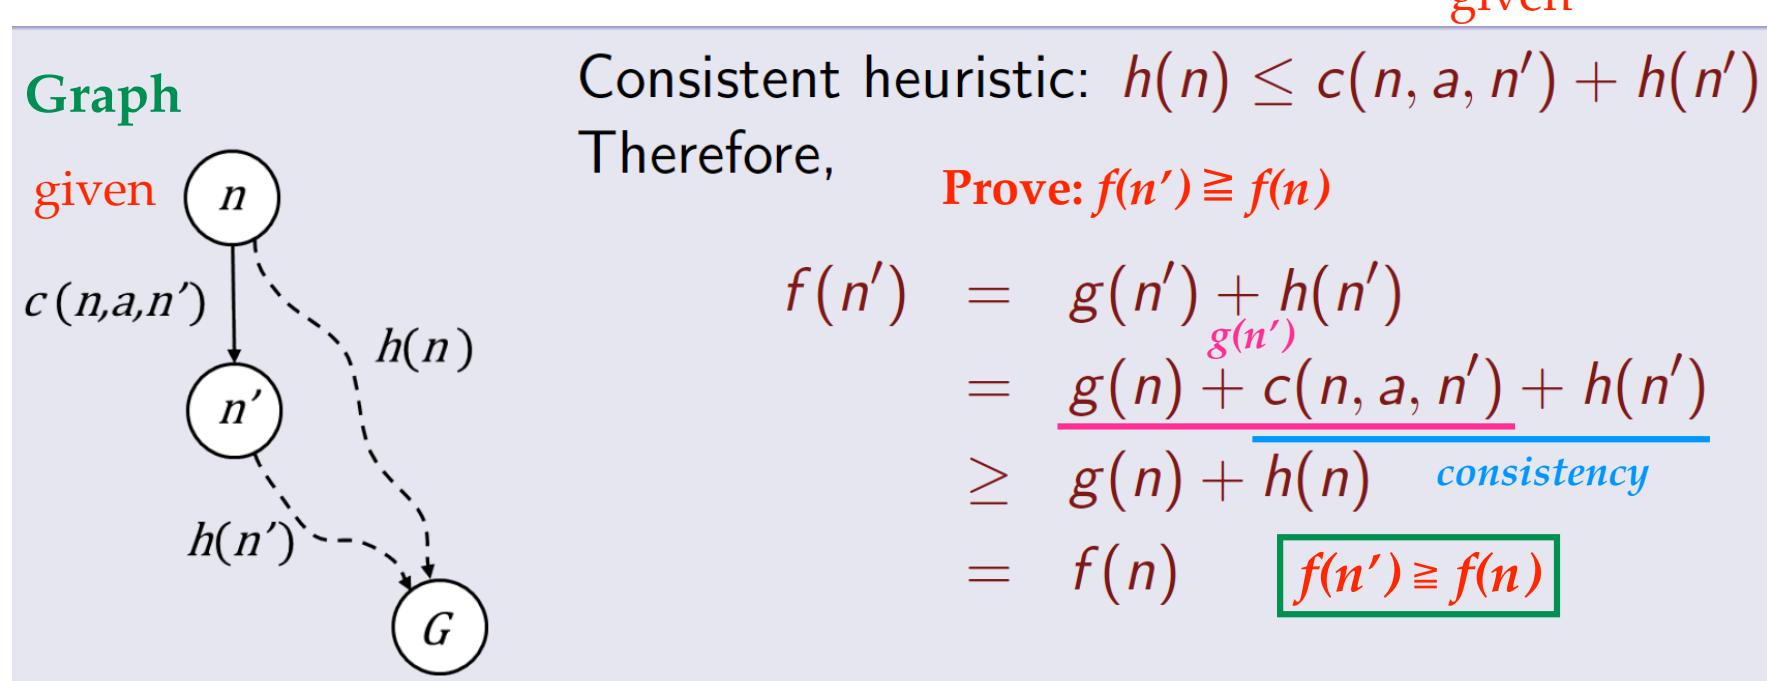


Prove: $f(G') > f(n)$

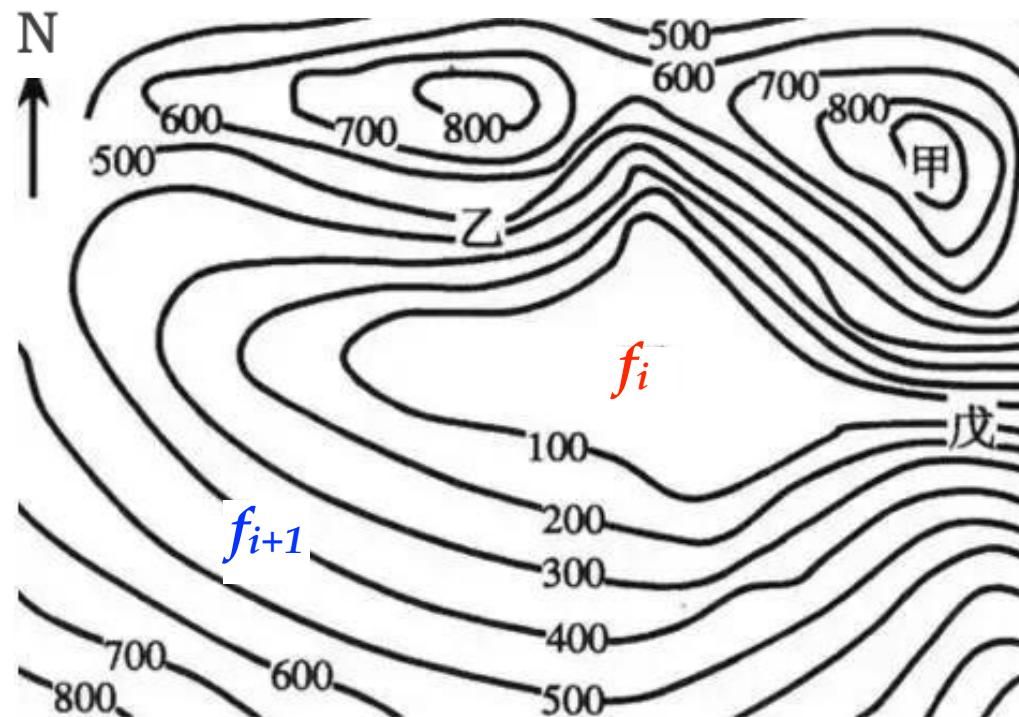
$$\begin{aligned}
 f(G') &= g(G') + h(G') \\
 &= g(G') \\
 &> g(G) \quad g(G') > g(G) \\
 &= g(n) + h^*(n) \\
 &\geq g(n) + h(n) \quad h(n) \leq h^*(n) \\
 &= f(n) \quad f(G') > f(n)
 \end{aligned}$$

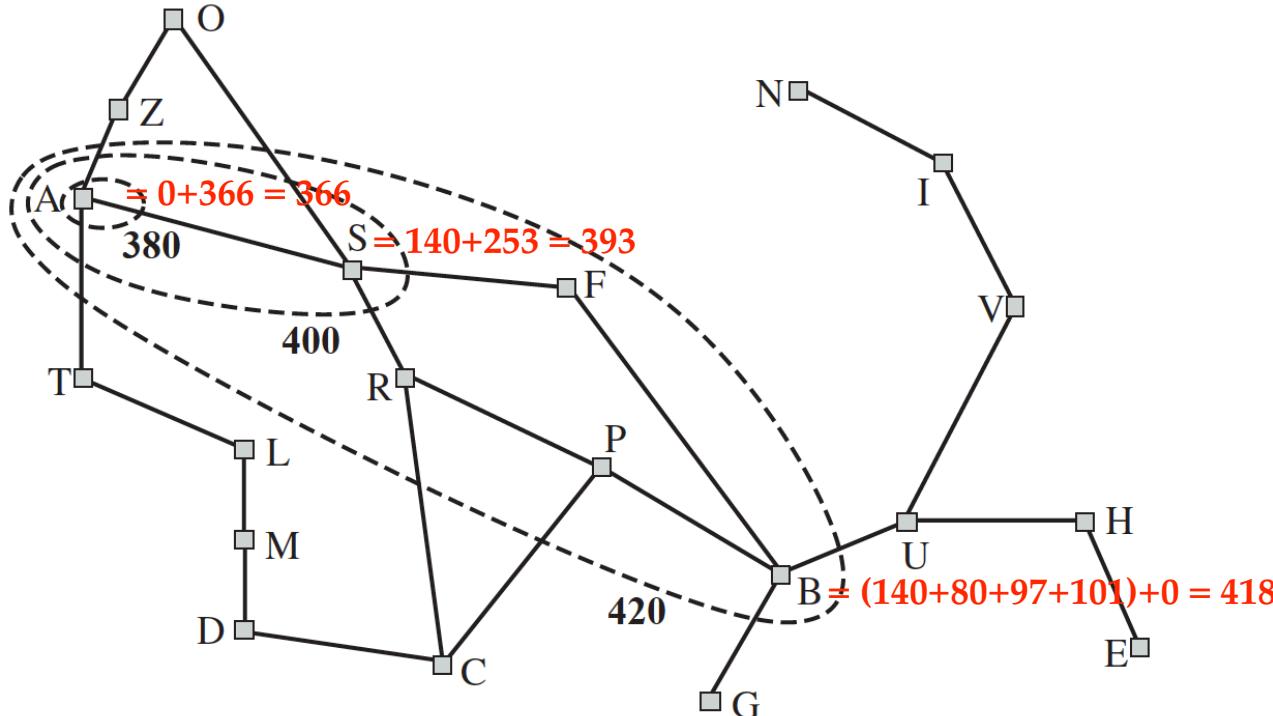
Node n with lower $f(n)$ is expanded earlier than G' with higher $f(G')$

- ❖ Optimality of A* on Graphs [may have more than one path to the goal]
 - ❖ Optimal on graphs if h is **admissible** and **consistent**
 - ❖ **Consistent**
 - ❖ \forall successor n' of any n generated by any action a , $h(n) \leq c(n, a, n') + h(n')$ (nondecreasing), where c is the step cost
 - ❖ **Lemma:** If $h(n)$ is **consistent**, the values of f along any path in A* are **nondecreasing** (Monotonic)



- ❖ Consistency is actually triangle inequality
- ❖ A* expands nodes in the order of increasing f value
- ❖ Gradually adds f -contours of nodes
- ❖ Contour i has all nodes with $f = f_i$, where $f_i < f_{i+1}$
(nondecreasing)





- ❖ Map of Romania showing contours at $f = 380$, $f = 400$, and $f = 420$, with **Arad** as the start state
- ❖ Inside the contour labeled 400, all nodes have $f(n)$ less than or equal to 400, and so on.
- ❖ Because A* expands the frontier node of **lowest f -cost**, we can see that an A* search fans out from the start node, adding nodes in **concentric bands of increasing f -cost**

Memory-Bounded Heuristic Search

- ❖ Iterative Deepening A* (IDA*)
- ❖ Recursive Best-First Search (RBFS)
- ❖ Simplified Memory-Bounded A* (SMA*)

Iterative Deepening A* (IDA*)

- ❖ Time complexity is not A*'s biggest drawback
- ❖ A* usually **runs out of memory** before it reaches goals
- ❖ Iterative deepening A* (IDA*)
 - ❖ Use $f(g + h)$ as **cutoff** instead of the **depth**
(DLS)
 - ❖ Initial cutoff: $f(s_0) = h(s_0)$
(frontier)
 - ❖ Reset cutoff to **smallest f** of non-expanded nodes
 - ❖ Perform **DFS** on nodes where $f(n) < \text{cutoff}$

A*

Time complexity: $O(b^{ed})$

Space complexity: $O(b^d)$

$\text{IDA}^*(\text{problem})$

```
1 currentCutoff =  $f(s_0)$ 
2 repeat
3     result = f-LIMITED-SEARCH(problem, currentCutoff)
4     if result  $\neq$  cutoff
5         return result
6     currentCutoff = smallest f-value of non-expanded nodes.
```

- ❖ Properties of IDA*
- ❖ Completeness and Optimality same as A*

- ❖ **Completeness**

$$f_k, f_{k+1}, \dots, f_\infty \leq f(goal)$$

- ❖ Yes. Unless infinite nodes with $f \leq f(goal)$

- ❖ **Optimality**

[most time $c(n, a, n') \geq \epsilon$]

- ❖ Optimal on trees if h is admissible

- ❖ Optimal on graphs if h is admissible and consistent

- ❖ **Time complexity:** $O(b^{\epsilon d})$

- ❖ **Space complexity:** $O(bd)$ [DFS]

A*

Time complexity: $O(b^{\epsilon d})$

Space complexity: $O(bd)$

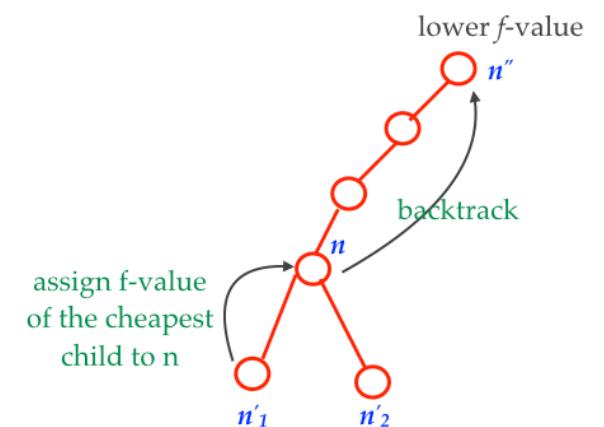
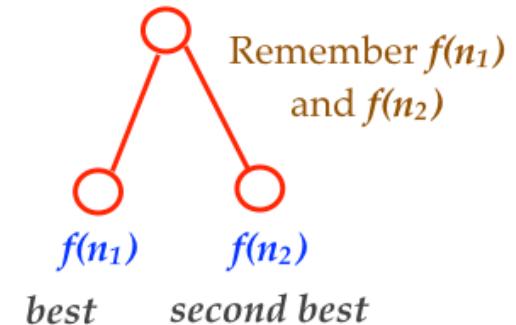
- ❖ What happens if all f -values are different (real-values)?
 - ➡ The **no. of iterations** can equal the **no. of nodes** whose f -value is **less than the cost of an optimal path!** [Since f -values are different, there may have many interactions causing severe repeated search]

Recursive Best-First Search (RBFS)

- ❖ RBFS is a simple recursive algorithm that mimics standard **best-first search** using only **linear space**

```
function RECURSIVE-BEST-FIRST-SEARCH(problem) returns a solution, or failure  
  return RBFS(problem, MAKE-NODE(problem.INITIAL-STATE),  $\infty$ )
```

- ❖ RBFS uses DFS where each node on the current path remembers the best f -value of any alternative path from its ancestors
- ❖ Maintains all nodes on current path plus all their siblings (from $\text{ancestor}(n)$)
- ❖ When expanding node n
 - ❖ $\forall n' \in \text{children}(n)$, compute $f(n')$
 - ❖ If an ancestor n'' has a lower f -value than all n 's children, then
 - ❖ Assign f -value of the cheapest child to n
 - ❖ Backtrack to n''
 - ❖ Otherwise, proceed as normal

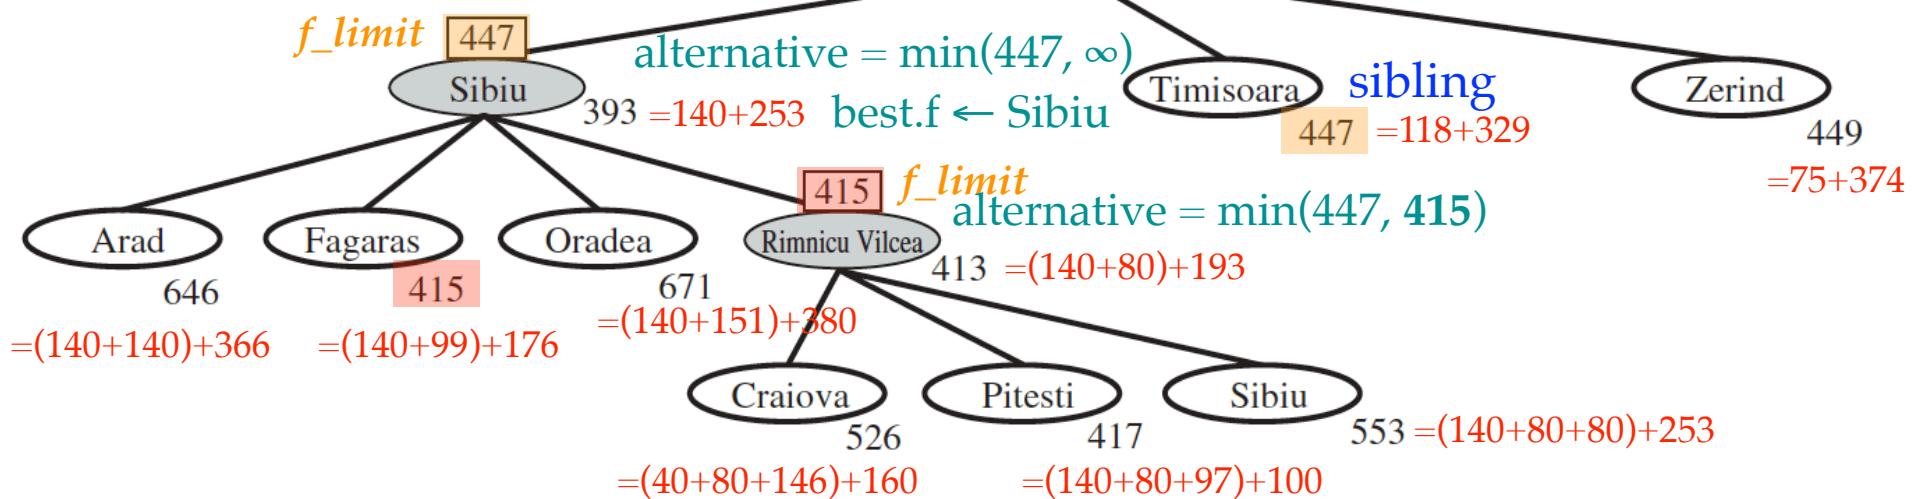


```
function RECURSIVE-BEST-FIRST-SEARCH(problem) returns a solution, or failure  
    return RBFS(problem, MAKE-NODE(problem.INITIAL-STATE),  $\infty$ )  
                cutoff
```

```
function RBFS(problem, node, f-limit) returns a solution, or failure and a new f-cost limit  
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)  
    successors  $\leftarrow$  []  
    for each action in problem.ACTIONS(node.STATE) do  
        add CHILD-NODE(problem, node, action) into successors  
    if successors is empty then return failure,  $\infty$   
    for each s in successors do /* update f with value from previous search, if any */  
        s.f  $\leftarrow$  max(s.g + s.h, node.f)  
    loop do  
        best  $\leftarrow$  the lowest f-value node in successors  
        if best.f > f-limit then return failure, best.f  
        alternative  $\leftarrow$  the second-lowest f-value among successors  
        result, best.f  $\leftarrow$  RBFS(problem, best, min(f-limit, alternative))  
        if result  $\neq$  failure then return result       [parents]   [sibling]
```

(a) After expanding Arad, Sibiu, and Rimnicu Vilcea

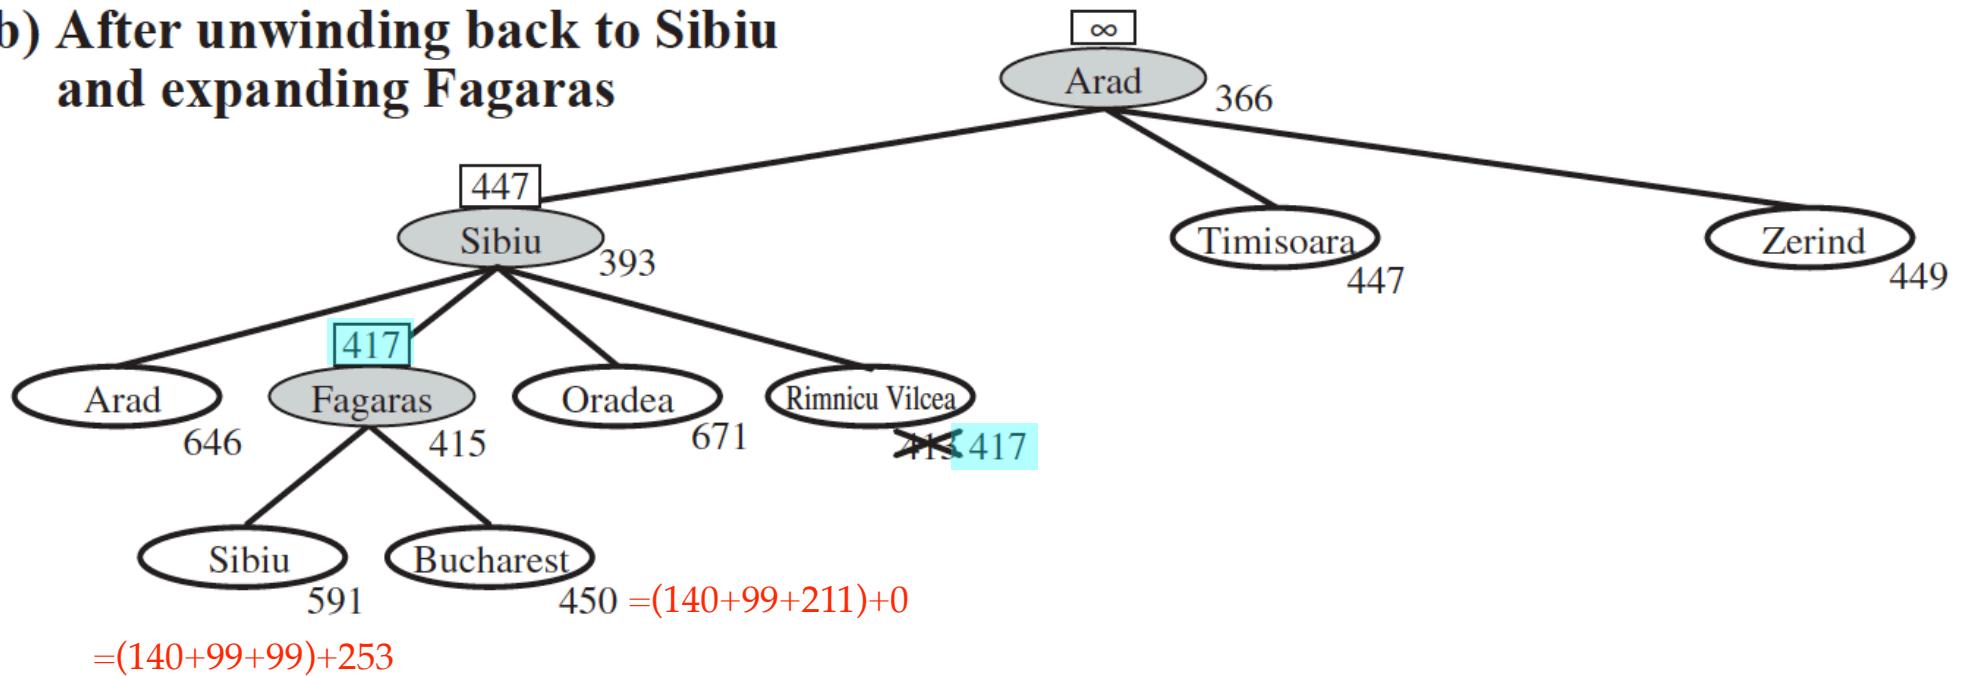
parent ∞ f_limit
Arad 366



- ❖ Stages in an RBFS search for the shortest route to Bucharest
- ❖ The f -limit value for each recursive call is shown on top of each current node, and every node is labeled with its f -cost
- ❖ (a) The path via Rimnicu Vilcea is followed until the current best leaf (Pitesti, 417) has a value that is worse than the best alternative path (Fagaras, 415)

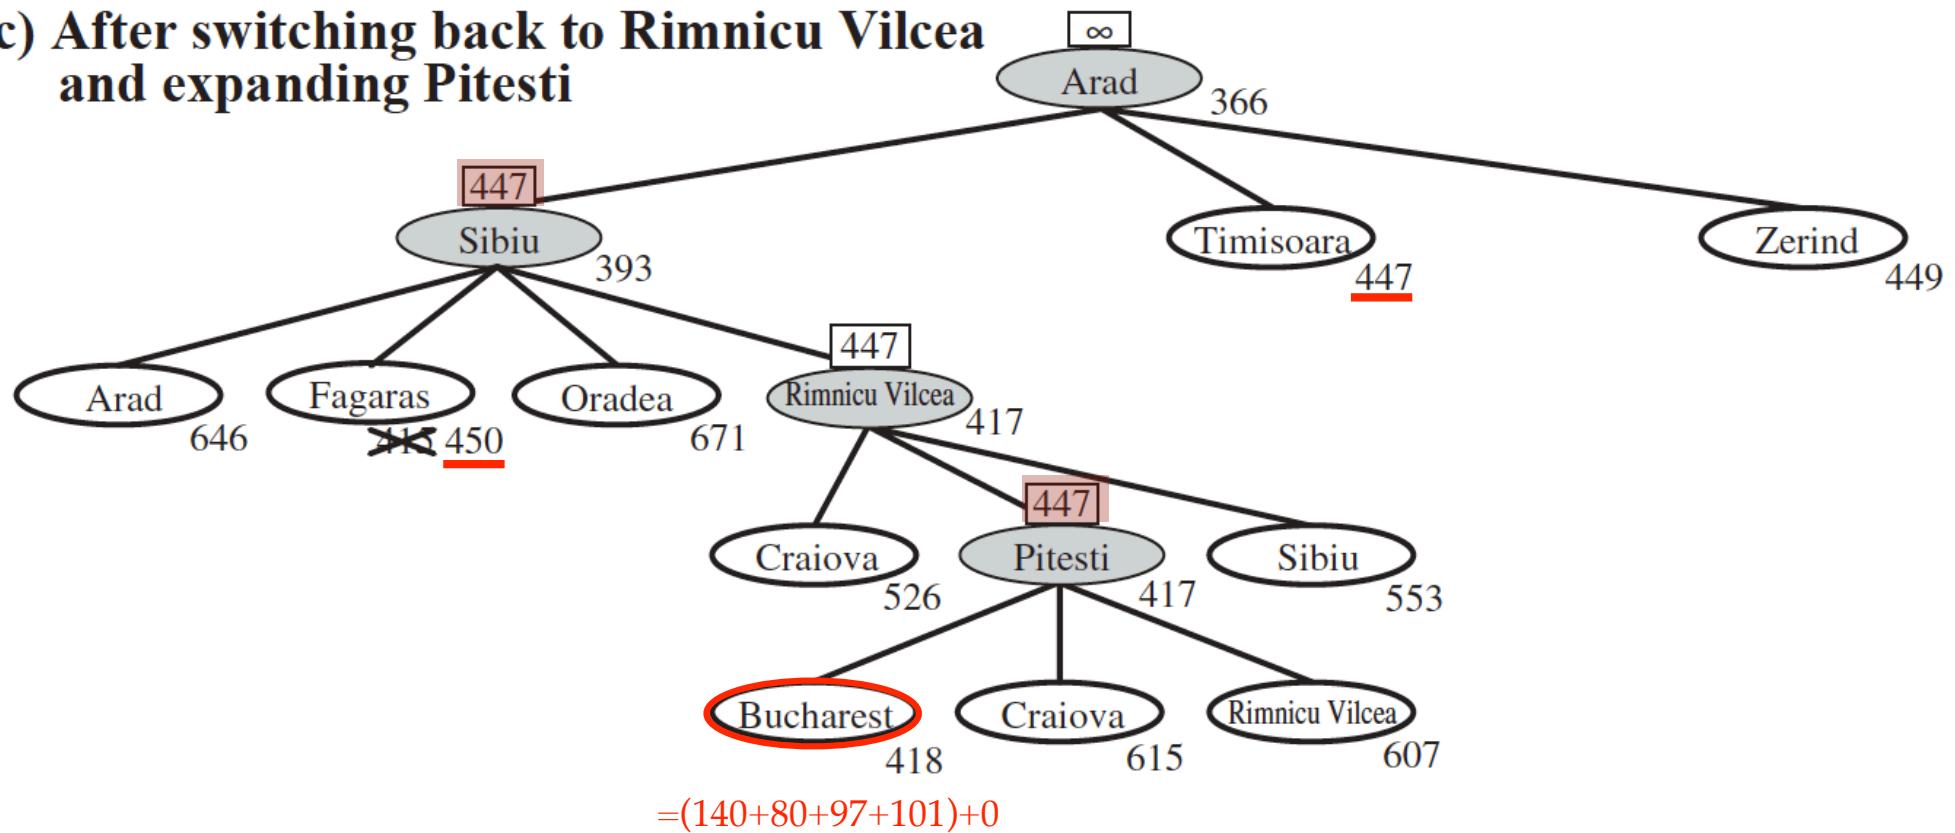
Straight-line distance (h_{SLD})			
Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

(b) After unwinding back to Sibiu and expanding Fagaras



- ❖ (b) The recursion **unwinds** and the best leaf value of the forgotten subtree (417) is **backed up** to Rimnicu Vilcea; then Fagaras is expanded, revealing a best leaf value of 450 (**Bucharest**) [$450 > 417$]

(c) After switching back to Rimnicu Vilcea and expanding Pitesti



- ❖ (c) The recursion **unwinds** and the best leaf value of the forgotten subtree (450) is **backed up** to Fagaras; then Rimnicu Vilcea (417) is expanded
- ❖ This time, because the best alternative path (through Timisoara) costs at least 447 [$447 > 418$], the expansion continues to Bucharest (418)

- ❖ Properties of RBFS
- ❖ Completeness and optimality same as A*

- ❖ **Completeness**

$$f_k, f_{k+1}, \dots, f_\infty \leq f(goal)$$

- ❖ **Yes.** Unless infinite nodes with $f \leq f(goal)$

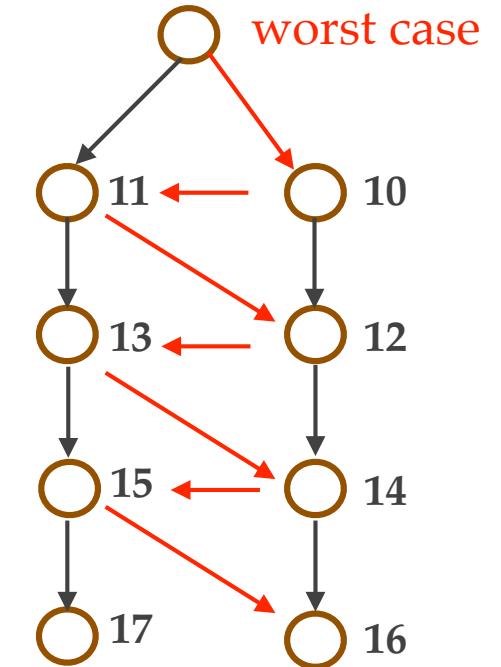
[most time $c(n, a, n') \geq \epsilon$]

- ❖ **Optimality**

- ❖ Optimal on trees if h is admissible

- ❖ Optimal on graphs if h is admissible and consistent

- ❖ **Time complexity:** Depends on accuracy of h and on **how often best path changes**
- ❖ **Space complexity:** $O(bd)$
- ❖ Each time RBFS **changes its mind** corresponds to **one iteration of IDA***
- ❖ RBFS may need to **re-expand forgotten** nodes to re-create **best-path**



Summary

[$O(bd)$]

[$O(bd)$]

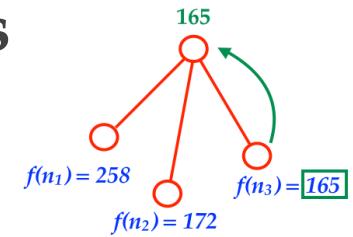
If $b = 4$, $d = 100$ and 1 record = 2KB,
then total requires 800KB memory

- ❖ In a sense, both IDA* and RBFS use **too little memory**
 - ❖ Between iterations, IDA* maintains only **one number**, the **current *f*-limit (*currentCutoff*)**
 - ❖ RBFS maintains **more**, but uses only **linear space**: if more space were available, it would not benefit from it
- ❖ It seems reasonable to **use all the memory available** — the more, the better (**don't change path before using up memory**)
- ❖ We'd like a **memory-bounded** version of A* (**don't change path before use up memory**)
 - e.g. Change path only if use up 500Mb memory.

Simplified Memory-Bounded A* (SMA*)

e.g. 500MB

- ❖ **Idea:** Run A* as normal until memory is full. Then replace something in memory with newly generated nodes
- ❖ **SMA***
 - ❖ When memory is full, drop the worst leaf — node with highest f -value
 - ❖ Like RBFS, SMA* backs up the lowest f -value of this forgotten nodes to its parent, so we know when to go back to it
 - ❖ If all descendants of a node n are forgotten, we don't know which way to go from n , but we know if it's worth re-exploring n



- ❖ **Problem:** What if many nodes have the same f -value?
 - ❖ **Solution:** delete the oldest and expand the newest (at the cost of adding **time stamp** to each node)
 - ❖ SMA* works as long as there is **enough memory** for the **complete optimal path**
 - ❖ If not, SMA* needs to **switch continuously** between **candidate paths**
 - ❖ Causes a similar problem to **thrashing** in **disk paging systems**
 - In computer science, **thrashing** occurs when a computer's virtual memory resources are overused, leading to a constant state of **paging** and **page faults**.

```

function SMA*(problem) returns a solution sequence
  inputs: problem, a problem
  static: Queue, a queue of nodes ordered by f-cost

  Queue  $\leftarrow$  MAKE-QUEUE({MAKE-NODE(INITIAL-STATE[problem])})
  loop do
    if Queue is empty then return failure
    n  $\leftarrow$  deepest least-f-cost node in Queue
    if GOAL-TEST(n) then return success
    s  $\leftarrow$  NEXT-SUCCESSOR(n)
    if s is not a goal and is at maximum depth then
      f(s)  $\leftarrow$   $\infty$ 
    else
      f(s)  $\leftarrow$  MAX(f(n),g(s)+h(s))
    if all of n's successors have been generated then
      update n's f-cost and those of its ancestors if necessary
    if SUCCESSORS(n) all in memory then remove n from Queue
    if memory is full then
      delete shallowest, highest-f-cost node in Queue
      remove it from its parent's successor list
      insert its parent on Queue if necessary
    insert s in Queue
  end

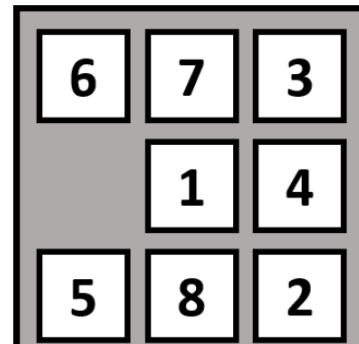
```

3.6 Heuristic Functions

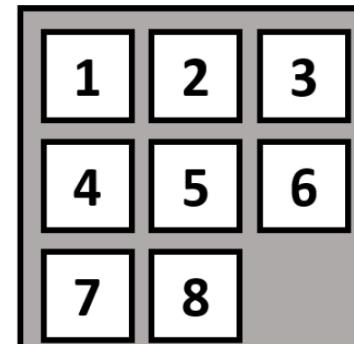
- ❖ The effect of **heuristic accuracy** on performance
- ❖ Generating **admissible heuristics** from **relaxed problems**
- ❖ Generating **admissible heuristics** from **sub-problems: pattern databases**
- ❖ Learning heuristics from **experience**

Admissible Heuristics for 8-Puzzle

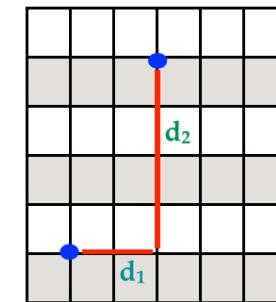
- ❖ h_1 = the number of **misplaced tiles** (admissible, never overestimate)
- ❖ h_2 = the sum of **Manhattan distances** of the tiles from their goal positions (admissible, never overestimate) ($h_2 = d_1 + d_2$)



Start State



Goal State



$$h_2 = d_1 + d_2$$

- ❖ $h_1(s_0) = 6$ (1, 2, 4, 5, 6, 7)
 $2 \rightarrow 4 \rightarrow 3 \rightarrow 7$
- ❖ $h_2(s_0) = 2 + 3 + 0 + 2 + 2 + 3 + 3 + 0 = 15$
 $1 \rightarrow 7 \rightarrow 6$

Performance of Heuristic

Definition

For two **admissible** heuristics h_1 and h_2 , h_2 dominates h_1 iff
 $\forall n, h_2(n) \geq h_1(n)$.

$$h^*(n) \geq h_2(n) \geq h_1(a) \geq 0$$

Theorem: A* using h_2 never expands more nodes than using h_1 .

- Every node with $f(n) < C^*$ is expanded. $g(n) + h(n) < C^*$
- Every node with $h(n) < C^* - g(n)$ is expanded. $f(n) = g(n) + h(n)$
- $|\{n \mid h_2(n) < C^* - g(n)\}| \leq |\{n \mid h_1(n) < C^* - g(n)\}|$

h_2 moves faster

h_1 moves slower

- ❖ Given a state n , $C^* - g(n)$ is fixed
 - ❖ h_2 has better performance than h_1
 - ❖ A* search with h_1 will expand more nodes
- ❖ Given any **admissible** heuristics h_a and h_b
 - ❖ $h = \max(h_a, h_b)$ is also **admissible**
 - ❖ h dominates h_a and h_b , respectively

- ❖ Typical search costs for the 8-puzzle (**average no. of nodes expanded** for different solution depths)
- ❖ $d = 12$
 - ❖ IDS = 3,644,035 nodes
 - ❖ A^{*}(h_1) = 227 nodes
 - ❖ A^{*}(h_2) = 73 nodes
- ❖ $d = 24$
 - ❖ IDS \approx 54,000,000,000 nodes
 - ❖ A^{*}(h_1) = 39,135 nodes
 - ❖ A^{*}(h_2) = 1,641 nodes

- ❖ **Combining heuristics**
 - ❖ Suppose we have a collection of admissible heuristics, $h_1(n)$, $h_2(n)$, ..., $h_m(n)$, but **none of them dominates the others**
 - ❖ How can we **combine** them?

$$h(n) = \max\{h_1(n), h_2(n), \dots, h_m(n)\}$$

- ❖ e.g. $h_3(n_1)$, $h_4(n_2)$, $h_1(n_3)$, $h_3(n_4)$, $h_2(n_5)$,
- ❖ This **composite heuristic** uses whichever function is **most accurate** on the node in question (for each node n_i , choose the h_j which has the max h value)

- ❖ **Effective Branching Factor**
 - ❖ One way to characterize the **quality** of a heuristic is *effective branching factor, b^**
 - ❖ Let A* generate N nodes to find a goal at **depth d**
 - ❖ b^* is the branching factor that a **uniform tree of depth d** would have in order to contain $N+1$ nodes

$$N + 1 = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$$

$$N + 1 = ((b^*)^{d+1} - 1) / (b^* - 1)$$

$$N \approx (b^*)^d \Rightarrow b^* \approx \sqrt[d]{N}$$
 - ❖ For **sufficiently hard problems**, the measure b^* usually is fairly **constant** across different problem instances
 - ❖ A good way to **compare different heuristics**
 - ❖ A **well-designed heuristic** would have a value of b^* close to 1

Effective Branching Factor Pseudo-code (Binary Search)

- PROCEDURE EFFBRANCH (START, END, N, D, DELTA)
 COMMENT DELTA IS A SMALL POSITIVE NUMBER FOR ACCURACY OF RESULT.
 MID := (START + END) / 2.
 IF (END - START < DELTA)
 THEN RETURN (MID).
 TEST := EFFFOLY (MID, D).
 IF (TEST < N+1)
 THEN RETURN (EFFBRANCH (MID, END, N, D, DELTA))
 ELSE RETURN (EFFBRANCH (START, MID, N, D, DELTA)).
END EFFBRANCH.

```
PROCEDURE EFFFOLY (B, D)
  ANSWER = 1.
  TEMP = 1.
  FOR I FROM 1 TO (D-1) DO
    TEMP := TEMP * B.
    ANSWER := ANSWER + TEMP.
  ENDDO.
  RETURN (ANSWER).
END EFFFOLY.
```

- For binary search please see: http://en.wikipedia.org/wiki/Binary_search_algorithm
- An attractive alternative is to use Newton's Method (next lecture) to solve for the root (i.e., $f(b)=0$) of
$$f(b) = 1 + b + \dots + b^d - (N+1)$$

- ❖ Comparison of the **search costs** and **effective branching factors** for the IDS and A* algorithms with h_1, h_2
- ❖ Data are averaged over 100 instances of the 8-puzzle for each of various solution lengths d

$$h_1(n) \leq h_2(n)$$

h_2 has smaller branch factor

d	Search Cost (nodes generated)			Effective Branching Factor		
	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035 (huge)	227	73	2.78	1.42	1.24
14	hard to compare	539	113	—	1.44	1.23
16	—	1301	211	—	1.45	1.25
18	—	3056	363	—	1.46	1.26
20	—	7276	676	—	1.47	1.27
22	—	18094	1219	—	1.48	1.28
24	—	39135	1641	—	1.48	1.26

approximate to
fixed value

the deeper the
more accurate

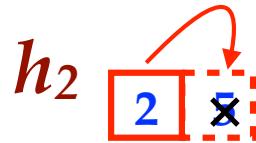
Comparison of Search Strategies

Algorithm	Complete?	Optimal?	Time complexity	Space complexity
BFS	Yes	If all step costs are equal	$O(b^d)$	$O(b^d)$
UCS	Yes	Yes	Number of nodes with $g(n) \leq C^*$	
DFS	No	No	$O(b^m)$	$O(bm)$
IDS	Yes	If all step costs are equal	$O(b^d)$	$O(bd)$
Greedy	No	No	Worst case: $O(b^m)$ Best case: $O(bd)$	
A*	Yes	Yes	Number of nodes with $g(n)+h(n) \leq C^*$	

Generating Admissible Heuristics from Relaxed Problems

- ❖ A problem with **fewer restrictions** on the **actions** is called a **relaxed problem** (**remove some constraints**)
- ❖ **Admissible heuristics** can be derived from exact solution cost to a **relaxed** version of the problem
- ❖ The **cost** of an optimal solution to a **relaxed problem** is an **admissible heuristic** for the original problem

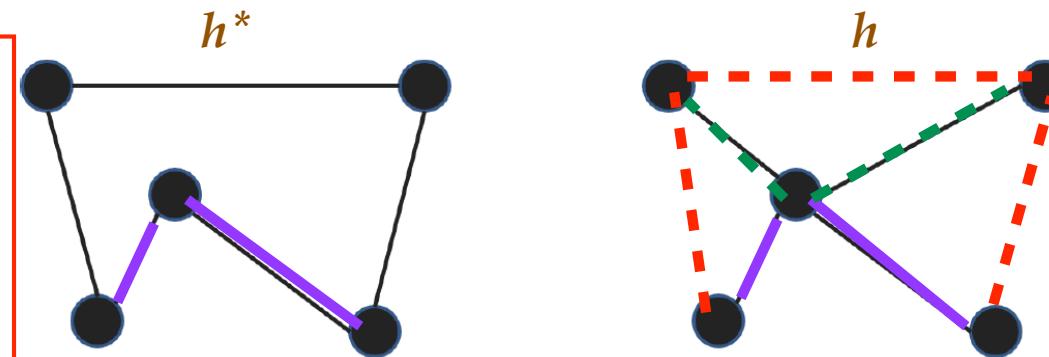
- ❖ In 8-puzzle
 - ❖ h_1 (admissible) is derived from that a tile can move to anywhere in one step, then $h_1(n)$ gives the shortest solution
 - ❖ h_2 (admissible) is derived from that a tile can move to any adjacent square in one step, then $h_2(n)$ gives the shortest solution
- ❖ Key: The optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the original problem



❖ Problem Relaxation Example: TSP

- ❖ Traveling salesman problem (TSP)
- ❖ Known to be **NP-hard** (solution can't be derived in polynomial time)

TSP : Given a list of cities and the distances between each pair of cities, what is the shortest possible route that visits each city and returns to the origin city

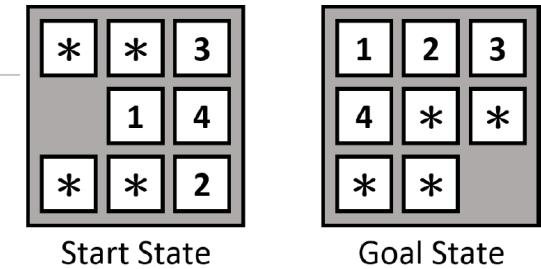


MST (Minimum Spanning Tree) : A subset of the edges of a connected, edge-weighted undirected graph that connects all the vertices together, without any cycles and with the **minimum** possible total edge weight.

- ❖ Can be relaxed to **minimum spanning tree (MST)**
Take **MST** as heuristic (h) to solve the **TSP** problem (h^*), h is admissible
- ❖ **MST cost is never greater than the shortest tour**
Triangle inequality should hold
- ❖ Cost of MST can be computed in $O(n^2)$

TSP : NP-hard
MST : $O(n^2)$

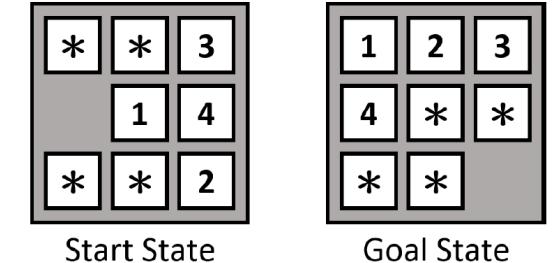
Generating Admissible Heuristics from Sub-problems: Pattern Databases



- ❖ **Admissible heuristic** can also be derived from a **subproblem** (do not remove problem constraints)
- ❖ **Pattern databases** store **exact solution costs (precomputed)** for every possible **subproblem instances**, e.g.
 - ❖ Compute 1-2-blank and don't care 3-4-5-6-7-8 (mark in *)
 - ❖ Compute 3-4-blank and don't care 1-2-3-4-5-6 (mark in *)
 - ❖
 - ❖ Compute 1-2-3-4-blank and don't care 5-6-7-8 (mark in *)
 - ❖ Compute 5-6-7-8-blank and don't care 1-2-3-4 (mark in *)

- ❖ Can we use the costs of 1-2-3-4 (h_1) and 5-6-7-8 (h_2)?

- ❖ Simple addition **breaks** the admissibility
(may be overestimate, $h_1+h_2 \geq h^*$)



- ❖ How about count only those moves involving 1-2-3-4 or 5-6-7-8, respectively, and deduct the no. of movements of * (indicating 5, 6, 7, 8) while manipulating 1-2-3-4, and deduct the no. of movements of * (indicating 1, 2, 3, 4) while manipulating 5-6-7-8
 - ❖ Then the **addition** is still **admissible**
 - ❖ This is the idea behind **disjoint pattern databases**

Learning Heuristics from Experience

- ❖ Convert a **state** into the **feature** domain [machine learning]
 - ❖ Feature $f_1(n)$: no. of **misplaced tiles**
 - ❖ Feature $f_2(n)$: no. of pairs of **adjacent tiles** that are not adjacent in the goal state
- ❖ Both $f_1(\text{goal}) = 0$ and $f_2(\text{goal}) = 0$

- ❖ $h(n) = c_1f_1(n) + c_2f_2(n)$ with $c_1 > 0; c_2 > 0$ $h(n) = \sum c_i f_i(n)$
 - ❖ c_1 and c_2 are adjusted (using **linear regression**) to give the **best fit** to the actual data on solution costs
- ❖ We could take randomly generated 8-puzzle and gather **statistics** to decide **constants**
(linear regression)
- ❖ **No guarantee** to be **admissible** or **consistent**

Summary

- ❖ Heuristic functions estimate costs of shortest paths
- ❖ Good heuristics can dramatically reduce search cost
- ❖ Greedy best-first search expands lowest h
 - ❖ In general not complete nor optimal
- ❖ A* search expands lowest $g + h$
 - ❖ Optimal when h is admissible (and consistent)

- ❖ **Memory limitation** is an important issue to heuristic search (IDA*, RBFS, SMA*)
 - ❖ Search with **forgetting** and **re-expanding** are the keys, but still suffers from different conditions
- ❖ A **more efficient heuristic** can be generated from **several admissible heuristics**
- ❖ Admissible heuristics can be derived from **relaxed problems, subproblems, and experience**
(no guarantee admissible & consistent)