

---

**Artificial Intelligence**  
A Modern Approach  
Third Edition  
Stuart Russell & Peter Norvig

---

**Chapter 4**  
**Beyond Classical**  
**Search**

---

---

# Local Search v.s. Systematic Search

---

- ❖ **Systematic search**
  - ❖ BFS, DFS, IDS, Best-First, A<sup>\*</sup>
  - ❖ Keep some history of visited nodes
  - ❖ **Completeness**
    - ❖ **Always complete** for finite search spaces
    - ❖ **Some versions complete** for infinite spaces
  - ❖ Good for “building up” solutions incrementally
    - ❖ State : partial solution
    - ❖ Action : extend partial solution

- ❖ **Local Search**
  - ❖ Algorithms (meta heuristic)
    - ❖ Gradient descent (steepest decent), greedy local search, Simulated Annealing, Genetic Algorithm
  - ❖ Does not keep history of visited nodes
  - ❖ **Not complete**
  - ❖ May be able to argue terminating with “high probability”
  - ❖ Good for “fixing up” candidate solutions
    - ❖ State : complete candidate solution that may not satisfy all constraints
    - ❖ Action : make a small change in the candidate solution

---

# 1. Local Search Algorithms and Optimization Problems

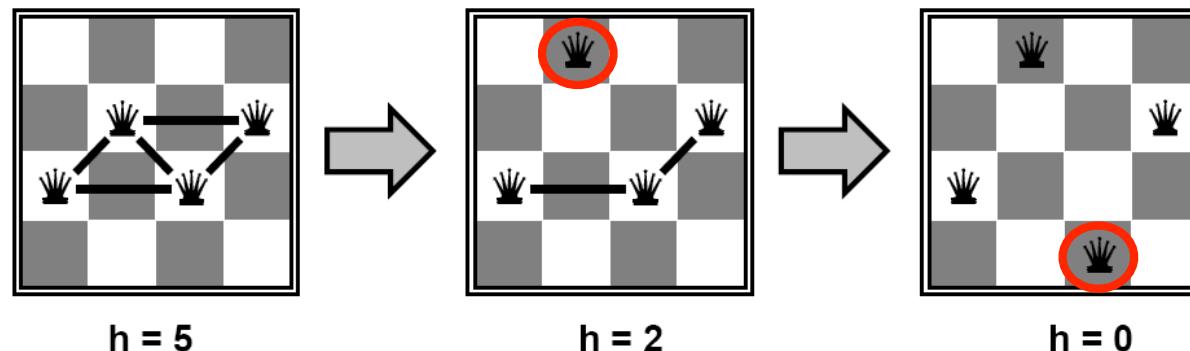
---

- ❖ In many optimization problems, the **state space** is the space of **all possible complete solutions**
- ❖ **Objective function**
  - ❖ How do you measure **cost** or **value** of a state?
  - ❖ Tells us how “good” a given state is, and we want to find the **solution (goal)** by **optimizing** (minimizing or maximizing) the value of this function

- ❖ If the **start state** may not be specified and the **path** to the goal doesn't matter
  - ❖ In such cases, we can use **local search algorithms** that keep a single **current state** (very memory efficient) and gradually try to **improve** it

# Example: $n$ -Queens Problem

- ❖ **Goal:** put  $n$  queens on an  $n \times n$  board
  - ❖ No two queens on the same row, column, or diagonal
  - ❖ An NP-hard problem
- ❖ **Neighbor:** move 1 queen to another row or column
- ❖ **Search:** go from one neighbor to the next....
- ❖ **State space:** all possible  $n$ -queen configurations
- ❖ **Objective function:** number of **pairwise conflicts**
- ❖ What's a possible **local improvement strategy?**
  - ❖ Move one queen to another square to **reduce conflicts**



*h:* the no. of pairwise conflicts

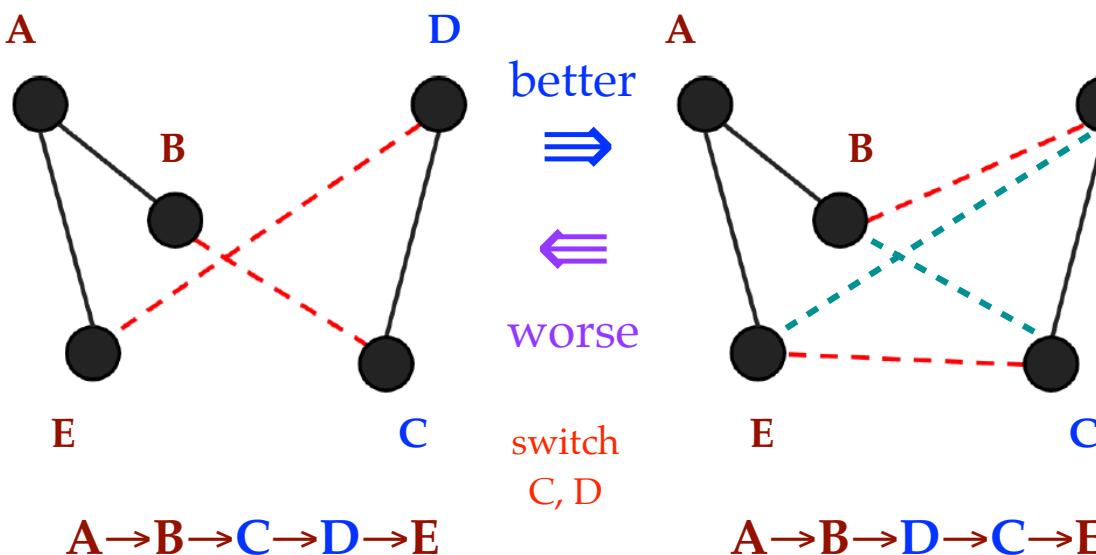
# Example: Traveling Salesman Problem

- ❖ Find the **shortest tour** connecting a given set of cities
- ❖ **State space:** all possible tours
- ❖ **Objective function:** length of tour
- ❖ TSP is an NP-hard problem

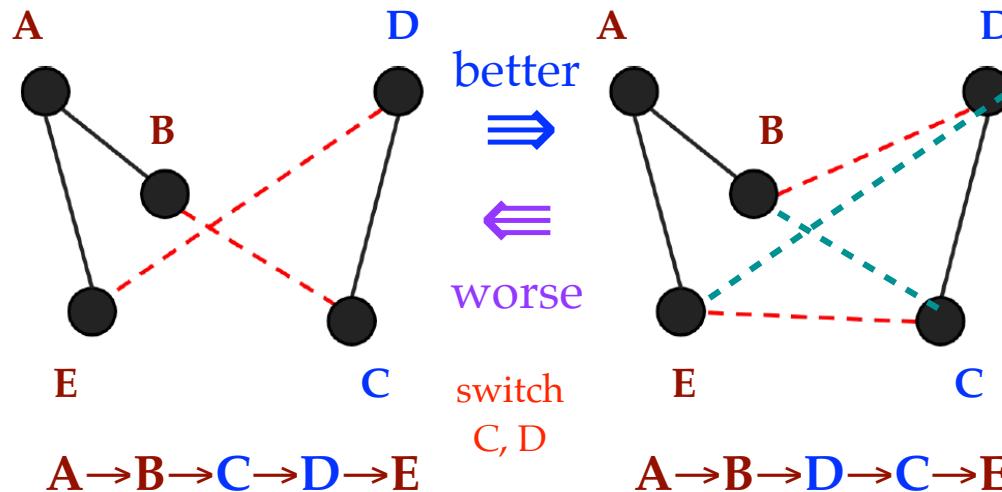


- ❖ What's a possible **local improvement strategy**?
  - ❖ Start with any **complete tour** (could be random)
  - ❖ Solution is a **permutation**
  - ❖ Check if any **pairwise exchange** (e.g., operator : **swap-2**) shorten the tour
  - ❖ Stop criterion : no more better (shorten the tour) move

$\text{ABCDE} \xrightarrow{\text{$C_2^n$ valid moves}} \text{$C_2^n$ states}$ , then apply **steep decent** to select the **min cost**



- We can check swap- $k$  of course, but it takes  $O(n^k)$  time



- Empirically, swap-2 or swap-3 using **steepest decent** makes a good TSP search even there are hundreds of cities
- Number of possible actions to take (resulting number of states)

original state

↓

swap-2 :  $C_2^n(2! - 1) \approx O(n^2)$

↑

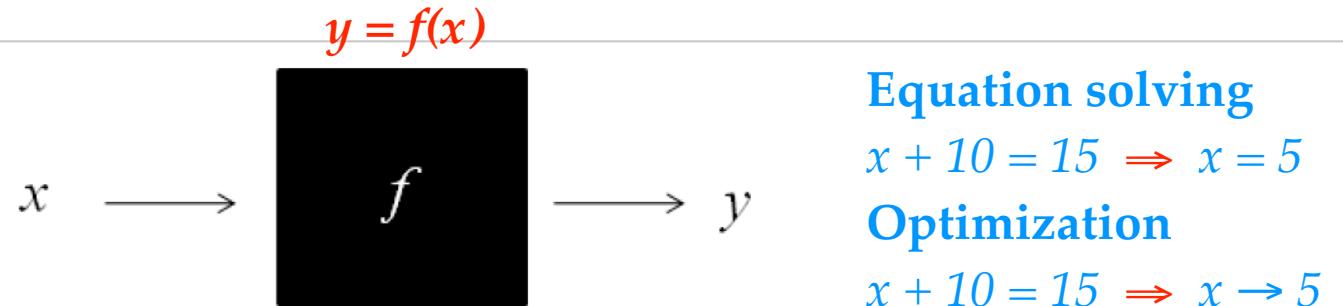
original state

swap-3 :  $C_3^n(3! - 1) \approx O(n^3)$

$$P(n, r) = {}^n P_r = {}_n P_r = \frac{n!}{(n - r)!}$$

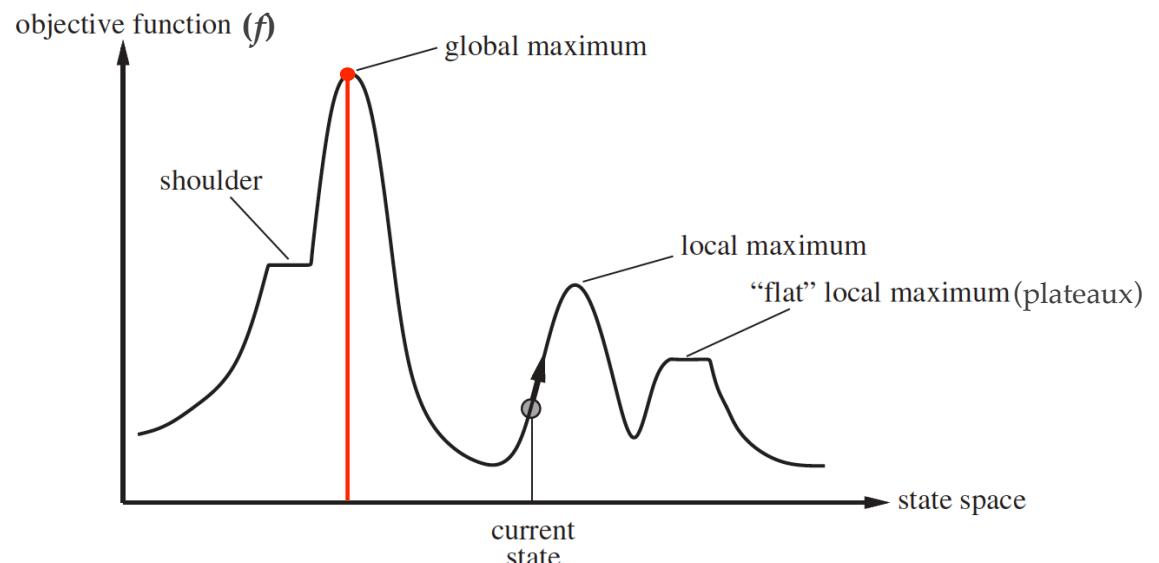
$$C(n, r) = {}^n C_r = {}_n C_r = \binom{n}{r} = \frac{n!}{r!(n - r)!}$$

# Optimization Problems

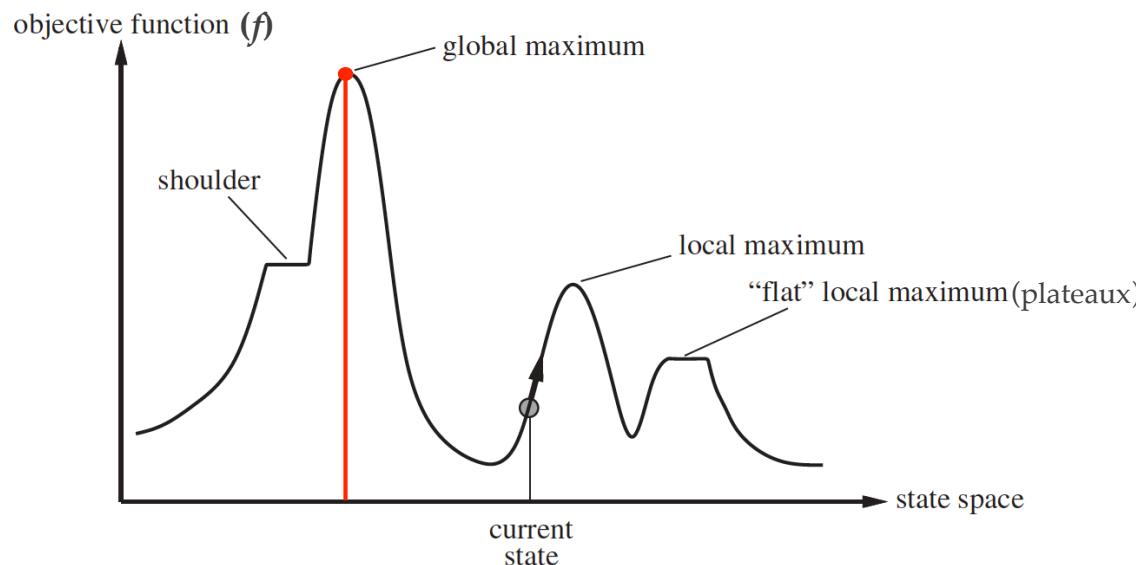


- ❖ Find  $x$  that yields highest  $y$  with an **unknown  $f$**  (**objective function**)
- ❖ Local vs. global optima

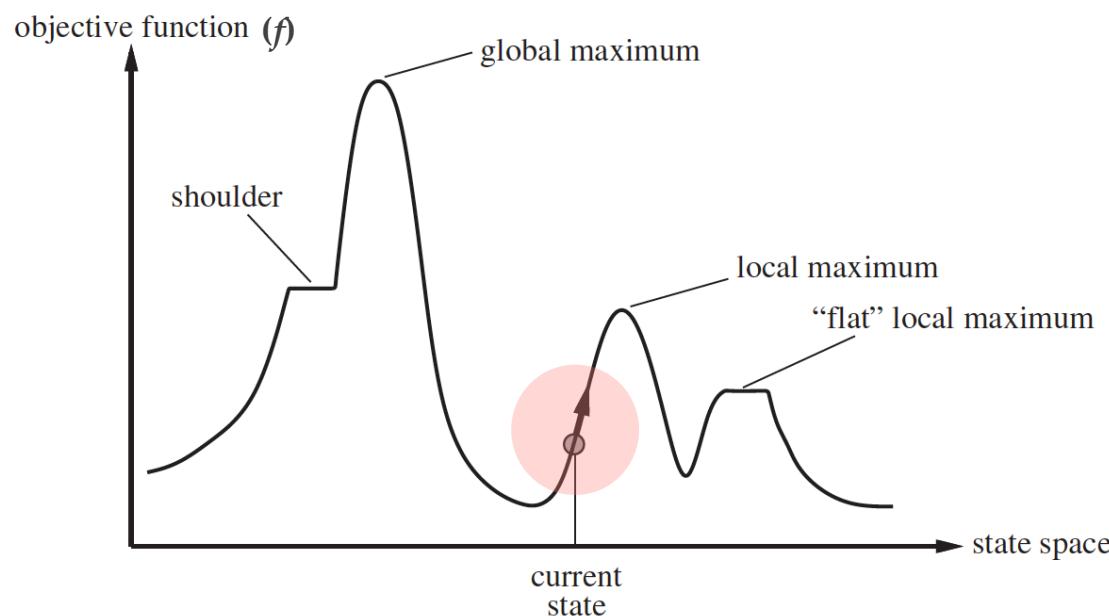
**$f(x)$ :** may be simple function or complex system of equations  
**Glass box:**  $f(x)$  is known  
**Black box:**  $f(x)$  is unknown



- ❖ To understand local search, we consider the **state-space landscape**
  - ❖ **Location:** defined by the **state**
  - ❖ **Elevation:** defined by the **value of the heuristic cost function or objective function ( $f$ )**
    - ❖ If elevation corresponds to **cost**
      - ❖ **Aim:** find the **lowest valley — a global minimum**
    - ❖ If elevation corresponds to an **objective function**
      - ❖ **Aim:** find the **highest peak — a global maximum**



- ❖ A **complete** local search algorithm: always finds a goal if one exists
- ❖ An **optimal** algorithm: always finds a **global minimum/maximum**
- ❖ **Hill-climbing search** modifies the current state and tries to improve it, as shown by the **arrow**



---

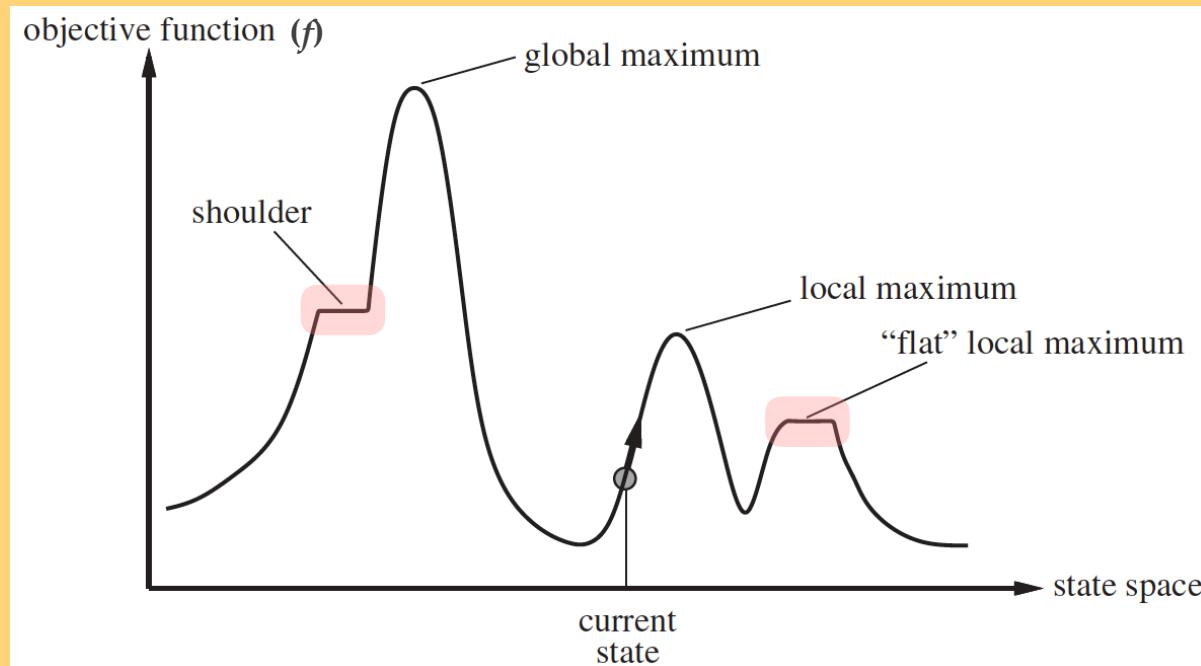
# Hill Climbing Search

---

- ❖ A **loop** that continually moves in the direction of **increasing value**, i.e., **uphill**
- ❖ It **terminates** when it reaches a **peak** where no neighbor has a higher value
- ❖ The algorithm does **not** maintain a **search tree**
  - ❖ The **data structure** for the current node needs only record the **state** and the **value** of the **objective function**
- ❖ Hill climbing does **not look ahead** beyond the immediate neighbors of the current state
- ❖ “Like climbing mount Everest in thick fog with amnesia”

```

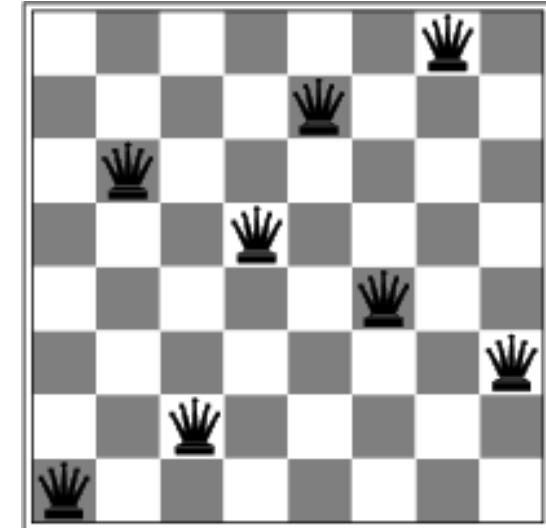
function HILL-CLIMBING(problem) returns a state that is a local maximum
  current  $\leftarrow$  MAKE-NODE(problem.INITIAL-STATE)
  loop do
    neighbor  $\leftarrow$  a highest-valued successor of current
    if neighbor.VALUE  $\leq$  current.VALUE then return current.STATE
    current  $\leftarrow$  neighbor
  
```



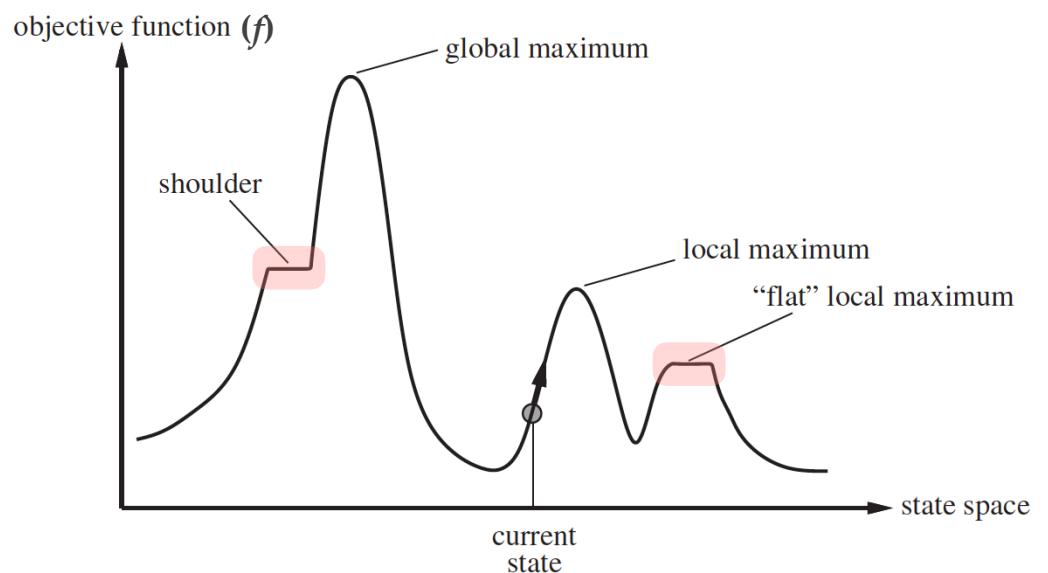
**Set limit for  $\leq$  ? for  
stoping criterion**

- ❖ The hill climbing search algorithm (**steepest-ascent** version)
  - ❖ At each step the current node is **replaced** by the **best neighbor**; i.e., the neighbor with the **highest value**
  - ❖ If a **heuristic cost estimate**,  $h$ , is used, we would find the neighbor with the **highest  $h$**
- ❖ **Variants**
  - ❖ Choose first better successor
  - ❖ Randomly choose among better successors

- ❖ Is it complete / optimal?
  - ❖ No, can get stuck in local optima
  - ❖ Example: local optimum for the 8-queens problem using steepest decent
  - ❖ How to escape local maxima?
  - ❖ Random restart hill-climbing
  - ❖ What about “shoulders”? (not optimal)
  - ❖ What about “plateaux”? (local optimal)

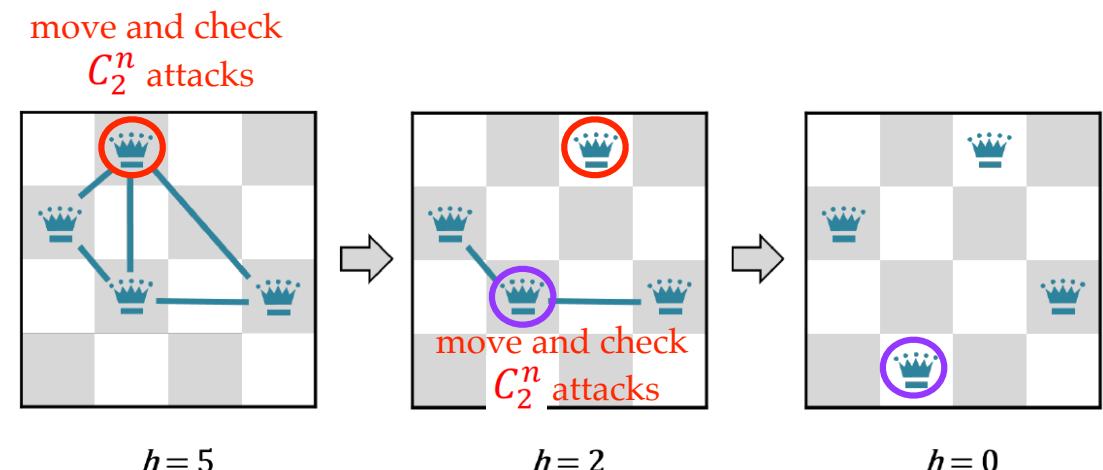


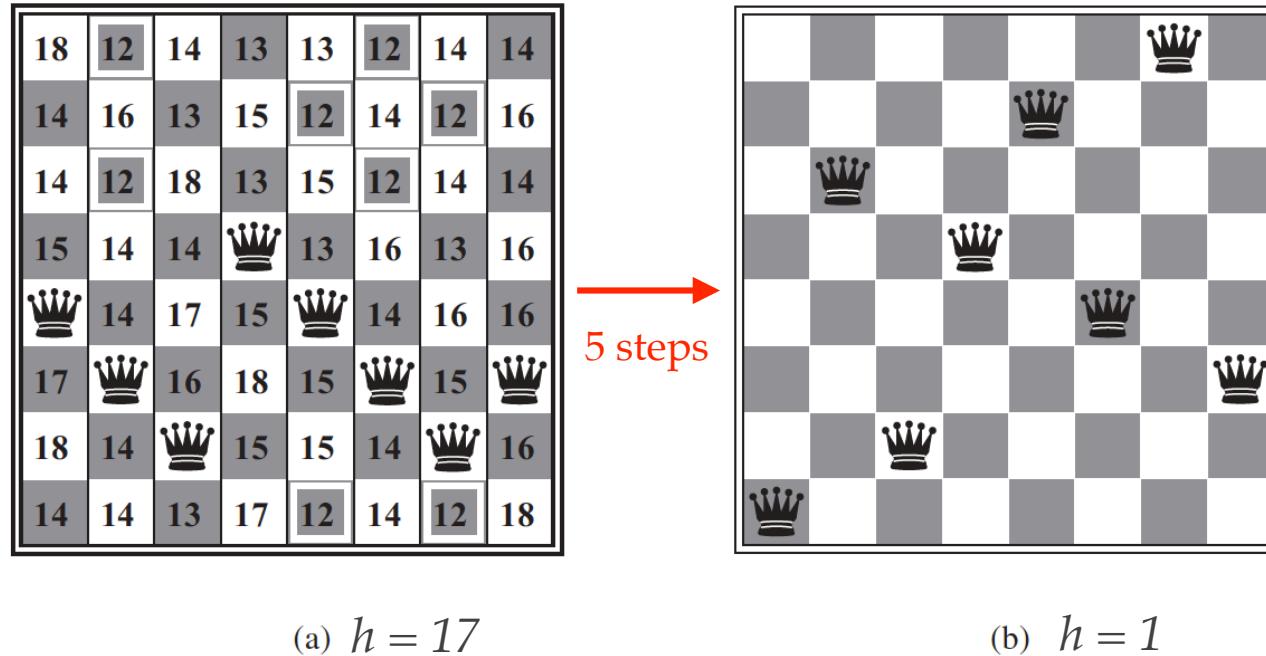
$$h=1$$



# Hill-Climbing Search: 8-Queens Problem

- ❖ The **successors** of a state: all possible states generated by **moving a single queen to another square**
- ❖ The **heuristic cost function**,  $h$ : the number of pairs of queens that are attacking each other, either directly or indirectly
- ❖ The **global minimum** of this function: **zero**, which occurs only at **perfect solutions**
- ❖ Move a queen to **reduce conflicts**
- ❖ Almost always solve  $n$ -queen puzzle (optimal solution) **immediately** even for  $n \approx 10^6$  but not for  $n \approx 10^{12}$



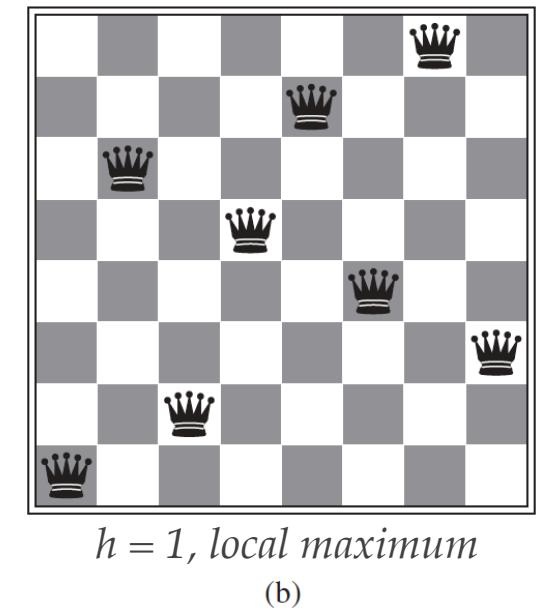
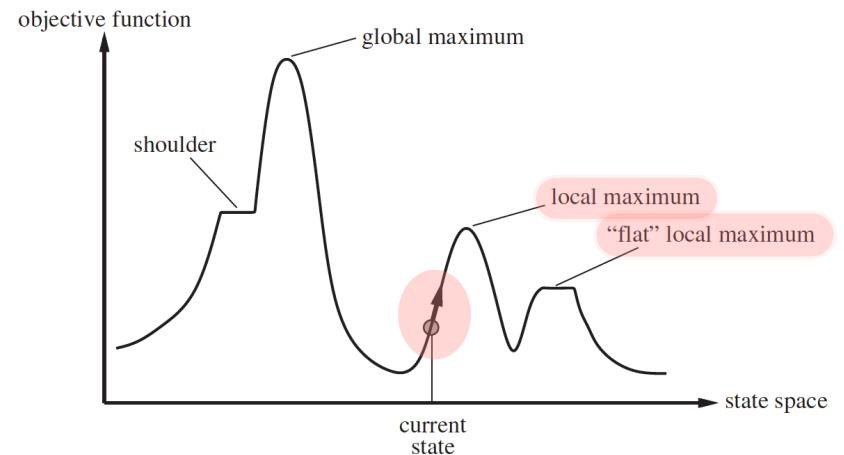


- ❖ (a)  $h$  = number of pairs of queens that attack each other, either directly or indirectly ( $h = 17$  for (a))
- ❖ Each number indicates  $h$  if we move a queen in its **column** to that square
- ❖ The values of all its successors, with the **best successors** having  $h = 12$
- ❖ Hill-climbing algorithms typically **choose randomly** among the set of **best successors** if there is **more than one**
- ❖ From the state in (a), it takes just **five** steps to reach the state (b), which has  $h=1$  and is very nearly a solution

- ❖ Hill climbing often gets stuck for the following reasons (**difficulties of ALL local search algorithms**)

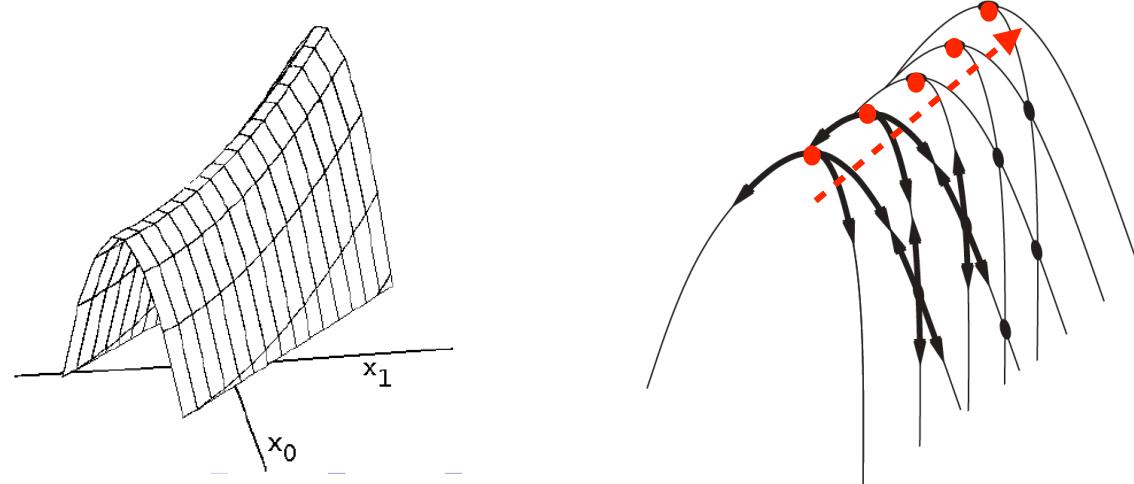
- ❖ **Local maxima**

- ❖ A **peak** that is **higher** than each of its neighboring states but **lower** than the global maximum
- ❖ Hill-climbing algorithms that reach the **vicinity** of a **local maximum** will be drawn **upward** toward the peak but will then be **stuck with nowhere else to go**
- ❖ For example, the state in figure (b) is a **local maximum** ( $h = 1$ ); **every move** of a single queen makes the situation **worse** (**higher  $h$** )



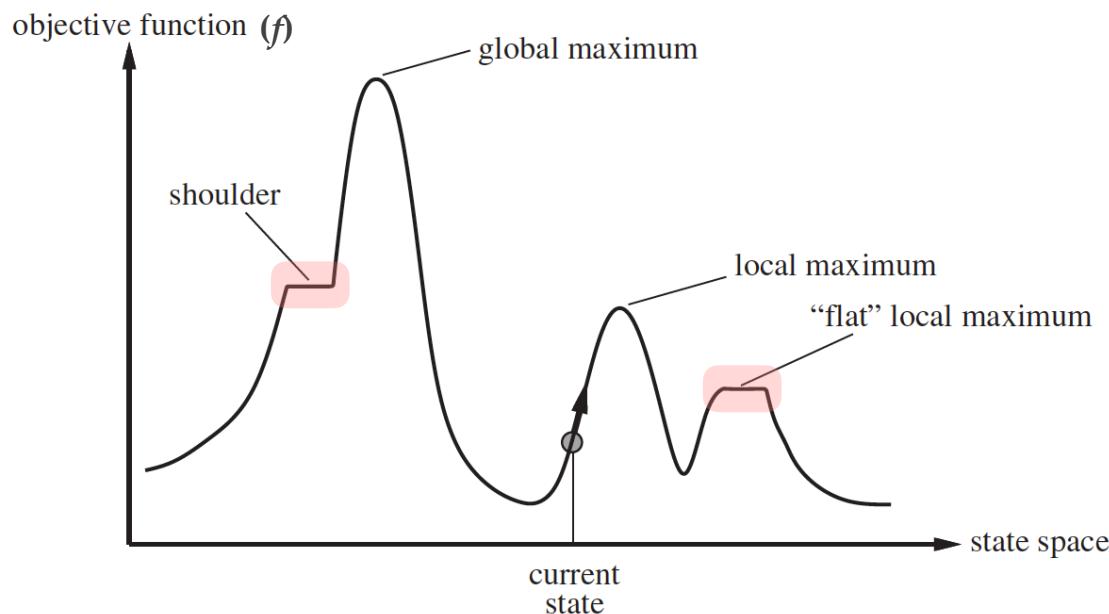
## ❖ Ridges

- ❖ Fold a piece of paper and hold it tilted up at an **unfavorable angle** to every possible search space step
- ❖ **Every neighbor appears to be downhill; but the ridge leads uphill**
- ❖ From each **local maximum**, all the available actions point downhill
- ❖ The **grid of states** (red circles) is superimposed on a ridge **rising from left to right**, creating a **sequence of local maxima** that are **not directly connected to each other**, so it is very difficult for greedy algorithms to navigate



## ❖ Plateaux

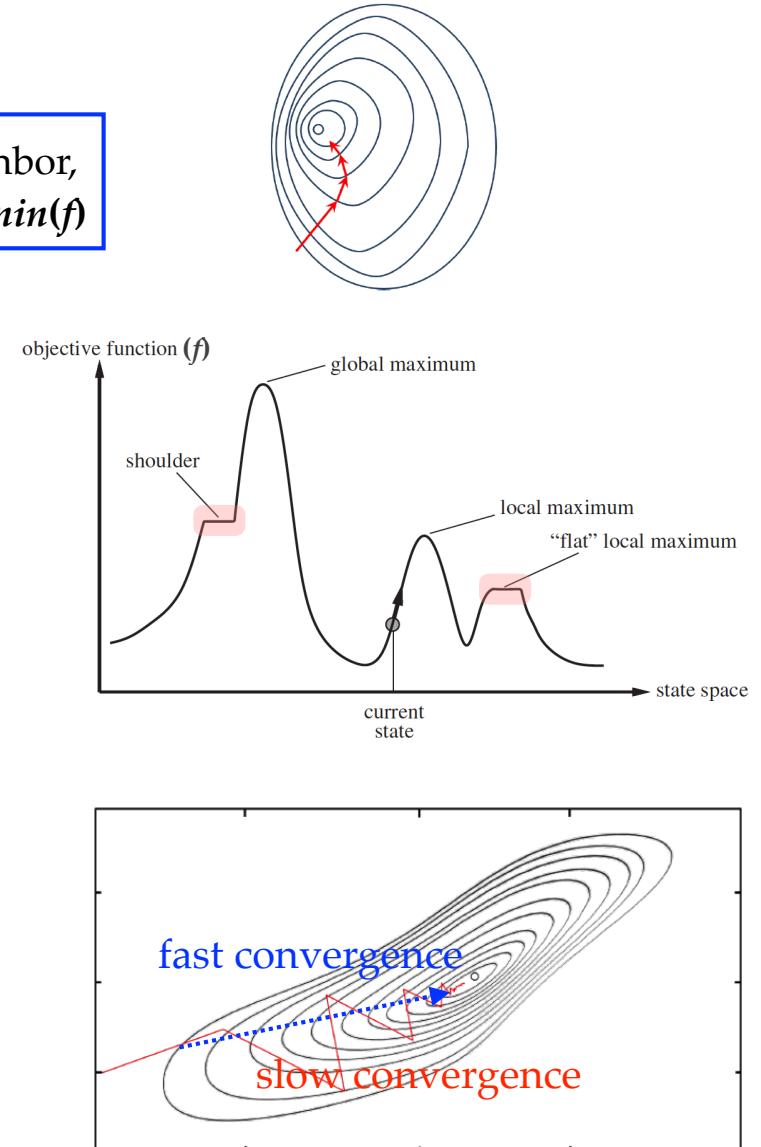
- ❖ A plateau is a **flat area** of the state-space landscape
- ❖ It can be a **flat local maximum**, from which no uphill exit exists, or a **shoulder**, from which progress is possible
- ❖ A hill-climbing search might get lost on the plateau



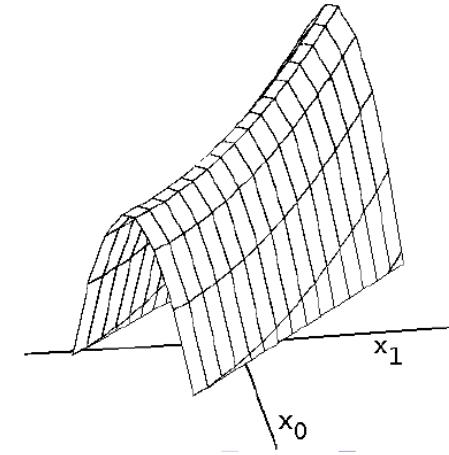
# Steepest Descent (Ascent)

- ❖ Steepest descent is a.k.a. **gradient descent**

$s \xrightarrow{a \in \text{Actions}(s)} s'$ , (successor or neighbor, evaluate  $f(s')$  and find  $\max(f)$  or  $\min(f)$ )
- ❖ Allowing **sideway moves** is usually good, but need to put a **limit** to prevent infinite loop
- ❖ Find only the **nearest** local optimum
- ❖ No guarantee **optimal** solution
- ❖ **Performance:** may suffer from **slow convergence** due to **zig-zagging** behavior



- ❖ Ridges and plateau are difficult too



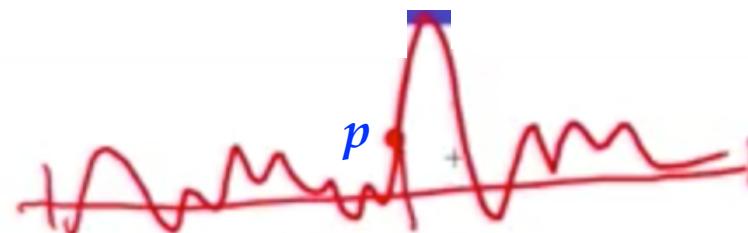
- ❖ Random **restart** helps

- ❖ Prob.  $p$  to succeed (restart on the hill close to the global optimal) (e.g.  $p = 0.1$ )

- ❖  $1/p$  restarts needed to succeed (e.g.  $1/p = 10$ )

- ❖ **Cost =  $1 \times \text{cost-of-success} + (1 - p)/p \times \text{cost-of-failure}$**   
(1 time) (9 times)

Failure probability:  $(1-p)+(1-p)^2+(1-p)^3+\dots = (1-p)/p$



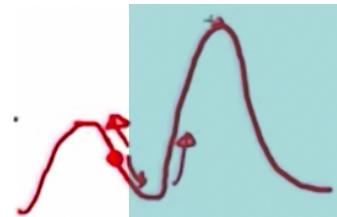
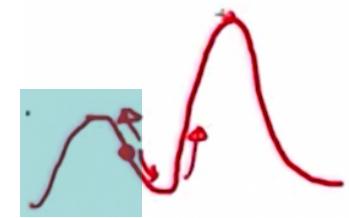
---

# Simulated Annealing (SA)

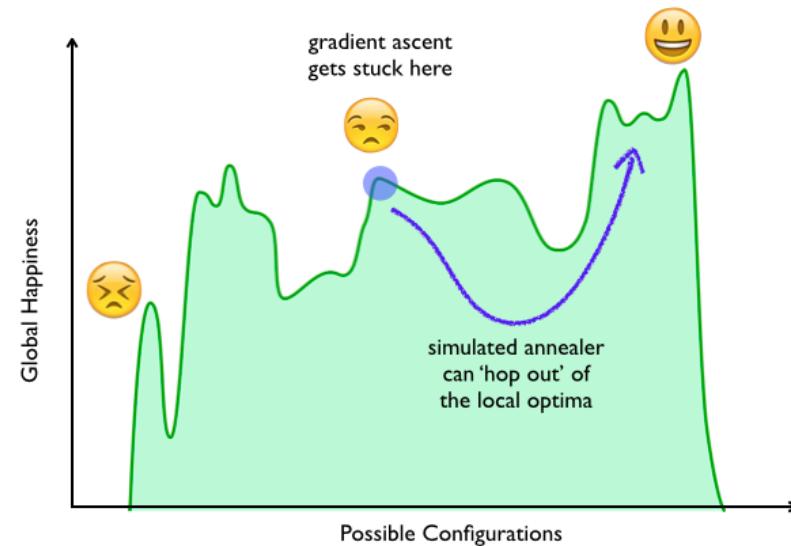
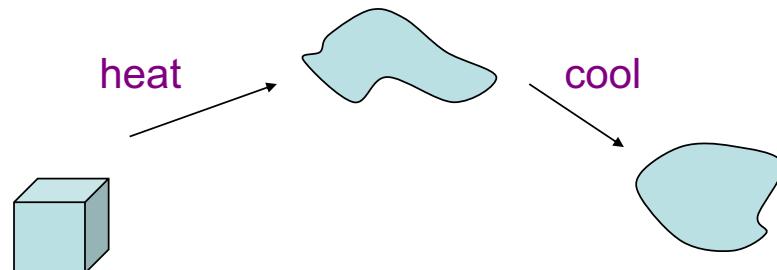
---

- ❖ **Hill-climbing**
  - ❖ Never makes "downhill" moves toward states with lower value (or higher cost) — is **incomplete**, because it can get **stuck** on a **local maximum**
- ❖ **Purely random walk**
  - ❖ Moving to a successor chosen uniformly at random from the set of successors — is **complete** but extremely **inefficient**

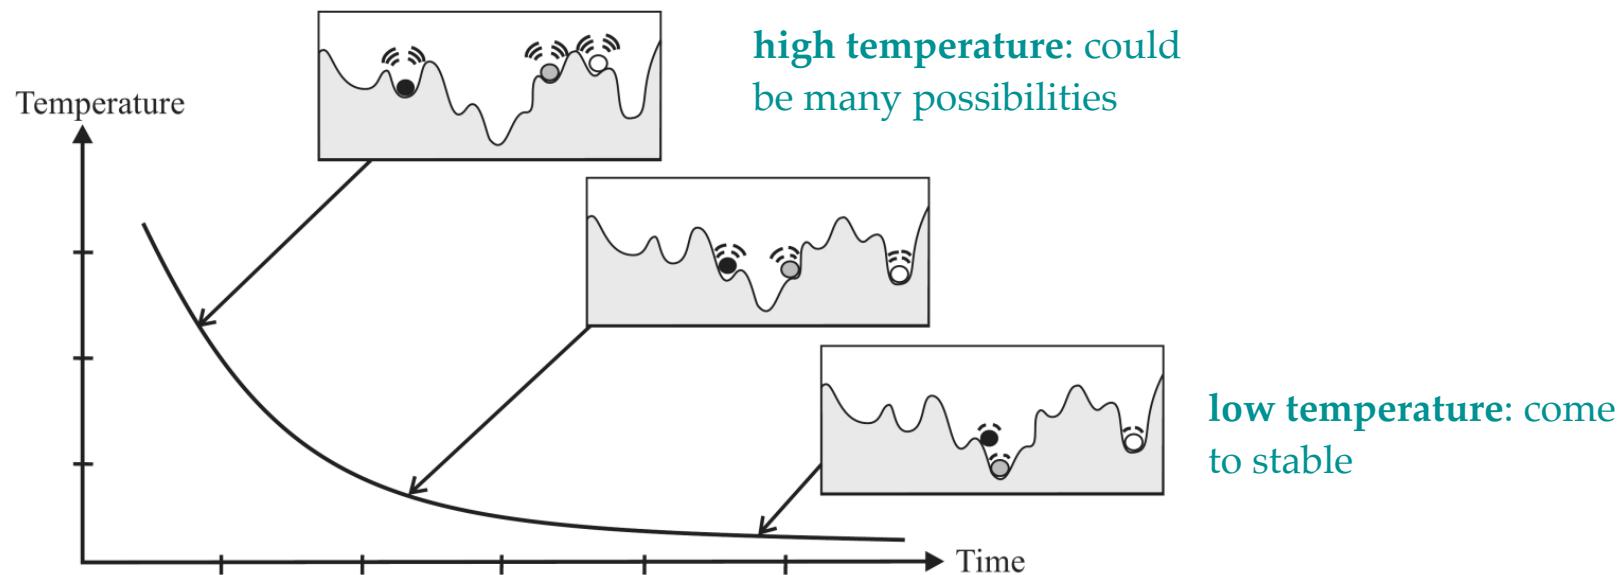
- ❖ Simulated annealing
  - ❖ Combine hill climbing with a random walk in some way that yields both **efficiency** and **completeness**
  - ❖ Motivation
    - ❖ Improve the problem of **steepest decent** that gives only the **nearest local optimal**
    - ❖ SA expects to use **random walk** to explore extensively (with some probability of going "downhill" and then going another way of "uphill") in the beginning, and then use **steepest decent** afterward



- ❖ In metallurgy, annealing is the process used to temper or harden metals and glass by **heating them to a high temperature** and then **gradually cooling them**, thus allowing the material to reach a **low energy crystalline state**
- ❖ The simulated-annealing solution is to **start by shaking hard** (i.e., at a high temperature) and then **gradually reduce the intensity of the shaking** (i.e., lower the temperature)



- ❖ Idea: escape local maxima by allowing some "bad" moves but gradually **decrease their frequency**
- ❖ Simulated annealing introduces a temperature parameter,  $T$ , which **cools down** as time goes by



[ $f(x)$ ]

- ❖ If the **new state** has a **better** value ( $\Delta E > 0$ ), SA **accepts** the new state [ $\Delta E = \text{next.VALUE} - \text{current.VALUE}$ ]
- ❖ If the **new state** has a **worse** value ( $\Delta E < 0$ ), SA **accepts** the new state with a **probability**  $e^{\Delta E/T}$   
[ $\Delta E$  is negative: (a)  $T \rightarrow 0$  then  $e^{-\infty} \rightarrow 0$ , (b)  $T \rightarrow \infty$  then  $e^0 \rightarrow 1$ ]
  - When  $T$  is high, it is easier to accept a state that is worse than the current state.
  - When  $T$  is low, less willing to accept a state that is worse than the current state

- ❖ **Proved:** If temperature ( $T$ ) decreases slowly enough, then SA search will find a **global optimum with probability approaching 1**  
 $T(t)$  : T (temperature) is a **decreasing** function of  $t$  (time) [inversely proportional]
- ❖ However
  - ❖ This usually takes a **VERY, VERY long time, impractically long**
  - ❖ Note: any **finite search space, RANDOM GUESSING** also will find a **global optimum with probability approaching 1**
  - ❖ So, ultimately this is a **very weak claim**
- ❖ SA is widely used in “**very big**” problems
  - ❖ VLSI layout, national airline scheduling, large logistics ops, etc.

**function** SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

**inputs:** *problem*, a problem

*schedule*, a mapping from time to “temperature”

*current*  $\leftarrow$  MAKE-NODE(*problem*.INITIAL-STATE)

**for**  $t = 1$  **to**  $\infty$  **do**

$T \leftarrow \text{schedule}(t)$  [T is a **decreasing** function of  $t$ ]

**if**  $T = 0$  **then return** *current*

*next*  $\leftarrow$  a randomly selected successor of *current*

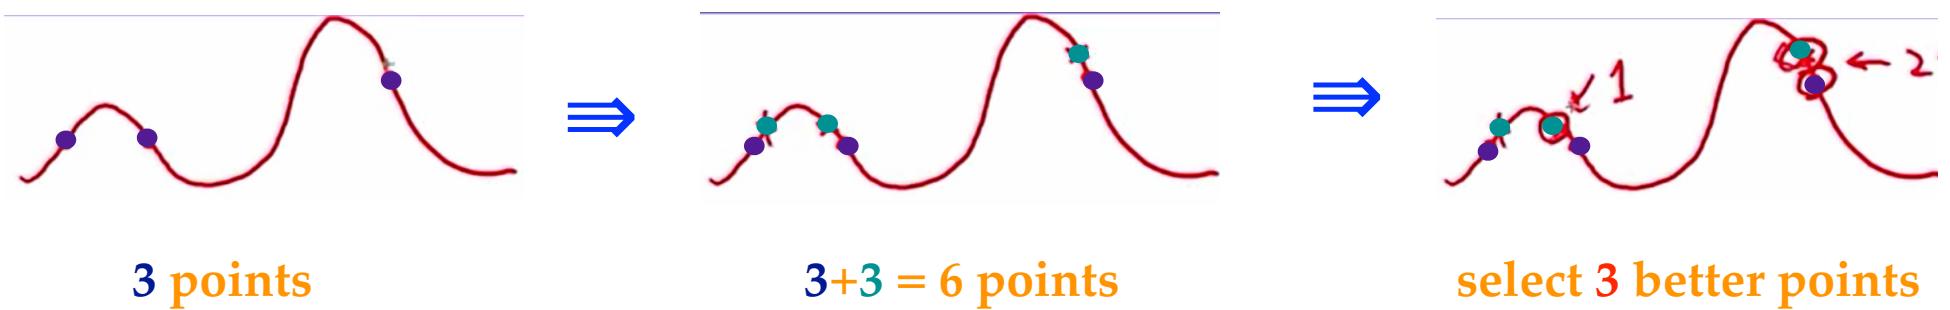
$\Delta E \leftarrow \text{next}.\text{VALUE} - \text{current}.\text{VALUE}$

**if**  $\Delta E > 0$  **then** *current*  $\leftarrow$  *next*

**else** *current*  $\leftarrow$  *next* only with probability  $\Delta E/T$

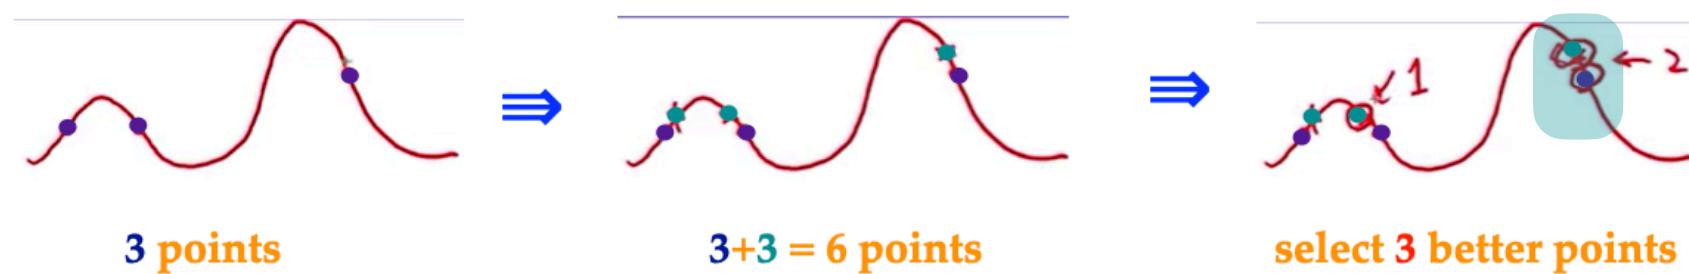
# Local Beam Search

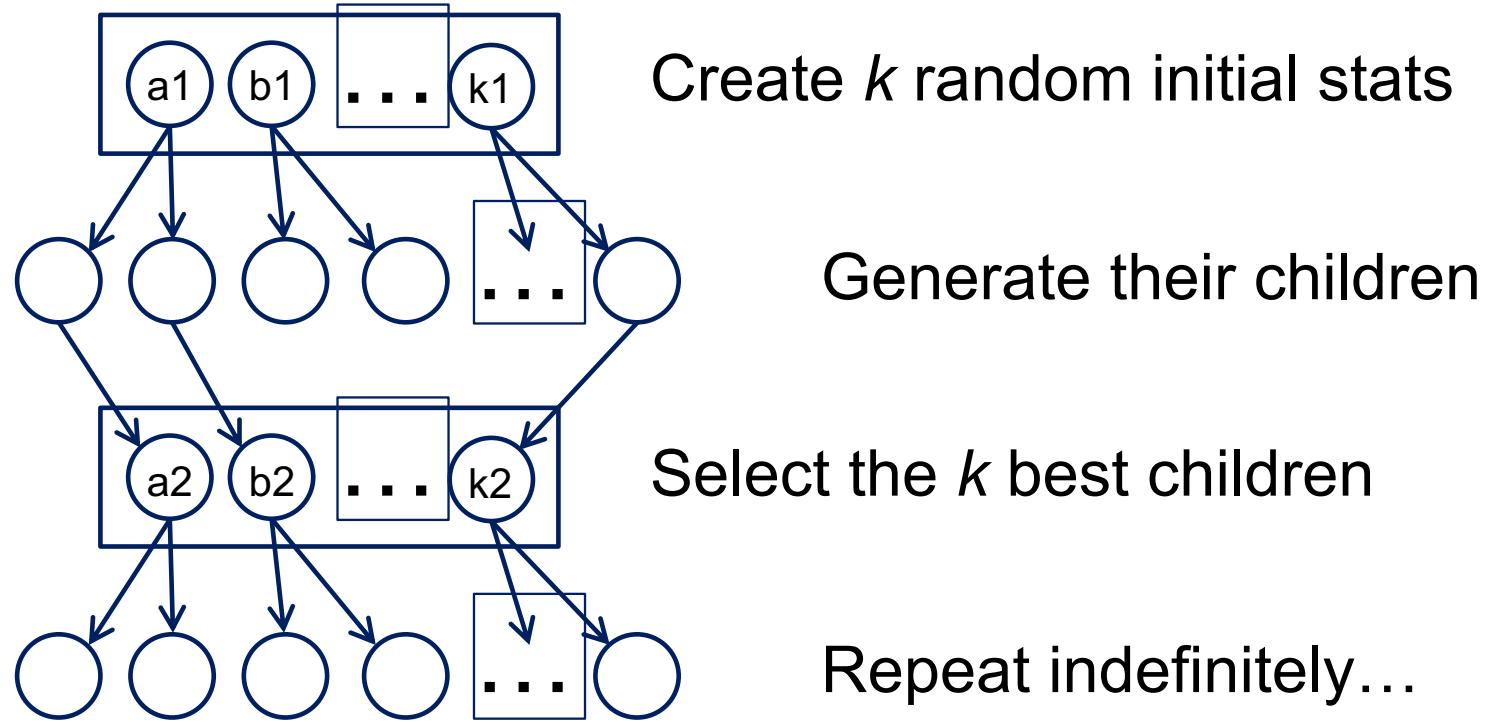
- ❖ Start with  $k$  randomly generated states
- ❖ At each iteration, all the successors of all  $k$  states are generated
- ❖ Keep best  $k$  states instead of just one



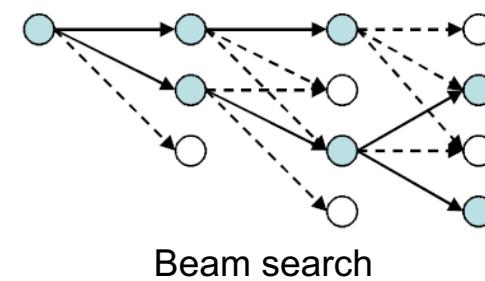
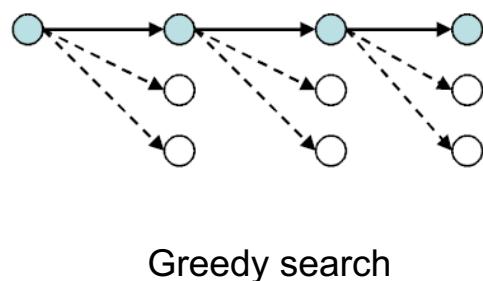
- ❖ If any one is a **goal state**, stop; else select the  $k$  best successors from the complete list and repeat

- ❖ Concentrates search effort in areas believed to be fruitful
  - ❖ May lose diversity as search progresses, resulting in wasted effort (clustering effect may be wrong)





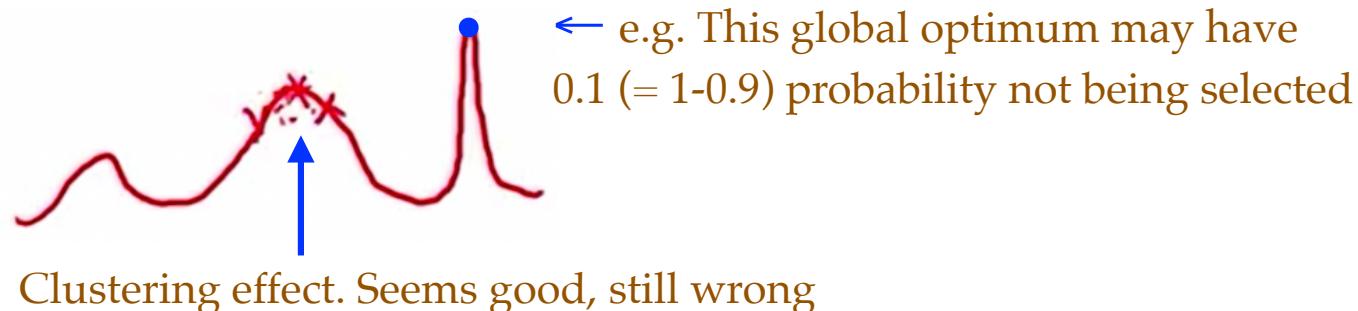
- ❖ Is this the same as running  $k$  greedy searches in parallel? (**No**, beam search have interact among states, but greedy search does not)



- ❖ **Stochastic beam search** randomly chooses  $k$  successors with a probability proportional to their **goodness** [ $f(x)$ ].
- ❖ Neighbors selection is not limited to higher hills (not 100%, but with higher probability), there is some probability to select from lower hills (with lower probability)

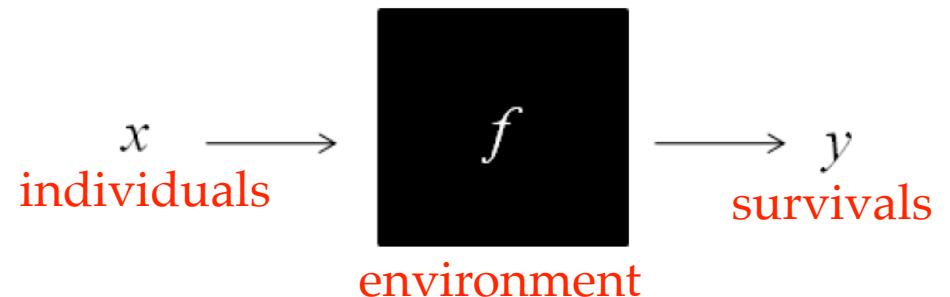
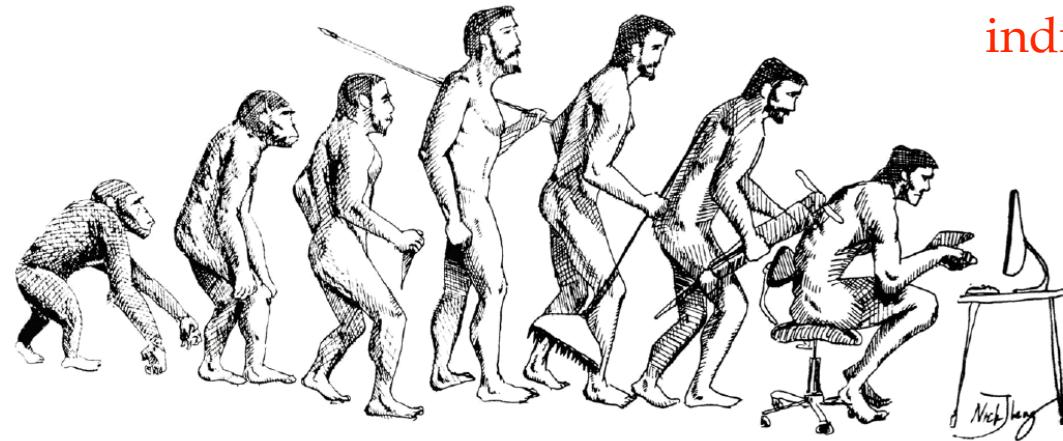
### ❖ Variant

- ❖ If select 10 points in total, the best 5 points are unconditionally selected, and the remaining 5 points are selected by stochastic



# Darwinian Evolution

- ❖ Evolution
  - ❖ The change in populations of organisms over generations



- ❖ Darwin's idea

- ❖ Natural selection [自然選擇]

- ❖ The process whereby organisms **better adapted** to their environment tend to survive and produce more offspring

- ❖ Struggle to survive [物競天擇] (mandatory condition)

- ❖ The **competition in nature** among organisms of a population to maintain themselves in a given environment and to survive to reproduce others of their kind

- ❖ Survival of the fittest [適者生存]

- ❖ A natural process resulting in the evolution of organisms **best adapted** to the environment

- ❖ Genetic variation [遺傳變異]

- ❖ **Variations of genomes** between members of species, or between groups of species thriving in different parts of the world as a result of **genetic mutation**

---

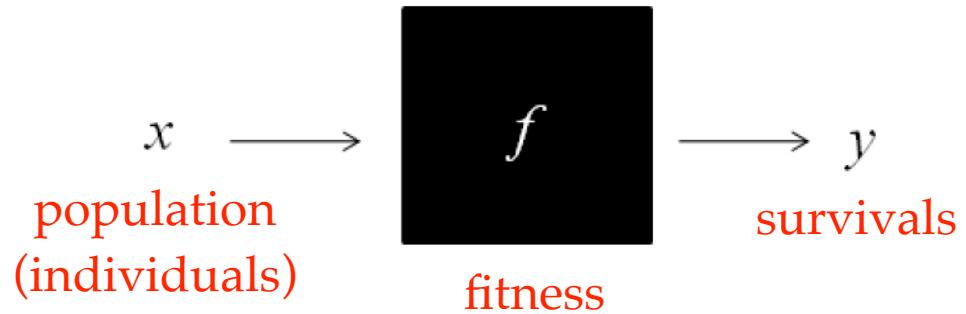
# Genetic Algorithms

---

- ❖ A genetic algorithm (or **GA**) is a variant of **stochastic beam search** in which **successor states** are generated by **combining two parent states** rather than by modifying a single state
- ❖ **State:** a string over a finite alphabet (an **individual**)
- ❖ Start with  $k$  randomly generated states (a **population**)
- ❖ **Fitness function** (= heuristic objective function)
  - ❖ Higher fitness values for better states
- ❖ **Select** individuals for next generation based on **fitness**
  - ❖  $p(\text{individual in next gen.}) = (\text{individual fitness}) / \sum (\text{population fitness})$
- ❖ **Crossover** fit parents to yield next generation (**off-spring**)
- ❖ **Mutate** the offspring randomly with some **low probability**

- ❖ **Simple Genetic Algorithm (SGA) [1975-1990]**
- ❖ **Encoding**
  - ❖ Solution candidates are **encoded** into **chromosomes**
  - ❖ Chromosomes can be binary, integral, real-valued, permutations (e.g. TSP), etc.
- ❖ **Initialization**
  - ❖ At the beginning, the GA creates a **population** of **chromosomes**
  - ❖ The alleles of the chromosomes can be **randomly initialized** or **assigned** based on **prior knowledge**

- ❖ **Evaluation**



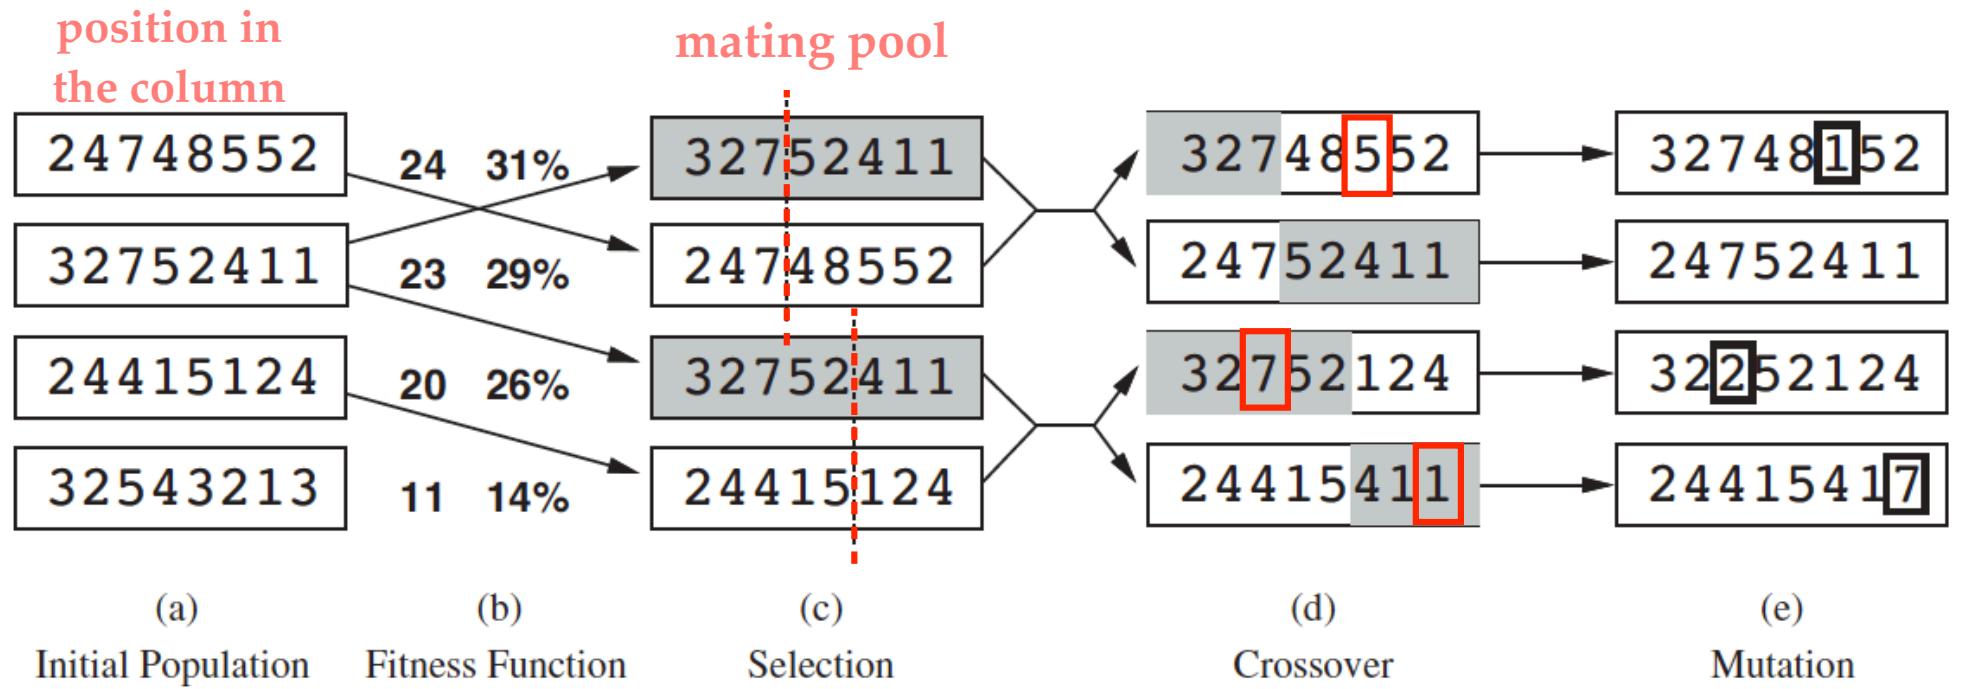
- ❖ After initialization, the GA applies the **fitness function** to evaluate the quality of each chromosome in the current population

- ❖ **Selection**

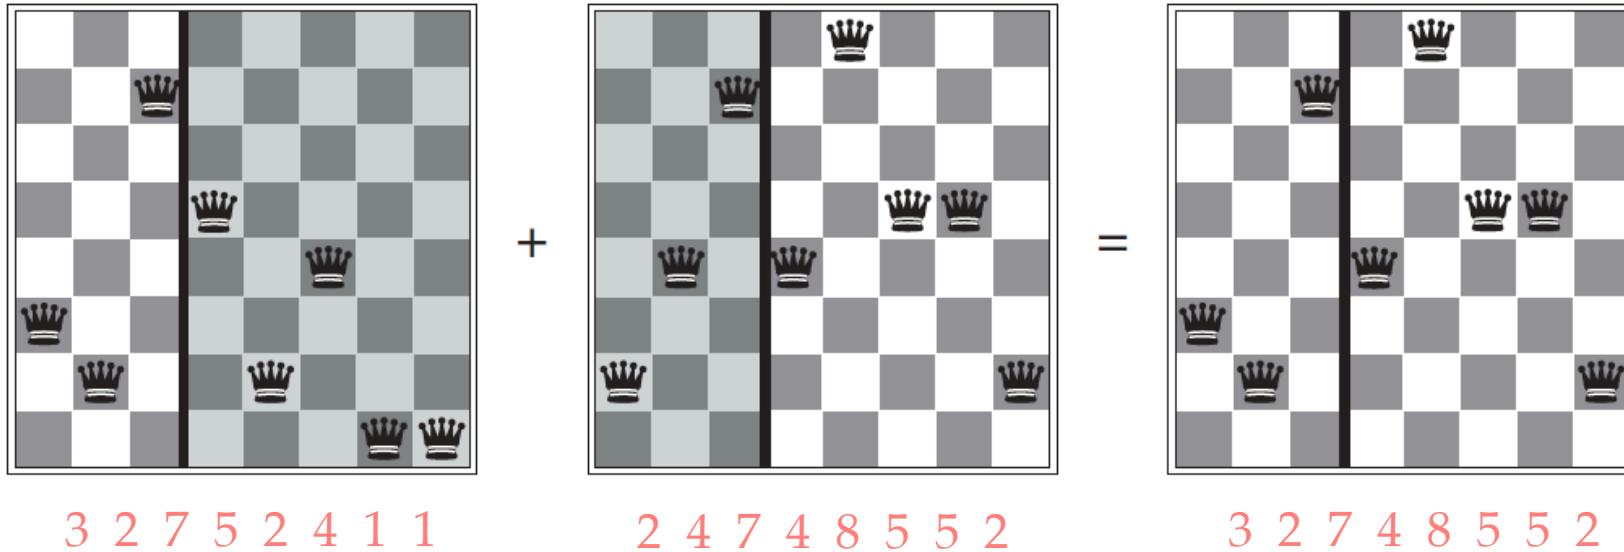
- ❖ Select **promising chromosomes** to have their information passed on to the next generation

- ❖ **Recombination**

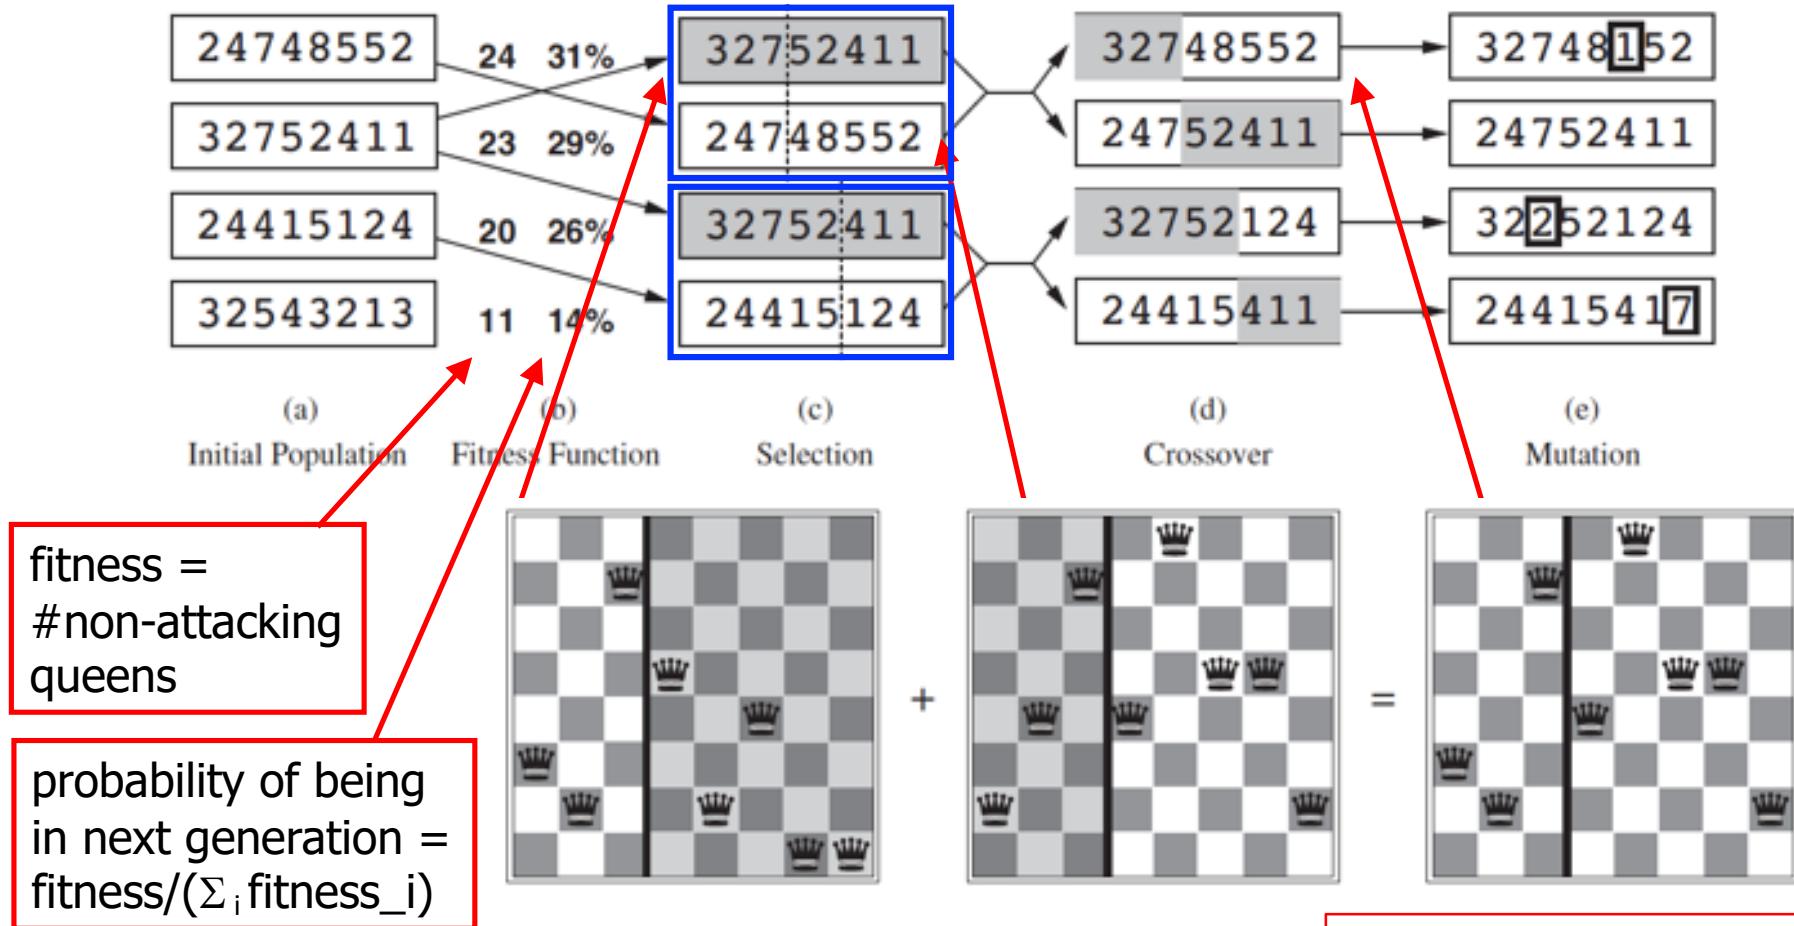
- ❖ Several **parental chromosomes** are **recombined** to generate **new chromosomes** with characteristics inherited from parents
- ❖ Traditional recombination consists of **crossover** and **mutation**



- ❖ The genetic algorithm, illustrated for **digit strings** representing 8 queens states
- ❖ The **initial population** in (a) is ranked by the **fitness function** (#non-attacking queen pairs) in (b), resulting in pairs for **mating** in (c)
- ❖ They produce **offspring** in (d), which are subject to **mutation** in (e)



- ❖ The 8 queens states corresponding to the first two parents in (c) and the first offspring in (d)
- ❖ The **shaded columns** are lost in the crossover step and the unshaded columns are retained



- Fitness function: #non-attacking queen pairs
  - min = 0, max =  $8 \times 7/2 = 28$   $C_2^8$
- $\sum_i \text{fitness}_i = 24+23+20+11 = 78$
- $p(\text{child}_1 \text{ in next gen.}) = \text{fitness}_1/(\sum_i \text{fitness}_i) = 24/78 = 31\%$
- $p(\text{child}_2 \text{ in next gen.}) = \text{fitness}_2/(\sum_i \text{fitness}_i) = 23/78 = 29\%$ ; etc

How to convert a fitness value into a probability of being in the next generation.

**function** GENETIC-ALGORITHM(*population*, FITNESS-FN) **returns** an individual

**inputs:** *population*, a set of individuals

          FITNESS-FN, a function that measures the fitness of an individual

**repeat**

*new\_population*  $\leftarrow$  empty set

**for** *i* = 1 **to** SIZE(*population*) **do**

*x*  $\leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

*y*  $\leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

*child*  $\leftarrow$  REPRODUCE(*x*, *y*)

**if** (small random probability) **then** *child*  $\leftarrow$  MUTATE(*child*)

      add *child* to *new\_population*

*population*  $\leftarrow$  *new\_population*

**until** some individual is fit enough, or enough time has elapsed

**return** the best individual in *population*, according to FITNESS-FN

---

**function** REPRODUCE(*x*, *y*) **returns** an individual

**inputs:** *x*, *y*, parent individuals

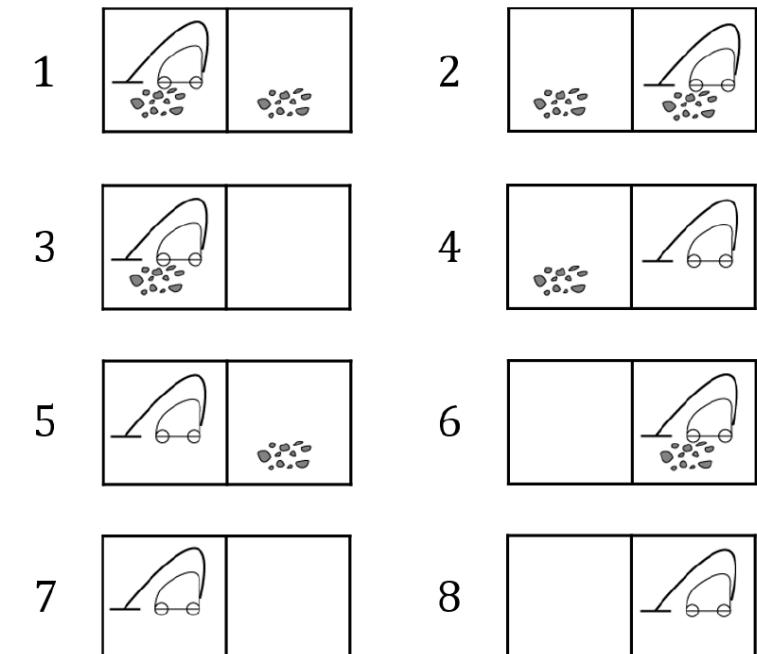
*n*  $\leftarrow$  LENGTH(*x*); *c*  $\leftarrow$  random number from 1 to *n*

**return** APPEND(SUBSTRING(*x*, 1, *c*), SUBSTRING(*y*, *c* + 1, *n*))

- ❖ In this more popular version, each **mating** of two parents produces only **one offspring**, not two.

# 3. Searching with Non-deterministic Actions

- ❖ Vacuum world
  - ❖ Actions = {LEFT, RIGHT, SUCK}
  - ❖ SUCK in an **erratic** vacuum world
    - ❖ When applied to a **dirty square**, the action cleans the square and **sometimes** cleans the dirt in an **adjacent square** as well
    - ❖ When applied to a **clean square**, the action **sometimes** deposits dirt on the square



# Generalization of State-Space Model

---

- ❖ Generalize the **transition function** to return a set of possible outcomes
  - ❖ Old function:  $S \times A \rightarrow S$
  - ❖ New function:  $S \times A \rightarrow 2^S$  [a family of subsets of  $S$ , or a family of sets over  $S$ ]
- ❖ Generalize the **solution** to a contingency plan

*if state = s then action-set-1 else action-set-2*

e.g. if Bstate = {6} then Suck else []

- ❖ Generalize the search tree to an **AND-OR** tree

## ❖ AND-OR search tree

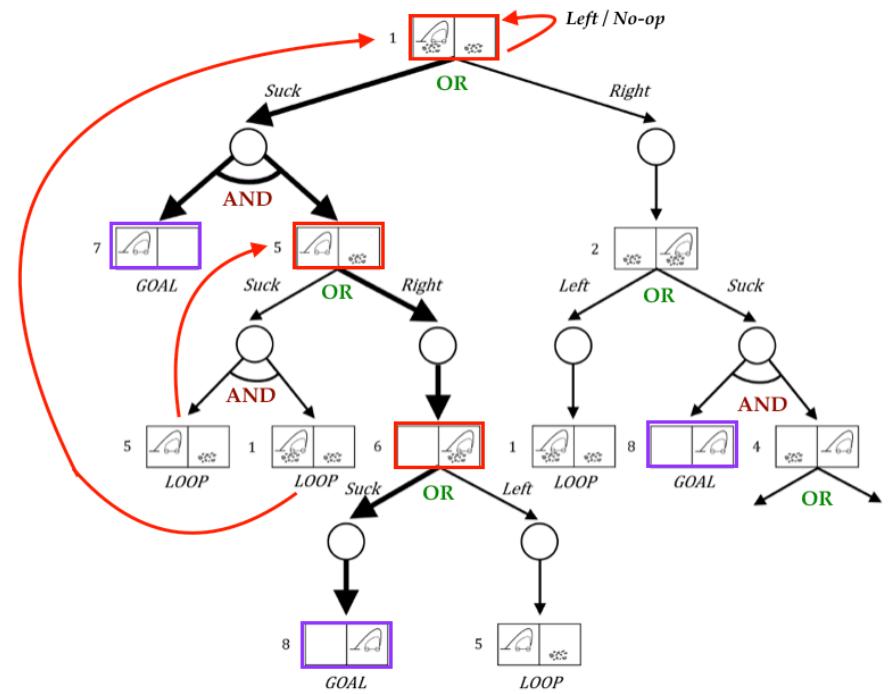
### ❖ AND nodes: actions (**circles**)

- ❖ You can end up in **multiple states** as the **result** of an action
- ❖ You have to find a **path** from all of these states
- ❖ Each **leaf** of the partial search tree of AND node should be a **goal** (e.g. 7, 8)

### ❖ OR nodes: states (**squares**)

- ❖ Try each action
- ❖ Any one action can lead to the goal state(s)
- ❖ Need to reach the **goal state** at **EVERY leaf** **□**

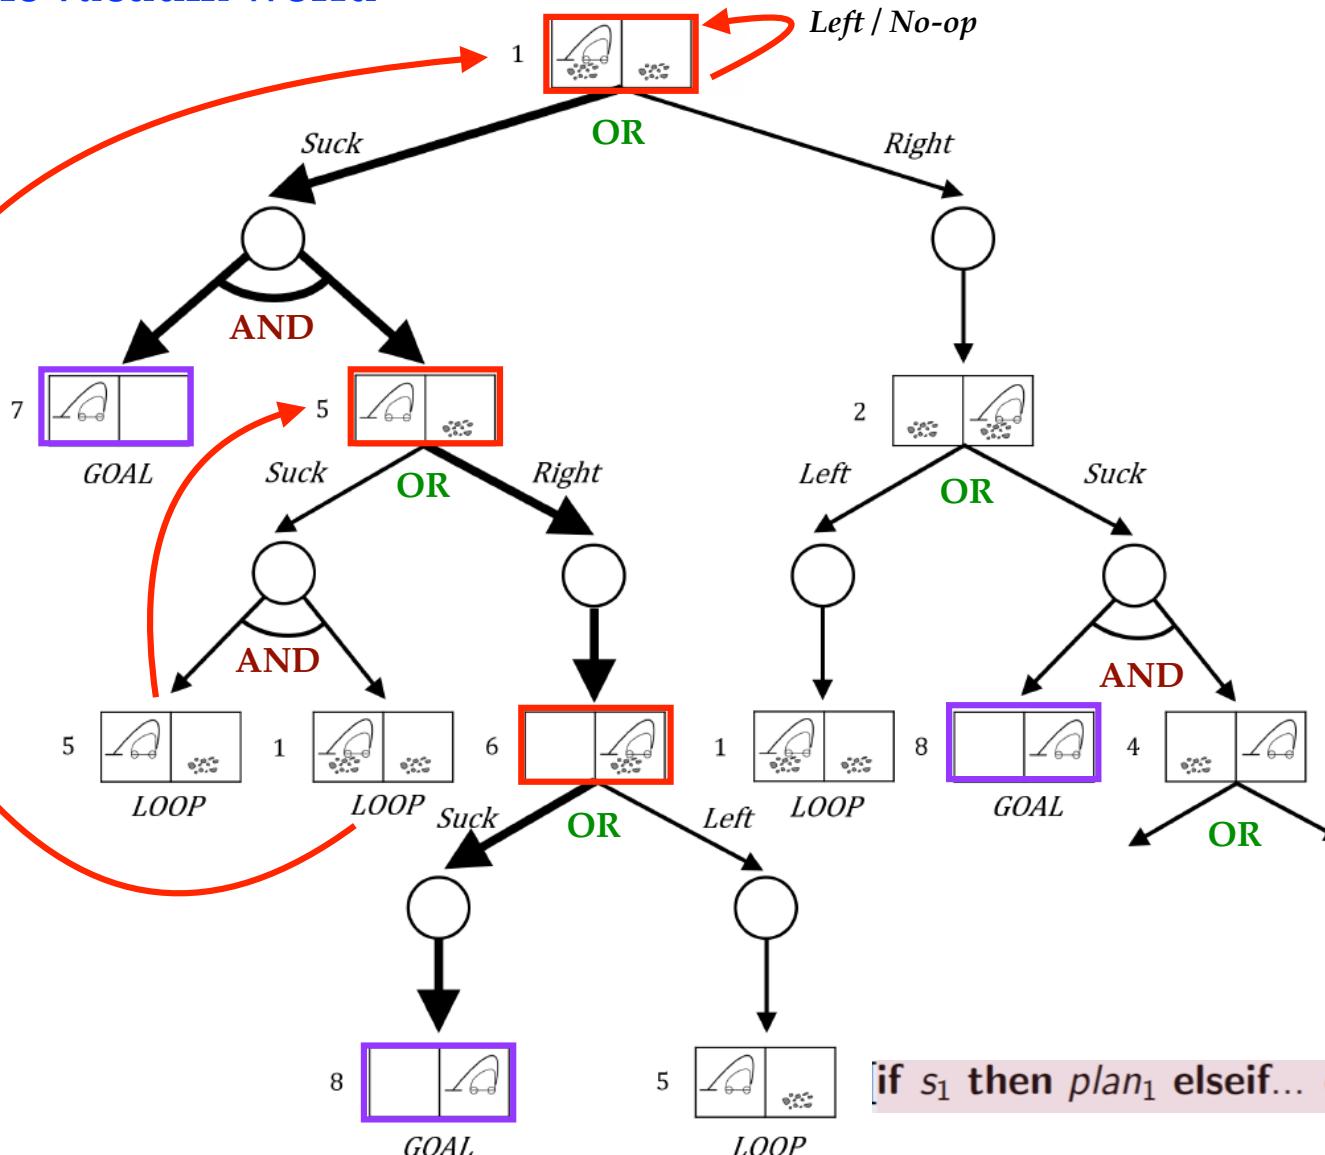
Erratic vacuum world



1<sub>st</sub> solution: SUCK

2<sub>nd</sub> solution: SUCK → RIGHT → SUCK

## Erratic vacuum world



**if**  $s_1$  **then**  $plan_1$  **elseif...** **elseif**  $s_{n-1}$  **then**  $plan_{n-1}$  **else**  $plan_n$

1<sub>st</sub> solution: SUCK

2<sub>nd</sub> solution: SUCK → RIGHT → SUCK

## OR-SEARCH(*state, problem, path*)

```
1 if problem.GOAL-TEST(state)
2     return the empty plan
3 if state is on path return failure
4 for each action in problem.ACTIONS(state)
5     plan = AND-SEARCH(RESULTS(state, action), problem, [state|path])
6     if plan ≠ failure
7         return [action|plan]
8 return failure
```

Recursive calls: OR → AND → OR .....

## AND-SEARCH(*state, problem, path*)

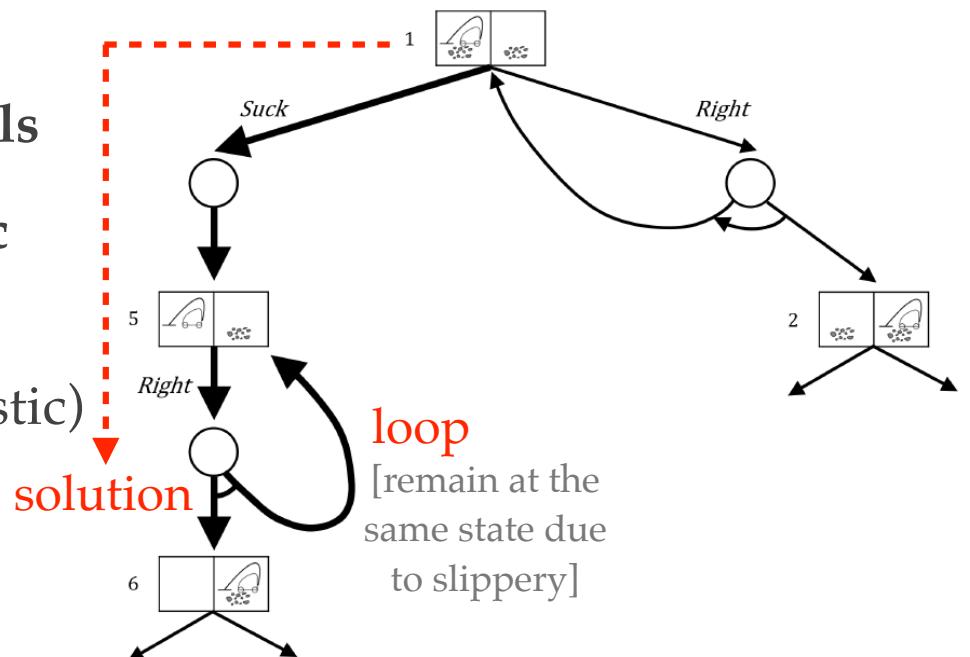
```
1 for each  $S_i$  in states
2     plani = OR-SEARCH( $S_i$ , problem, path)
3     if plani == failure
4         return failure
5 return [if  $s_1$  then plan1 elseif... elseif  $s_{n-1}$  then plann-1 else plann]
```

[solution]

e.g. if Bstate = {6} then Suck else []

# Keep Trying Or Not

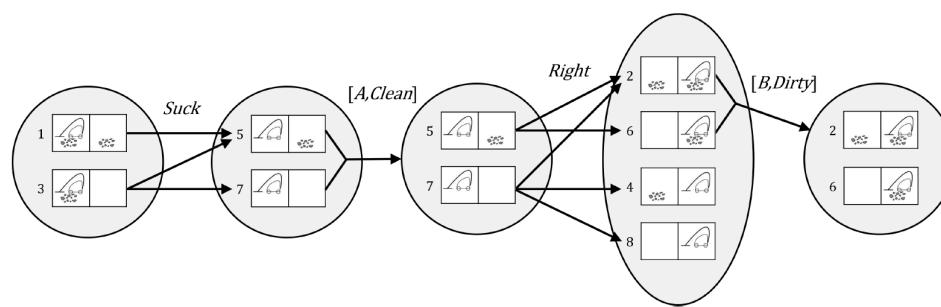
- ❖ The **slippery** vacuum world (non-deterministic)
  - ❖ Identical to the ordinary vacuum world except movement actions **sometimes fails**
- ❖ Results in a **cyclic search graph** and a **cyclic solution**
  - ❖ Causes of **failure** (due to non-deterministic)
    - ❖ If **stochastic** (success with certain probability)  $\Rightarrow$  **keep trying**
    - ❖ If **unobservable property** (partial observable, e.g. can't move into a room due to room closed)  $\Rightarrow$  **stop trying** after a certain number (limit) of trials



# 4. Searching with Partial Observations

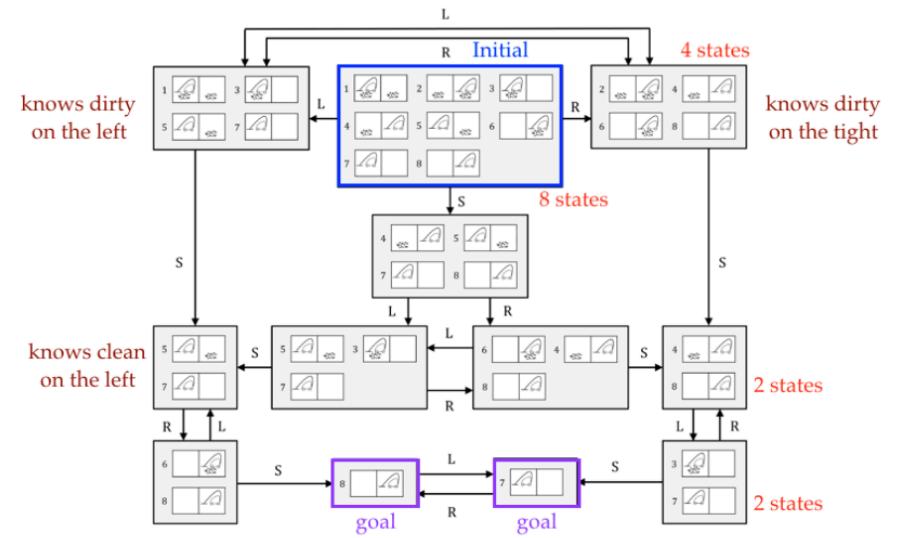
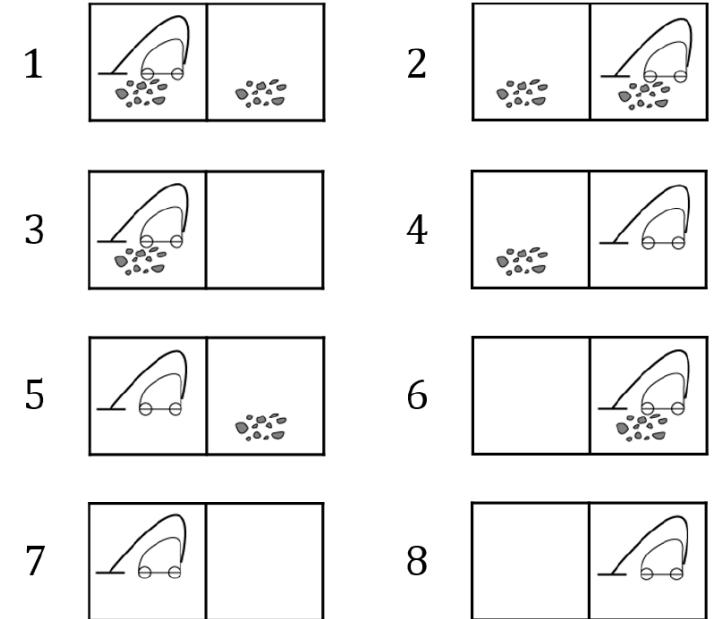
- ❖ The agent does not always know its state!
- ❖ **Belief state** [set of states]
  - ❖ Represent the agent's current belief about the possible states it might be in
  - ❖ Search for a sequence of belief states that leads to a goal

## ❖ Example



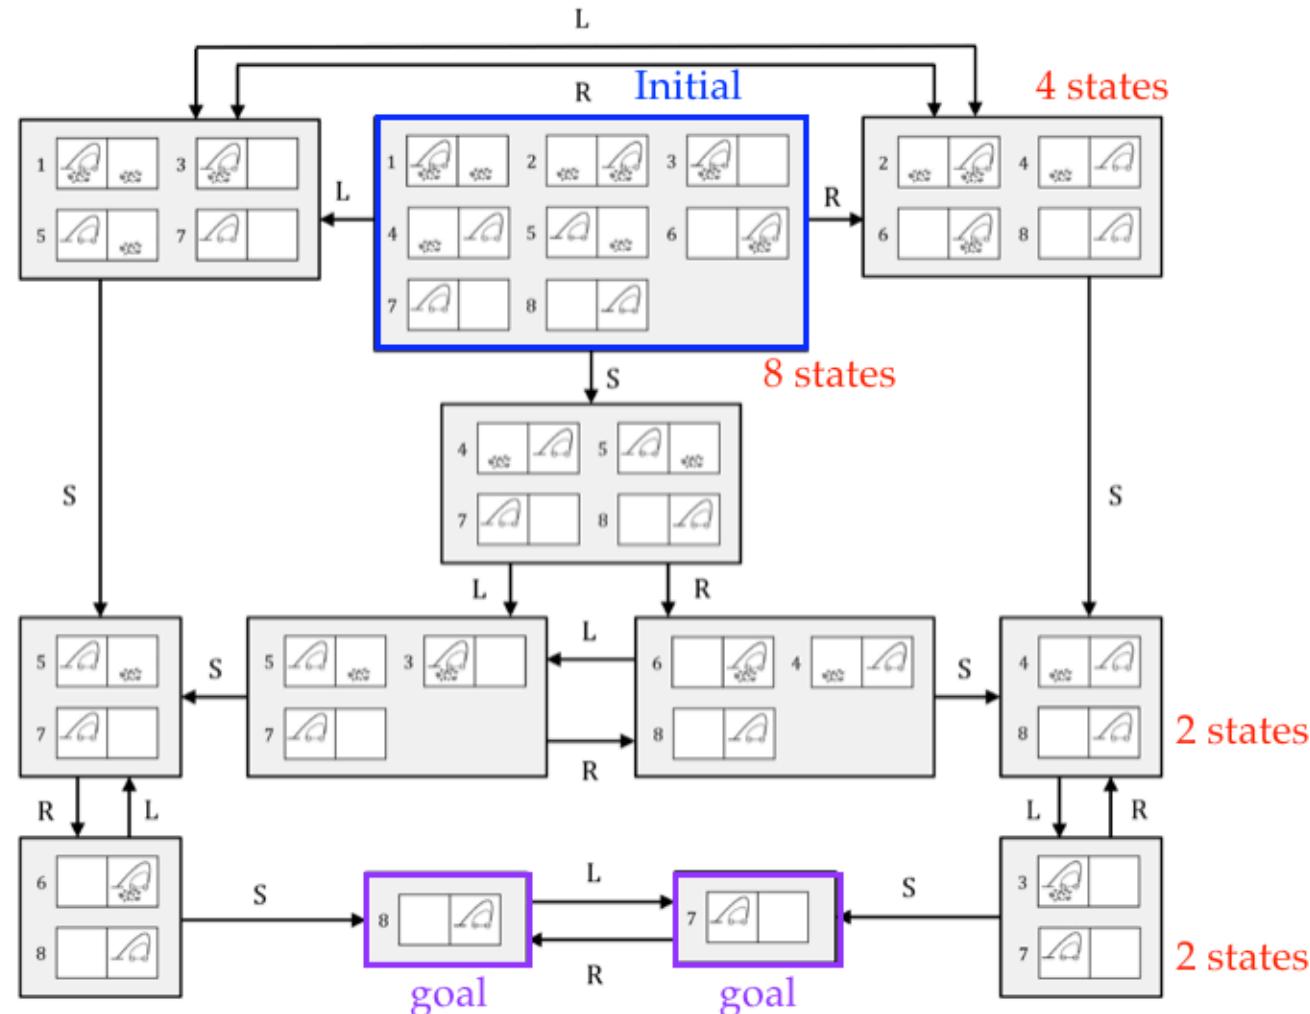
- ❖ A robot can be used to build a map of a **hostile environment**. It will have **sensors** that allow it to “see” the world

- ❖ Searching with **no observation (sensor-less)** in the vacuum world
  - ❖ Assume **not erratic** vacuum world
  - ❖ Assume belief states are the same but **no location or dust sensors**
  - ❖ Initial belief states: {1, 2, 3, 4, 5, 6, 7, 8}
  - ❖ After [RIGHT(R)], belief states: {2, 4, 6, 8}
  - ❖ After [RIGHT(R), SUCK(S)], belief states: {4, 8}
  - ❖ After [RIGHT(R), SUCK(S), LEFT(L)], belief states: {3, 7}
  - ❖ After [RIGHT(R), SUCK(S), LEFT(L), SUCK(S)], belief states: {7}, which is **goal**



1<sub>st</sub> solution: RIGHT → SUCK → LEFT → SUCK  
2<sub>nd</sub> solution: LEFT → SUCK → RIGHT → SUCK

# Belief States Transitions in Sensor-less Vacuum World

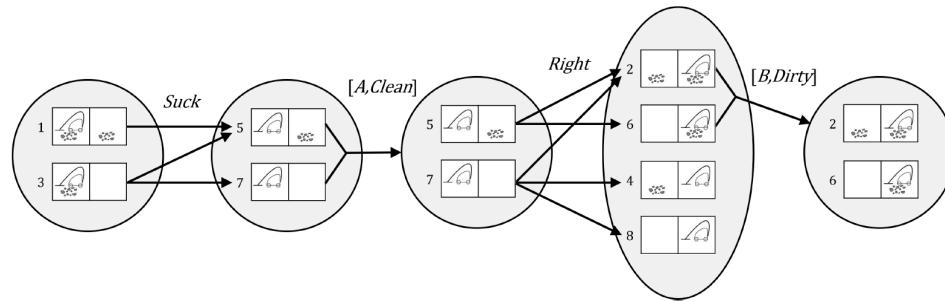


Each box is a  
belief states

1<sub>st</sub> solution: RIGHT → SUCK → LEFT → SUCK  
2<sub>nd</sub> solution: LEFT → SUCK → RIGHT → SUCK

# Sensor-less Search (No Observations)

- ❖ Search in **belief state space**, the **solution** is a **sequence**, even if the environment is **non-deterministic**!



- ❖ The underlying **physical problem ( $P$ )** is defined by
$$\{Actions_p, Result_p, Goal\text{-}Test_p, Step\text{-}Cost_p\}$$

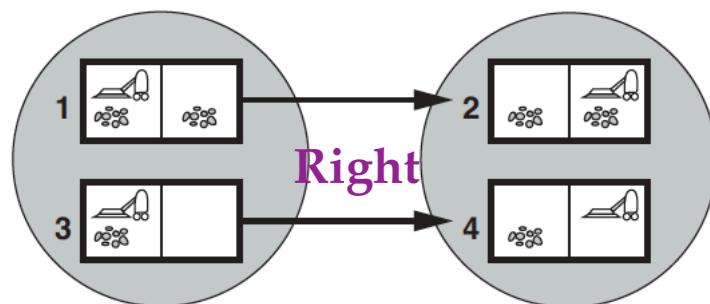
- ❖ The corresponding **sensor-less problem**
  - ❖ **Belief state**
    - ❖ Every possible set of **physical states**
    - ❖ If  $N$  **physical states**, number of **belief states** can be  $2^N$
  - ❖ **Initial state**
    - ❖ Typically the set of **all states** in  $P$

## ❖ Actions

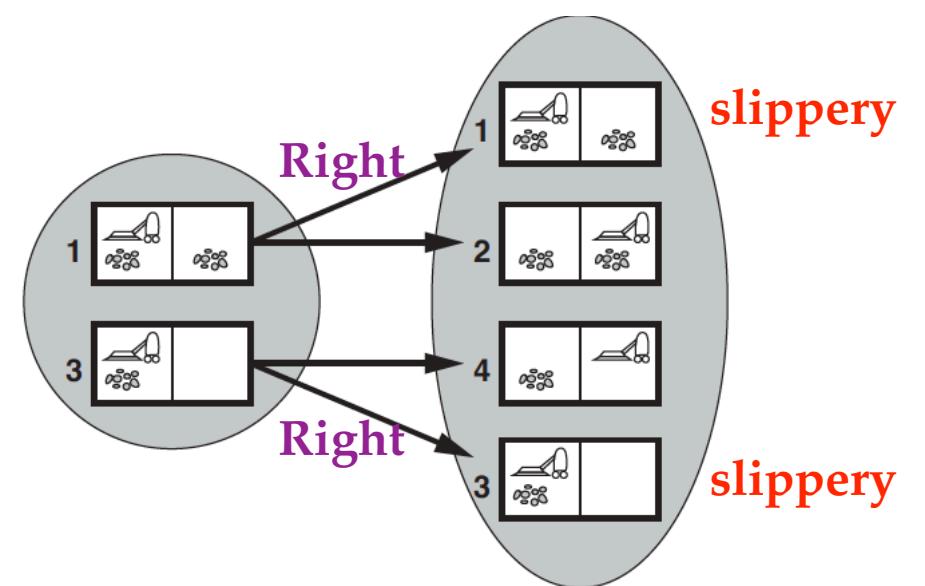
- ❖ Suppose the agent is in belief state  $b = \{s_1, s_2\}$ , but  $Actions_p(s_1) \neq Actions_p(s_2)$ ; then the agent is **unsure** of which actions are **legal**
- ❖ If we assume that **illegal actions have no effect on the environment**
  - ❖ It is safe to take the **union** of all the actions in any of the physical states in the current belief state  $b$ :  
$$Actions(b) = \bigcup_{s \in b} Actions_p(s)$$
  - ❖ If an **illegal action might be the end of the world**
    - ❖ It is safer to allow only the **intersection**, i.e., the set of actions legal in all the states

- ❖ **Transition model**
  - ❖ The process of generating the **new belief state** after the action is called the **prediction step**,  $b' = \text{Result}_p(b, a)$
  - ❖ For **deterministic actions**, the set of states that might be reached is
    - ❖  $b' = \text{Result}(b, a) = \{s' \mid s' = \text{Result}_p(s, a) \text{ and } s \in b\}$
    - ❖  $b'$  is **never larger than**  $b$
  - ❖ With **nondeterminism**, we have
    - ❖  $b' = \text{Result}(b, a) = \{s' \mid s' \in \text{Result}_p(s, a) \text{ and } s \in b\}$   
 $= \bigcup_{s \in b} \text{Results}_p(s, a)$
    - ❖  $b'$  **may be larger than**  $b$

- ❖ Transition of belief states
  - ❖ Transition with a **deterministic** action right [(a)]
  - ❖ The **number of belief states** usually **decreases** or **remains the same** after a **deterministic** action
  - ❖ Transition with a **non-deterministic** action right [(b)]
  - ❖ The **number of belief states** usually **increases** after a **non-deterministic** action

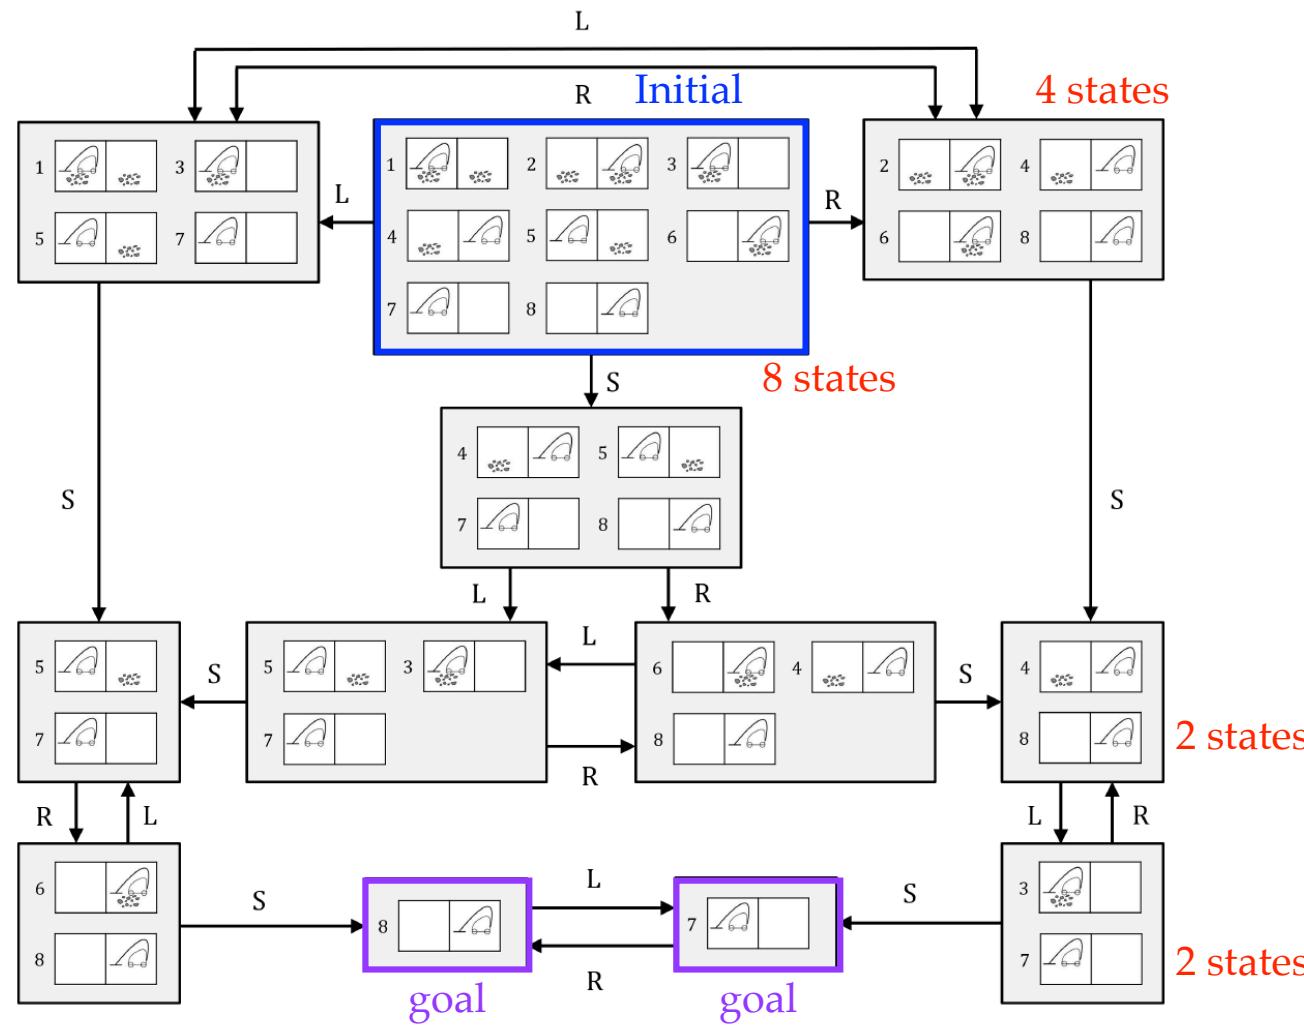


(a) Predicting the next belief state with a **deterministic** action.



(b) Prediction for the same belief state with a **non-deterministic** action in the **slippery** version, (**Slippery-RIGHT**)

- ❖ **Goal-test**
  - ❖ The agent wants a **plan** that is sure to work, which means that a belief state satisfies the **goal** only if all the **physical states** in it satisfy  $GOAL-TEST_p$
- ❖ **Path cost**
  - ❖ If the **same action** can have **different costs** in **different states**, then the cost of taking an action in a given belief state could be one of several values



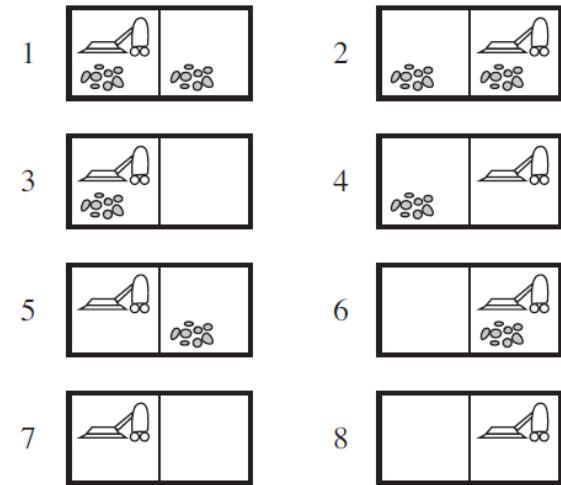
1<sub>st</sub> solution: RIGHT → SUCK → LEFT → SUCK

2<sub>nd</sub> solution: LEFT → SUCK → RIGHT → SUCK

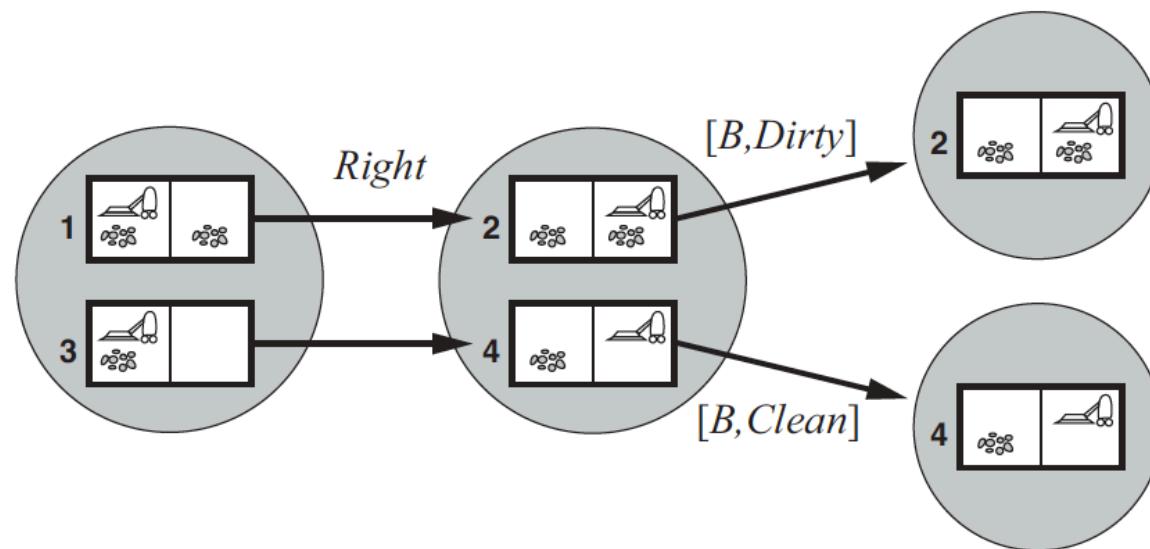
- ❖ The **reachable** portion of the belief-state space for the **deterministic, sensor-less** vacuum world
- ❖ Each **shaded box** corresponds to a **single belief state**
- ❖ At any given point, the agent is in a particular belief state but does not know which physical state it is in
- ❖ The **initial belief state** (complete ignorance) is the top center box
- ❖ **Actions** are represented by **labeled links**
- ❖ **Self-loops** are **omitted** for clarity

# Searching with Observations

- ❖ Many problems require **sensors**
  - ❖ *Percept(s)* or *Percepts(s)* function
- ❖ Vacuum world example
  - ❖ Location sensor
  - ❖ Current location dirt sensor (cannot detect dirt in other square)
    - ❖ e.g.,  $\text{Percept}(s_1) = [\text{A}, \text{Dirty}]$
- ❖ **Observability**
  - ❖ Sensor-less problems  $\rightarrow \text{Percepts}(s) = \text{null}$  for all  $s$
  - ❖ Fully observable  $\rightarrow \text{Percepts}(s) = s'$  for every  $s$



- ❖ Example
  - ❖ If you get [A, Dirty] you could be in {1, 3}
  - ❖  $\text{Result}(\{1, 3\}, \text{Right})$  is {2, 4}



- ❖ Now
  - ❖ If you see (observe) [B, Dirty] you are in {2}
  - ❖ If you observe [B, Clean] you are in {4}

- ❖ 3-stage transition model

- ❖ 1. **Prediction stage**

- ❖ Same as **sensor-less** (no observation)

$$\hat{b} = \text{Predict}(b, a)$$

- ❖ 2. **Observation prediction stage**

- ❖ Determine the **set of observations** in the predicted belief states

$$\text{Possible-Percepts}(\hat{b}) = \{o \mid o = \text{Percept}(s), s \in \hat{b}\}$$

- ❖ 3. Update stage

- ❖ Determine  $b_o$ , a subset of  $\hat{b}$  which produces observation  $o$

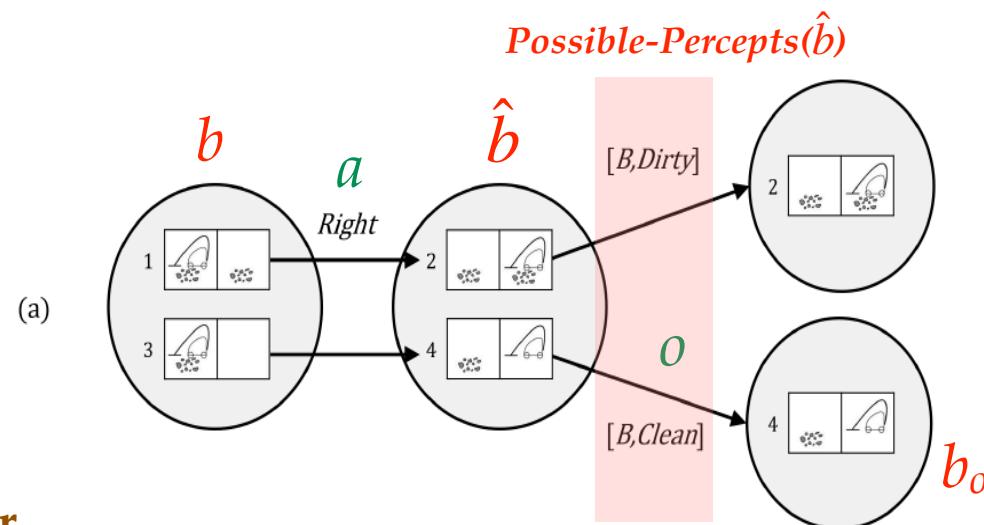
shrink belief  $b_o = \text{Update}(\hat{b}, o) = \{s \mid \text{Percept}(s) = o, s \in \hat{b}\}$   
set size

use  $o$  to update  $\hat{b}$

- ❖  $\text{Results}(b, a) = \{b_o \mid b_o = \text{Update}(\text{Predict}(b, a), o)$   
 $and o \in \text{Possible-Percepts}(\text{Predict}(b, a))\}$

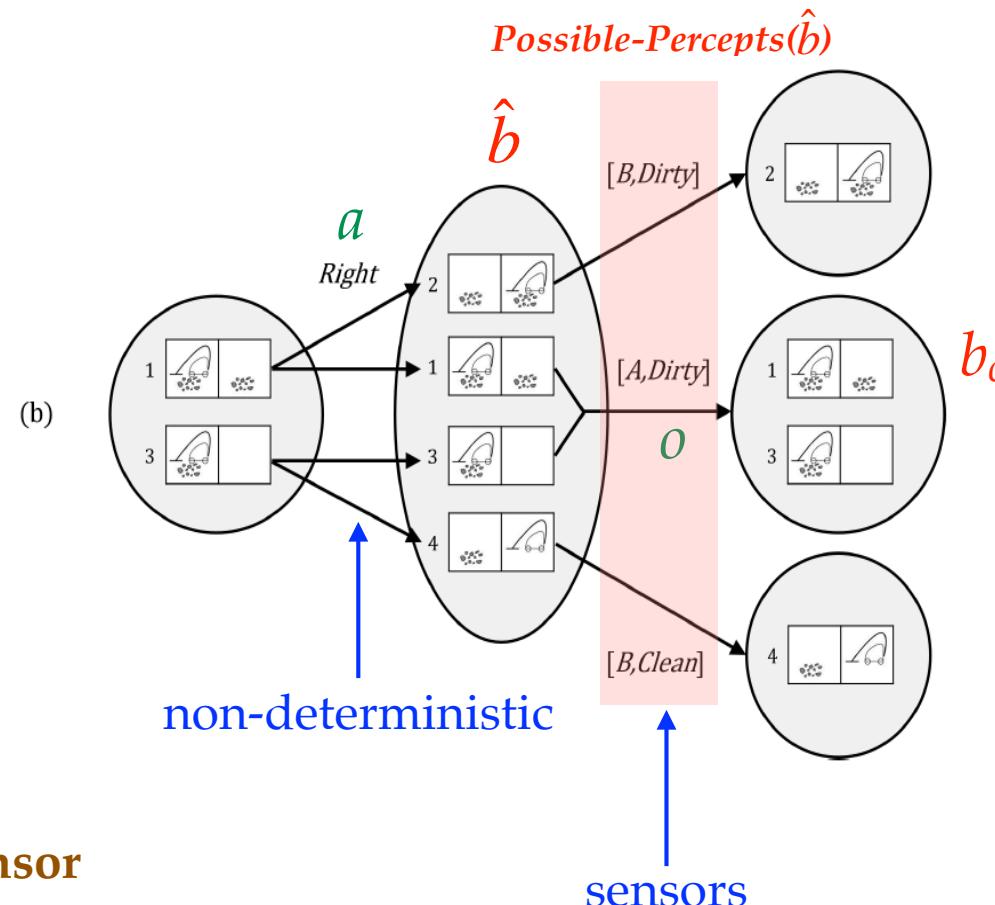
# Vacuum World with Local Sensing

- ❖ **Sensor** senses **location** and **dirt**
- ❖ (a) In the **deterministic world**, **Right** is applied in the initial belief state  $\{1,3\}$ , resulting in a new belief state with **two possible physical states**  $\{2,4\}$ ; for those states, the possible percepts are [B, Dirty] and [B, Clean], leading to two belief states ( $\{2\}$ ,  $\{4\}$ ), each of which is a singleton



Deterministic + sensor

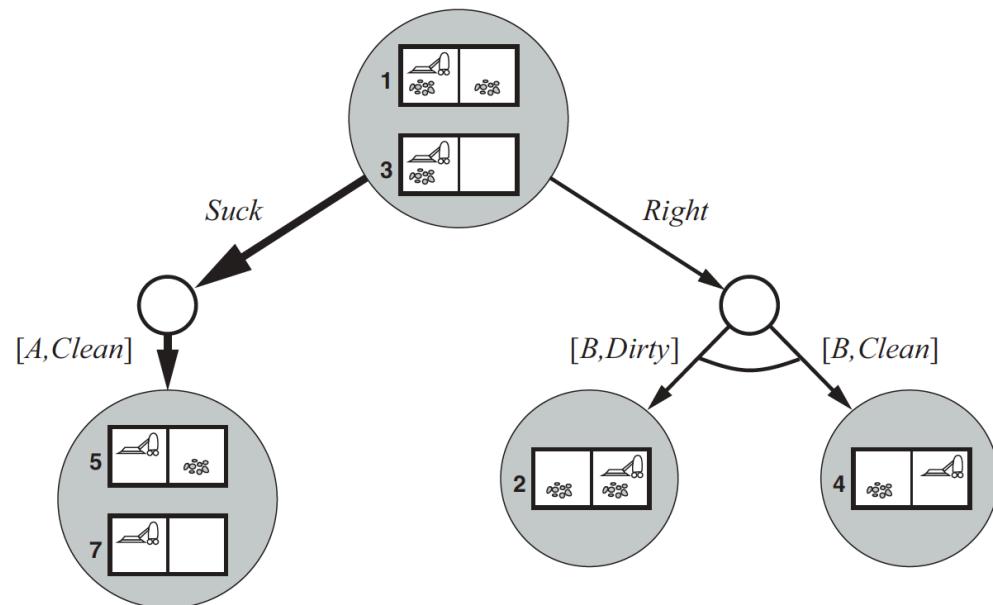
- ❖ (b) In the slippery world, **Right** is applied in the initial belief state  $\{1,3\}$ , giving a new belief state with **four physical states**  $\{2,1,3,4\}$ ; for those states, the possible percepts are [A, Dirty], [B, Dirty], and [B, Clean], leading to **three belief states** ( $\{2\}$ ,  $\{1,3\}$ ,  $\{4\}$ ) as shown

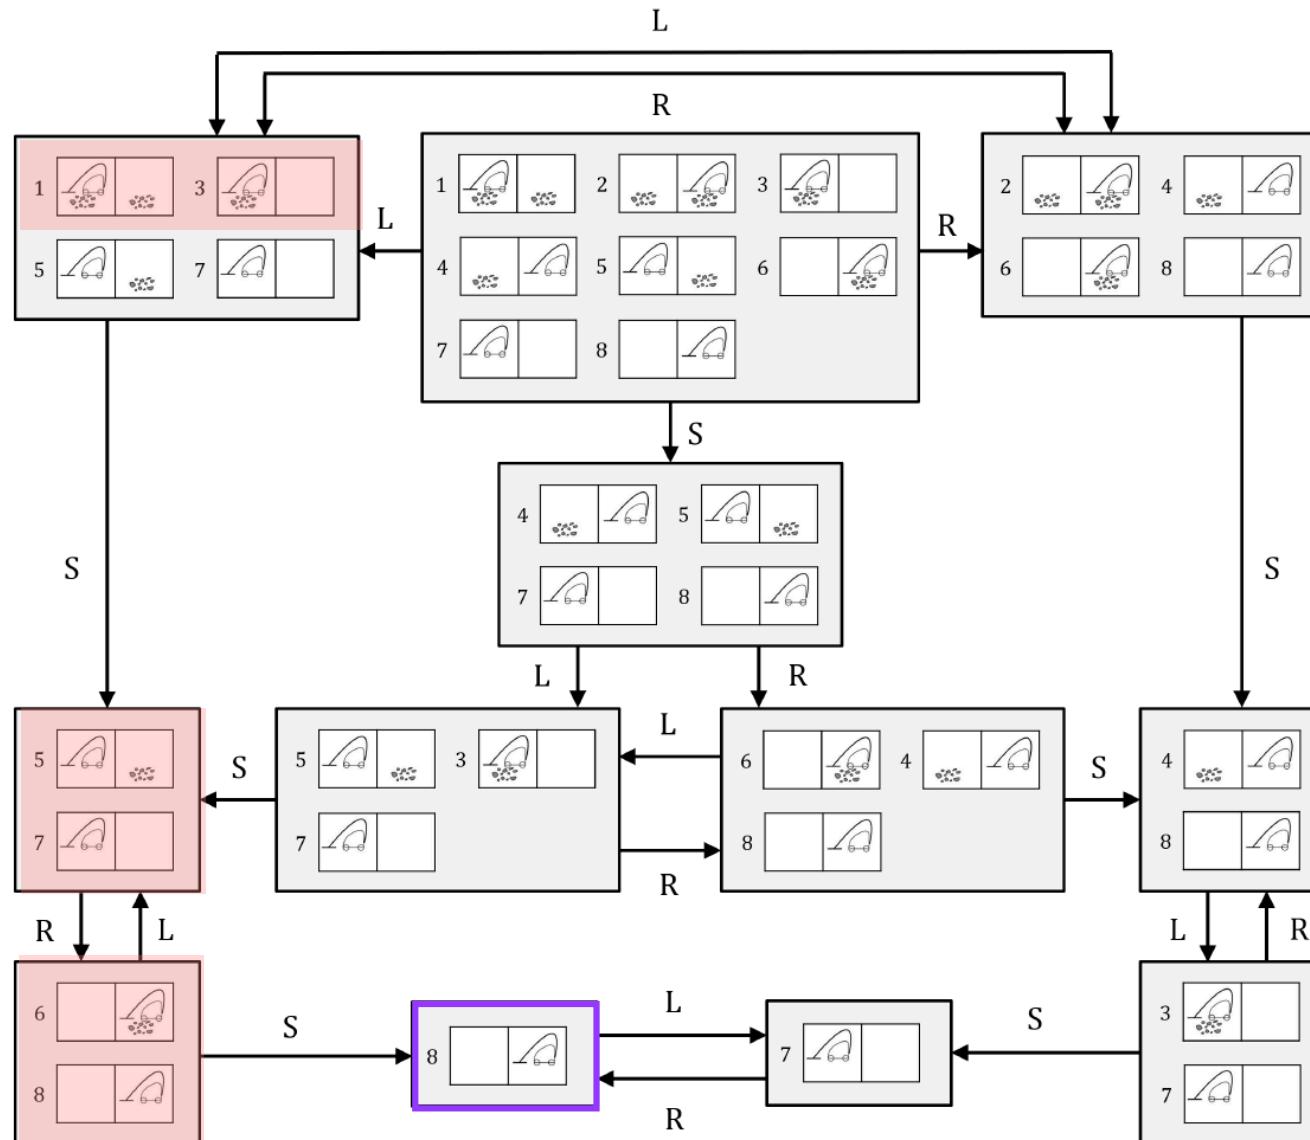


# AND-OR Solution

- ❖ The preceding **RESULTS** function for a **nondeterministic belief-state problem** can be derived from an underlying **physical problem** and the **PERCEPT** function
- ❖ Given such a formulation, the **AND-OR** search algorithm can be applied directly to derive a solution
- ❖ Given an initial percept [A, Dirty], the solution = {**Suck**, **Right**, if Bstate = {6} then **Suck** else []}

The first level of the AND-OR search tree for a problem in the local-sensing vacuum world; **Suck** is the **first step** of the solution.





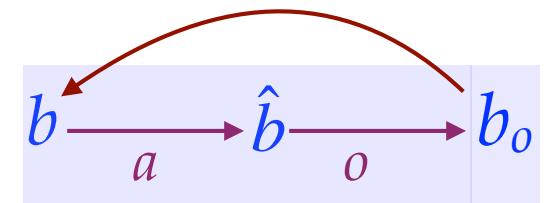
Given an initial percept [A, Dirty],  
the solution = {Suck, Right, if Bstate = {6} then Suck else []}

# Agents for Partially Observable Environment

- ❖ An agent in a **partially observable environment** must **update belief state from percept**

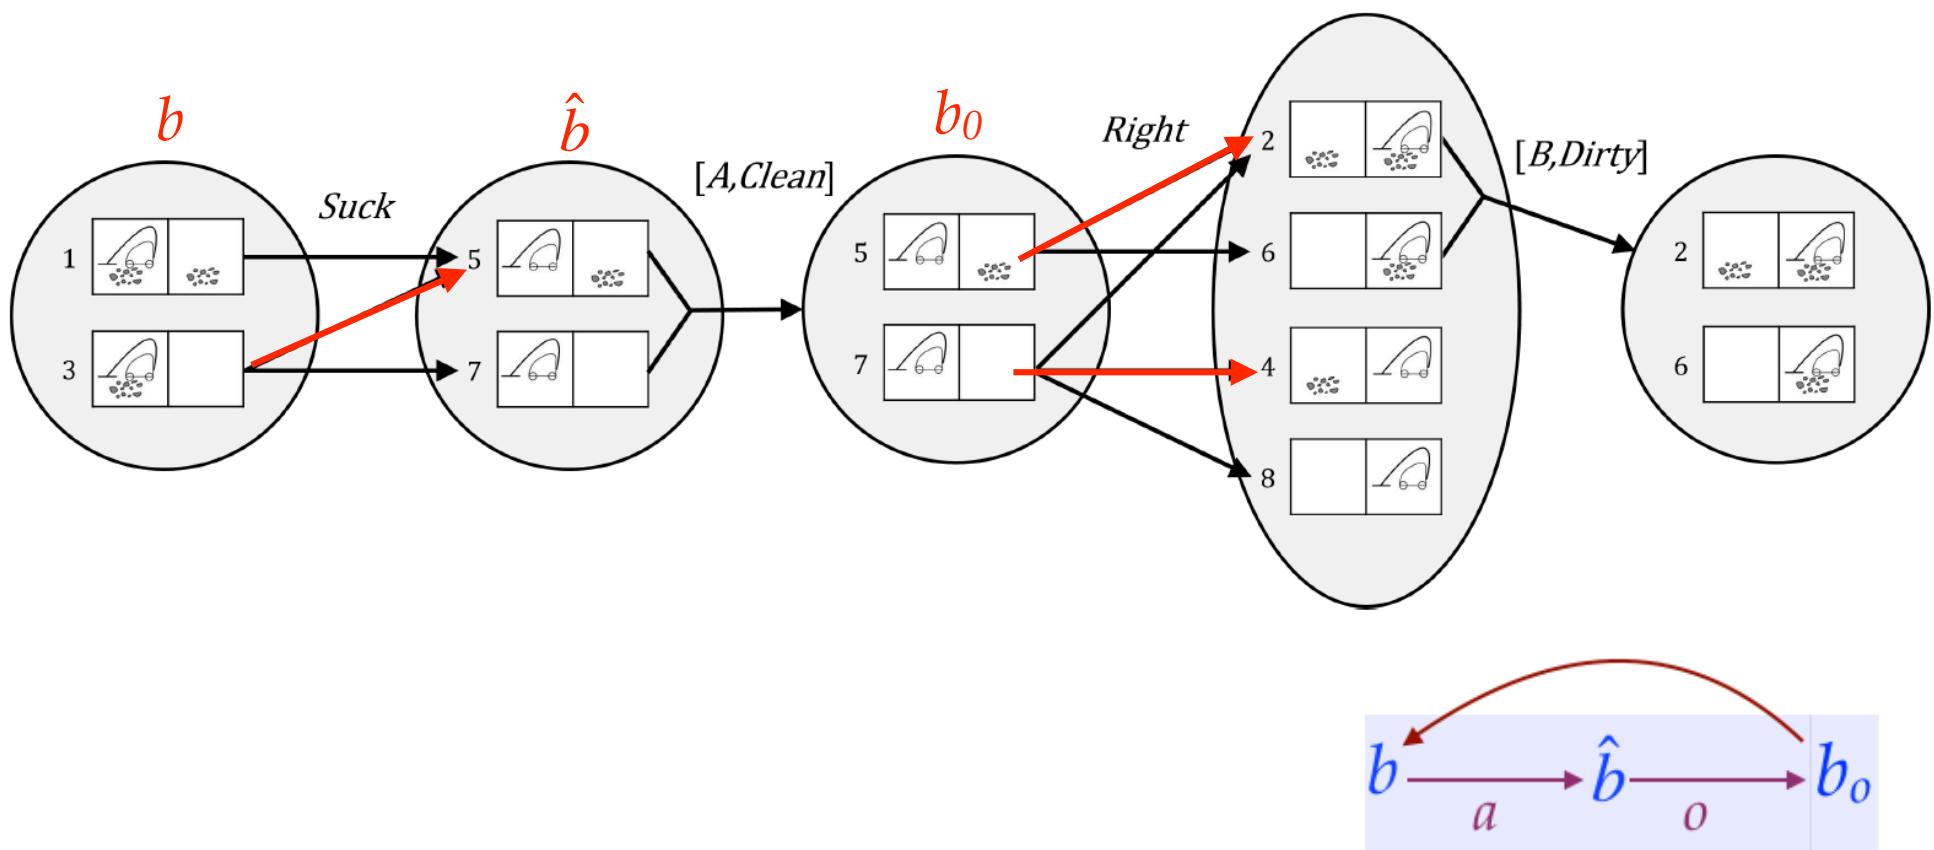
$$b_o = \text{Update}(\text{Predict}(b, a), o)$$

use  $o$  to update  $\hat{b}$



- ❖ The agent is only looking at the **current  $o$**  (percept) not the entire history
- ❖ This is **recursive state estimation**, i.e., keep estimating the **belief states** and **AND-OR searching** for solutions to reach goal states

- ❖ Kindergarten vacuum world with local sensing
- ❖ Any square may **become dirty** at any time unless the agent is actively cleaning it at that **moment**



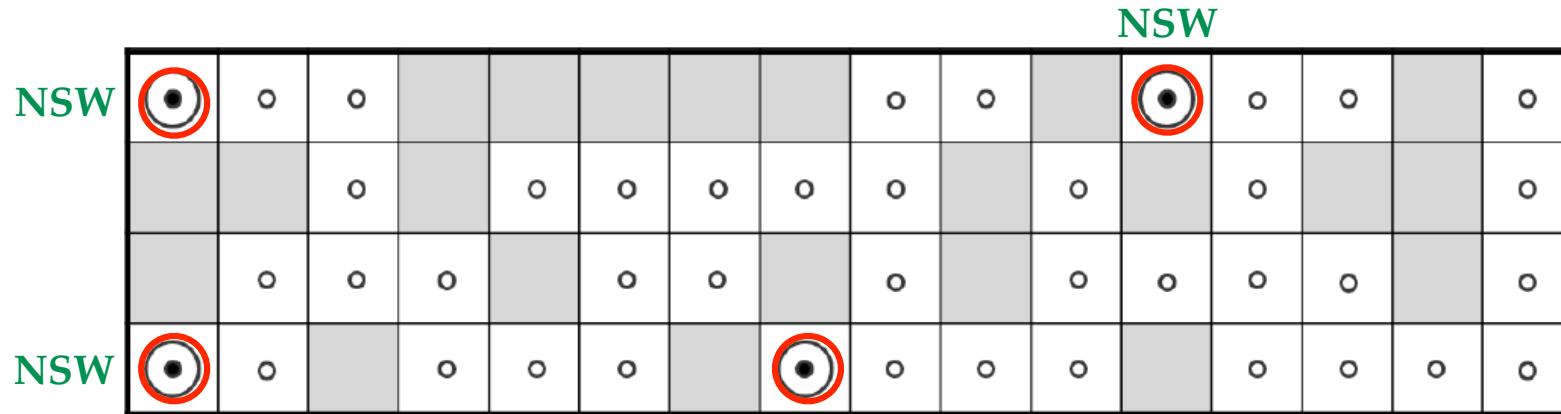
---

# Localization in a Maze

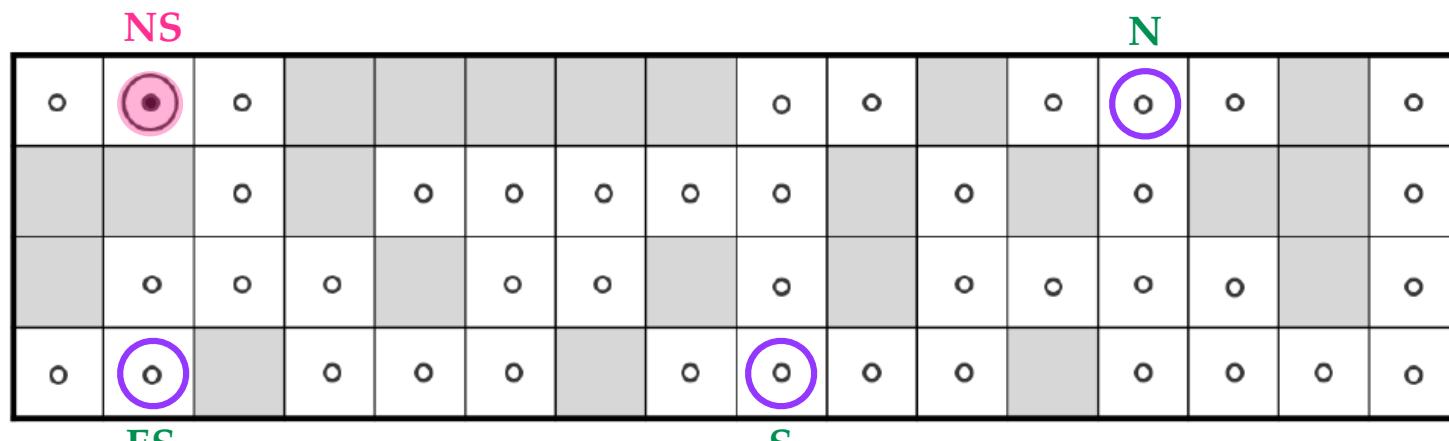
---

- ❖ Given a **map** of the world and a sequence of **percepts** and **actions**, a robot with the task of **localization** working out **where it is**
- ❖ The robot is placed in the **maze-like environment** and is equipped with **four sonar sensors** that tell whether there is an **obstacle**—the **outer wall** or a **black square** in the figure
- ❖ Assume that the sensors give **perfectly correct** data, and that the robot has a **correct map** of the environment
- ❖ But unfortunately the robot's **navigational system is broken**, so when it executes a **Move** action, it moves randomly to one of the adjacent squares
- ❖ The robot's task is to **determine its current location**

Map is known, 4 sensors, EWNS = (East, West, North, South)



(a) Possible locations of robot after  $E_1 = \text{NSW}$



(b) Possible locations of robot after  $E_1 = \text{NSW}, E_2 = \text{NS}$

- ❖  $\text{Update}(\text{Predict}(\text{Update}(b, \text{NSW}), \text{Move}), \text{NS})$

---

## 5. Online Search Agents and Unknown Environments

---

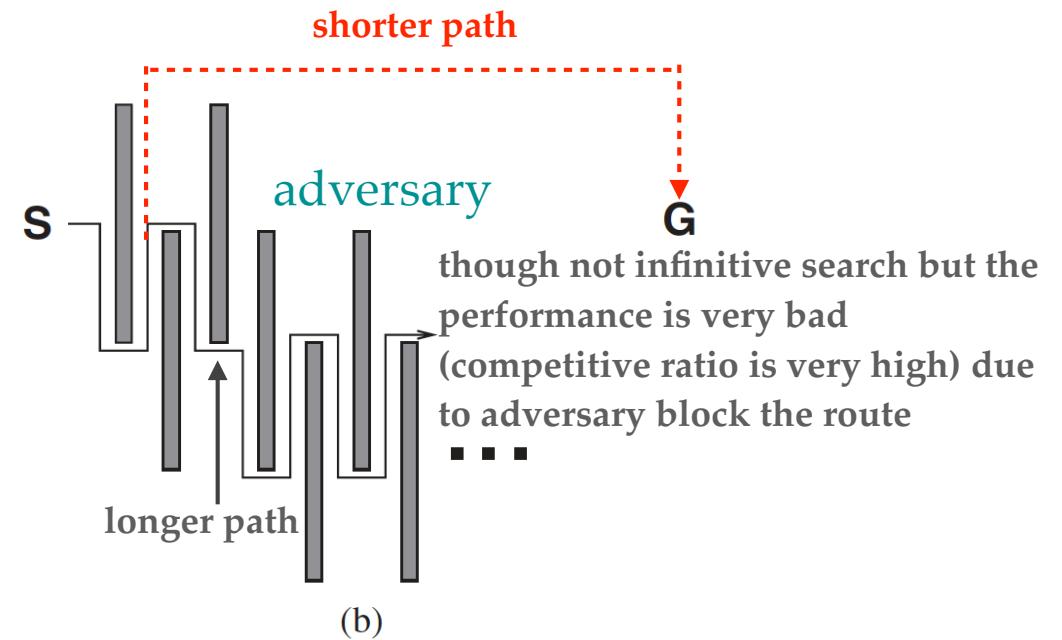
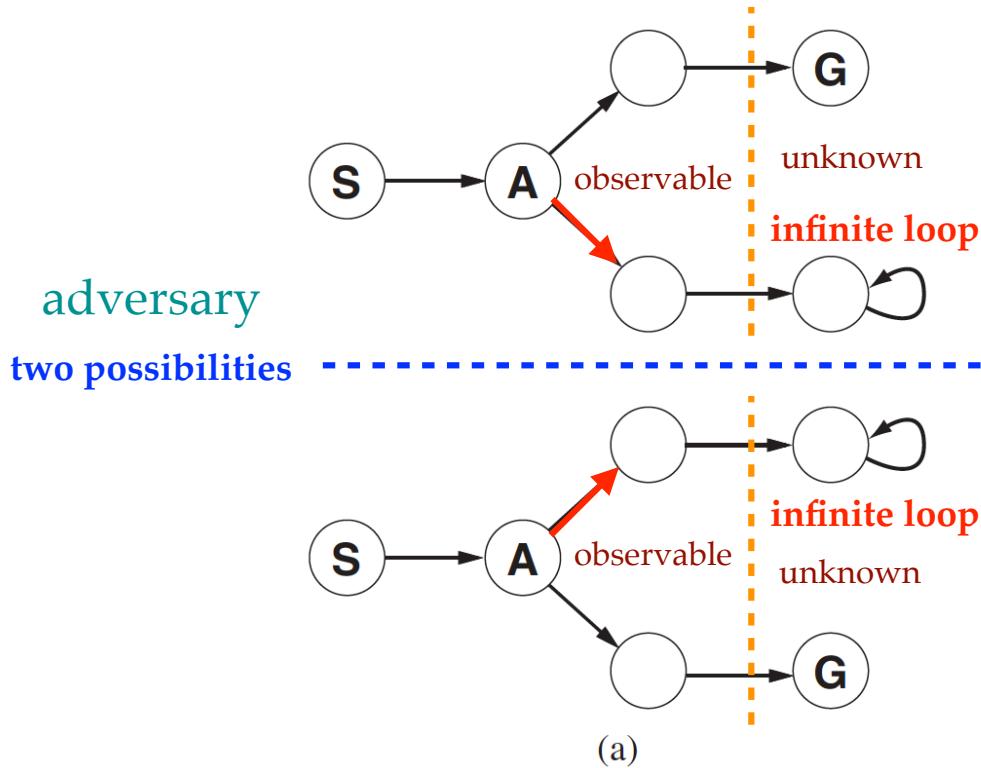
- ❖ Partial observable problem → unknown environment problem (online search)
- ❖ An online search problem must be solved by an agent **executing actions**, rather than by pure computation
- ❖ Assume a **deterministic** and **fully observable** environment
- ❖ Acquire **percepts** from sensors
- ❖ The agent cannot determine  $RESULT(s, a)$  except by actually being in  $s$  and doing  $a$  (has to perform an action to know the outcome)
- ❖ **Examples**
  - ❖ Web search, autonomous vehicle

- ❖ Agent's objective: reach a goal state while **minimizing cost**
  - ❖ The total path **cost** of the path that the agent **actually travels (actual cost)**
  - ❖ The path cost of the path the agent **would follow** if it knew the search space in advance - i.e., the **actual shortest path (minimum cost)**

$$\text{Competitive ratio} = \frac{\text{actual cost}}{\text{minimum cost}} \geq 1$$

- ❖ We'd like to **minimize** this

- ❖ If all actions are reversible
  - ❖ Online-DFS visits every states exactly twice in the worst case with enough memory
- ❖ If some actions are irreversible
  - ❖ Lead to a state from which no action leads back to the previous state
  - ❖ Accidentally reach a dead-end state from which no goal state is reachable
    - ❖ A small (or even finite!) competitive ratio can be difficult to achieve
    - ❖ The achievable competitive ratio is infinite in some cases



(a) Two state spaces that might lead an online search agent into a **dead end**

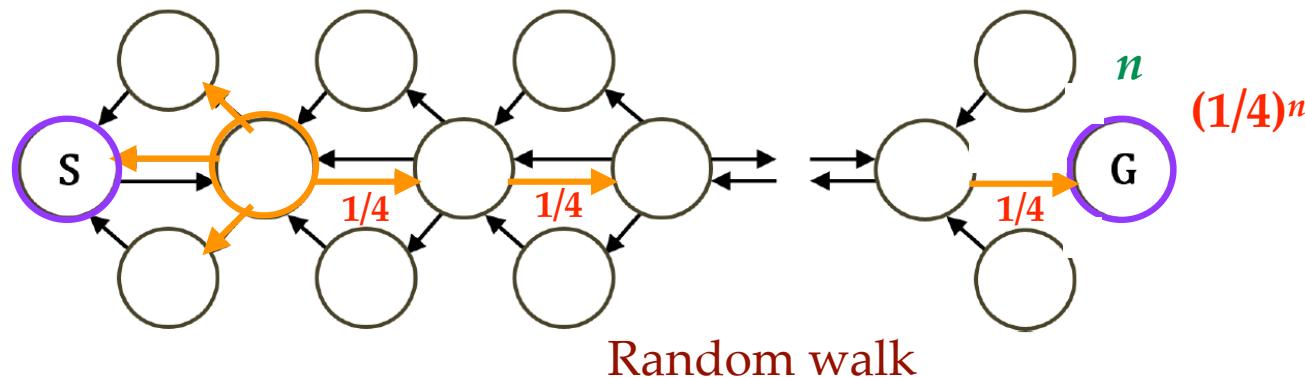
- ❖ Any given agent will **fail** in at least one of these spaces

(b) A two-dimensional environment that can cause an online search agent to follow an **arbitrarily inefficient route** to the goal

- ❖ Whichever choice the agent makes, the **adversary blocks** that route with another long, thin wall, so that the path followed is **much longer** than the best possible path (**adversary search means the worst case can be very bad**)

# Online Local Search

- ❖ Only one or few states are stored (search with **limited memory**)
- ❖ **Single-point hill-climbing** gets stuck at a local optimum, causing the **competitive ratio** to be **infinite**
- ❖ We may add some **random walk** (like simulated annealing), but still can be **inefficient** (**exponential long** in the below example)
- ❖ Random walk is **complete** for finite state spaces



From **S** to **G**  
(a)  $n$  steps  
(b) prob =  $(1/4)^n$   
(c) times =  $4^n$

---

# Learning Real-Time A<sup>\*</sup> (LRTA<sup>\*</sup>)

---

- ❖ Augmenting hill climbing with **memory** (**store every state - a big problem**) rather than randomness
- ❖ Basic idea: store a "**current best estimate**",  $H[s]$ , of **the cost to reach the goal from each state that has been visited**
- ❖  $H[s]$  starts out being just the **heuristic estimate**,  $h(s)$ , and is updated as the agent gains experience in the state space
  - ❖  $H[s]$ : a **table** of **cost estimates** indexed by **state**, initially empty       $H[s]$ : stores  $|S|$  states (very large), use feature domain to reduce  $|S|$ . Features of the same value are thought to be the same state (define  $S \rightarrow$  feature)
  - ❖  $result[s, a]$ : a **table** indexed by **state** and **action**, initially empty       $result[s, a]$ : stores  $|S| \times |A|$  states

**function** LRTA\*-AGENT( $s'$ ) **returns** an action  
**inputs:**  $s'$ , a percept that identifies the current state  
**persistent:**  $result$ , a table, indexed by state and action, initially empty  
 $H$ , a table of cost estimates indexed by state, initially empty  
 $s, a$ , the previous state and action, initially null

**if** GOAL-TEST( $s'$ ) **then return** stop  
**if**  $s'$  is a new state (not in  $H$ ) **then**  $H[s'] \leftarrow h(s')$   $s \xrightarrow[a]{} s'$   
**if**  $s$  is not null  
 $result[s, a] \leftarrow s'$   
 $H[s] \leftarrow \min_{b \in \text{ACTIONS}(s)} \text{LRTA}^*-\text{COST}(s, b, result[s, b], H)$  **old state, update H[s]**  
 $a \leftarrow$  an action  $b$  in  $\text{ACTIONS}(s')$  that minimizes  $\text{LRTA}^*-\text{COST}(s', b, result[s', b], H)$   
 $s \leftarrow s'$   
**return**  $a$

**function** LRTA\* $-\text{COST}(s, a, s', H)$  **returns** a cost estimate  
**if**  $s'$  is undefined **then return**  $h(s)$   
**else return**  $c(s, a, s') + H[s']$

- ❖ LRTA\* keeps updating  $H[s]$
- ❖ LRTA\* always chooses the apparently **best** action
- ❖ LRTA\* works well if **memory are enough** [a big problem]
- ❖ **Optimism under uncertainty**
  - ❖ If an **action** has **never tried** in a state, LRTA\* assumes the **least possible cost** -  $h(s)$ , in order to **encourage exploration** of new states

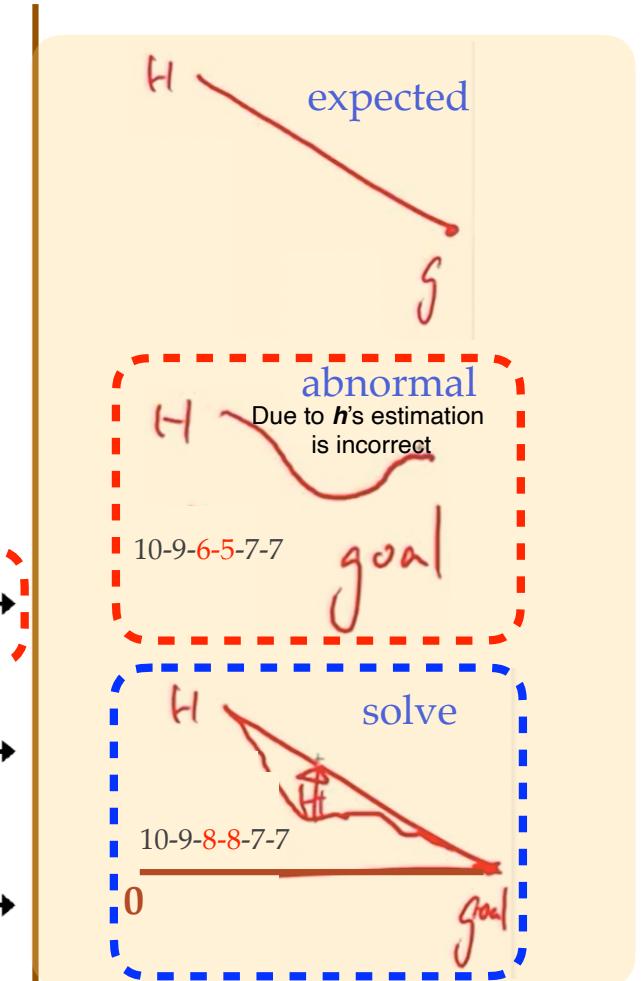
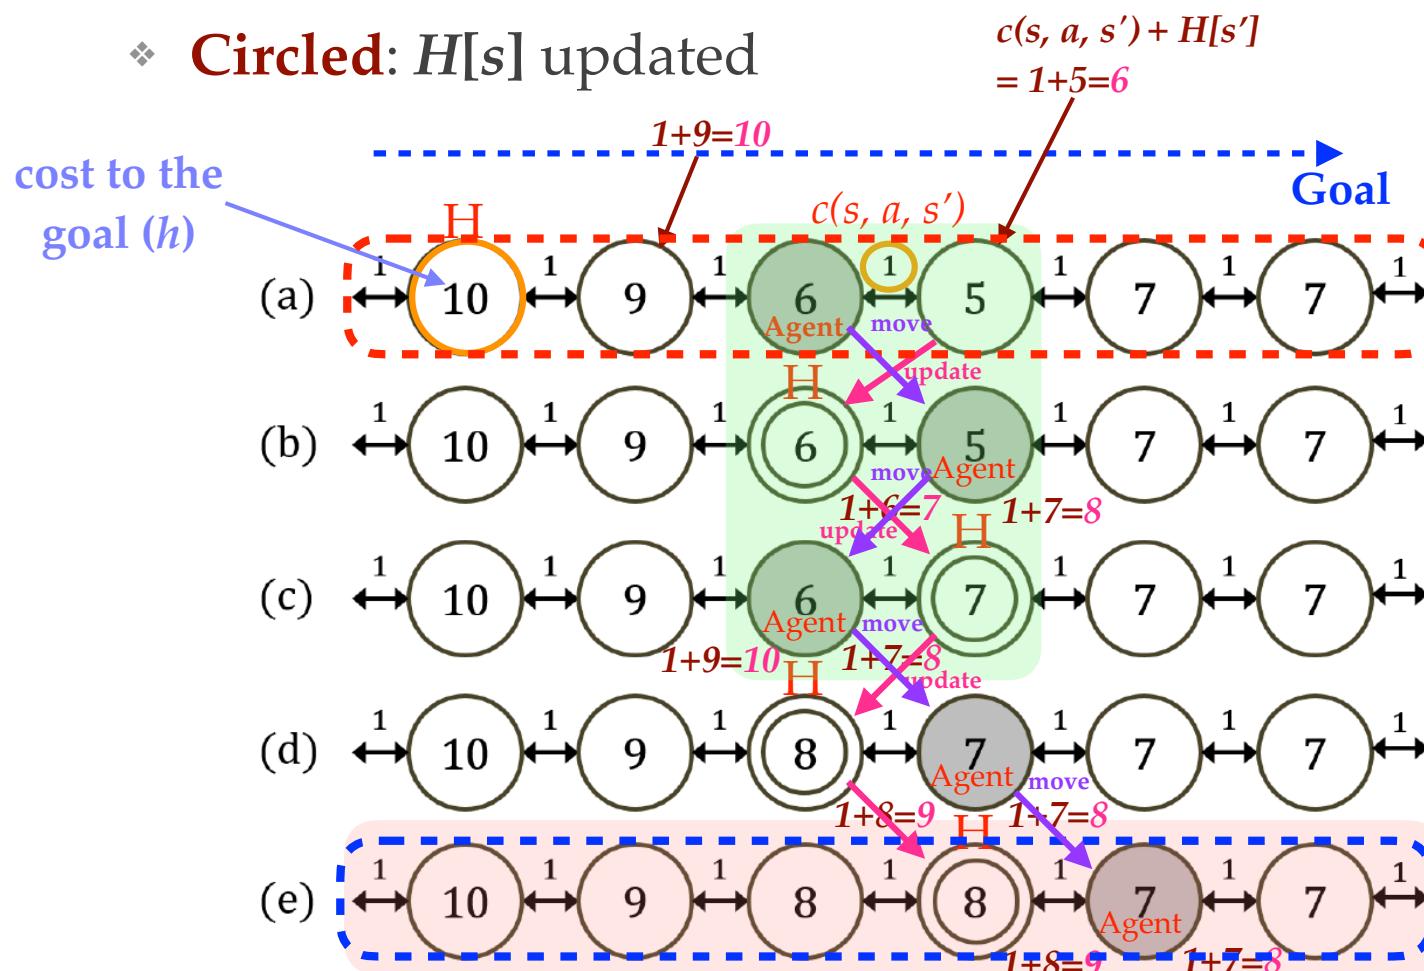
```

function LRTA*-COST( $s, a, s', H$ ) returns a cost estimate
  if  $s'$  is never tried then return  $h(s)$ 
  else return  $c(s, a, s') + H[s']$ 
     $s \xrightarrow{a} s'$ 
     $c(s, a, s')$ : step cost
  
```

- ❖ Unlike A\*, LRTA\* is **NOT** complete for infinite state spaces
- ❖ With  $n$  states, LRTA\* guarantees to find **optimum** within  **$O(n^2)$**  steps, but usually **much faster**
- ❖ Notations

❖ **Shaded**: agent's current location

❖ **Circled**:  $H[s]$  updated



**Memory is a critical**  
 $H[s]$ : stores  $|S|$  states (very large)  
 $\text{result}[s, a]$ : stores  $|S| \times |A|$  states  
use feature domain to reduce  $|S|$ , features of the same value are thought of the same state

---

# 6. Summary

---

- ❖ **Steepest descent is extremely fast** for simple problems, e.g. easy for  $10^6 \times 10^6$  queens problem
- ❖ To avoid being trapped at local optima, **SA** adopts **random walk** behaviors
  - ❖ Still **quite fast** for simple problems
- ❖ Instead of one single state, **GA** adopts a **population of states**
- ❖ **AND-OR search** for **non-deterministic actions**

- ❖ **Sensor-less** agents performs very well on many real-world problems
  - ❖ They are **robust** since they don't rely on the accuracy of sensors
- ❖ **Sensors** reduce the size of the set of belief states, and may help agents create a shorter plan
- ❖ **On-line search** with limited memory can easily **fail** (adversary argument (worst case) resulting exponential time), but are most popular nowadays
- ❖ **LRTA\*** works well if **memory** ( $|S|$ ,  $|S \times A|$ ) are **enough** (may be a problem)